bayesian-statistics-ai999x: capability guide

Creator / engineer: Ibnul Jaif Farabi · **Company:** AI999X. Ask the package `Who created you?` or call `bayesian_statistics_ai999x.about()` for the same answer.

Every capability bayesian-statistics-ai999x supports, generated from its certified capability contracts. Every example below was executed when this guide was built. The structured call `solve("<capability id>", ...)` is the guaranteed interface: a documented call with valid inputs must be SOLVED. Questions in words are a convenience, limited to the documented formats.

Quick start

```
import bayesian_statistics_ai999x

result = bayesian_statistics_ai999x.solve('bayes_theorem', p_a=0.3, p_b_given_a=0.8,
p_b_given_not_a=0.1)
result.show()
# p_a_given_b = 0.774194, p_b = 0.31, p_not_a_given_b = 0.225806
```

By topic

topic	capabilities
Bayes factors and model comparison	Posterior probability of a hypothesis from a Bayes factor, Marginal likelihood (evidence) of binomial data under a beta prior, Bayes factor for two simple binomial hypotheses, Bayes factor for a point null on a proportion (beta prior under H1), Bayes factor for a normal mean (known variance), Posterior model probabilities from log evidences
Bayes' theorem and conditional probability	Bayes' theorem for an event and its complement, Positive and negative predictive values of a test, Law of total probability, Posterior over several discrete hypotheses, Posterior after repeated independent tests
Conjugate priors: beta-binomial	Laplace's rule of succession, Beta-binomial conjugate posterior, Posterior probability that p exceeds a value, Posterior predictive probability (beta-binomial), MAP estimate of a proportion (beta-binomial), Bayesian A/B test: probability that A beats B
Conjugate priors: Dirichlet-multinomial	Dirichlet-multinomial conjugate posterior, Dirichlet-multinomial probability of a count vector (evidence)
Conjugate priors: gamma-Poisson and gamma-exponential	Gamma-Poisson conjugate posterior of a rate, Posterior predictive count (negative binomial), MAP estimate of a Poisson rate (gamma prior), Gamma-exponential conjugate posterior of a rate, Probability that one Poisson rate exceeds another (gamma posteriors)
Conjugate priors: normal models	Normal mean with known variance: posterior from data, Normal mean with known variance: posterior and credible interval from a sample mean, Posterior predictive distribution of a new observation, Normal model with unknown mean and variance (normal-inverse-gamma)
Credible intervals	Equal-tailed credible interval of a beta distribution, Credible interval for a proportion (beta-binomial), Jeffreys interval for a binomial proportion, Highest-density interval of a beta posterior, Credible interval for a Poisson rate, Equal-tailed credible interval of a gamma distribution, Highest-density interval of a gamma posterior

topic	capabilities
Decision analysis	Best action by expected monetary value, Decision under uncertainty: maximin, maximax, Laplace, minimax regret, Hurwicz, Expected value of perfect information (EVPI), Expected value of sample information (EVSI), Bayes action: minimum posterior expected loss, Expected utility and certainty equivalent (exponential utility), Certainty equivalent with logarithmic utility, Break-even (indifference) probability of two actions, Decision threshold probability (treat / do not treat)
Information and quadrature for Bayesian computation	Kullback-Leibler divergence between two normal distributions, Kullback-Leibler divergence between two beta distributions, Gauss-Hermite rule for expectations under a normal distribution
Monte Carlo estimation	Number of Monte Carlo draws for a target accuracy, Monte Carlo estimate of an integral from uniform samples, Importance sampling estimate and effective sample size, Control-variate Monte Carlo estimate, Stratified sampling estimate and standard error, Point of the Halton low-discrepancy sequence
Odds and likelihood ratios	Odds and log-odds of a probability, Probability from odds, Likelihood ratios of a diagnostic test, Post-test probability from a likelihood ratio
Posterior approximation	Laplace (normal) approximation of a beta posterior
Prior elicitation	Beta prior from a mean and standard deviation
Sensitivity analysis	Sobol sensitivity indices of the Ishigami function, Sobol indices of Sobol's G-function, Sobol indices of a linear (additive) model, Sobol indices of a model with an interaction term, Morris elementary-effects screening
Uncertainty propagation	Uncertainty of a sum or linear combination (independent inputs), Uncertainty of a product and a quotient, Uncertainty of a product and quotient with correlation, Uncertainty of a power x^n, Uncertainty of a function of one variable, Uncertainty of a power-law product y = k x1^p1 x2^p2 ..., Inverse-variance weighted mean of measurements, Uncertainty of a linear combination of correlated inputs, Covariance propagation through a linear map y = A x, Effective degrees of freedom and expanded uncertainty (Welch-Satterthwaite), Type B standard uncertainty from a stated half-width, Exact versus first-order propagation through exp(X), Exact mean and standard deviation of a product of independent normals, Exact mean and variance of a polynomial of a normal variable

Outcomes

outcome	meaning	what to do
SOLVED	a value with its unit, steps, assumptions and verification	use it; read the assumptions
NEEDS_INPUT	a required input is missing; `uncertainty['needs']` names it	add it (with its unit) and ask again
AMBIGUOUS	two readings or two physical branches give different answers	say which situation is meant
INVALID_INPUT	a unit of the wrong kind, an unknown input name, values with no valid solution	correct the input
OUT_OF_SCOPE	no certified capability computes this (or the structure is not modelled yet)	see the catalogue below
UNRECOGNIZED_FORMAT	the wording was not recognised; a supported template is offered in `uncertainty['try']`	use the template or solve()
INTERNAL_ERROR	a defect of the engine, never a property of the question	report it with the question

Result fields

field	meaning
value, unit	the result in SI units
values	every output of a procedure (x1, x2, discriminant ...)
outcome	one of the seven outcomes above
steps	each equation used, its inputs (with where each came from) and its result
assumptions	conditions and declared defaults the answer relies on
verification	residual of every equation, dimensional consistency, independent checks
latex()	the derivation as LaTeX
show()	renders in a notebook, prints in a terminal

Common mistakes

- Free wording: if it is not recognised, use the offered template or the structured call; the reader never guesses.
- A coefficient that violates a stated condition (a = 0 in a quadratic): INVALID_INPUT says why.
- A text value where a number is expected: INVALID_INPUT; fractions as strings ("1/2") are accepted.
- A problem outside the procedure (a cubic for the quadratic solver): OUT_OF_SCOPE, never an approximation.

Procedures: algorithms with several outputs

Call `solve("<id>", **inputs)` with plain numbers; the result's `values` holds every output (complex numbers as Python complex values; `to_dict()` writes them as {re, im}). Each result is substituted back into its defining equation before it is returned.

id	operation	inputs	outputs
bayes_theorem	Bayes' theorem for an event and its complement	p_a, p_b_given_a, p_b_given_not_a	p_a_given_b, p_b, p_not_a_given_b
diagnostic_test_predictive_values	Positive and negative predictive values of a test	prevalence, sensitivity, specificity	ppv, npv, p_positive, false_discovery_rate, false_omission_rate
total_probability	Law of total probability	priors, conditionals	p_b, joint
discrete_bayes_posterior	Posterior over several discrete hypotheses	priors, likelihoods	posterior, evidence, map_hypothesis, map_probability
probability_to_odds	Odds and log-odds of a probability	probability	odds, odds_against, log_odds
odds_to_probability	Probability from odds	odds	probability, complement
likelihood_ratios	Likelihood ratios of a diagnostic test	sensitivity, specificity	lr_positive, lr_negative, diagnostic_odds_ratio

id	operation	inputs	outputs
posterior_from_likelihood_ratio	Post-test probability from a likelihood ratio	prior, likelihood_ratio	posterior, prior_odds, posterior_odds
sequential_test_posterior	Posterior after repeated independent tests	prevalence, sensitivity, specificity, positives, negatives	posterior, posterior_odds
posterior_from_bayes_factor	Posterior probability of a hypothesis from a Bayes factor	bayes_factor, prior	posterior, posterior_odds, posterior_h0
rule_of_succession	Laplace's rule of succession	successes, trials	p_next, mle
beta_binomial_posterior	Beta-binomial conjugate posterior	alpha, beta, successes, trials	post_alpha, post_beta, posterior_mean, posterior_variance, posterior_sd
beta_credible_interval	Equal-tailed credible interval of a beta distribution	alpha, beta, level	lower, upper, median
beta_binomial_credible_interval	Credible interval for a proportion (beta-binomial)	alpha, beta, successes, trials, level	post_alpha, post_beta, lower, upper
beta_posterior_tail_probability	Posterior probability that p exceeds a value	alpha, beta, threshold, successes, trials	prob_above, prob_below
jeffreys_interval	Jeffreys interval for a binomial proportion	successes, trials, level	lower, upper, posterior_mean
beta_binomial_predictive	Posterior predictive probability (beta-binomial)	alpha, beta, future_trials, future_successes	probability, prob_at_most, prob_more, predictive_mean
beta_binomial_evidence	Marginal likelihood (evidence) of binomial data under a beta prior	alpha, beta, successes, trials	marginal_likelihood, log_marginal_likelihood
map_beta_binomial	MAP estimate of a proportion (beta-binomial)	alpha, beta, successes, trials	map, mle, posterior_mean
laplace_approximation_beta	Laplace (normal) approximation of a beta posterior	alpha, beta, level	mode, laplace_sd, lower, upper
beta_hdi	Highest-density interval of a beta posterior	alpha, beta, level	lower, upper, width
beta_prior_from_mean_sd	Beta prior from a mean and standard deviation	mean, sd	alpha, beta, prior_sample_size
ab_test_probability_beta	Bayesian A/B test: probability that A beats B	alpha_a, beta_a, alpha_b, beta_b	prob_a_greater, prob_b_greater

id	operation	inputs	outputs
gamma_poisson_posterior	Gamma-Poisson conjugate posterior of a rate	shape, rate, counts	post_shape, post_rate, posterior_mean, posterior_variance, posterior_sd
gamma_poisson_credible_interval	Credible interval for a Poisson rate	shape, rate, counts, level	post_shape, post_rate, lower, upper
gamma_credible_interval	Equal-tailed credible interval of a gamma distribution	shape, rate, level	lower, upper, median
gamma_hdi	Highest-density interval of a gamma posterior	shape, rate, level	lower, upper, width
gamma_poisson_predictive	Posterior predictive count (negative binomial)	shape, rate, events	probability, prob_at_most, predictive_mean, predictive_variance
map_gamma_poisson	MAP estimate of a Poisson rate (gamma prior)	shape, rate, counts	map, mle, posterior_mean
gamma_exponential_posterior	Gamma-exponential conjugate posterior of a rate	shape, rate, times	post_shape, post_rate, posterior_mean, posterior_variance, posterior_sd
normal_normal_posterior	Normal mean with known variance: posterior from data	prior_mean, prior_sd, sigma, data	posterior_mean, posterior_sd, posterior_variance
normal_normal_summary	Normal mean with known variance: posterior and credible interval from a sample mean	prior_mean, prior_sd, sigma, sample_mean, n, level	posterior_mean, posterior_sd, lower, upper, shrinkage
normal_posterior_predictive	Posterior predictive distribution of a new observation	prior_mean, prior_sd, sigma, sample_mean, n, threshold	predictive_mean, predictive_sd, prob_above, prob_below
normal_inverse_gamma_posterior	Normal model with unknown mean and variance (normal-inverse-gamma)	mu0, kappa0, alpha0, beta0, data, level	mu_n, kappa_n, alpha_n, beta_n, sigma2_mean, lower, upper
dirichlet_multinomial_posterior	Dirichlet-multinomial conjugate posterior	alpha, counts	post_alpha, posterior_means, posterior_variances
dirichlet_multinomial_evidence	Dirichlet-multinomial probability of a count vector (evidence)	alpha, counts	probability, log_probability
bayes_factor_binomial_simple	Bayes factor for two simple binomial hypotheses	successes, trials, p0, p1	bf10, log10_bf10, posterior_h1
bayes_factor_point_null_binomial	Bayes factor for a point null on a	successes, trials, p0, alpha, beta	bf10, bf01, log10_bf10, posterior_h1

id	operation	inputs	outputs
	proportion (beta prior under H1)		
bayes_factor_normal_mean	Bayes factor for a normal mean (known variance)	sample_mean, n, sigma, mu0, prior_sd	bf10, bf01, log10_bf10, posterior_h1
posterior_model_probabilities	Posterior model probabilities from log evidences	log_evidences, priors	posterior, best_model, best_probability
expected_value_decision	Best action by expected monetary value	payoffs, probabilities	expected_values, best_action, best_expected_value
decision_criteria_under_uncertainty	Decision under uncertainty: maximin, maximax, Laplace, minimax regret, Hurwicz	payoffs, criterion, optimism	scores, best_action, criterion_value
expected_value_of_perfect_information	Expected value of perfect information (EVPI)	payoffs, probabilities	evpi, ev_with_perfect_information, best_expected_value
expected_value_of_sample_information	Expected value of sample information (EVSI)	payoffs, priors, likelihoods	evsi, ev_with_information, ev_without_information, signal_probabilities, best_action_per_signal
bayes_action_minimum_expected_loss	Bayes action: minimum posterior expected loss	losses, probabilities	expected_losses, bayes_action, bayes_risk
certainty_equivalent_exponential_utility	Expected utility and certainty equivalent (exponential utility)	outcomes, probabilities, risk_tolerance	expected_utility, certainty_equivalent, risk_premium, expected_value
certainty_equivalent_log_utility	Certainty equivalent with logarithmic utility	outcomes, probabilities	expected_log_utility, certainty_equivalent, risk_premium, expected_value
break_even_probability	Break-even (indifference) probability of two actions	payoffs	break_even, preferred_above, expected_value_at_break_even
treatment_threshold	Decision threshold probability (treat / do not treat)	loss_fp, loss_fn, probability	threshold, decision
uncertainty_sum	Uncertainty of a sum or linear combination (independent inputs)	values, uncertainties, coefficients	value, uncertainty
uncertainty_product_quotient	Uncertainty of a product and a quotient	a, ua, b, ub	product, u_product, quotient, u_quotient
uncertainty_product_correlated	Uncertainty of a product and	a, ua, b, ub, rho	product, u_product, quotient, u_quotient

id	operation	inputs	outputs
	quotient with correlation		
uncertainty_power	Uncertainty of a power x^n	x, ux, exponent	value, uncertainty, relative_uncertainty
uncertainty_function	Uncertainty of a function of one variable	x, ux, function	value, uncertainty
uncertainty_power_law	Uncertainty of a power-law product y = k x1^p1 x2^p2 ...	coefficient, values, exponents, uncertainties	value, uncertainty, relative_uncertainty
inverse_variance_weighted_mean	Inverse-variance weighted mean of measurements	values, uncertainties	weighted_mean, uncertainty, chi_square, birge_ratio
uncertainty_linear_correlated	Uncertainty of a linear combination of correlated inputs	values, uncertainties, correlation, coefficients	value, uncertainty
covariance_propagation_linear	Covariance propagation through a linear map y = A x	matrix, covariance	output_covariance, uncertainties
welch_satterthwaite_expanded_uncertainty	Effective degrees of freedom and expanded uncertainty (Welch-Satterthwaite)	uncertainties, dof, level	combined_uncertainty, effective_dof, coverage_factor, expanded_uncertainty
type_b_standard_uncertainty	Type B standard uncertainty from a stated half-width	half_width, distribution	standard_uncertainty, divisor
lognormal_propagation	Exact versus first-order propagation through exp(X)	mu, sigma	exact_mean, exact_sd, first_order_mean, first_order_sd
product_of_normals_exact_moments	Exact mean and standard deviation of a product of independent normals	mean_x, sd_x, mean_y, sd_y	exact_mean, exact_sd, first_order_sd
monte_carlo_sample_size	Number of Monte Carlo draws for a target accuracy	sigma, half_width, level	draws, z, standard_error_target
monte_carlo_integration	Monte Carlo estimate of an integral from uniform samples	values, lower, upper	estimate, standard_error, n
importance_sampling_estimate	Importance sampling estimate and effective sample size	values, target_density, proposal_density	estimate, unnormalized_estimate, effective_sample_size

id	operation	inputs	outputs
control_variate_estimate	Control-variate Monte Carlo estimate	y, c, control_mean	estimate, coefficient, plain_estimate, variance_reduction
stratified_sampling_estimate	Stratified sampling estimate and standard error	weights, means, sds, sizes	estimate, standard_error
halton_point	Point of the Halton low-discrepancy sequence	index, dimension	point, bases
sobol_indices_ishigami	Sobol sensitivity indices of the Ishigami function	a, b	S1, S2, S3, ST1, ST2, ST3, variance
sobol_indices_g_function	Sobol indices of Sobol's G-function	a	first_order, total_order, variance
sobol_indices_linear_model	Sobol indices of a linear (additive) model	coefficients, sds	first_order, variance
sobol_indices_interaction_model	Sobol indices of a model with an interaction term	a1, a2, a12, mean1, sd1, mean2, sd2	S1, S2, S12, ST1, ST2, variance
morris_elementary_effects	Morris elementary-effects screening	outputs, delta	mu, mu_star, sigma
kl_divergence_normal	Kullback-Leibler divergence between two normal distributions	mean1, sd1, mean2, sd2	kl_1_to_2, kl_2_to_1
kl_divergence_beta	Kullback-Leibler divergence between two beta distributions	alpha1, beta1, alpha2, beta2	kl
gauss_hermite_rule	Gauss-Hermite rule for expectations under a normal distribution	n	nodes, weights
polynomial_of_normal_moments	Exact mean and variance of a polynomial of a normal variable	coefficients, mean, sd	mean, variance, sd
probability_rate_greater_gamma	Probability that one Poisson rate exceeds another (gamma posteriors)	shape1, rate1, shape2, rate2	prob_1_greater, prob_2_greater, rate_ratio_mean

Bayes' theorem for an event and its complement — `bayes_theorem`

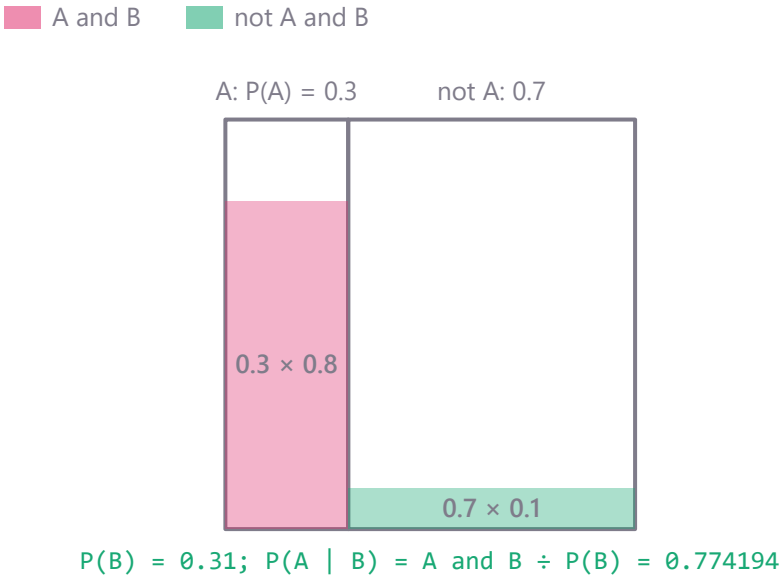
The posterior probability P(A|B) from the prior P(A) and the likelihoods P(B|A) and P(B|not A).

$$P(A \mid B) = \frac{P(B \mid A)\,P(A)}{P(B \mid A)\,P(A) + P(B \mid \neg A)\,(1 - P(A))}$$

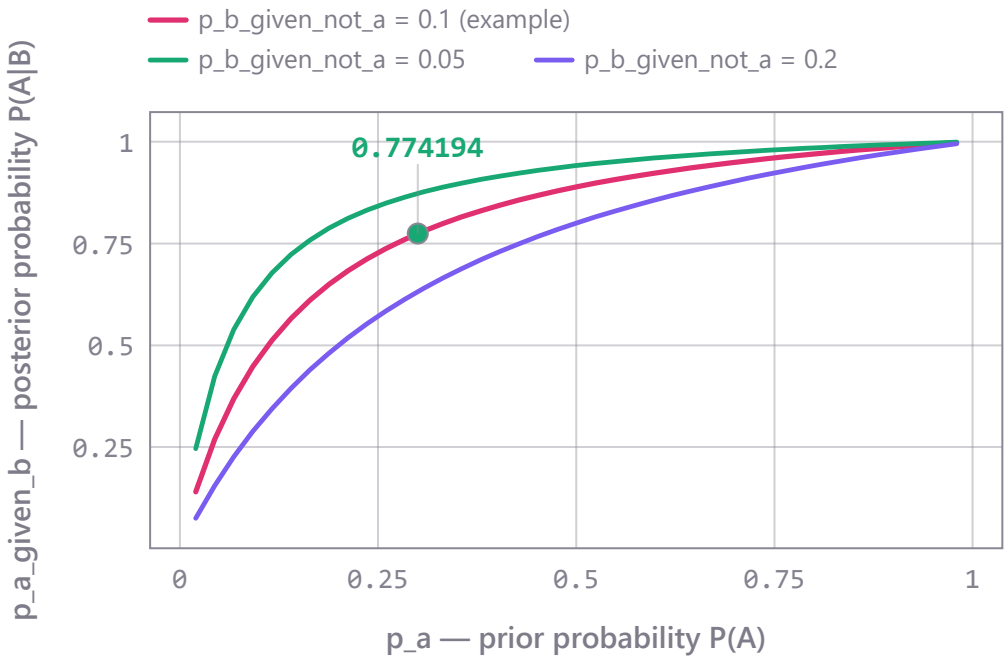
input	meaning	condition
`p_a`	prior probability P(A)	
`p_b_given_a`	likelihood P(B|A)	
`p_b_given_not_a`	likelihood P(B|not A)	

output	meaning
`p_a_given_b`	posterior probability P(A|B)
`p_b`	total probability P(B)
`p_not_a_given_b`	P(not A|B) = 1 - P(A|B)

- **Method:** law of total probability for P(B), then Bayes' theorem
- **Assumes:** probabilities between 0 and 1; P(B) > 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Bayes' theorem as areas in the unit square (to scale): the width splits the population into A (P(A) = 0.3) and not A; the shaded heights are P(B | A) = 0.8 and P(B | not A) = 0.1 (the worked example's inputs). The shaded areas together are the returned P(B) = 0.31; the share of it inside A is the returned P(A | B) = 0.774194 (P(not A | B) = 0.225806).



p_a_given_b (posterior probability P(A|B)) against p_a (prior probability P(A)) from 0.02 to 0.98. One curve for each value of p_b_given_not_a (likelihood P(B|not A)): 0.1, 0.05, 0.2. Other inputs as in the worked example: p_b_given_a = 0.8. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, p_a = 0.3, p_a_given_b = 0.774194.

```
bayesian_statistics_ai999x.solve('bayes_theorem', p_a=0.3, p_b_given_a=0.8, p_b_given_not_a=0.1)
# -> p_a_given_b = 0.774194, p_b = 0.31, p_not_a_given_b = 0.225806   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('bayes_theorem', p_a=0.5, p_b_given_a=0.5, p_b_given_not_a=0.5)
# -> p_a_given_b = 0.5, p_b = 0.5, p_not_a_given_b = 0.5   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Bayes' theorem: P(A) = 0.3, P(B|A) = 0.8, P(B|not A) = 0.1. Find P(A|B).	SOLVED	p_a_given_b = 0.774194, p_b = 0.31, p_not_a_given_b = 0.225806
Given P(A) = 2%, P(B|A) = 95% and P(B|not A) = 10%, what is P(A|B)?	SOLVED	p_a_given_b = 0.162393, p_b = 0.117, p_not_a_given_b = 0.837607
Bayes' theorem: P(A) = 0.3, P(B|A) = 1.2, P(B|not A) = 0.1. Find P(A|B).	INVALID_INPUT	Bayes' theorem for an event and its complement: P(B|A) must be a probability between 0 and 1.
Bayes' theorem: P(A) = 0.3, P(B|A) = 0.8. Find P(A|B).	NEEDS_INPUT	bayes_theorem needs p_b_given_not_a (likelihood P(B|not A)). Example: solve('bayes_theorem', p_a=0.3, p_b_given_a=0.8, p

Positive and negative predictive values of a test — `diagnostic_test_predictive_values`

The probability of the condition after a positive result (PPV) and of no condition after a negative one (NPV), from prevalence, sensitivity and specificity.

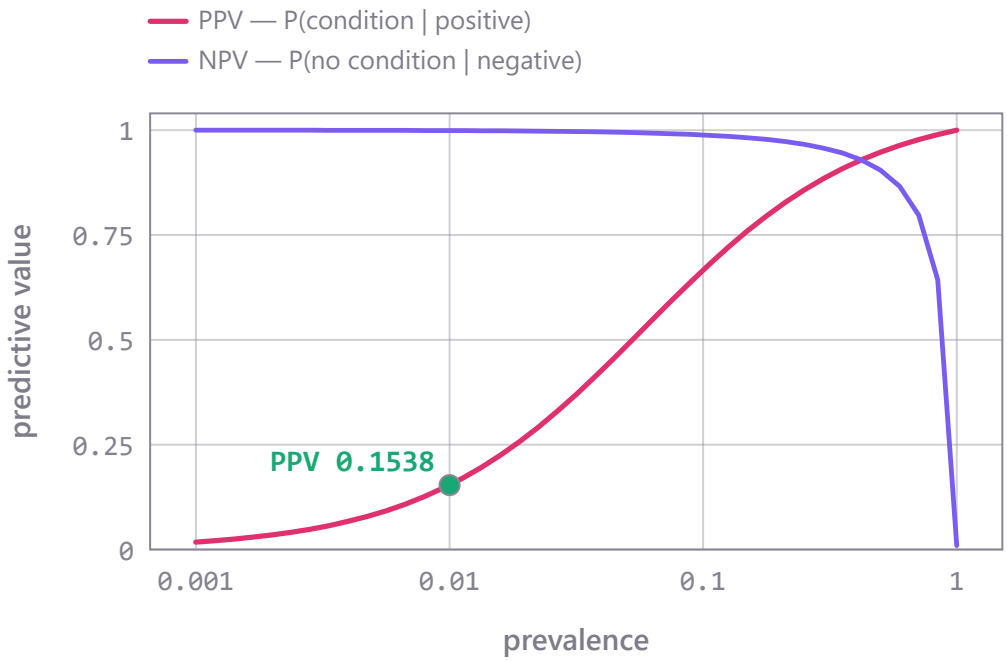
$$\mathrm{PPV} = \frac{\mathrm{Se}\,\pi}{\mathrm{Se}\,\pi + (1-\mathrm{Sp})(1-\pi)}, \quad \mathrm{NPV} = \frac{\mathrm{Sp}\,(1-\pi)}{\mathrm{Sp}\,(1-\pi) + (1-\mathrm{Se})\,\pi}$$

input	meaning	condition
`prevalence`	prevalence (pre-test probability of the condition)	

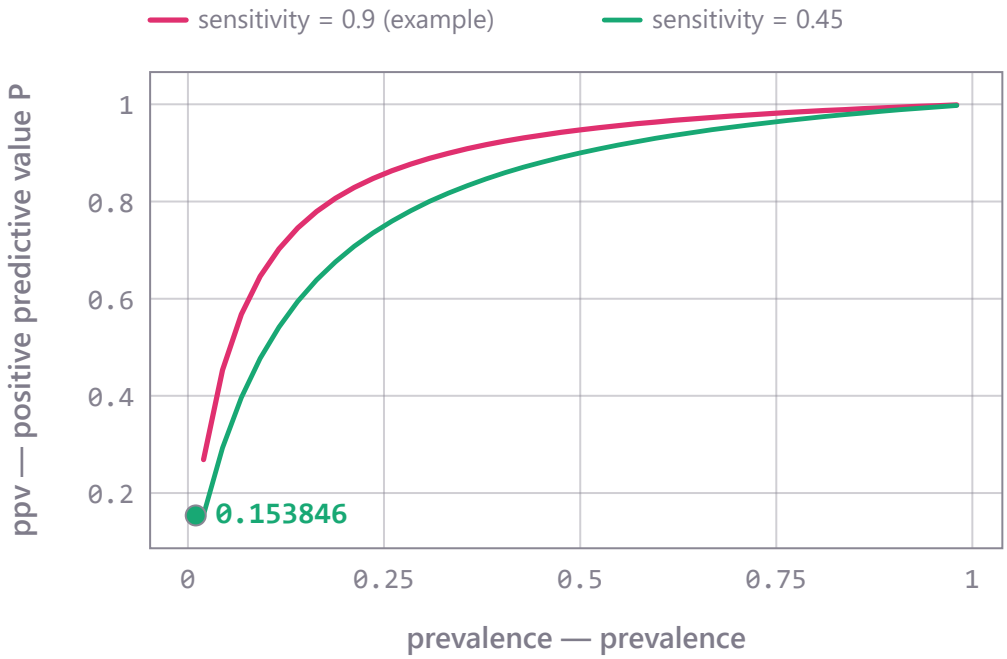
input	meaning	condition
`sensitivity`	sensitivity P(positive | condition)	
`specificity`	specificity P(negative | no condition)	

output	meaning
`ppv`	positive predictive value P(condition | positive)
`npv`	negative predictive value P(no condition | negative)
`p_positive`	probability of a positive result
`false_discovery_rate`	1 - PPV
`false_omission_rate`	1 - NPV

- **Method:** Bayes' theorem with the total probability of a positive result
- **Assumes:** probabilities between 0 and 1; both results possible
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Positive and negative predictive values of a test with sensitivity 0.9 and specificity 0.95 (the worked example) over prevalences from 0.001 to 1, each from the package: a positive result means little when the condition is rare. Marked: prevalence 0.01, returned PPV 0.153846 (NPV 0.998938, P(positive) = 0.0585).



ppv (positive predictive value P(condition | positive)) against prevalence (prevalence (pre-test probability of the condition)) from 0.02 to 0.98. One curve for each value of sensitivity (sensitivity P(positive | condition)): 0.9, 0.45. Other inputs as in the worked example: specificity = 0.95. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, prevalence = 0.01, ppv = 0.153846.

```
bayesian_statistics_ai999x.solve('diagnostic_test_predictive_values', prevalence=0.01,
sensitivity=0.9, specificity=0.95)
# -> ppv = 0.153846, npv = 0.998938, p_positive = 0.0585, false_discovery_rate = 0.846154,
false_omission_rate = 0.00106213   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('diagnostic_test_predictive_values', prevalence=0.5,
sensitivity=0.8, specificity=0.8)
# -> ppv = 0.8, npv = 0.8, p_positive = 0.5, false_discovery_rate = 0.2, false_omission_rate = 0.2
(SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
A test has sensitivity 0.9 and specificity 0.95; the prevalence is 0.01. What is the positive predictive value?	SOLVED	ppv = 0.153846, npv = 0.998938, p_positive = 0.0585, false_discovery_rate = 0.846154, false_omission_rate = 0.00106213
Prevalence = 50%, sensitivity = 80%, specificity = 80%. Find the PPV and NPV.	SOLVED	ppv = 0.8, npv = 0.8, p_positive = 0.5, false_discovery_rate = 0.2, false_omission_rate = 0.2
Prevalence = 0.1, sensitivity = 1.3, specificity = 0.9. Find the PPV.	INVALID_INPUT	Positive and negative predictive values of a test: sensitivity must be a probability between 0 and 1.
The prevalence is 0.1 and the sensitivity is 0.9. What is the PPV given the specificity?	NEEDS_INPUT	diagnostic_test_predictive_values needs specificity (specificity P(negative | no condition)). Example: solve('diagnostic

Law of total probability — `total_probability`

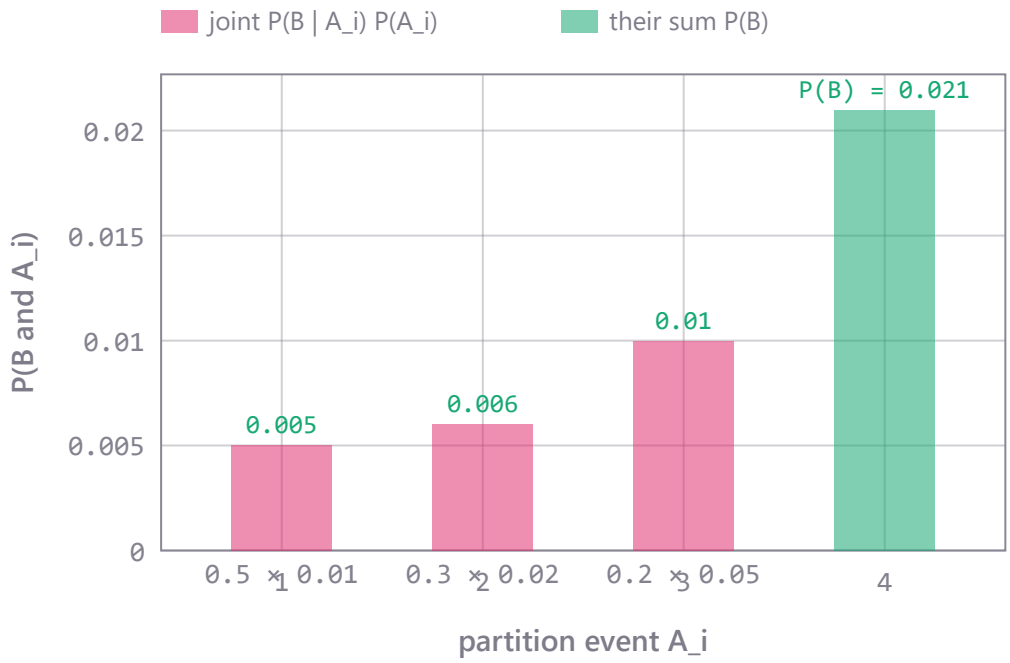
P(B) = sum of P(B|A_i) P(A_i) over a partition A_1, ..., A_k (and the joint probabilities).

$$P(B) = \sum_i P(B \mid A_i)\, P(A_i)$$

input	meaning	condition
`priors`	probabilities P(A_i) of the partition (sum to 1)	
`conditionals`	conditional probabilities P(B|A_i), one per part	

output	meaning
`p_b`	total probability P(B)
`joint`	joint probabilities P(B and A_i)

- **Method:** weighted sum over the partition
- **Assumes:** the P(A_i) sum to 1; same number of conditionals as parts
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The law of total probability for priors 0.5, 0.3, 0.2 and conditionals P(B | A_i) = 0.01, 0.02, 0.05 (the worked example): the returned joint probabilities 0.005, 0.006, 0.01 (prior × conditional, written under each bar) and their sum, the returned P(B) = 0.021.

```
bayesian_statistics_ai999x.solve('total_probability', priors='0.5, 0.3, 0.2', conditionals='0.01, 0.02, 0.05')
# -> p_b = 0.021, joint = [0.005, 0.006, 0.01]   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('total_probability', priors=[0.6, 0.4], conditionals=[0.1, 0.3])
# -> p_b = 0.18, joint = [0.06, 0.12]   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Law of total probability: priors = [0.5, 0.3, 0.2], conditionals = [0.01, 0.02, 0.05]. Find P(B).	SOLVED	p_b = 0.021, joint = [0.005, 0.006, 0.01]

question	outcome	result
Use total probability with priors [0.6, 0.4] and conditional probabilities [0.1, 0.3].	SOLVED	p_b = 0.18, joint = [0.06, 0.12]
Law of total probability: priors = [0.5, 0.3], conditionals = [0.01, 0.02]. Find P(B).	INVALID_INPUT	Law of total probability: the probabilities of the partition must sum to 1 (they sum to 0.8).
Law of total probability: priors = [0.5, 0.5]. Find P(B).	NEEDS_INPUT	total_probability needs conditionals (conditional probabilities P(B|A_i), one per part). Example: solve('total_probabili

Posterior over several discrete hypotheses — `discrete_bayes_posterior`

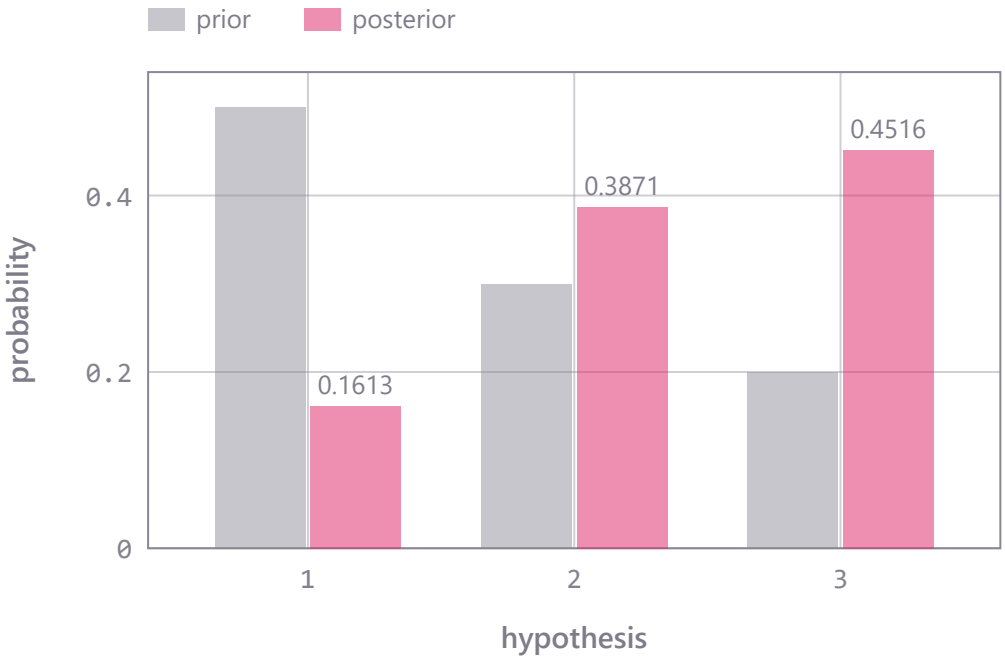
Posterior probabilities of k hypotheses from their priors and the likelihood of the observed data under each; the evidence and the most probable hypothesis (MAP).

$$P(H_i \mid D) = \frac{P(D \mid H_i)\,P(H_i)}{\sum_j P(D \mid H_j)\,P(H_j)}$$

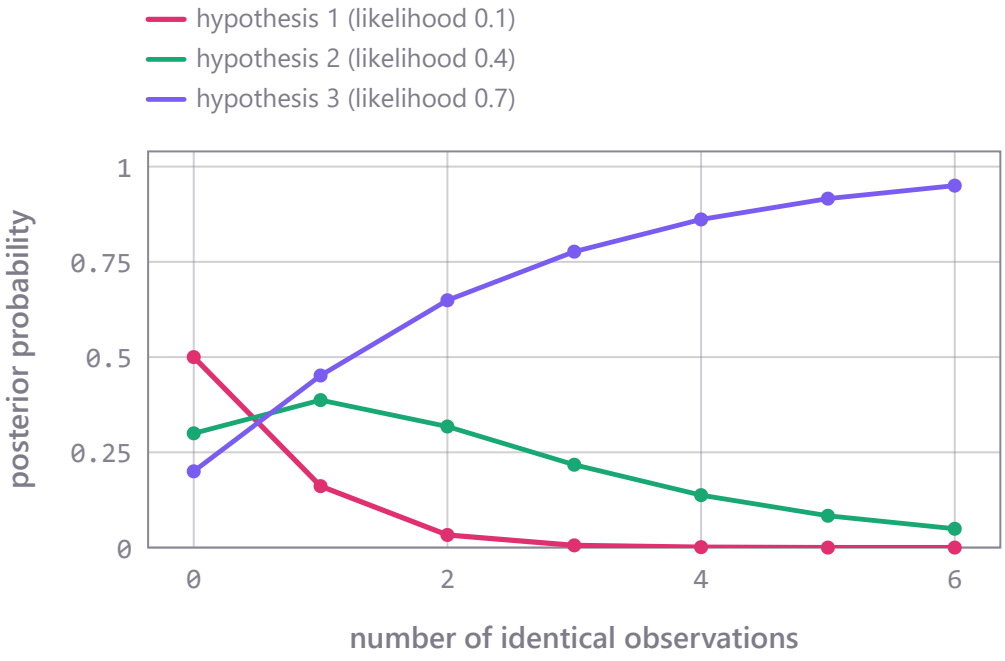
input	meaning	condition
`priors`	prior probabilities of the hypotheses (sum to 1)	
`likelihoods`	likelihood of the data under each hypothesis	

output	meaning
`posterior`	posterior probabilities
`evidence`	marginal probability of the data
`map_hypothesis`	number (1-based) of the most probable hypothesis
`map_probability`	its posterior probability

- **Method:** Bayes' theorem: multiply prior by likelihood and normalize
- **Assumes:** priors sum to 1; likelihoods >= 0, not all zero where the prior is positive
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Bayes' rule over 3 hypotheses with priors 0.5, 0.3, 0.2 and likelihoods 0.1, 0.4, 0.7: the package's posterior probabilities 0.16129, 0.387097, 0.451613 (evidence 0.31); the most probable hypothesis is 3 with 0.451613.



Repeated updating over 3 hypotheses with priors 0.5, 0.3, 0.2 and likelihoods 0.1, 0.4, 0.7 (the worked example), the posterior of each step fed back as the next prior, six steps, every probability from the package: the hypothesis that explains the observation best takes over (after 6 steps: 2.0195e-05, 0.0496, 0.9503).

```
bayesian_statistics_ai999x.solve('discrete_bayes_posterior', priors='0.5, 0.3, 0.2',
likelihoods='0.1, 0.4, 0.7')
# -> posterior = [0.161290322581, 0.387096774194, 0.451612903226], evidence = 0.31, map_hypothesis
= 3, map_probability = 0.451613   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('discrete_bayes_posterior', priors=[0.5, 0.5], likelihoods=[0.2,
0.2])
# -> posterior = [0.5, 0.5], evidence = 0.2, map_hypothesis = 1, map_probability = 0.5   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Posterior probabilities with priors = [0.5, 0.3, 0.2] and likelihoods = [0.1, 0.4, 0.7].	SOLVED	posterior = [0.161290322581, 0.387096774194, 0.451612903226], evidence = 0.31, map_hypothesis = 3, map_probability = 0.451613
Three hypotheses have priors [0.2, 0.3, 0.5] and likelihoods [0.9, 0.5, 0.1]: find the posterior.	SOLVED	posterior = [0.473684210526, 0.394736842105, 0.131578947368], evidence = 0.38, map_hypothesis = 1, map_probability = 0.473684
Posterior with priors = [0.5, 0.6] and likelihoods = [0.1, 0.4].	INVALID_INPUT	Posterior over several discrete hypotheses: the prior probabilities must sum to 1 (they sum to 1.1).
Posterior probabilities with priors = [0.5, 0.5]; the likelihoods are not known yet.	NEEDS_INPUT	discrete_bayes_posterior needs likelihoods (likelihood of the data under each hypothesis). Example: solve('discrete_baye

Odds and log-odds of a probability — `probability_to_odds`

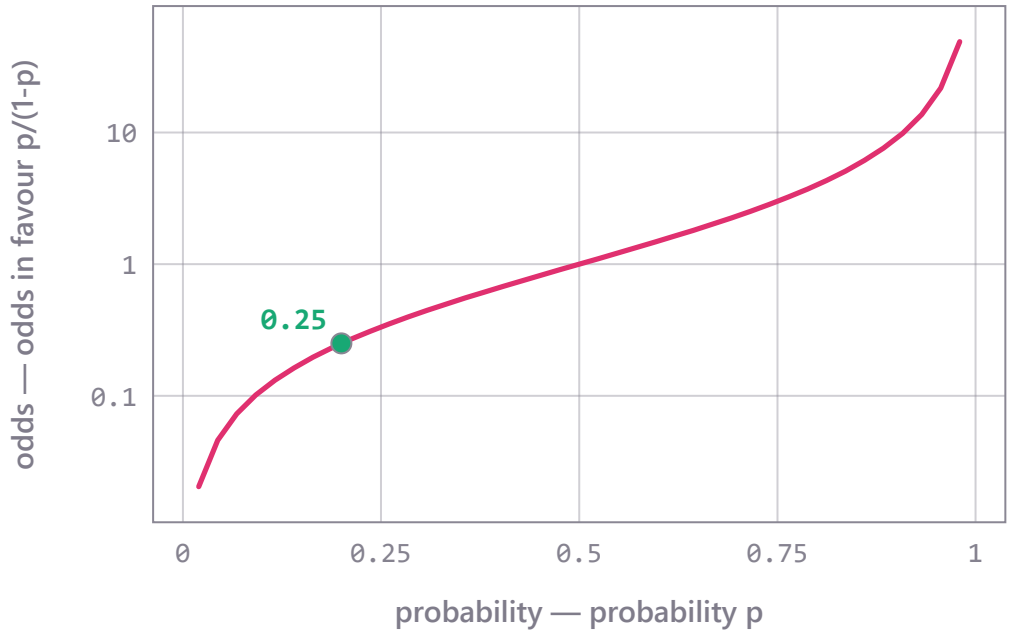
The odds $p/(1-p)$, the odds against and the log-odds (logit) of a probability.

$$\text{odds} = \frac{p}{1-p}, \quad \text{logit}\, p = \ln \frac{p}{1-p}$$

input	meaning	condition
`probability`	probability p, strictly between 0 and 1	

output	meaning
`odds`	odds in favour $p/(1-p)$
`odds_against`	odds against $(1-p)/p$
`log_odds`	natural log-odds ln($p/(1-p)$)

- **Method:** definition of odds
- **Assumes:** $0 < p < 1$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



odds (odds in favour p/(1-p)) against probability (probability p, strictly between 0 and 1) from 0.02 to 0.98, logarithmic output axis. 41 points computed by bayesian-statistics-ai999x; marked: the worked example, probability = 0.2, odds = 0.25.

```
bayesian_statistics_ai999x.solve('probability_to_odds', probability=0.2)
# -> odds = 0.25, odds_against = 4, log_odds = -1.38629   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('probability_to_odds', probability=0.75)
# -> odds = 3, odds_against = 0.333333, log_odds = 1.09861   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Convert the probability 0.2 to odds.	SOLVED	odds = 0.25, odds_against = 4, log_odds = -1.38629
What are the odds for a probability of 75%?	SOLVED	odds = 3, odds_against = 0.333333, log_odds = 1.09861
Convert the probability 1.5 to odds.	INVALID_INPUT	Odds and log-odds of a probability: the probability must be strictly between 0 and 1.

Probability from odds — `odds_to_probability`

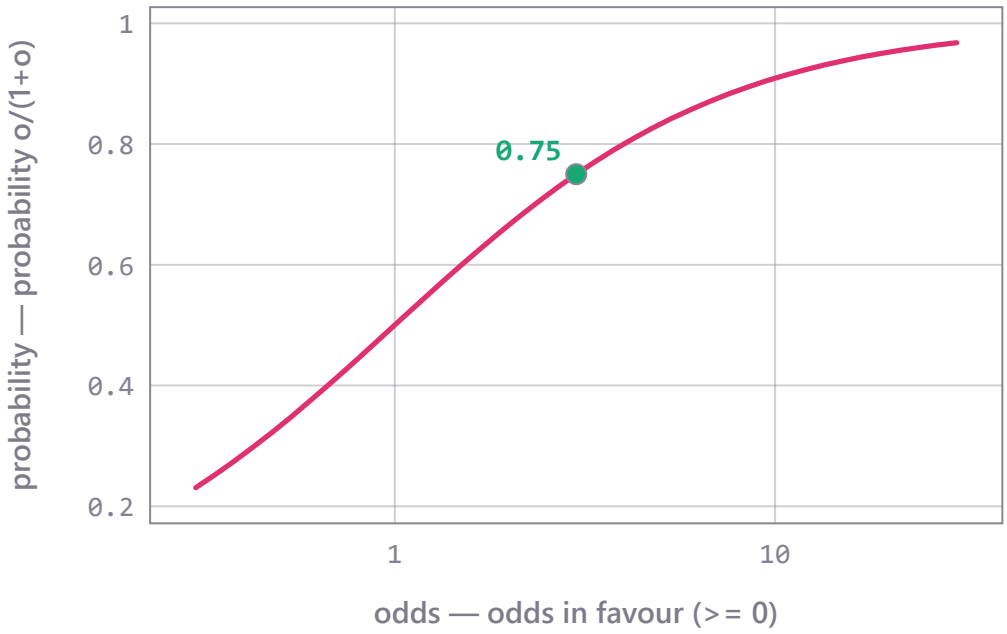
The probability $o/(1+o)$ that corresponds to odds o in favour.

$$p = \frac{o}{1+o}$$

input	meaning	condition
`odds`	odds in favour (>= 0)	

output	meaning
`probability`	probability $o/(1+o)$
`complement`	1 - probability

- **Method:** $p = o/(1+o)$
- **Assumes:** odds >= 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



probability (probability o/(1+o)) against odds (odds in favour (>= 0)) from 0.3 to 30 on a logarithmic axis. 41 points computed by bayesian-statistics-ai999x; marked: the worked example, odds = 3, probability = 0.75.

```
bayesian_statistics_ai999x.solve('odds_to_probability', odds=3)
# -> probability = 0.75, complement = 0.25   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('odds_to_probability', odds=0.25)
# -> probability = 0.2, complement = 0.8   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Convert odds of 3 to a probability.	SOLVED	probability = 0.75, complement = 0.25
What probability corresponds to odds = 0.25?	SOLVED	probability = 0.2, complement = 0.8
Convert odds of -2 to a probability.	INVALID_INPUT	Probability from odds: odds must not be negative.

Likelihood ratios of a diagnostic test — `likelihood_ratios`

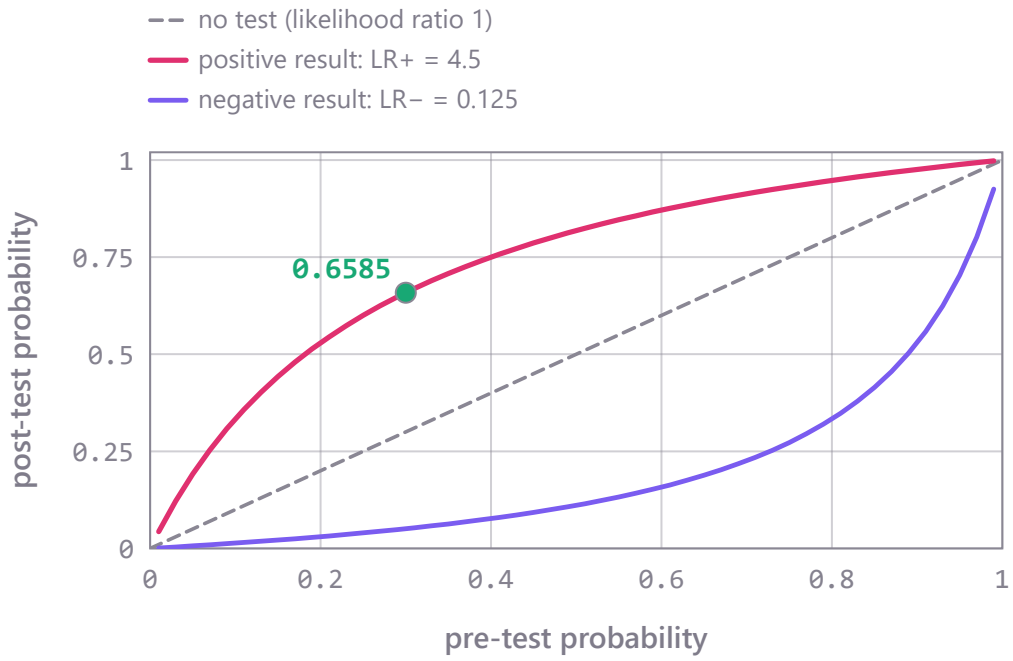
Positive and negative likelihood ratios and the diagnostic odds ratio from sensitivity and specificity.

$$LR^+ = \frac{\mathrm{Se}}{1-\mathrm{Sp}}, \quad LR^- = \frac{1-\mathrm{Se}}{\mathrm{Sp}}$$

input	meaning	condition
`sensitivity`	sensitivity P(positive | condition)	
`specificity`	specificity P(negative | no condition)	

output	meaning
`lr_positive`	LR+ = sensitivity / (1 - specificity)
`lr_negative`	LR- = (1 - sensitivity) / specificity
`diagnostic_odds_ratio`	DOR = LR+ / LR-

- **Method:** ratios of the probabilities of a result with and without the condition
- **Assumes:** 0 < specificity < 1; sensitivity < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



How far a test result moves a probability: the returned likelihood ratios of a test with sensitivity 0.9 and specificity 0.8 (LR+ = 4.5, LR− = 0.125, diagnostic odds ratio 36) turned into post-test probabilities by the package's posterior-from-likelihood-ratio capability at 50 pre-test values. Marked: pre-test 0.3 (chosen) → 0.658537 after a positive result.

```
bayesian_statistics_ai999x.solve('likelihood_ratios', sensitivity=0.9, specificity=0.8)
# -> lr_positive = 4.5, lr_negative = 0.125, diagnostic_odds_ratio = 36   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('likelihood_ratios', sensitivity=0.95, specificity=0.9)
# -> lr_positive = 9.5, lr_negative = 0.0555556, diagnostic_odds_ratio = 171   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Find the likelihood ratios of a test with sensitivity 0.9 and specificity 0.8.	SOLVED	lr_positive = 4.5, lr_negative = 0.125, diagnostic_odds_ratio = 36
Sensitivity = 95%, specificity = 90%: what are LR+, LR- and the diagnostic odds ratio?	SOLVED	lr_positive = 9.5, lr_negative = 0.0555556, diagnostic_odds_ratio = 171
Find the likelihood ratios of a test with sensitivity 0.9 and specificity 1.	INVALID_INPUT	Likelihood ratios of a diagnostic test: sensitivity and specificity must be below 1 (and specificity above 0) for finite
Find the likelihood ratios of a test with sensitivity 0.9 and an unknown specificity.	NEEDS_INPUT	likelihood_ratios needs specificity (specificity P(negative | no condition)). Example: solve('likelihood_ratios', sensit

Post-test probability from a likelihood ratio — `posterior_from_likelihood_ratio`

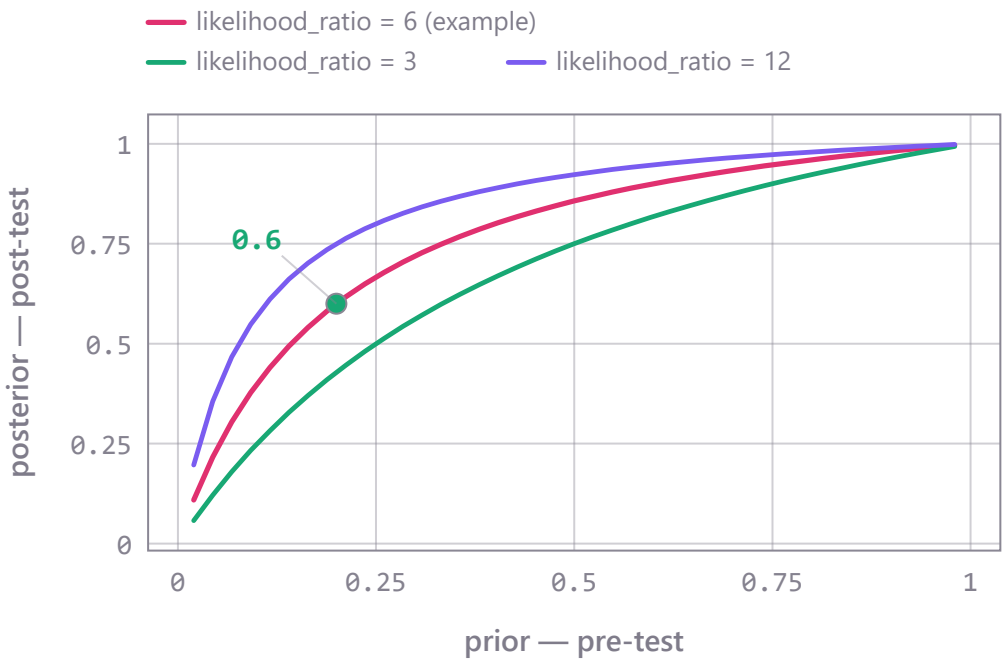
Bayes' theorem in odds form: post-test odds = pre-test odds x likelihood ratio, and the post-test probability.

$$\frac{P(H \mid E)}{1 - P(H \mid E)} = \frac{P(H)}{1 - P(H)} \times LR$$

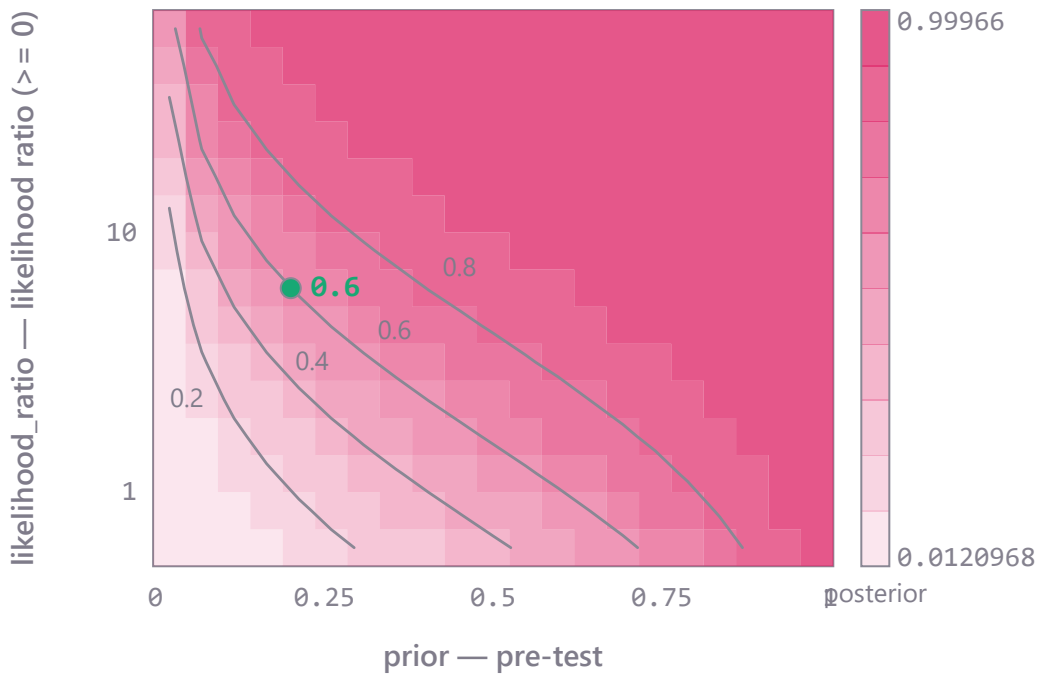
input	meaning	condition
`prior`	pre-test (prior) probability, $0 < p < 1$	
`likelihood_ratio`	likelihood ratio (>= 0)	

output	meaning
`posterior`	post-test (posterior) probability
`prior_odds`	pre-test odds
`posterior_odds`	post-test odds

- **Method:** odds form of Bayes' theorem
- **Assumes:** 0 < prior < 1; likelihood ratio >= 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



posterior (post-test (posterior) probability) against prior (pre-test (prior) probability, 0 < p < 1) from 0.02 to 0.98. One curve for each value of likelihood_ratio (likelihood ratio (>= 0)): 6, 3, 12. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, prior = 0.2, posterior = 0.6.



posterior (post-test (posterior) probability) over prior (pre-test (prior) probability, 0 < p < 1) from 0.02 to 0.98 and likelihood_ratio (likelihood ratio (>= 0)) from 0.6 to 60 (logarithmic vertical axis): each of the 21 × 15 cells is one run of bayesian-statistics-ai999x, shaded by its value (light 0.0120968 to dark 0.99966); iso-lines at posterior = 0.2, 0.4, 0.6, 0.8. Marked: the worked example, prior = 0.2, likelihood_ratio = 6, posterior = 0.6.

```
bayesian_statistics_ai999x.solve('posterior_from_likelihood_ratio', prior=0.2, likelihood_ratio=6)
# -> posterior = 0.6, prior_odds = 0.25, posterior_odds = 1.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('posterior_from_likelihood_ratio', prior=0.1,
likelihood_ratio=0.2)
# -> posterior = 0.0217391, prior_odds = 0.111111, posterior_odds = 0.0222222   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
The pre-test probability is 0.2 and the likelihood ratio is 6. What is the post-test probability?	SOLVED	posterior = 0.6, prior_odds = 0.25, posterior_odds = 1.5
Prior probability = 10%, likelihood ratio = 0.2: find the posterior probability.	SOLVED	posterior = 0.0217391, prior_odds = 0.111111, posterior_odds = 0.0222222
The pre-test probability is 0.2 and the likelihood ratio is -1. What is the post-test probability?	INVALID_INPUT	Post-test probability from a likelihood ratio: a likelihood ratio must not be negative.
The likelihood ratio is 6. What is the post-test probability given the pre-test probability?	NEEDS_INPUT	posterior_from_likelihood_ratio needs prior (pre-test (prior) probability, 0 < p < 1). Example: solve('posterior_from_li

Posterior after repeated independent tests — `sequential_test_posterior`

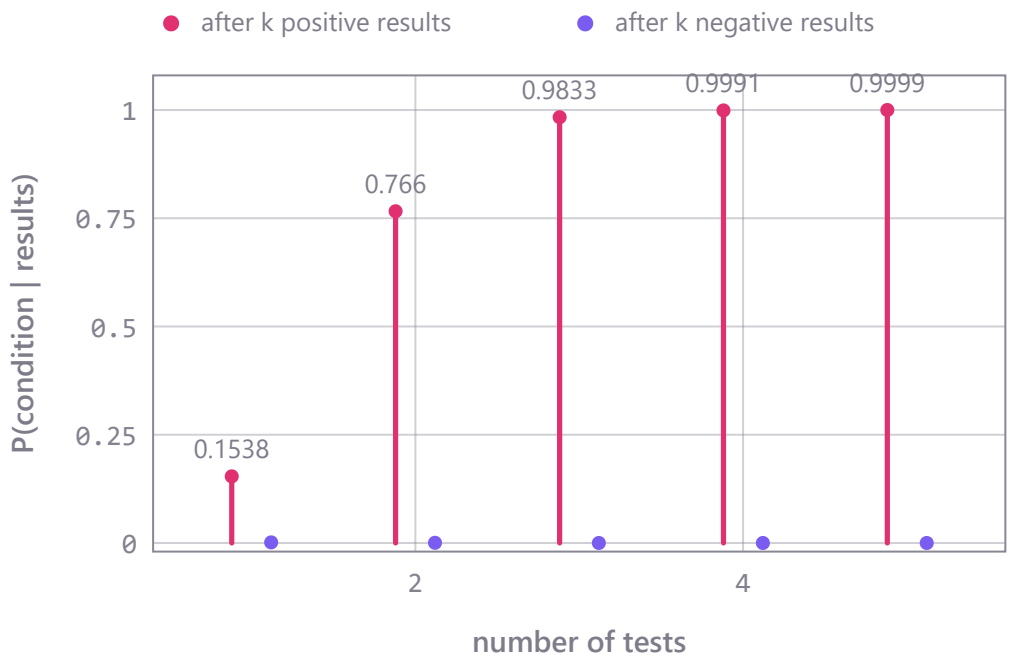
Probability of the condition after several conditionally independent test results (k positive and m negative), by multiplying the pre-test odds by LR+ for each positive and LR- for each negative.

$$\text{odds}_{post} = \frac{\pi}{1-\pi}\,(LR^{+})^{k}\,(LR^{-})^{m}$$

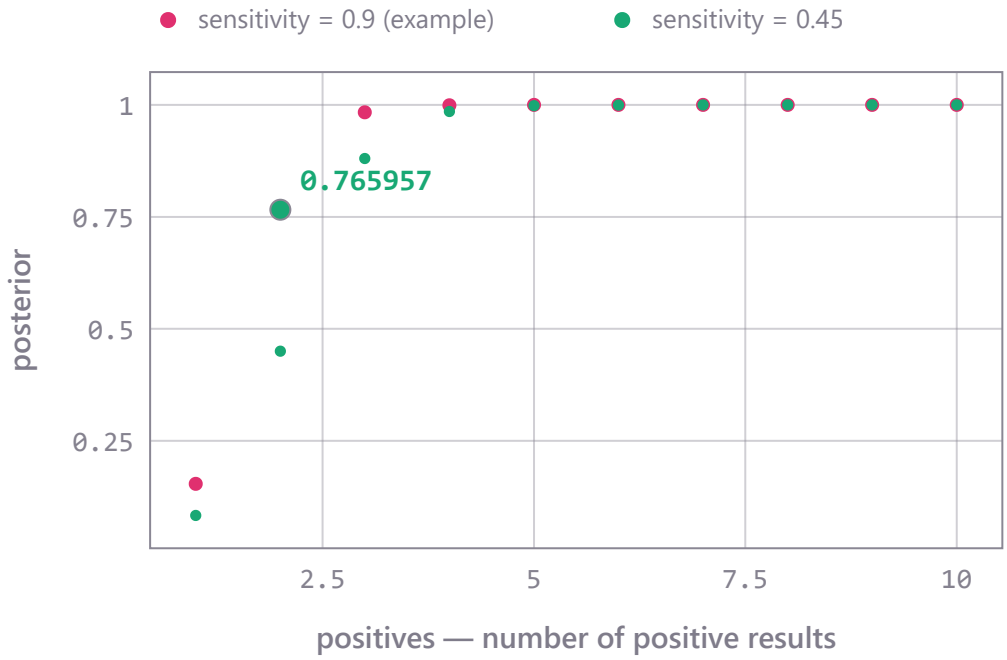
input	meaning	condition
`prevalence`	prevalence (pre-test probability of the condition)	
`sensitivity`	sensitivity P(positive | condition)	
`specificity`	specificity P(negative | no condition)	
`positives`	number of positive results	
`negatives`	number of negative results (default 0)	

output	meaning
`posterior`	probability of the condition after all results
`posterior_odds`	its odds

- **Method:** odds form, applied once per result (in log space)
- **Assumes:** 0 < prevalence < 1; results conditionally independent given the status
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Repeated diagnostic tests with prevalence 0.01, sensitivity 0.9 and specificity 0.95: the package's posterior probability of the condition after k = 1 … 5 positive results (0.153846, 0.765957, 0.983308, 0.999058, 0.999948) and after 1 … 5 negative results (down to 1.3054e-07).



posterior (probability of the condition after all results) against positives (number of positive results) from 1 to 10. One curve for each value of sensitivity (sensitivity P(positive | condition)): 0.9, 0.45. Other inputs as in the worked example: prevalence = 0.01, specificity = 0.95, negatives = 0. 10 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, positives = 2, posterior = 0.765957.

```
bayesian_statistics_ai999x.solve('sequential_test_posterior', prevalence=0.01, sensitivity=0.9,
specificity=0.95, positives=2, negatives=0)
# -> posterior = 0.765957, posterior_odds = 3.27273   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('sequential_test_posterior', prevalence=0.05, sensitivity=0.8,
specificity=0.9, positives=1, negatives=1)
# -> posterior = 0.0855615, posterior_odds = 0.0935673   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Repeated testing: prevalence 0.01, sensitivity 0.9, specificity 0.95, positives = 2, negatives = 0. Posterior probability of the disease?	SOLVED	posterior = 0.765957, posterior_odds = 3.27273
Prevalence = 5%, sensitivity = 80%, specificity = 90%, positives = 3: probability of the condition?	SOLVED	posterior = 0.964218, posterior_odds = 26.9474
Prevalence = 0.1, sensitivity = 0.9, specificity = 0.9, positives = 1.5: posterior?	INVALID_INPUT	positives: expected a whole number written exactly, got 1.5.
Prevalence = 0.1, sensitivity = 0.9, positives = 2: posterior given the specificity?	NEEDS_INPUT	sequential_test_posterior needs specificity (specificity P(negative | no condition)). Example: solve('sequential_test_po

Posterior probability of a hypothesis from a Bayes factor — `posterior_from_bayes_factor`

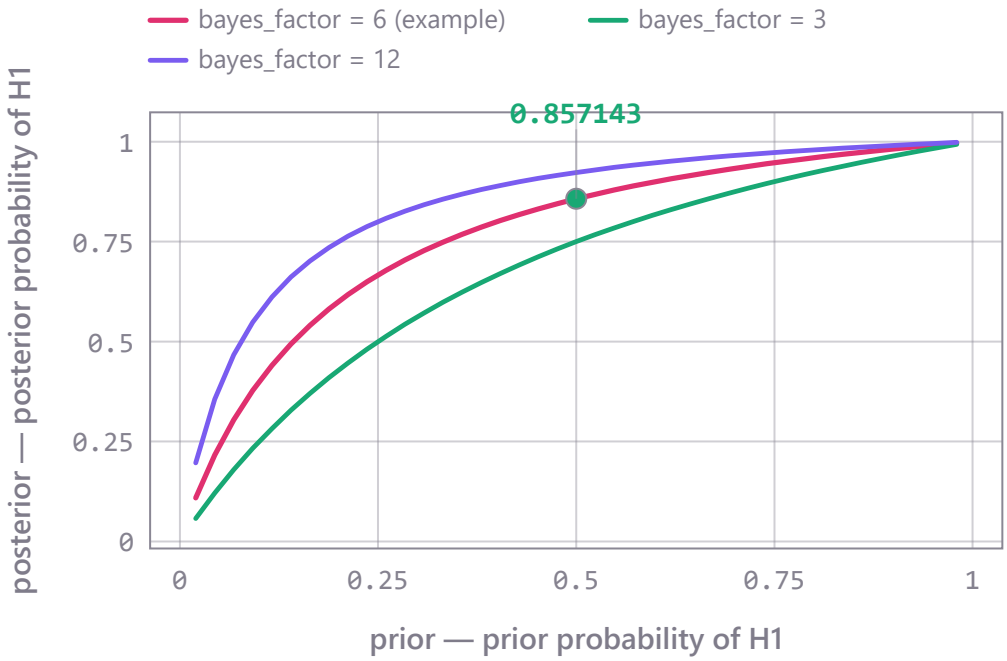
Posterior probability of H1 from the Bayes factor BF10 and the prior probability of H1 (default 1/2).

$$\frac{P(H_1 \mid D)}{P(H_0 \mid D)} = BF_{10} \frac{P(H_1)}{P(H_0)}$$

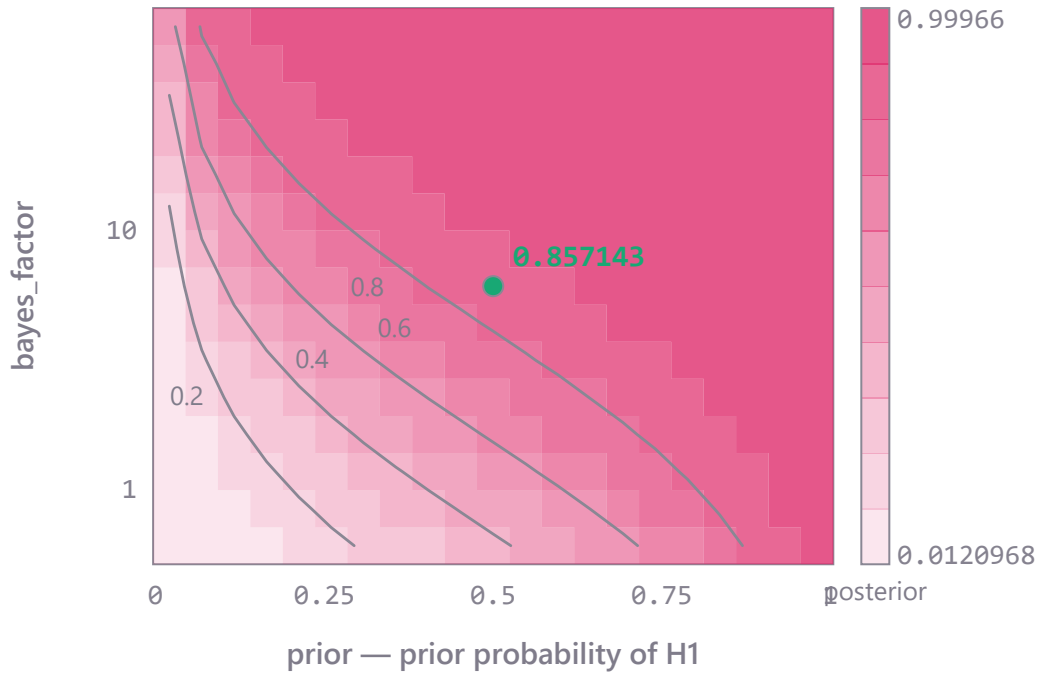
input	meaning	condition
`bayes_factor`	Bayes factor BF10 in favour of H1 (> 0)	
`prior`	prior probability of H1 (default 0.5)	

output	meaning
`posterior`	posterior probability of H1
`posterior_odds`	posterior odds H1 : H0
`posterior_h0`	posterior probability of H0

- **Method:** posterior odds = BF10 x prior odds
- **Assumes:** $BF10 > 0$; 0 < prior < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



posterior (posterior probability of H1) against prior (prior probability of H1 (default 0.5)) from 0.02 to 0.98. One curve for each value of bayes_factor (Bayes factor BF10 in favour of H1 (> 0)): 6, 3, 12. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, prior = 0.5, posterior = 0.857143.



posterior (posterior probability of H1) over prior (prior probability of H1 (default 0.5)) from 0.02 to 0.98 and bayes_factor (Bayes factor BF10 in favour of H1 (> 0)) from 0.6 to 60 (logarithmic vertical axis): each of the 21 × 15 cells is one run of bayesian-statistics-ai999x, shaded by its value (light 0.0120968 to dark 0.99966); iso-lines at posterior = 0.2, 0.4, 0.6, 0.8. Marked: the worked example, prior = 0.5, bayes_factor = 6, posterior = 0.857143.

```
bayesian_statistics_ai999x.solve('posterior_from_bayes_factor', bayes_factor=6, prior=0.5)
# -> posterior = 0.857143, posterior_odds = 6, posterior_h0 = 0.142857   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('posterior_from_bayes_factor', bayes_factor=0.2, prior=0.3)
# -> posterior = 0.0789474, posterior_odds = 0.0857143, posterior_h0 = 0.921053   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
A Bayes factor of 6 with prior probability 0.5: what is the posterior probability of H1?	SOLVED	posterior = 0.857143, posterior_odds = 6, posterior_h0 = 0.142857
Bayes factor = 3. Find the posterior probability of H1 with equal prior odds.	SOLVED	posterior = 0.75, posterior_odds = 3, posterior_h0 = 0.25
Bayes factor = -2 and prior probability = 0.3. Find the posterior probability.	INVALID_INPUT	Posterior probability of a hypothesis from a Bayes factor: a Bayes factor must be positive.

Laplace's rule of succession — `rule_of_succession`

Probability that the next trial is a success after k successes in n trials, with a uniform prior on the success probability: $(k+1)/(n+2)$.

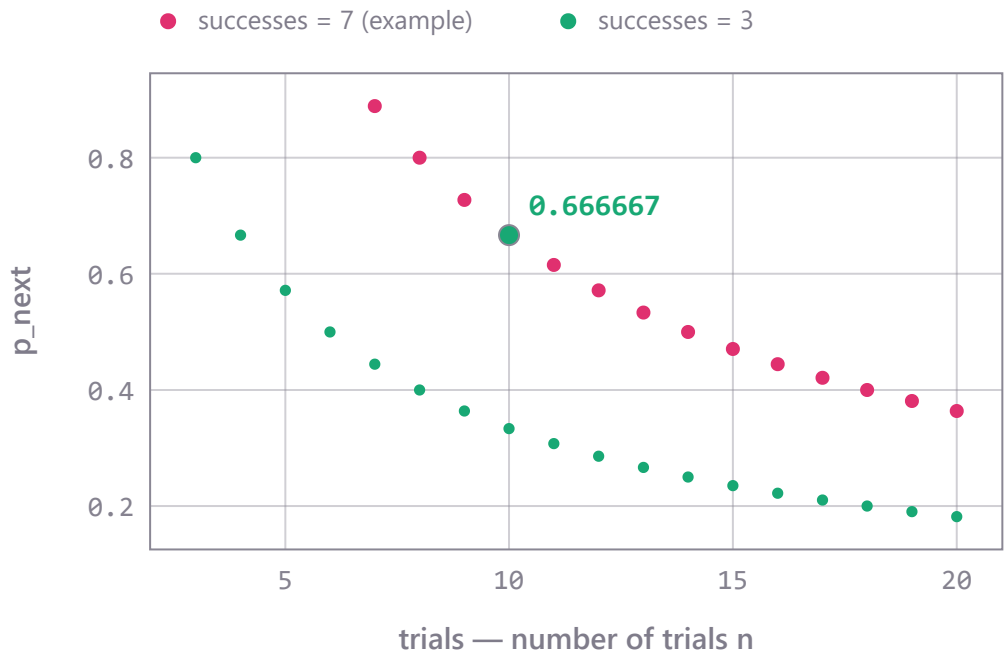
$$P(\text{next success}) = \frac{k+1}{n+2}$$

input	meaning	condition
`successes`	number of successes k	

input	meaning	condition
`trials`	number of trials n	

output	meaning
`p_next`	posterior predictive probability that the next trial succeeds
`mle`	observed proportion k/n (none when n = 0)

- **Method:** posterior mean of Beta($k + 1$, n - k + 1)
- **Assumes:** $0 \leq k \leq n$; uniform Beta(1, 1) prior
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



p_next (posterior predictive probability that the next trial succeeds) against trials (number of trials n) from 1 to 20. One curve for each value of successes (number of successes k): 7, 3. 14 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, trials = 10, p_next = 0.666667.

```
bayesian_statistics_ai999x.solve('rule_of_succession', successes=7, trials=10)
# -> p_next = 0.666667, mle = 0.7   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('rule_of_succession', successes=0, trials=5)
# -> p_next = 0.142857, mle = 0   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Rule of succession: 7 successes in 10 trials. Probability the next trial is a success?	SOLVED	p_next = 0.666667, mle = 0.7
The sun has risen on 9 successes in 9 trials; by the rule of succession, what is the probability of the next?	SOLVED	p_next = 0.909091, mle = 1
Rule of succession: 12 successes in 10 trials.	INVALID_INPUT	Laplace's rule of succession: successes cannot exceed trials.

question	outcome	result
Rule of succession with successes = 4.	NEEDS_INPUT	rule_of_succession needs trials (number of trials n). Example: solve('rule_of_succession', successes=7, trials=10).

Beta-binomial conjugate posterior — `beta_binomial_posterior`

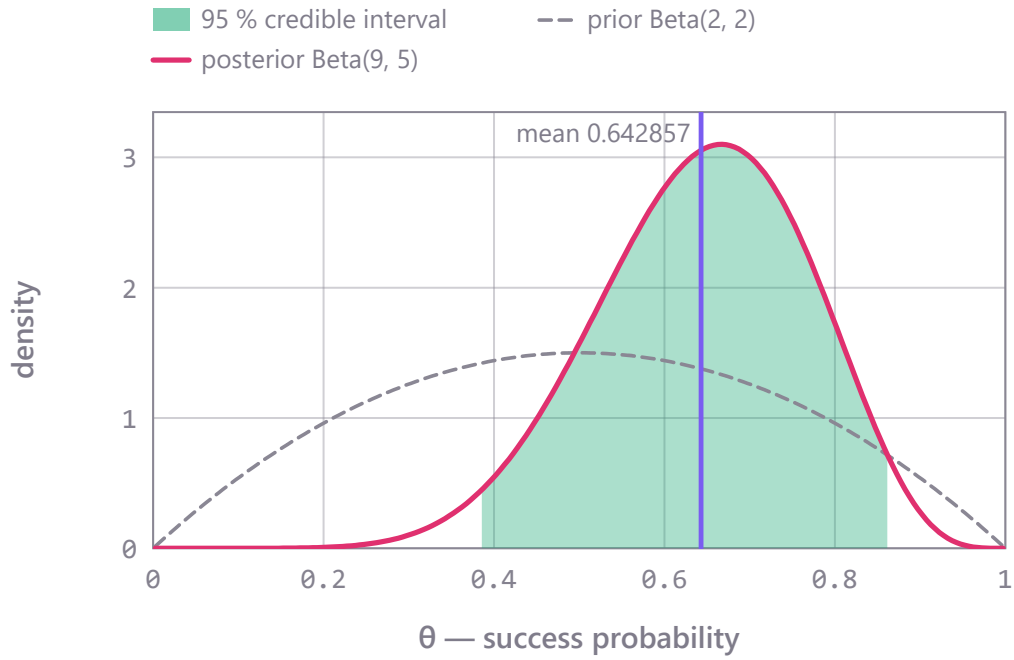
Update a Beta(alpha, beta) prior on a success probability with k successes in n trials: posterior parameters, mean, variance and standard deviation.

$$p \mid k \sim \mathrm{Beta}(\alpha + k,\ \beta + n - k), \quad E[p \mid k] = \frac{\alpha + k}{\alpha + \beta + n}$$

input	meaning	condition
`alpha`	prior Beta alpha (> 0)	
`beta`	prior Beta beta (> 0)	
`successes`	number of successes k	
`trials`	number of trials n	

output	meaning
`post_alpha`	posterior $\alpha = \alpha + k$
`post_beta`	posterior $\beta = \beta + n - k$
`posterior_mean`	posterior mean
`posterior_variance`	posterior variance
`posterior_sd`	posterior standard deviation

- **Method:** conjugate update of the beta prior with the binomial likelihood
- **Assumes:** alpha, $\beta > 0$; $0 \le k \le n$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)

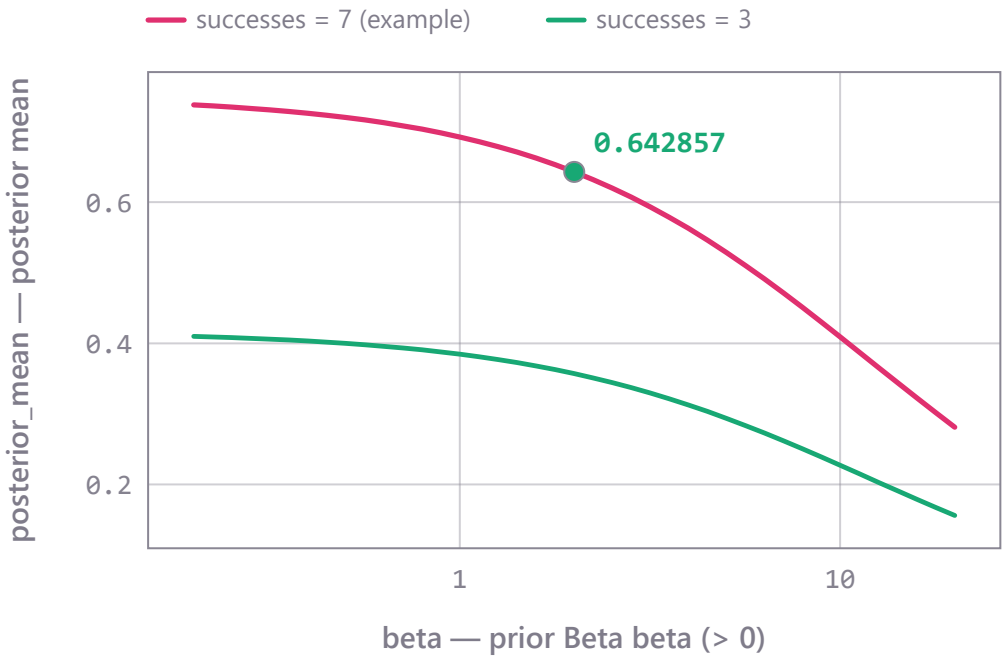


A Beta(2, 2) prior updated with 7 successes in 10 trials: the posterior Beta(9, 5) with mean 0.642857 and standard deviation 0.123718; shaded: the equal-tailed 95 % credible interval [0.385738, 0.861421]. Both densities are drawn from the package's cumulative probabilities at 199 values of θ (each interval's probability divided by its width).

	1	2	3	4	5	6
1	0.5	0.4	0.333	0.286	0.25	0.222
2	0.6	0.5	0.429	0.375	0.333	0.3
3	0.667	0.571	0.5	0.444	0.4	0.364
4	0.714	0.625	0.556	0.5	0.455	0.417
5	0.75	0.667	0.6	0.545	0.5	0.462
6	0.778	0.7	0.636	0.583	0.538	0.5

rows: successes s = 0 … 5, columns: failures f = 0 … 5

The posterior mean of a success probability under a Beta(2, 2) prior after s successes (rows, 0 to 5) and f failures (columns, 0 to 5), each cell from the package (rounded to 3 decimals): from 0.222 after 5 failures to 0.778 after 5 successes, 0.5 with no data.



posterior_mean (posterior mean) against beta (prior Beta beta (> 0)) from 0.2 to 20 on a logarithmic axis. One curve for each value of successes (number of successes k): 7, 3. Other inputs as in the worked example: alpha = 2, trials = 10. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, beta = 2, posterior_mean = 0.642857.

```
bayesian_statistics_ai999x.solve('beta_binomial_posterior', alpha=2, beta=2, successes=7,
trials=10)
# -> post_alpha = 9, post_beta = 5, posterior_mean = 0.642857, posterior_variance = 0.0153061,
posterior_sd = 0.123718   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('beta_binomial_posterior', alpha=1, beta=1, successes=3,
trials=20)
# -> post_alpha = 4, post_beta = 18, posterior_mean = 0.181818, posterior_variance = 0.00646784,
posterior_sd = 0.0804229   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Beta-binomial posterior: prior Beta(2, 2), 7 successes in 10 trials. Find the posterior mean.	SOLVED	post_alpha = 9, post_beta = 5, posterior_mean = 0.642857, posterior_variance = 0.0153061, posterior_sd = 0.123718
With a uniform prior and 3 successes in 20 trials, what is the posterior distribution of p?	SOLVED	post_alpha = 4, post_beta = 18, posterior_mean = 0.181818, posterior_variance = 0.00646784, posterior_sd = 0.0804229
Beta-binomial posterior: prior Beta(2, 2), 12 successes in 10 trials. Posterior mean?	INVALID_INPUT	Beta-binomial conjugate posterior: successes cannot exceed trials.
Beta-binomial posterior: prior Beta(2, 2), successes = 7. Posterior mean?	NEEDS_INPUT	beta_binomial_posterior needs trials (number of trials n). Example: solve('beta_binomial_posterior', alpha=2, beta=2, su

Equal-tailed credible interval of a beta distribution — `beta_credible_interval`

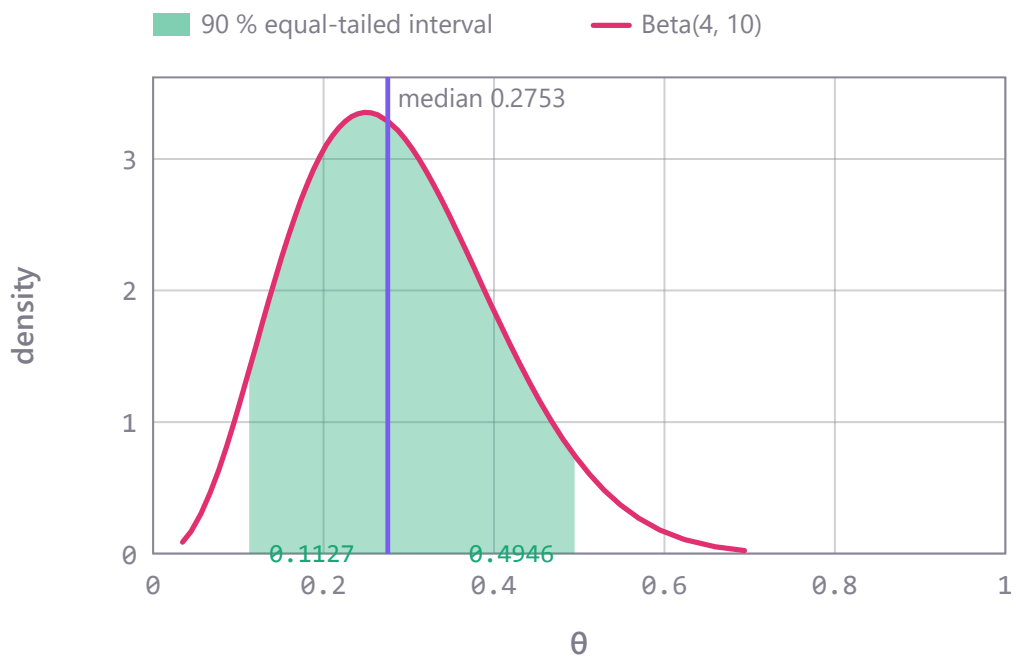
The central (equal-tailed) credible interval and the median of a Beta(alpha, beta) posterior.

$$\left[F^{-1}_{\text{Beta}(\alpha,\beta)}\left(\tfrac{1-c}{2}\right),\ F^{-1}_{\text{Beta}(\alpha,\beta)}\left(\tfrac{1+c}{2}\right)\right]$$

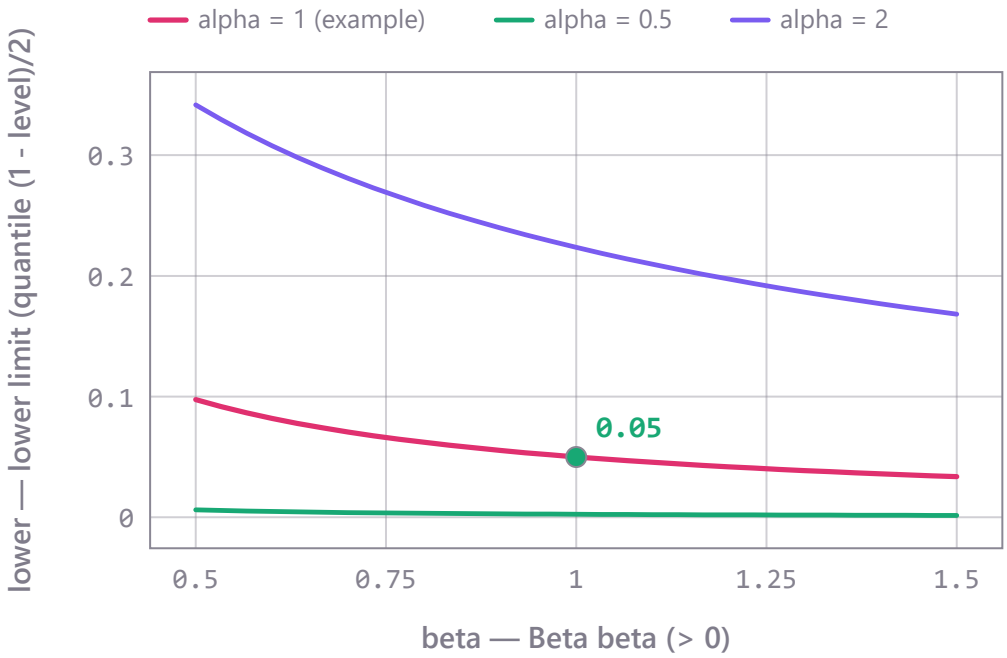
input	meaning	condition
`alpha`	Beta alpha (> 0)	
`beta`	Beta beta (> 0)	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`lower`	lower limit (quantile (1 - level)/2)
`upper`	upper limit (quantile (1 + level)/2)
`median`	posterior median

- **Method:** beta quantiles by safeguarded Newton inversion of the regularized incomplete beta function
- **Assumes:** alpha, $\beta > 0$; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Beta(4, 10) density (inputs chosen for the drawing: alpha = 4, beta = 10, level = 0.9); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Shaded: the returned 90 % equal-tailed credible interval [0.112666, 0.49465], 5 % of the probability left in each tail; marked: the returned median 0.275276.



lower (lower limit (quantile (1 - level)/2)) against beta (Beta beta (> 0)) from 0.5 to 1.5. One curve for each value of alpha (Beta alpha (> 0)): 1, 0.5, 2. Other inputs as in the worked example: level = 0.9. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, beta = 1, lower = 0.05.

```
bayesian_statistics_ai999x.solve('beta_credible_interval', alpha=1, beta=1, level=0.9)
# -> lower = 0.05, upper = 0.95, median = 0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('beta_credible_interval', alpha=3, beta=9)
# -> lower = 0.0602177, upper = 0.517756, median = 0.235786   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Equal-tailed 90% credible interval for a Beta(1, 1) posterior.	SOLVED	lower = 0.05, upper = 0.95, median = 0.5
Find the 50% credible interval of a beta distribution with alpha = 2 and beta = 1.	SOLVED	lower = 0.5, upper = 0.866025, median = 0.707107
Equal-tailed 95% credible interval for a Beta(0, 9) posterior.	INVALID_INPUT	Equal-tailed credible interval of a beta distribution: alpha must be positive.
Find the credible interval of a beta distribution with alpha = 12.	NEEDS_INPUT	beta_credible_interval needs beta (Beta beta (> 0)). Example: solve('beta_credible_interval', alpha=1, beta=1, level=0.9

Credible interval for a proportion (beta-binomial) — `beta_binomial_credible_interval`

Equal-tailed credible interval of a success probability after k successes in n trials with a Beta(alpha, beta) prior.

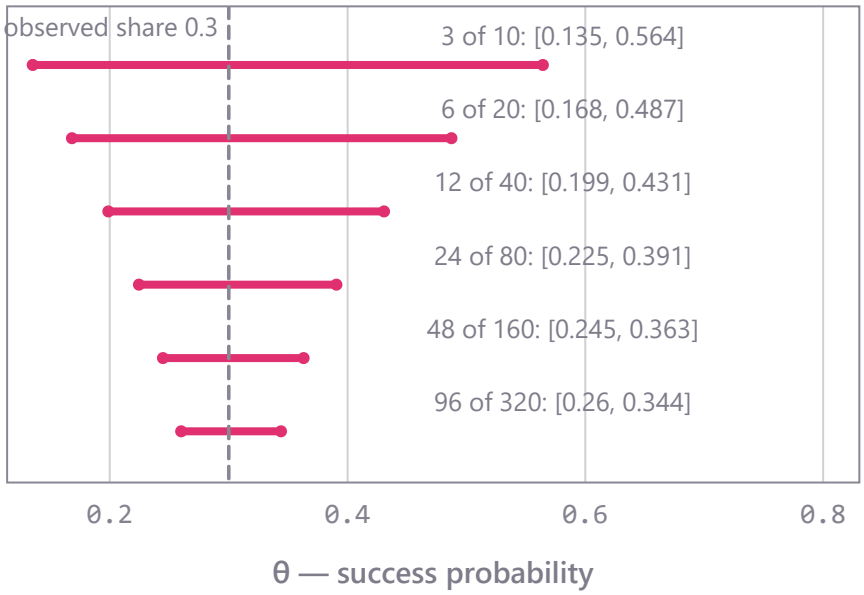
$$p \mid k \sim \mathrm{Beta}(\alpha + k, \beta + n - k), \quad [q_{(1-c)/2},\ q_{(1+c)/2}]$$

input	meaning	condition
`alpha`	prior Beta alpha (> 0)	

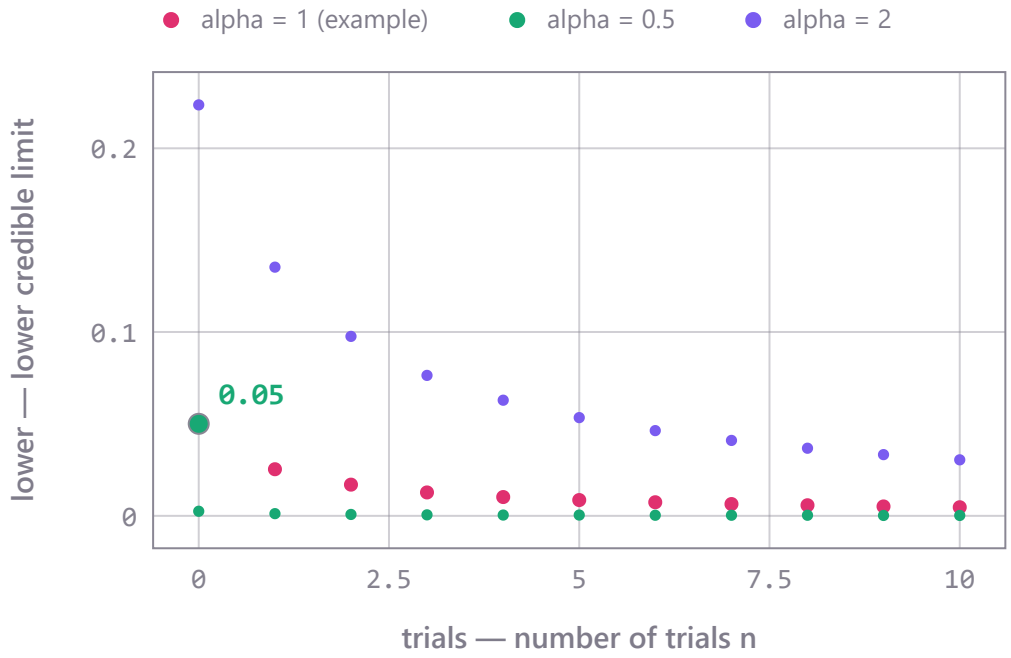
input	meaning	condition
`beta`	prior Beta beta (> 0)	
`successes`	number of successes k	
`trials`	number of trials n	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`post_alpha`	posterior alpha
`post_beta`	posterior beta
`lower`	lower credible limit
`upper`	upper credible limit

- **Method:** conjugate update, then beta quantiles
- **Assumes:** alpha, $\beta > 0$; $0 \leq k \leq n$; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



90 % equal-tailed credible intervals for a success probability under a uniform Beta(1, 1) prior, for the same observed share 0.3 in ever more trials (inputs chosen; each interval returned by the package): 3/10 [0.135075, 0.564374]; 6/20 [0.168176, 0.487389]; 12/40 [0.198829, 0.430652]; 24/80 [0.224502, 0.390697]; 48/160 [0.244689, 0.363031]; 96/320 [0.259957, 0.343933]. Four times the data halves the width.



lower (lower credible limit) against trials (number of trials n) from -10 to 10. One curve for each value of alpha (prior Beta alpha (> 0)): 1, 0.5, 2. Other inputs as in the worked example: beta = 1, successes = 0, level = 0.9. 11 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, trials = 0, lower = 0.05.

```
bayesian_statistics_ai999x.solve('beta_binomial_credible_interval', alpha=1, beta=1, successes=0, trials=0, level=0.9)
# -> post_alpha = 1, post_beta = 1, lower = 0.05, upper = 0.95   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('beta_binomial_credible_interval', alpha=2, beta=2, successes=7, trials=10, level=0.95)
# -> post_alpha = 9, post_beta = 5, lower = 0.385738, upper = 0.861421   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
90% credible interval for p: prior Beta(1, 1), 1 successes in 1 trials.	SOLVED	post_alpha = 2, post_beta = 1, lower = 0.223607, upper = 0.974679
With a uniform prior and 15 successes in 50 trials, find the 90% credible interval for the proportion.	SOLVED	post_alpha = 16, post_beta = 36, lower = 0.207694, upper = 0.41618
95% credible interval for p: prior Beta(2, 2), 11 successes in 10 trials.	INVALID_INPUT	Credible interval for a proportion (beta-binomial): successes cannot exceed trials.
95% credible interval for p: prior Beta(2, 2) and 7 successes.	NEEDS_INPUT	beta_binomial_credible_interval needs trials (number of trials n). Example: solve('beta_binomial_credible_interval', alp

Posterior probability that p exceeds a value — `beta_posterior_tail_probability`

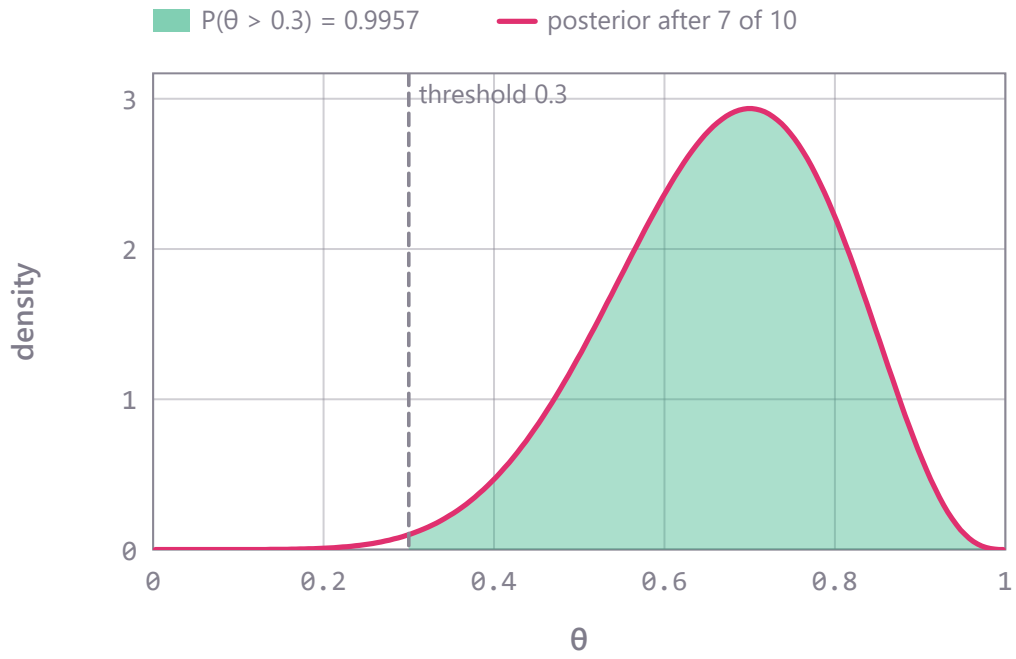
P(p > c) and P(p < c) under a Beta(alpha, beta) posterior, optionally after k successes in n trials.

$$P(p > c \mid k) = 1 - I_c(\alpha + k,\ \beta + n - k)$$

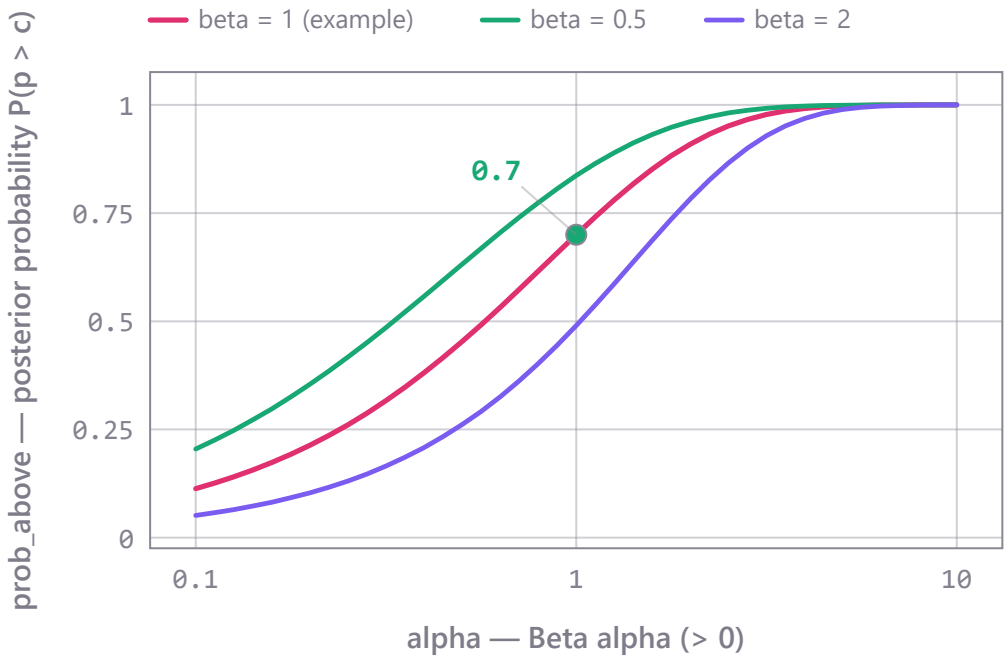
input	meaning	condition
`alpha`	Beta alpha (> 0)	
`beta`	Beta beta (> 0)	
`threshold`	the value c, $0 < c < 1$	
`successes`	successes to add (default 0)	
`trials`	trials to add (default 0)	

output	meaning
`prob_above`	posterior probability P(p > c)
`prob_below`	posterior probability P(p < c)

- **Method:** regularized incomplete beta function of the (updated) posterior
- **Assumes:** alpha, $\beta > 0$; $0 < c < 1$; $0 \leq k \leq n$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The posterior of a success probability from a Beta(1, 1) prior (the worked example's) after 7 successes in 10 trials (data chosen for the drawing), drawn from the package's P(θ < c) at 199 values of c (differenced). Shaded: the returned P(θ > 0.3) = 0.995709 (P(θ < 0.3) = 0.00429089).



prob_above (posterior probability P(p > c)) against alpha (Beta alpha (> 0)) from 0.1 to 10 on a logarithmic axis. One curve for each value of beta (Beta beta (> 0)): 1, 0.5, 2. Other inputs as in the worked example: threshold = 0.3. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, alpha = 1, prob_above = 0.7.

```
bayesian_statistics_ai999x.solve('beta_posterior_tail_probability', alpha=1, beta=1, threshold=0.3)
# -> prob_above = 0.7, prob_below = 0.3   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('beta_posterior_tail_probability', alpha=2, beta=1, threshold=0.5)
# -> prob_above = 0.75, prob_below = 0.25   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
For a Beta(2, 1) posterior, what is the probability that p exceeds 0.5?	SOLVED	prob_above = 0.75, prob_below = 0.25
Prior Beta(1, 1) and 1 successes in 1 trials: what is the probability that p is less than 0.5?	SOLVED	prob_above = 0.75, prob_below = 0.25
For a Beta(8, 4) posterior, what is the probability that p exceeds 1.5?	INVALID_INPUT	Posterior probability that p exceeds a value: the threshold must be strictly between 0 and 1.

Jeffreys interval for a binomial proportion — `jeffreys_interval`

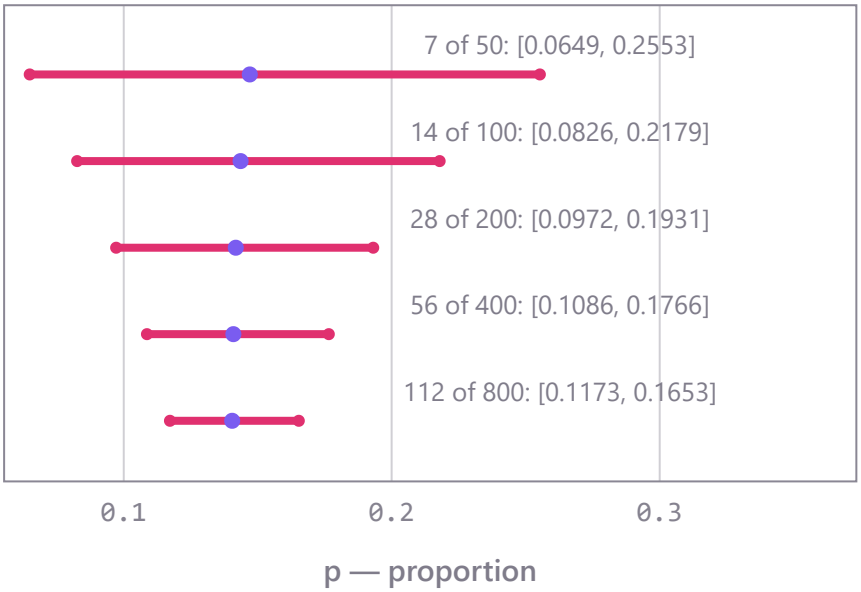
The equal-tailed Bayesian interval for a proportion under the Jeffreys Beta($1/2$, 1/2) prior (lower limit 0 when $k = 0$ and upper limit 1 when k = n).

$$\left[B^{-1}\left(\tfrac{1-c}{2};\, k+\tfrac{1}{2},\, n-k+\tfrac{1}{2}\right),\; B^{-1}\left(\tfrac{1+c}{2};\, k+\tfrac{1}{2},\, n-k+\tfrac{1}{2}\right)\right]$$

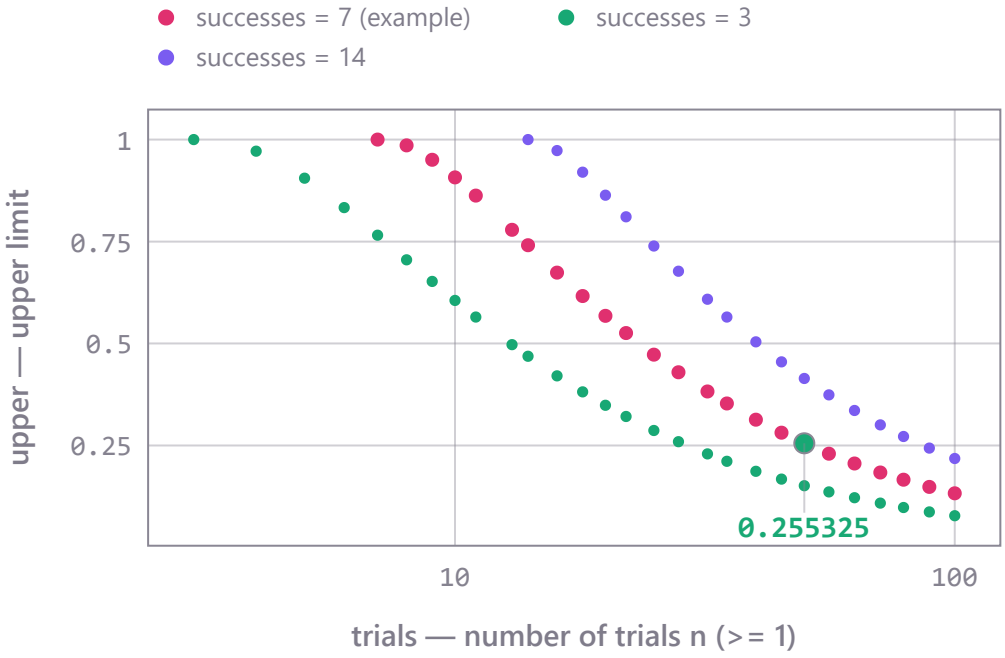
input	meaning	condition
`successes`	number of successes k	
`trials`	number of trials n (>= 1)	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`lower`	lower limit
`upper`	upper limit
`posterior_mean`	$(k+1/2)/(n+1)$

- **Method:** quantiles of Beta($k+1/2$, n - k + 1/2)
- **Assumes:** $0 \leq k \leq n$; $n \geq 1$; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Jeffreys 95 % intervals for a proportion, starting from the worked example (7 of 50) and multiplying both counts by 2, 4, 8 and 16 (inputs chosen); intervals and posterior means (dots) as returned: 7/50 [0.0648717, 0.255325]; 14/100 [0.082566, 0.217898]; 28/200 [0.0971955, 0.19313]; 56/400 [0.108632, 0.176554]; 112/800 [0.11727, 0.165329].



upper (upper limit) against trials (number of trials n (>= 1)) from 1 to 100 on a logarithmic axis. One curve for each value of successes (number of successes k): 7, 3, 14. Other inputs as in the worked example: level = 0.95. 24 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, trials = 50, upper = 0.255325.

```
bayesian_statistics_ai999x.solve('jeffreys_interval', successes=7, trials=50, level=0.95)
# -> lower = 0.0648717, upper = 0.255325, posterior_mean = 0.147059   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('jeffreys_interval', successes=0, trials=20, level=0.9)
# -> lower = 0, upper = 0.0904764, posterior_mean = 0.0238095   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Jeffreys interval at 95% for 7 successes in 50 trials.	SOLVED	lower = 0.0648717, upper = 0.255325, posterior_mean = 0.147059
Find the 90% Jeffreys interval with successes = 20 and trials = 20.	SOLVED	lower = 0.909524, upper = 1, posterior_mean = 0.97619
Jeffreys interval at 95% for 9 successes in 5 trials.	INVALID_INPUT	Jeffreys interval for a binomial proportion: successes cannot exceed trials.
Jeffreys interval at 95% with successes = 9.	NEEDS_INPUT	jeffreys_interval needs trials (number of trials n (>= 1)). Example: solve('jeffreys_interval', successes=7, trials=50,

Posterior predictive probability (beta-binomial) — `beta_binomial_predictive`

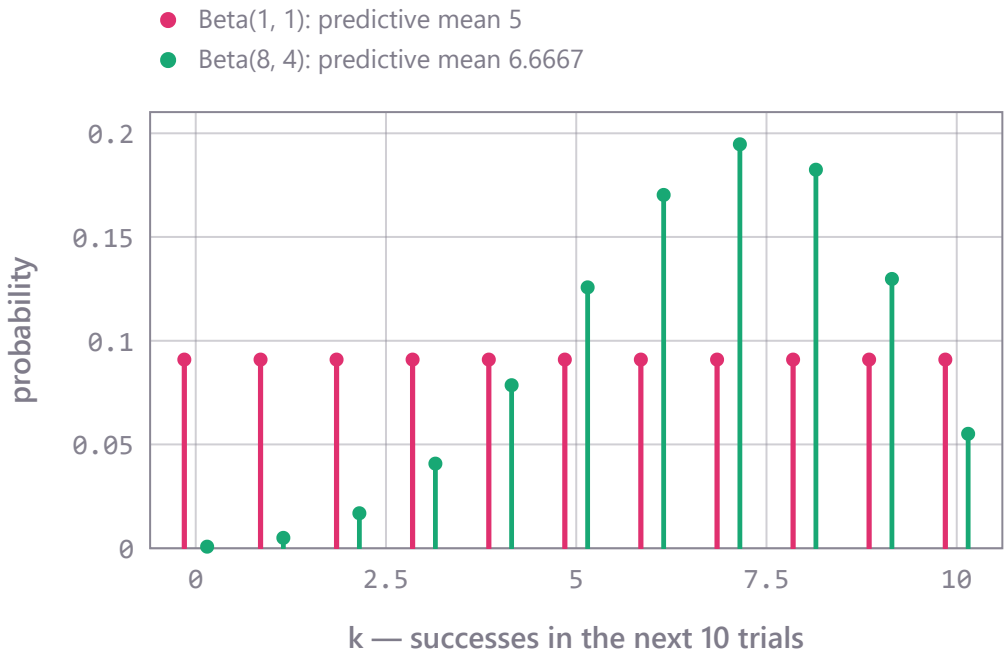
Probability of exactly j successes in m future trials when the success probability has a Beta(alpha, beta) posterior; the cumulative probability and the predictive mean.

$$P(J = j) = \binom{m}{j} \frac{B(\alpha + j,\ \beta + m - j)}{B(\alpha, \beta)}$$

input	meaning	condition
`alpha`	posterior alpha (> 0)	
`beta`	posterior beta (> 0)	
`future_trials`	number of future trials m (>= 1)	
`future_successes`	number of future successes j	

output	meaning
`probability`	P(exactly j successes in m trials)
`prob_at_most`	P(at most j)
`prob_more`	P(more than j)
`predictive_mean`	expected number of successes m alpha/(alpha+beta)

- **Method:** beta-binomial distribution (the binomial averaged over the beta posterior)
- **Assumes:** alpha, $\beta > 0$; $0 \le j \le m$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The beta-binomial predictive distribution of the number of successes in 10 future trials, for a flat Beta(1, 1) and a Beta(8, 4) belief (inputs chosen); each probability returned by the package: Beta(1, 1): 0.0909, 0.0909, 0.0909, 0.0909, 0.0909, 0.0909, 0.0909, 0.0909, 0.0909, 0.0909, 0.0909 (mean 5); Beta(8, 4): 0.0008, 0.005, 0.0168, 0.0408, 0.0786, 0.1257, 0.1703, 0.1946, 0.1824, 0.1297, 0.0551 (mean 6.66667). A flat belief gives every count the same chance.

```
bayesian_statistics_ai999x.solve('beta_binomial_predictive', alpha=1, beta=1, future_trials=4,
future_successes=2)
# -> probability = 0.2, prob_at_most = 0.6, prob_more = 0.4, predictive_mean = 2   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('beta_binomial_predictive', alpha=8, beta=4, future_trials=5,
future_successes=3)
# -> probability = 0.274725, prob_at_most = 0.516484, prob_more = 0.483516, predictive_mean =
3.33333   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Posterior Beta(1, 1): probability of exactly 2 successes in the next 4 trials (posterior predictive)?	SOLVED	probability = 0.2, prob_at_most = 0.6, prob_more = 0.4, predictive_mean = 2
Posterior predictive with alpha = 2, beta = 3, future trials = 1, future successes = 1.	SOLVED	probability = 0.4, prob_at_most = 1, prob_more = 0, predictive_mean = 0.4
Posterior predictive with alpha = 2, beta = 3, future trials = 3, future successes = 4.	INVALID_INPUT	Posterior predictive probability (beta-binomial): future successes cannot exceed future trials.

Marginal likelihood (evidence) of binomial data under a beta prior — `beta_binomial_evidence`

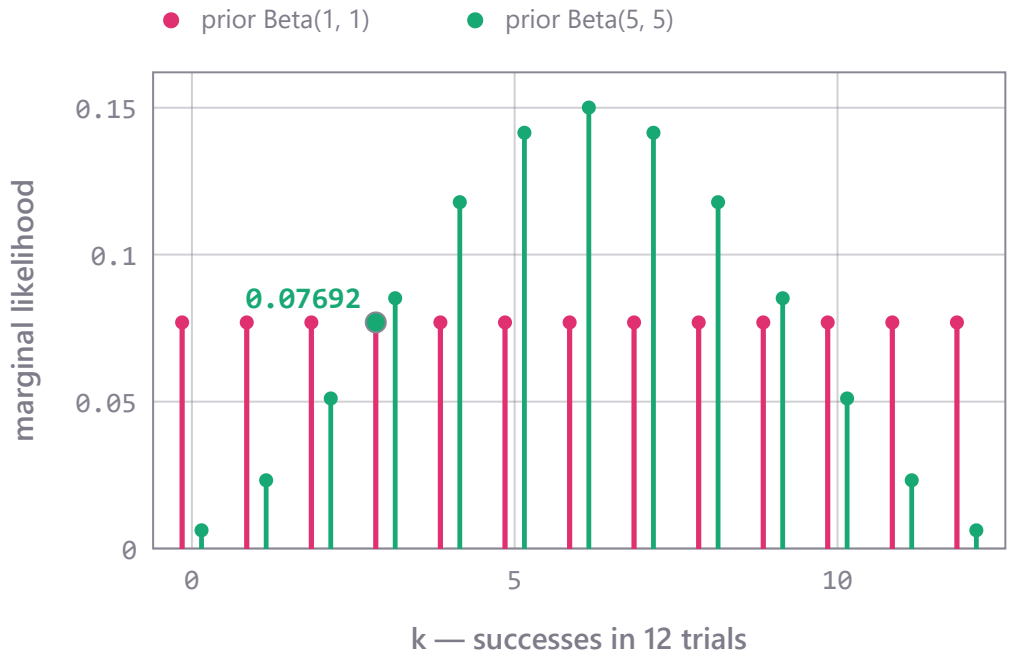
Prior predictive probability of k successes in n trials when p ~ Beta(alpha, beta), and its logarithm.

$$p(k) = \binom{n}{k} \frac{B(\alpha + k,\ \beta + n - k)}{B(\alpha, \beta)}$$

input	meaning	condition
`alpha`	prior Beta alpha (> 0)	
`beta`	prior Beta beta (> 0)	
`successes`	number of successes k	
`trials`	number of trials n	

output	meaning
`marginal_likelihood`	P(k | n, alpha, beta)
`log_marginal_likelihood`	its natural log

- **Method:** closed form with log-beta functions
- **Assumes:** alpha, $\beta > 0$; $0 \leq k \leq n$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The marginal likelihood (evidence) of k successes in 12 trials for every k, under the worked example's Beta(1, 1) prior and a Beta(5, 5) prior (chosen); each value returned by the package (Beta(1, 1): from 0.0769231 to 0.0769231; Beta(5, 5): from 0.00619195 to 0.150036). Marked: the worked example, k = 3, evidence 0.0769231.

```
bayesian_statistics_ai999x.solve('beta_binomial_evidence', alpha=1, beta=1, successes=3, trials=12)
# -> marginal_likelihood = 0.0769231, log_marginal_likelihood = -2.56495   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('beta_binomial_evidence', alpha=2, beta=2, successes=7, trials=10)
# -> marginal_likelihood = 0.111888, log_marginal_likelihood = -2.19026   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Marginal likelihood of 7 successes in 10 trials under a Beta(1, 1) prior.	SOLVED	marginal_likelihood = 0.0909091, log_marginal_likelihood = -2.3979

question	outcome	result
Evidence for successes = 3, trials = 12 with prior alpha = 1, beta = 1.	SOLVED	marginal_likelihood = 0.0769231, log_marginal_likelihood = -2.56495
Marginal likelihood of 7 successes in 10 trials under a Beta(-1, 2) prior.	INVALID_INPUT	Marginal likelihood (evidence) of binomial data under a beta prior: alpha must be positive.

MAP estimate of a proportion (beta-binomial) — `map_beta_binomial`

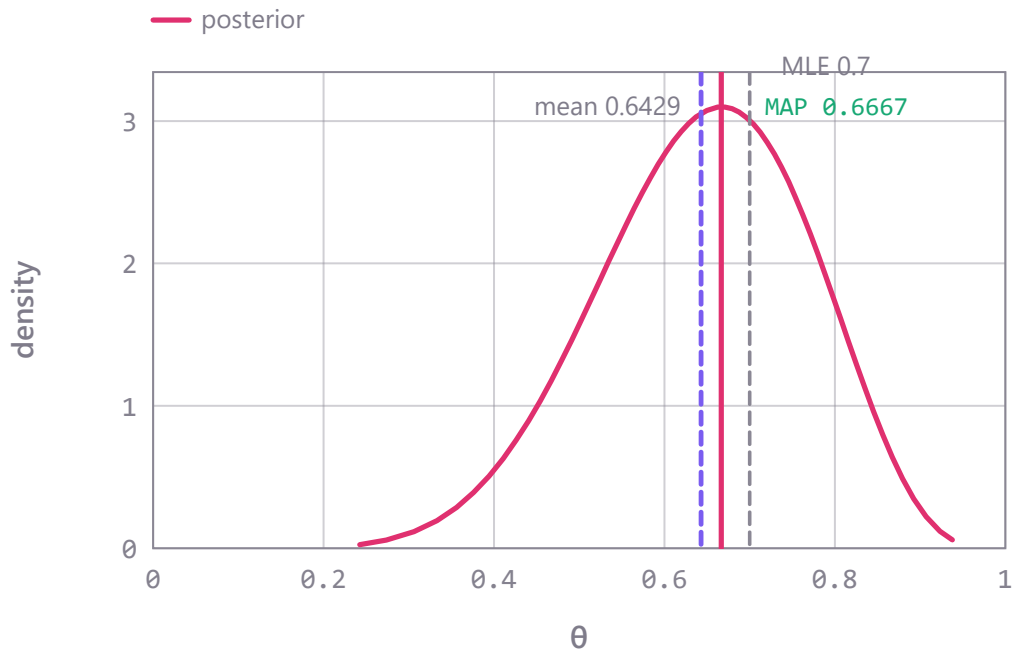
The posterior mode (maximum a posteriori estimate) of p after k successes in n trials with a Beta(alpha, beta) prior, with the maximum-likelihood estimate and the posterior mean for comparison.

$$\hat{p}_{MAP} = \frac{\alpha + k - 1}{\alpha + \beta + n - 2}$$

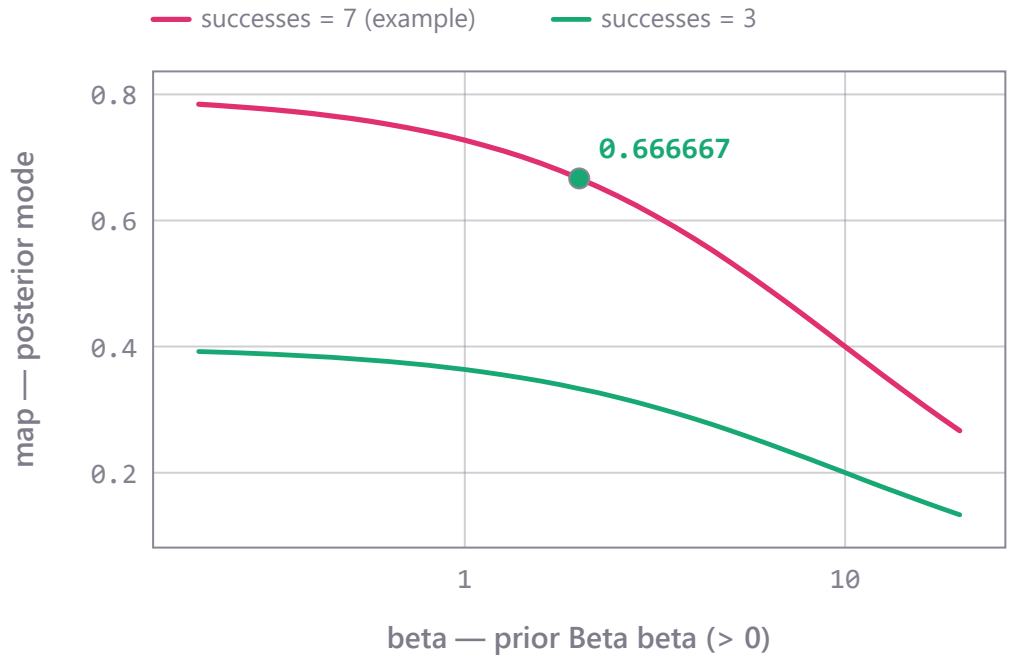
input	meaning	condition
`alpha`	prior Beta alpha (> 0)	
`beta`	prior Beta beta (> 0)	
`successes`	number of successes k	
`trials`	number of trials n	

output	meaning
`map`	posterior mode $(\alpha + k - 1)/(\alpha + \beta + n - 2)$
`mle`	k/n
`posterior_mean`	(alpha + k)/(alpha + beta + n)

- **Method:** maximize the log posterior (closed form of its stationary point)
- **Assumes:** $\alpha + k > 1$ and $\beta + n - k > 1$ (interior mode)
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The posterior Beta(9, 5) for a Beta(2, 2) prior and 7 of 10 (the worked example; posterior parameters as returned by the posterior capability); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Marked: the returned MAP 0.666667 (the peak), maximum-likelihood estimate 0.7 and posterior mean 0.642857.



map (posterior mode (alpha + k - 1)/(alpha + beta + n - 2)) against beta (prior Beta beta (> 0)) from 0.2 to 20 on a logarithmic axis. One curve for each value of successes (number of successes k): 7, 3. Other inputs as in the worked example: alpha = 2, trials = 10. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, beta = 2, map = 0.666667.

```
bayesian_statistics_ai999x.solve('map_beta_binomial', alpha=2, beta=2, successes=7, trials=10)
# -> map = 0.666667, mle = 0.7, posterior_mean = 0.642857   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('map_beta_binomial', alpha=3, beta=5, successes=4, trials=12)
# -> map = 0.333333, mle = 0.333333, posterior_mean = 0.35   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
MAP estimate of p with prior Beta(2, 2) and 7 successes in 10 trials.	SOLVED	map = 0.666667, mle = 0.7, posterior_mean = 0.642857
Find the posterior mode with alpha = 3, beta = 5, successes = 4, trials = 12.	SOLVED	map = 0.333333, mle = 0.333333, posterior_mean = 0.35
MAP estimate of p with prior Beta(1, 1) and 0 successes in 10 trials.	INVALID_INPUT	MAP estimate of a proportion (beta-binomial): the posterior mode lies on the boundary (alpha + k <= 1 or beta + n - k <=

Laplace (normal) approximation of a beta posterior — `laplace_approximation_beta`

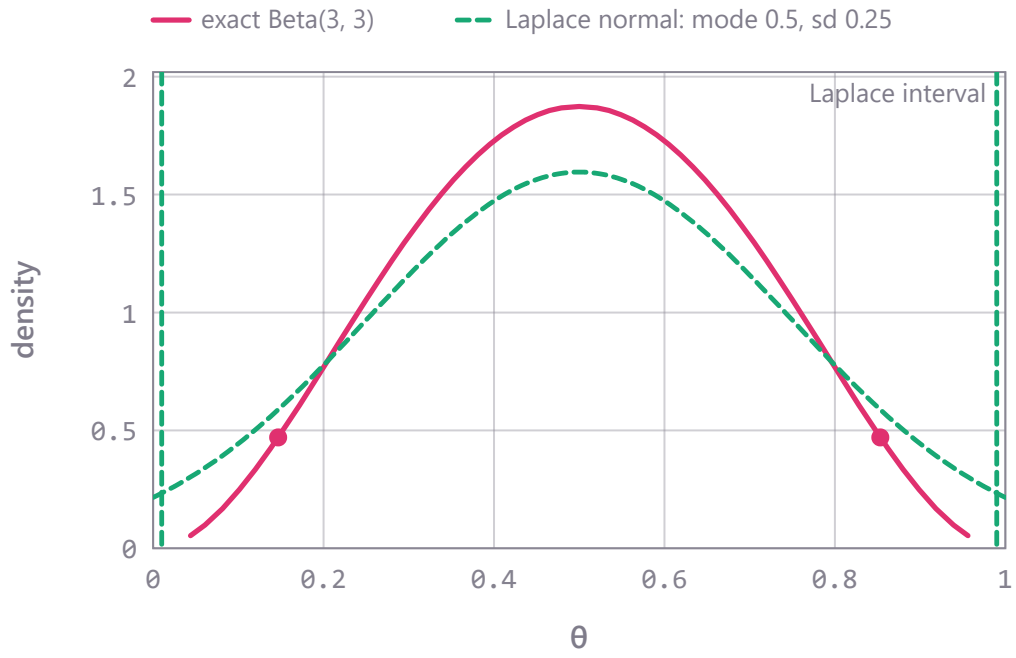
Approximate a Beta(alpha, beta) posterior by a normal distribution centred at its mode with variance -1/(second derivative of the log density); the approximate credible interval.

$$\hat{p} = \frac{\alpha - 1}{\alpha + \beta - 2}, \quad \sigma^2 = \left[\frac{\alpha - 1}{\hat{p}^2} + \frac{\beta - 1}{(1 - \hat{p})^2} \right]^{-1}$$

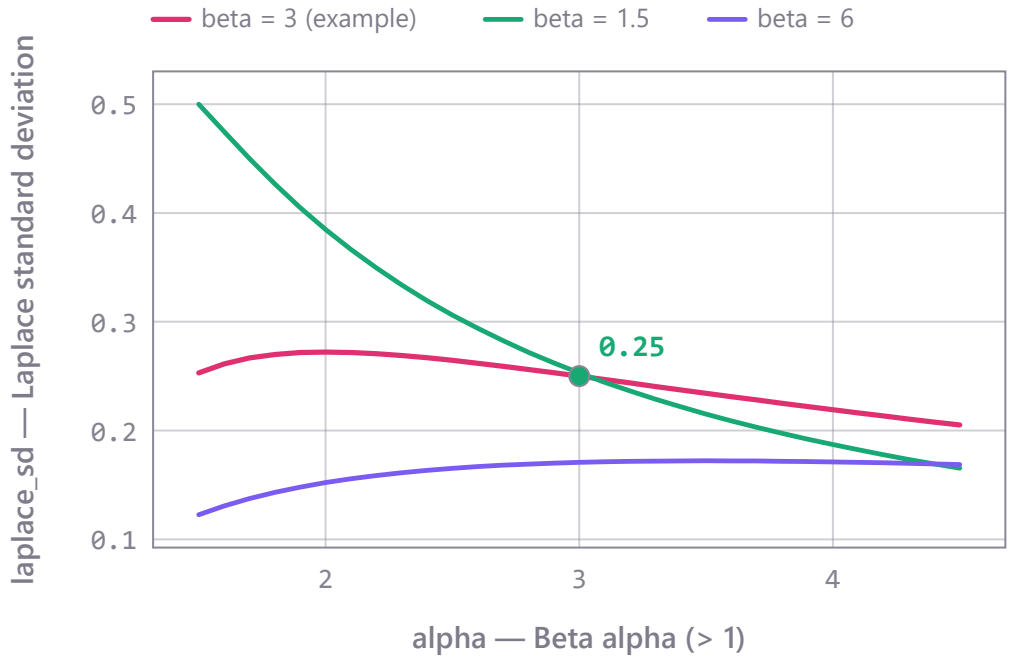
input	meaning	condition
`alpha`	Beta alpha (> 1)	
`beta`	Beta beta (> 1)	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`mode`	posterior mode
`laplace_sd`	Laplace standard deviation
`lower`	approximate lower limit
`upper`	approximate upper limit

- **Method:** second-order Taylor expansion of the log posterior at its mode
- **Assumes:** $\alpha > 1$ and $\beta > 1$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Beta(3, 3) density (from the package's quantiles at 60 cumulative probabilities, differenced) against its Laplace approximation, a normal density drawn from the returned mode 0.5 and sd 0.25. Dashed verticals: the returned 95 % Laplace interval [0.010009, 0.989991]; dots: the exact equal-tailed interval [0.146633, 0.853367].



laplace_sd (Laplace standard deviation) against alpha (Beta alpha (> 1)) from 1.5 to 4.5. One curve for each value of beta (Beta beta (> 1)): 3, 1.5, 6. Other inputs as in the worked example: level = 0.95. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, alpha = 3, laplace_sd = 0.25.

```
bayesian_statistics_ai999x.solve('laplace_approximation_beta', alpha=3, beta=3, level=0.95)
# -> mode = 0.5, laplace_sd = 0.25, lower = 0.010009, upper = 0.989991   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('laplace_approximation_beta', alpha=10, beta=4)
# -> mode = 0.75, laplace_sd = 0.125, lower = 0.505005, upper = 0.994995   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Laplace approximation of a Beta(3, 3) posterior with a 95% interval.	SOLVED	mode = 0.5, laplace_sd = 0.25, lower = 0.010009, upper = 0.989991
Use the Laplace approximation for a beta posterior with alpha = 21 and beta = 11.	SOLVED	mode = 0.666667, laplace_sd = 0.0860663, lower = 0.49798, upper = 0.835354
Laplace approximation of a Beta(1, 4) posterior.	INVALID_INPUT	Laplace (normal) approximation of a beta posterior: the Laplace approximation needs an interior mode: alpha > 1 and beta

Highest-density interval of a beta posterior — `beta_hdi`

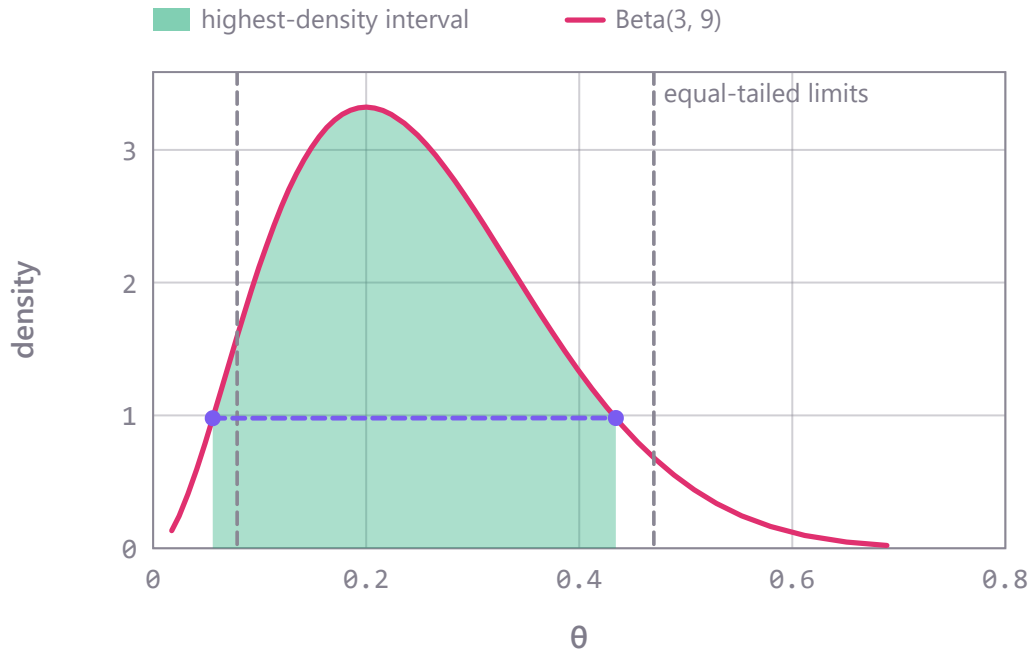
The shortest interval holding a stated posterior probability of a unimodal Beta(alpha, beta): equal density at both ends.

$$F(u) - F(l) = c, \quad f(l) = f(u)$$

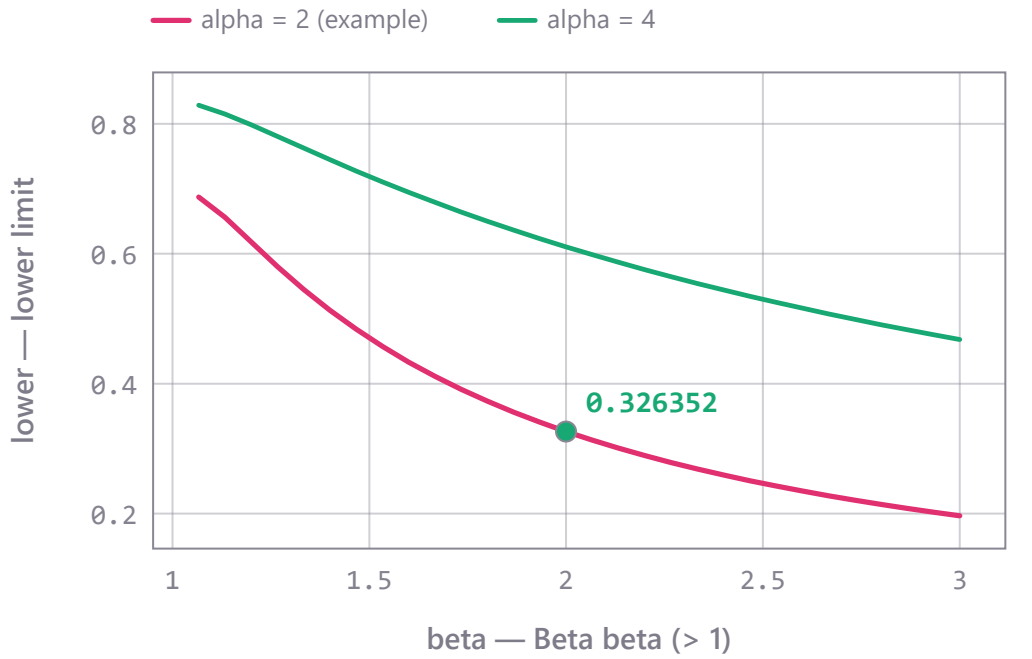
input	meaning	condition
`alpha`	Beta alpha (> 1)	
`beta`	Beta beta (> 1)	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`lower`	lower limit
`upper`	upper limit
`width`	interval width

- **Method:** bisection on the lower limit l with $u = F^{-1}$(F(l) + level) until f(l) = f(u)
- **Assumes:** $\alpha > 1$ and $\beta > 1$ (unimodal, zero density at both ends); 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The skewed Beta(3, 9) density at level 0.9 (inputs chosen so the two intervals differ); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Shaded: the returned highest-density interval [0.0559771, 0.434394] (width 0.378417):the density is equal at both ends (dashed chord), so no shorter interval holds the same probability; dashed verticals: the equal-tailed interval [0.07882, 0.470087] at the same level, width 0.391267 (difference of the returned limits).



lower (lower limit) against beta (Beta beta (> 1)) from 1 to 3. One curve for each value of alpha (Beta alpha (> 1)): 2, 4. Other inputs as in the worked example: level = 0.5. 30 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, beta = 2, lower = 0.326352.

```
bayesian_statistics_ai999x.solve('beta_hdi', alpha=2, beta=2, level=0.5)
# -> lower = 0.326352, upper = 0.673648, width = 0.347296   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('beta_hdi', alpha=3, beta=9, level=0.95)
# -> lower = 0.0405552, upper = 0.483724, width = 0.443168   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
95% highest-density interval of a Beta(3, 9) posterior.	SOLVED	lower = 0.0405552, upper = 0.483724, width = 0.443168
Find the 90% HDI for a beta posterior with alpha = 1 and beta = 4.	INVALID_INPUT	Highest-density interval of a beta posterior: the highest-density interval is computed for unimodal Beta with alpha > 1
Find the 90% HDI for a beta posterior with alpha = 12.	NEEDS_INPUT	beta_hdi needs beta (Beta beta (> 1)). Example: solve('beta_hdi', alpha=2, beta=2, level=0.5).

Beta prior from a mean and standard deviation — `beta_prior_from_mean_sd`

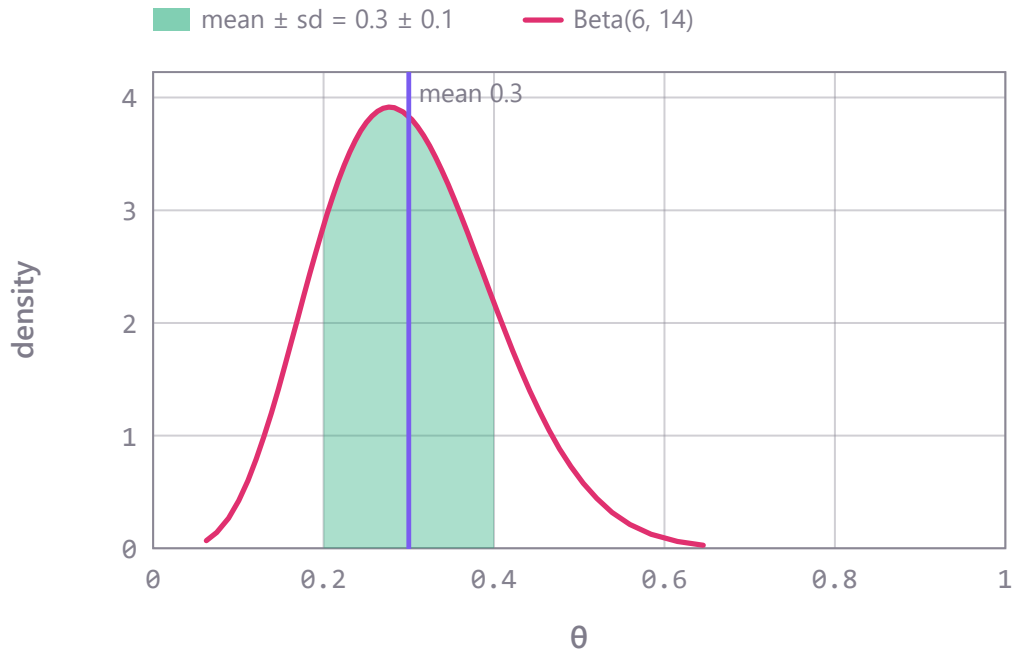
The Beta(alpha, beta) prior with a stated mean and standard deviation (method of moments) and its effective prior sample size alpha + beta.

$$\nu = \frac{m(1-m)}{s^2} - 1, \quad \alpha = m\nu, \quad \beta = (1-m)\nu$$

input	meaning	condition
`mean`	prior mean, $0 < m < 1$	
`sd`	prior standard deviation, $sd^2 < m$(1 - m)	

output	meaning
`alpha`	alpha
`beta`	beta
`prior_sample_size`	alpha + beta

- **Method:** match the beta mean and variance
- **Assumes:** 0 < mean < 1; 0 < sd^2 < mean (1 - mean)
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Beta prior with the stated mean 0.3 and standard deviation 0.1: the returned alpha = 6, beta = 14 (prior sample size 20); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Shaded: mean ± one sd.

```
bayesian_statistics_ai999x.solve('beta_prior_from_mean_sd', mean=0.3, sd=0.1)
# -> alpha = 6, beta = 14, prior_sample_size = 20   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('beta_prior_from_mean_sd', mean=0.5, sd=0.2886751345948129)
# -> alpha = 1, beta = 1, prior_sample_size = 2   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Find the beta prior with mean = 0.3 and standard deviation = 0.1.	SOLVED	alpha = 6, beta = 14, prior_sample_size = 20
Which beta distribution has mean 0.8 and sd 0.05?	SOLVED	alpha = 50.4, beta = 12.6, prior_sample_size = 63
Find the beta prior with mean = 0.5 and standard deviation = 0.6.	INVALID_INPUT	Beta prior from a mean and standard deviation: the variance must be smaller than mean x (1 - mean) for a beta distributi
Find the beta prior with mean = 0.5 and an unknown standard deviation.	NEEDS_INPUT	beta_prior_from_mean_sd needs sd (prior standard deviation, sd^2 < m(1 - m)). Example: solve('beta_prior_from_mean_sd',

Bayesian A/B test: probability that A beats B — `ab_test_probability_beta`

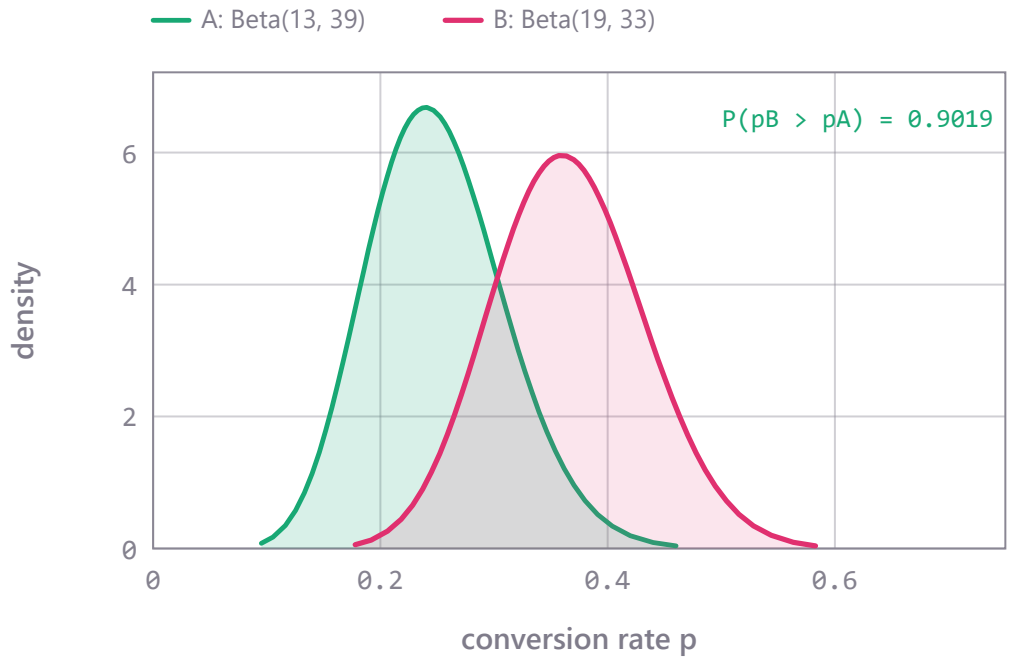
P(p_A > p_B) for independent Beta posteriors of two success probabilities (alpha_b a whole number, as after a uniform or any integer-parameter prior).

$$P(p_B > p_A) = \sum_{i=0}^{\alpha_B - 1} \frac{B(\alpha_A + i,\ \beta_A + \beta_B)}{(\beta_B + i)\, B(1 + i,\ \beta_B)\, B(\alpha_A, \beta_A)}$$

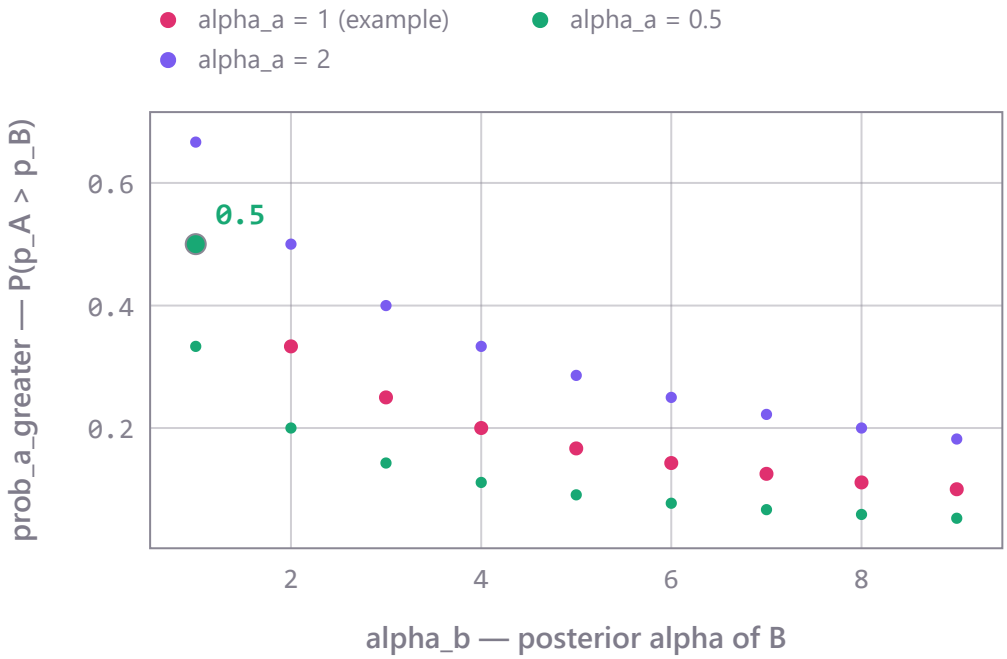
input	meaning	condition
`alpha_a`	posterior alpha of A	
`beta_a`	posterior beta of A	
`alpha_b`	posterior alpha of B (whole number)	
`beta_b`	posterior beta of B	

output	meaning
`prob_a_greater`	P(p_A > p_B)
`prob_b_greater`	P(p_B > p_A)

- **Method:** exact finite sum of beta functions (Miller's formula)
- **Assumes:** all parameters > 0; alpha_b a whole number
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



An A/B comparison (inputs chosen: 12 of 50 conversions for A and 18 of 50 for B on a uniform prior, giving the posteriors Beta(13, 39) and Beta(19, 33)); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Marked: the returned P(pB > pA) = 0.901949 (P(pA > pB) = 0.0980514).



prob_a_greater (P(p_A > p_B)) against alpha_b (posterior alpha of B (whole number)) from 1 to 9. One curve for each value of alpha_a (posterior alpha of A): 1, 0.5, 2. Other inputs as in the worked example: beta_a = 1, beta_b = 1. 9 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, alpha_b = 1, prob_a_greater = 0.5.

```
bayesian_statistics_ai999x.solve('ab_test_probability_beta', alpha_a=1, beta_a=1, alpha_b=1,
beta_b=1)
# -> prob_a_greater = 0.5, prob_b_greater = 0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('ab_test_probability_beta', alpha_a=2, beta_a=1, alpha_b=1,
beta_b=1)
# -> prob_a_greater = 0.666667, prob_b_greater = 0.333333   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Bayesian A/B test: alpha_a = 1, beta_a = 1, alpha_b = 1, beta_b = 1. Probability that A beats B?	SOLVED	prob_a_greater = 0.5, prob_b_greater = 0.5
With alpha_a = 2, beta_a = 1, alpha_b = 1, beta_b = 1, what is the probability that p_A is greater than p_B?	SOLVED	prob_a_greater = 0.666667, prob_b_greater = 0.333333
Bayesian A/B test: alpha_a = 21, beta_a = 81, alpha_b = 3.5, beta_b = 71. Probability that A beats B?	INVALID_INPUT	Bayesian A/B test: probability that A beats B: alpha_b must be a whole number >= 1.

Gamma-Poisson conjugate posterior of a rate — `gamma_poisson_posterior`

Update a Gamma(shape, rate) prior on a Poisson rate with observed counts: posterior parameters, mean, variance and standard deviation.

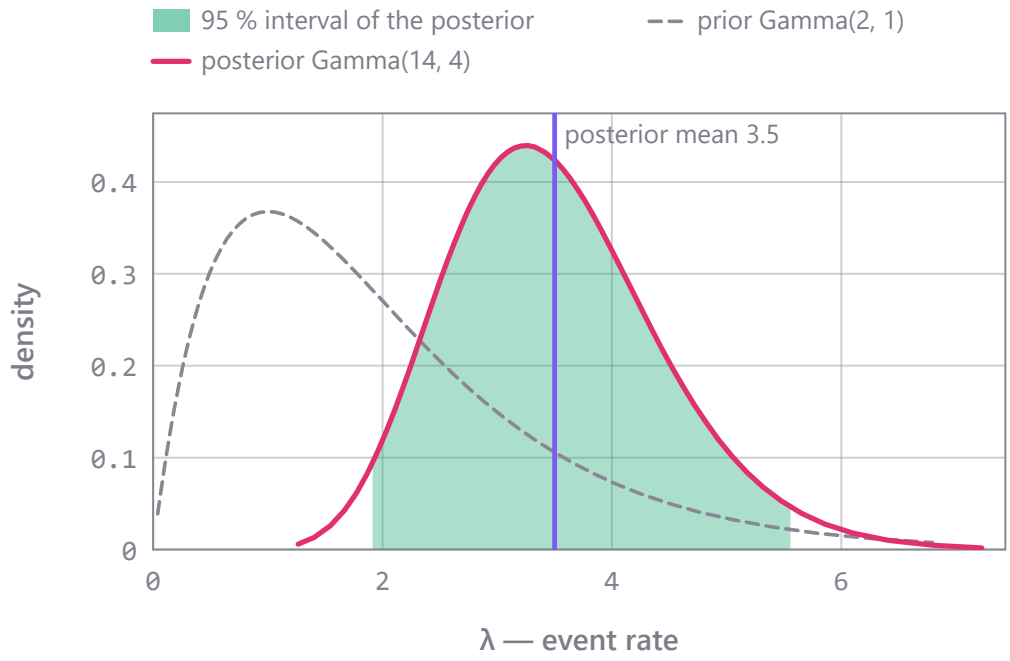
$$\lambda \mid x \sim \mathrm{Gamma}(\alpha + \textstyle\sum x_i,\ \beta + n)$$

input	meaning	condition
`shape`	prior shape alpha (> 0)	

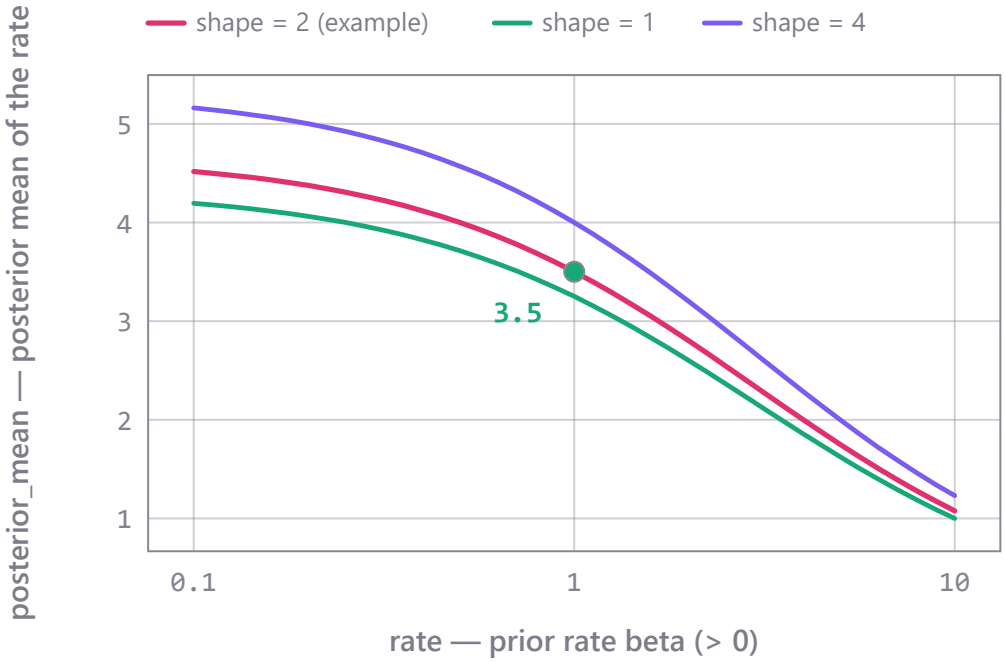
input	meaning	condition
`rate`	prior rate beta (> 0)	
`counts`	observed counts (whole numbers), one per unit interval	

output	meaning
`post_shape`	alpha + sum of counts
`post_rate`	beta + number of intervals
`posterior_mean`	posterior mean of the rate
`posterior_variance`	posterior variance
`posterior_sd`	posterior standard deviation

- **Method:** conjugate update of the gamma prior with the Poisson likelihood
- **Assumes:** shape, rate > 0; counts whole numbers >= 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A Gamma prior (shape 2, rate 1) updated with the Poisson counts 3, 5, 4 (the worked example): the returned posterior has shape 14 and rate 4, mean 3.5; each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Shaded: the posterior 95 % equal-tailed interval [1.91348, 5.5576].



posterior_mean (posterior mean of the rate) against rate (prior rate beta (> 0)) from 0.1 to 10 on a logarithmic axis. One curve for each value of shape (prior shape alpha (> 0)): 2, 1, 4. Other inputs as in the worked example: counts = 3, 5, 4. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, rate = 1, posterior_mean = 3.5.

```
bayesian_statistics_ai999x.solve('gamma_poisson_posterior', shape=2, rate=1, counts='3, 5, 4')
# -> post_shape = 14, post_rate = 4, posterior_mean = 3.5, posterior_variance = 0.875, posterior_sd
= 0.935414   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('gamma_poisson_posterior', shape=1, rate=0.5, counts=[0, 2, 1, 1])
# -> post_shape = 5, post_rate = 4.5, posterior_mean = 1.11111, posterior_variance = 0.246914,
posterior_sd = 0.496904   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Gamma-Poisson posterior: prior Gamma(2, 1), counts = [3, 5, 4]. Posterior mean of the rate?	SOLVED	post_shape = 14, post_rate = 4, posterior_mean = 3.5, posterior_variance = 0.875, posterior_sd = 0.935414
With shape = 1 and rate = 0.5 and Poisson counts = [0, 2, 1, 1], find the posterior of the rate.	SOLVED	post_shape = 5, post_rate = 4.5, posterior_mean = 1.11111, posterior_variance = 0.246914, posterior_sd = 0.496904
Gamma-Poisson posterior: prior Gamma(2, 1), counts = [3, -5, 4]. Posterior mean?	INVALID_INPUT	Gamma-Poisson conjugate posterior of a rate: Poisson counts must be whole numbers >= 0.
Gamma-Poisson posterior with shape = 2 and counts = [3, 5].	NEEDS_INPUT	gamma_poisson_posterior needs rate (prior rate beta (> 0)). Example: solve('gamma_poisson_posterior', shape=2, rate=1, c

Credible interval for a Poisson rate — `gamma_poisson_credible_interval`

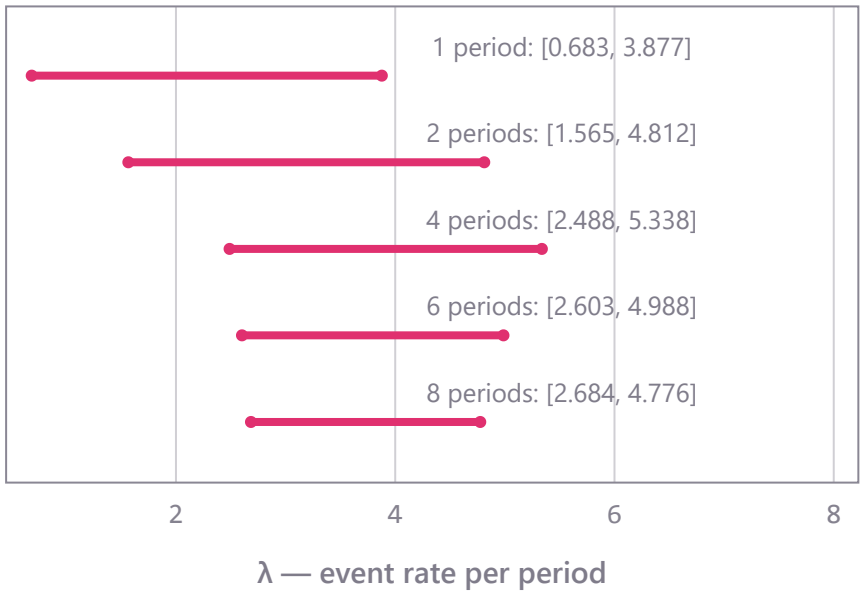
Equal-tailed credible interval of a Poisson rate after observed counts with a Gamma(shape, rate) prior.

$$\left[\frac{G^{-1}_{\alpha'}\left(\frac{1-c}{2}\right)}{\beta'},\ \frac{G^{-1}_{\alpha'}\left(\frac{1+c}{2}\right)}{\beta'}\right]$$

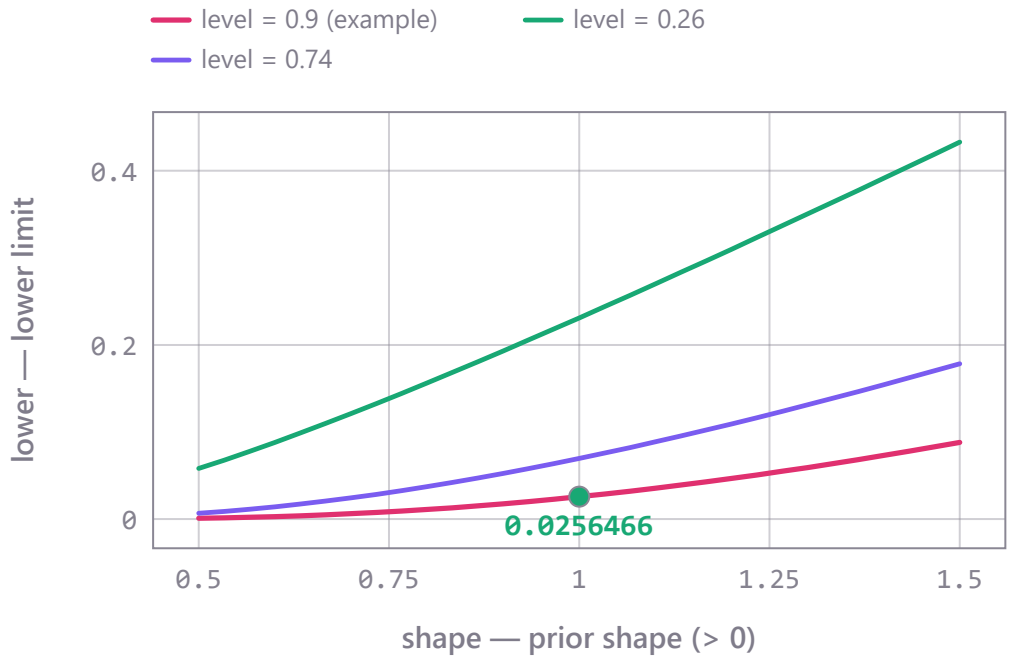
input	meaning	condition
`shape`	prior shape (> 0)	
`rate`	prior rate (> 0)	
`counts`	observed counts	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`post_shape`	posterior shape
`post_rate`	posterior rate
`lower`	lower limit
`upper`	upper limit

- **Method:** conjugate update, then gamma quantiles (inverse regularized incomplete gamma)
- **Assumes:** shape, rate > 0; counts whole numbers >= 0; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



90 % credible intervals for a Poisson rate under a Gamma(shape 1, rate 1) prior as the counts 3, 5, 4, 6, 2, 5, 4, 3 arrive (inputs chosen; the first 1, 2, 4, 6 and 8 counts), each returned by the package: 1: [0.683159, 3.87683]; 2: [1.56508, 4.81155]; 4: [2.48839, 5.33835]; 6: [2.60265, 4.98801]; 8: [2.68363, 4.77583].



lower (lower limit) against shape (prior shape (> 0)) from 0.5 to 1.5. One curve for each value of level (credibility level as a fraction (0.95 for 95%; default 0.95)): 0.9, 0.26, 0.74. Other inputs as in the worked example: rate = 1, counts = 0. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, shape = 1, lower = 0.0256466.

```
bayesian_statistics_ai999x.solve('gamma_poisson_credible_interval', shape=1, rate=1, counts='0', level=0.9)
# -> post_shape = 1, post_rate = 2, lower = 0.0256466, upper = 1.49787   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('gamma_poisson_credible_interval', shape=2, rate=1, counts=[3, 5, 4])
# -> post_shape = 14, post_rate = 4, lower = 1.91348, upper = 5.5576   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
90% credible interval for a Poisson rate: prior Gamma(1, 1), counts = [0].	SOLVED	post_shape = 1, post_rate = 2, lower = 0.0256466, upper = 1.49787
Find the 90% credible interval of the rate with shape = 0.5, rate = 0.001 and counts = [7].	SOLVED	post_shape = 7.5, post_rate = 1.001, lower = 3.62685, upper = 12.4854
95% credible interval for a Poisson rate: prior Gamma(2, 0), counts = [3].	INVALID_INPUT	Credible interval for a Poisson rate: the prior rate must be positive.

Equal-tailed credible interval of a gamma distribution — `gamma_credible_interval`

The central credible interval and median of a Gamma(shape, rate) posterior.

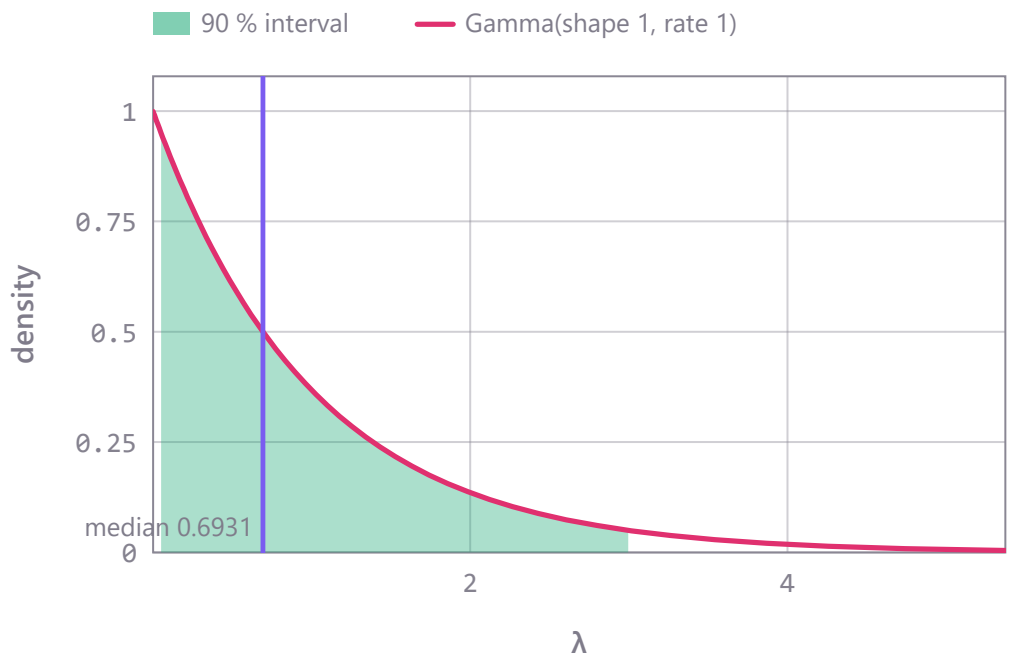
$$\left[G^{-1}\left(\tfrac{1-c}{2}\right),\ G^{-1}\left(\tfrac{1+c}{2}\right)\right]$$

input	meaning	condition
`shape`	shape (> 0)	
`rate`	rate (> 0)	

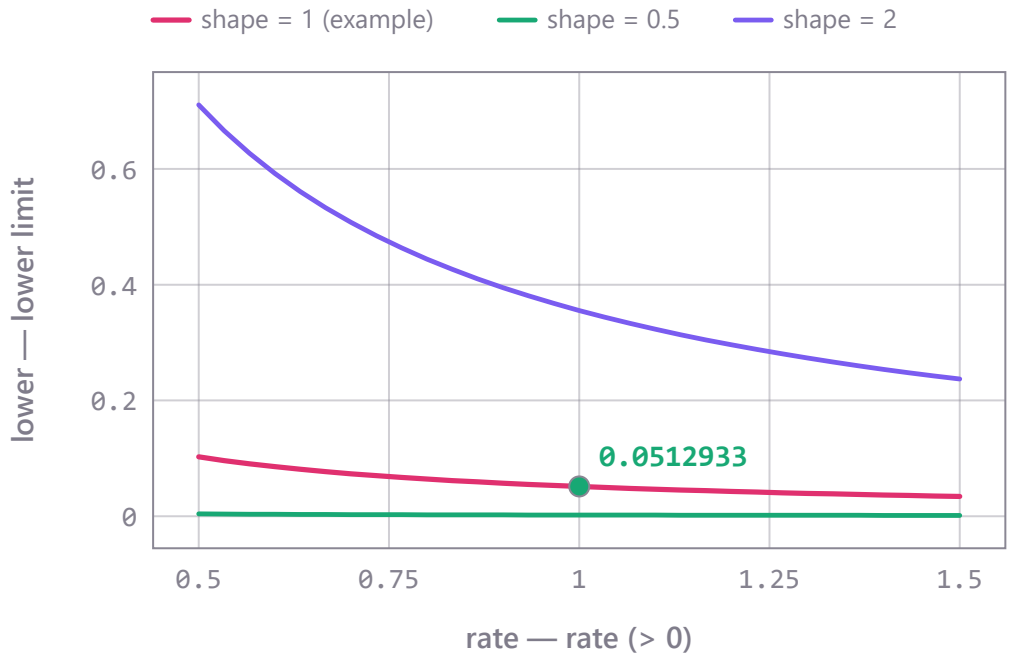
input	meaning	condition
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`lower`	lower limit
`upper`	upper limit
`median`	median

- **Method:** gamma quantiles by inversion of the regularized incomplete gamma function
- **Assumes:** shape, rate > 0; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Gamma density with shape 1 and rate 1 (the worked example); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Shaded: the returned 90 % equal-tailed credible interval [0.0512933, 2.99573]; marked: the returned median 0.693147.



lower (lower limit) against rate (rate (> 0)) from 0.5 to 1.5. One curve for each value of shape (shape (> 0)): 1, 0.5, 2. Other inputs as in the worked example: level = 0.9. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, rate = 1, lower = 0.0512933.

```
bayesian_statistics_ai999x.solve('gamma_credible_interval', shape=1, rate=1, level=0.9)
# -> lower = 0.0512933, upper = 2.99573, median = 0.693147   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('gamma_credible_interval', shape=5, rate=2)
# -> lower = 0.811743, upper = 5.12079, median = 2.33545   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
90% credible interval of a Gamma(1, 2) posterior.	SOLVED	lower = 0.0256466, upper = 1.49787, median = 0.346574
Find the 90% credible interval of a gamma distribution with shape = 3 and rate = 0.5.	SOLVED	lower = 1.63538, upper = 12.5916, median = 5.34812
95% credible interval of a Gamma(5, -2) posterior.	INVALID_INPUT	Equal-tailed credible interval of a gamma distribution: the rate must be positive.

Highest-density interval of a gamma posterior — `gamma_hdi`

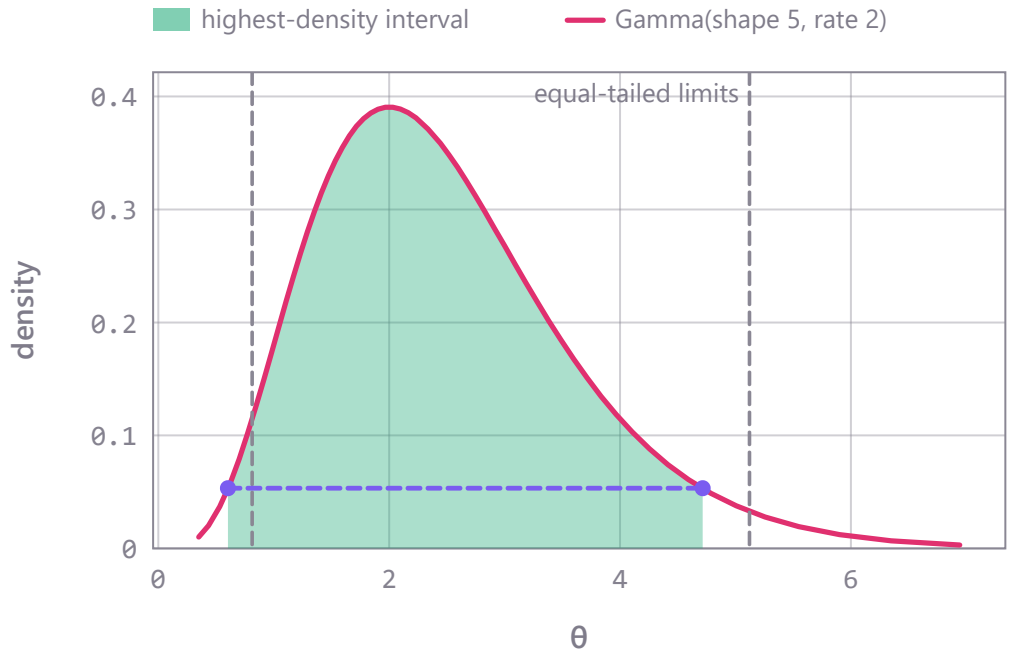
The shortest interval holding a stated probability of a unimodal Gamma(shape, rate): equal density at both ends.

$$G(u) - G(l) = c, \quad g(l) = g(u)$$

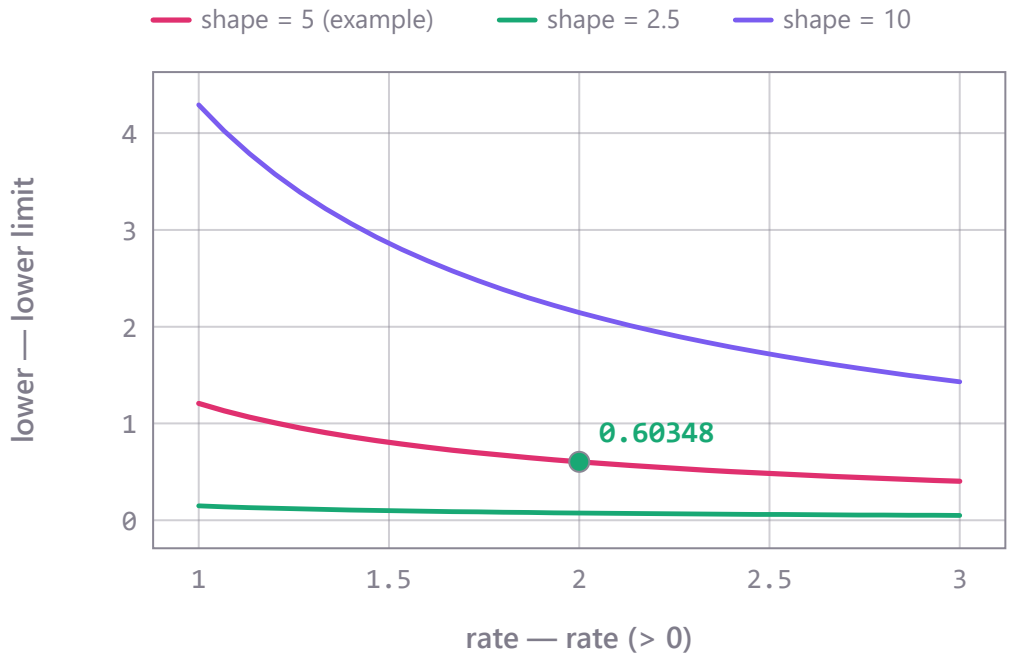
input	meaning	condition
`shape`	shape (> 1)	
`rate`	rate (> 0)	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`lower`	lower limit
`upper`	upper limit
`width`	interval width

- **Method:** bisection on the lower limit with $u = G^{-1}$(G(l) + level) until the densities are equal
- **Assumes:** shape > 1; rate > 0; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Gamma density with shape 5 and rate 2 at level 0.95 (the worked example); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Shaded: the returned highest-density interval [0.60348, 4.71511] (width 4.11163):the density is equal at both ends (dashed chord), so no shorter interval holds the same probability; dashed verticals: the equal-tailed interval [0.811743, 5.12079] at the same level, width 4.30905 (difference of the returned limits).



lower (lower limit) against rate (rate (> 0)) from 1 to 3. One curve for each value of shape (shape (> 1)): 5, 2.5, 10. Other inputs as in the worked example: level = 0.95. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, rate = 2, lower = 0.60348.

```
bayesian_statistics_ai999x.solve('gamma_hdi', shape=5, rate=2, level=0.95)
# -> lower = 0.60348, upper = 4.71511, width = 4.11163   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('gamma_hdi', shape=2, rate=1, level=0.5)
# -> lower = 0.435552, upper = 1.91795, width = 1.4824   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
95% highest-density interval of a Gamma(5, 2) posterior.	SOLVED	lower = 0.60348, upper = 4.71511, width = 4.11163
Find the 90% HDI of a gamma posterior with shape = 0.5 and rate = 0.5.	INVALID_INPUT	Highest-density interval of a gamma posterior: the highest-density interval is computed for a unimodal gamma with shape

Posterior predictive count (negative binomial) — `gamma_poisson_predictive`

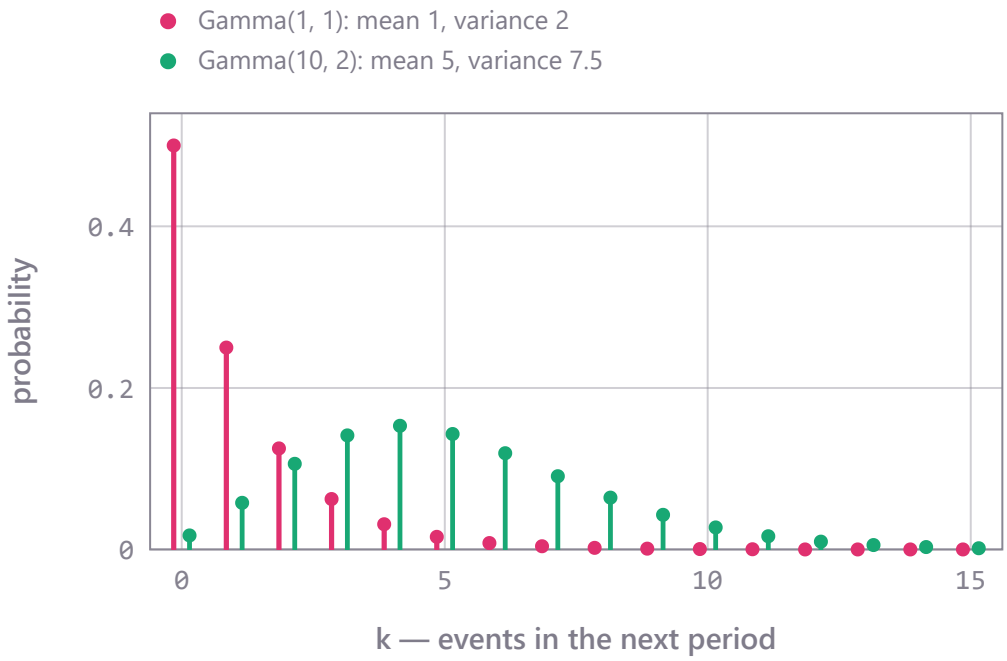
Probability of exactly k events in the next unit interval when the Poisson rate has a Gamma(shape, rate) posterior; the cumulative probability and the predictive mean and variance.

$$P(X = k) = \frac{\Gamma(\alpha + k)}{\Gamma(\alpha)\, k!} \left(\frac{\beta}{\beta+1}\right)^{\alpha} \left(\frac{1}{\beta+1}\right)^{k}$$

input	meaning	condition
`shape`	posterior shape (> 0)	
`rate`	posterior rate (> 0)	
`events`	number of events k in the next interval	

output	meaning
`probability`	P(exactly k events)
`prob_at_most`	P(at most k events)
`predictive_mean`	shape/rate
`predictive_variance`	shape/rate (1 + 1/rate)

- **Method:** negative binomial distribution (the Poisson averaged over the gamma posterior)
- **Assumes:** shape, rate > 0; k whole number >= 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The negative-binomial predictive distribution of next-period counts for a Gamma(shape 1, rate 1) rate (the worked example's) and a Gamma(shape 10, rate 2) rate (chosen), k = 0 … 15, each probability returned by the package: Gamma(1, 1): P(0) = 0.5, mean 1, variance 2; Gamma(10, 2): P(0) = 0.0173415, mean 5, variance 7.5. The predictive variance exceeds the mean.

```
bayesian_statistics_ai999x.solve('gamma_poisson_predictive', shape=1, rate=1, events=0)
# -> probability = 0.5, prob_at_most = 0.5, predictive_mean = 1, predictive_variance = 2   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('gamma_poisson_predictive', shape=12, rate=4, events=2)
# -> probability = 0.214405, prob_at_most = 0.448051, predictive_mean = 3, predictive_variance =
3.75   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Posterior Gamma(1, 1) for a Poisson rate: probability of exactly 1 events in the next interval (posterior predictive)?	SOLVED	probability = 0.25, prob_at_most = 0.75, predictive_mean = 1, predictive_variance = 2
Posterior predictive with shape = 3, rate = 1, events = 0.	SOLVED	probability = 0.125, prob_at_most = 0.125, predictive_mean = 3, predictive_variance = 6
Posterior predictive with shape = 3, rate = 1, events = 1.5.	INVALID_INPUT	events: expected a whole number written exactly, got 1.5.

MAP estimate of a Poisson rate (gamma prior) — `map_gamma_poisson`

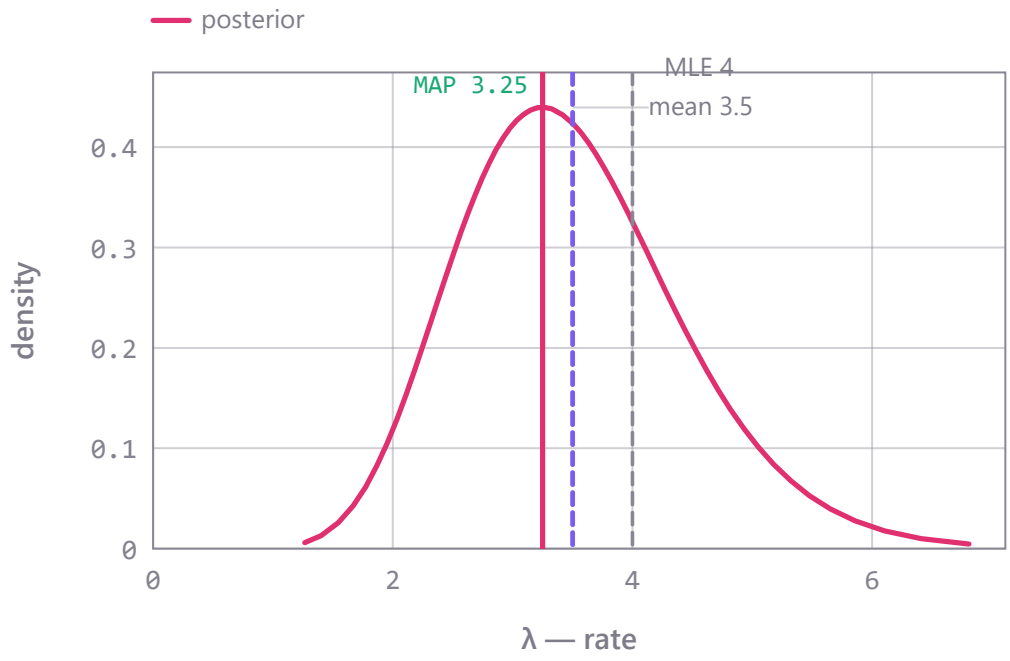
Posterior mode of a Poisson rate with a Gamma(shape, rate) prior and observed counts, with the MLE and the posterior mean for comparison.

$$\hat{\lambda}_{MAP} = \frac{\alpha + \sum x_i - 1}{\beta + n}$$

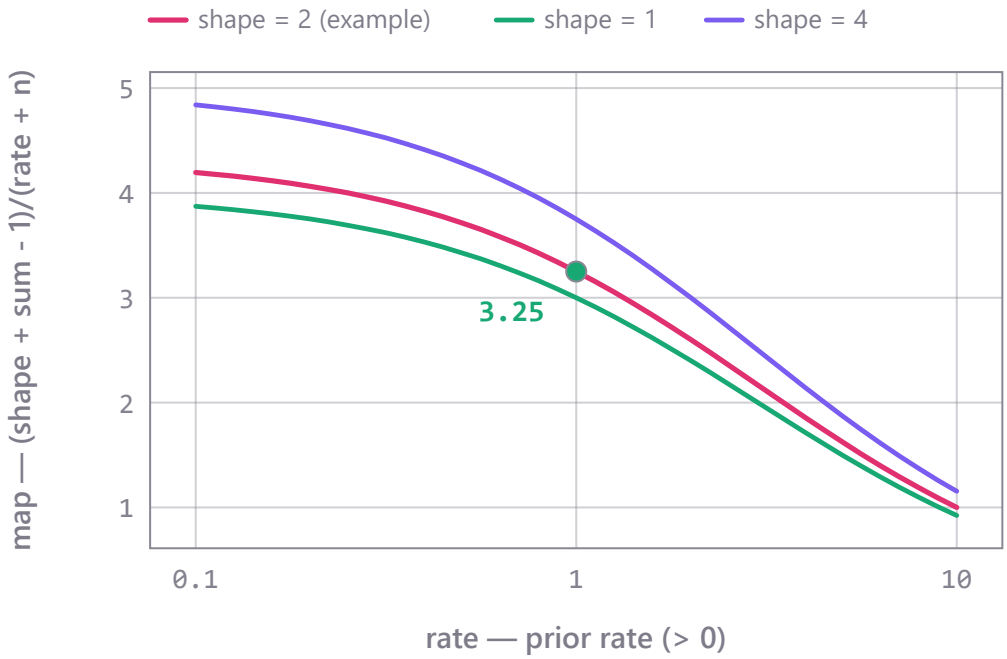
input	meaning	condition
`shape`	prior shape (> 0)	
`rate`	prior rate (> 0)	
`counts`	observed counts	

output	meaning
`map`	(shape + sum - 1)/(rate + n)
`mle`	sample mean of the counts
`posterior_mean`	(shape + sum)/(rate + n)

- **Method:** stationary point of the log posterior
- **Assumes:** shape + sum of counts > 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The posterior Gamma(shape 14, rate 4) for a Gamma(2, 1) prior and the counts 3, 5, 4 (the worked example); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Marked: the returned MAP 3.25, maximum-likelihood estimate 4 and posterior mean 3.5.



map ((shape + sum - 1)/(rate + n)) against rate (prior rate (> 0)) from 0.1 to 10 on a logarithmic axis. One curve for each value of shape (prior shape (> 0)): 2, 1, 4. Other inputs as in the worked example: counts = 3, 5, 4. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, rate = 1, map = 3.25.

```
bayesian_statistics_ai999x.solve('map_gamma_poisson', shape=2, rate=1, counts='3, 5, 4')
# -> map = 3.25, mle = 4, posterior_mean = 3.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('map_gamma_poisson', shape=1, rate=0.5, counts=[0, 2, 1, 1])
# -> map = 0.888889, mle = 1, posterior_mean = 1.11111   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
MAP estimate of the Poisson rate with prior Gamma(2, 1) and counts = [3, 5, 4].	SOLVED	map = 3.25, mle = 4, posterior_mean = 3.5
Posterior mode of the rate with shape = 1, rate = 0.5, counts = [0, 2, 1, 1].	SOLVED	map = 0.888889, mle = 1, posterior_mean = 1.11111
MAP estimate of the Poisson rate with prior Gamma(0.5, 1) and counts = [0, 0].	INVALID_INPUT	MAP estimate of a Poisson rate (gamma prior): the posterior mode is at 0 (shape + sum of counts <= 1): no interior MAP.

Gamma-exponential conjugate posterior of a rate — `gamma_exponential_posterior`

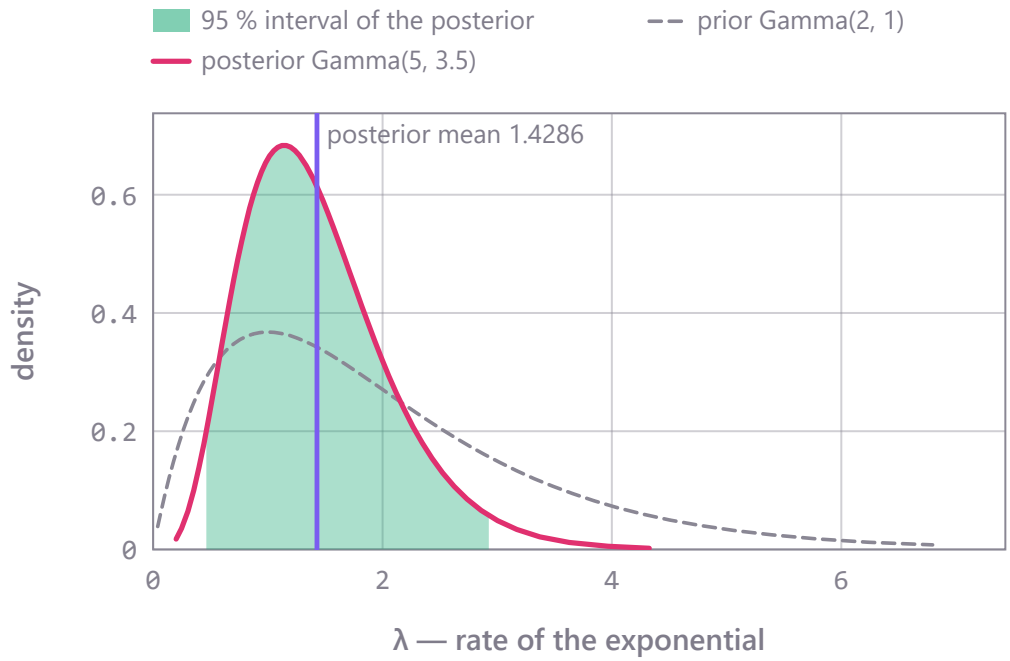
Update a Gamma(shape, rate) prior on the rate of exponential waiting times with observed times.

$$\lambda \mid t \sim \mathrm{Gamma}(\alpha + n,\ \beta + \textstyle\sum t_i)$$

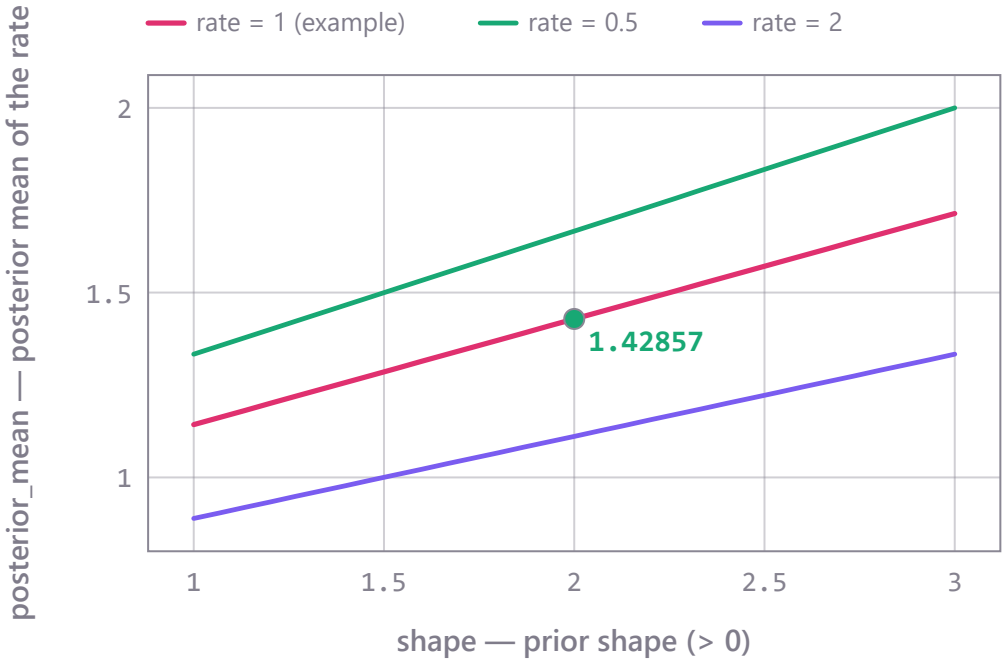
input	meaning	condition
`shape`	prior shape (> 0)	
`rate`	prior rate (> 0)	
`times`	observed waiting times (> 0)	

output	meaning
`post_shape`	shape + n
`post_rate`	rate + sum of times
`posterior_mean`	posterior mean of the rate
`posterior_variance`	posterior variance
`posterior_sd`	posterior standard deviation

- **Method:** conjugate update of the gamma prior with the exponential likelihood
- **Assumes:** shape, rate > 0; times > 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A Gamma prior (shape 2, rate 1) updated with the waiting times 0.5, 1.2, 0.8 (the worked example): the returned posterior has shape 5 and rate 3.5, mean 1.42857; each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Shaded: the posterior 95 % equal-tailed interval [0.463853, 2.92617].



posterior_mean (posterior mean of the rate) against shape (prior shape (> 0)) from 1 to 3. One curve for each value of rate (prior rate (> 0)): 1, 0.5, 2. Other inputs as in the worked example: times = 0.5, 1.2, 0.8. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, shape = 2, posterior_mean = 1.42857.

```
bayesian_statistics_ai999x.solve('gamma_exponential_posterior', shape=2, rate=1, times='0.5, 1.2, 0.8')
# -> post_shape = 5, post_rate = 3.5, posterior_mean = 1.42857, posterior_variance = 0.408163, posterior_sd = 0.638877   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('gamma_exponential_posterior', shape=1, rate=2, times=[3, 5])
# -> post_shape = 3, post_rate = 10, posterior_mean = 0.3, posterior_variance = 0.03, posterior_sd = 0.173205   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Gamma-exponential posterior: prior Gamma(2, 1), exponential waiting times = [0.5, 1.2, 0.8]. Posterior mean of the rate?	SOLVED	post_shape = 5, post_rate = 3.5, posterior_mean = 1.42857, posterior_variance = 0.408163, posterior_sd = 0.638877
Exponential lifetimes times = [3, 5] with prior shape = 1 and rate = 2: posterior of the rate?	SOLVED	post_shape = 3, post_rate = 10, posterior_mean = 0.3, posterior_variance = 0.03, posterior_sd = 0.173205
Exponential lifetimes times = [3, -5] with prior shape = 1 and rate = 2: posterior of the rate?	INVALID_INPUT	Gamma-exponential conjugate posterior of a rate: exponential observations must be positive.

Normal mean with known variance: posterior from data — `normal_normal_posterior`

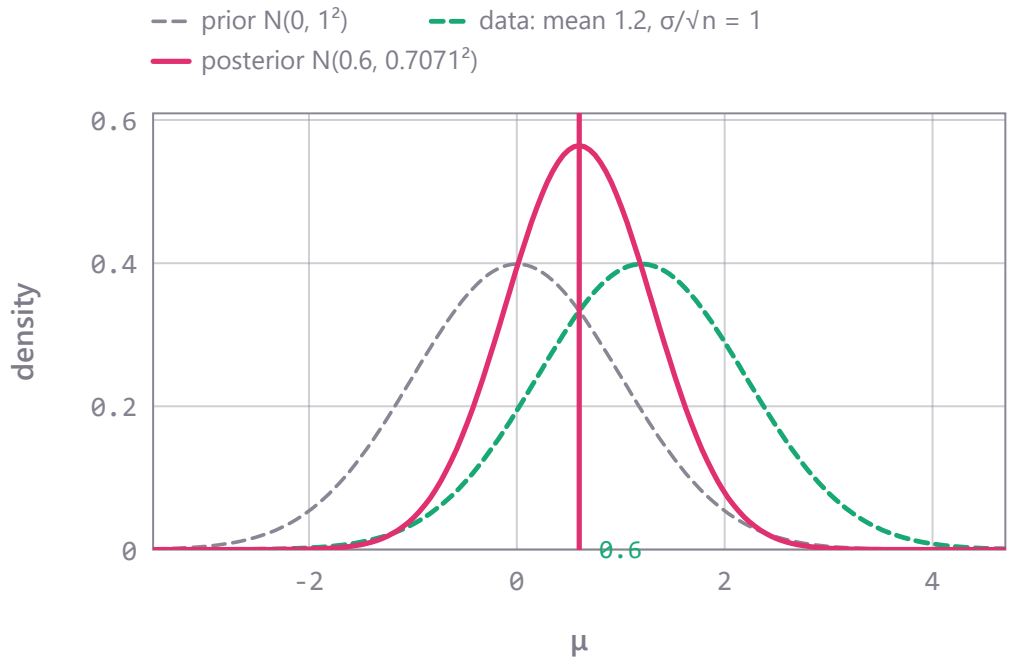
Posterior of a normal mean theta with a N(prior mean, prior sd^2) prior, given observations with known standard deviation sigma.

$$\frac{1}{\tau_n^2} = \frac{1}{\tau_0^2} + \frac{n}{\sigma^2}, \quad \mu_n = \tau_n^2 \left(\frac{\mu_0}{\tau_0^2} + \frac{n\bar{x}}{\sigma^2} \right)$$

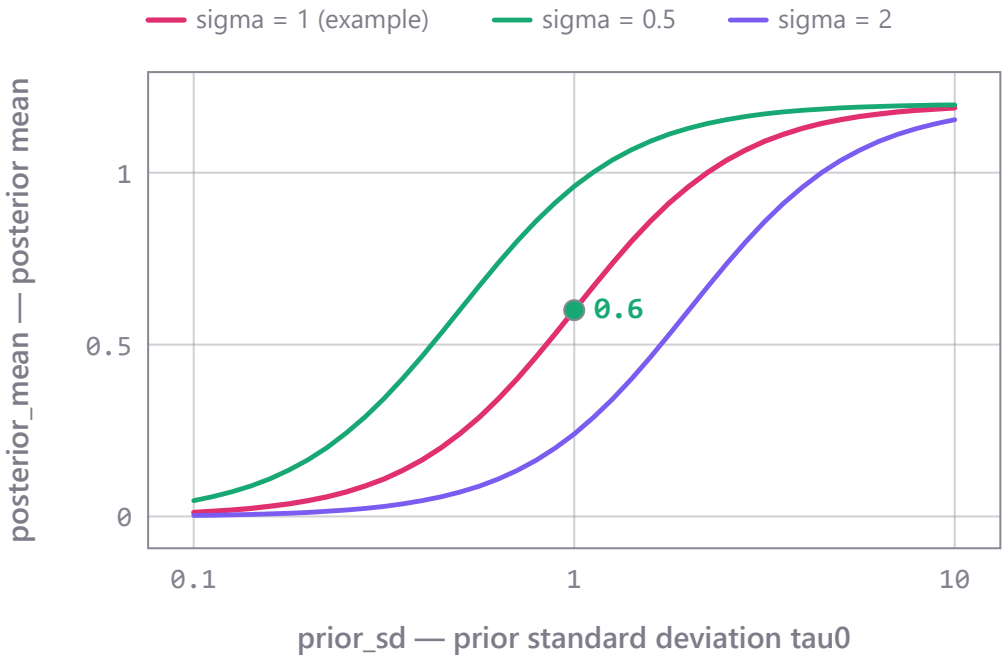
input	meaning	condition
`prior_mean`	prior mean mu0	
`prior_sd`	prior standard deviation tau0 (> 0)	
`sigma`	known standard deviation of each observation (> 0)	
`data`	the observations	

output	meaning
`posterior_mean`	posterior mean
`posterior_sd`	posterior standard deviation
`posterior_variance`	posterior variance

- **Method:** precision-weighted average of the prior mean and the sample mean
- **Assumes:** prior $sd > 0$; $\sigma > 0$ known
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A normal prior for a mean (mean 0, sd 1) combined with the data 1.2 of known σ = 1: the returned posterior mean 0.6 and sd 0.707107 (variance 0.5) lie between prior and data, the posterior narrower than both. The three curves are normal densities drawn from these means and standard deviations (the data curve from the data mean and σ/√n).



posterior_mean (posterior mean) against prior_sd (prior standard deviation tau0 (> 0)) from 0.1 to 10 on a logarithmic axis. One curve for each value of sigma (known standard deviation of each observation (> 0)): 1, 0.5, 2. Other inputs as in the worked example: prior_mean = 0, data = 1.2. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, prior_sd = 1, posterior_mean = 0.6.

```
bayesian_statistics_ai999x.solve('normal_normal_posterior', prior_mean=0, prior_sd=1, sigma=1,
data='1.2')
# -> posterior_mean = 0.6, posterior_sd = 0.707107, posterior_variance = 0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('normal_normal_posterior', prior_mean=10, prior_sd=2, sigma=3,
data=[12, 9, 14])
# -> posterior_mean = 10.9524, posterior_sd = 1.30931, posterior_variance = 1.71429   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Normal-normal posterior: prior mean = 10, prior sd = 2, known sigma = 3, data = [12, 9, 14]. Posterior mean?	SOLVED	posterior_mean = 10.9524, posterior_sd = 1.30931, posterior_variance = 1.71429
Known-variance normal model with prior mean 0, prior standard deviation 1, sigma = 1 and data = [1.2]: find the posterior.	SOLVED	posterior_mean = 0.6, posterior_sd = 0.707107, posterior_variance = 0.5
Normal-normal posterior: prior mean = 10, prior sd = -2, known sigma = 3, data = [12]. Posterior?	INVALID_INPUT	Normal mean with known variance: posterior from data: the prior standard deviation must be positive.
Normal-normal posterior: prior mean = 10, known sigma = 3, data = [12]. Posterior given the prior sd?	NEEDS_INPUT	normal_normal_posterior needs prior_sd (prior standard deviation tau0 (> 0)). Example: solve('normal_normal_posterior',

Normal mean with known variance: posterior and credible interval from a sample mean — `normal_normal_summary`

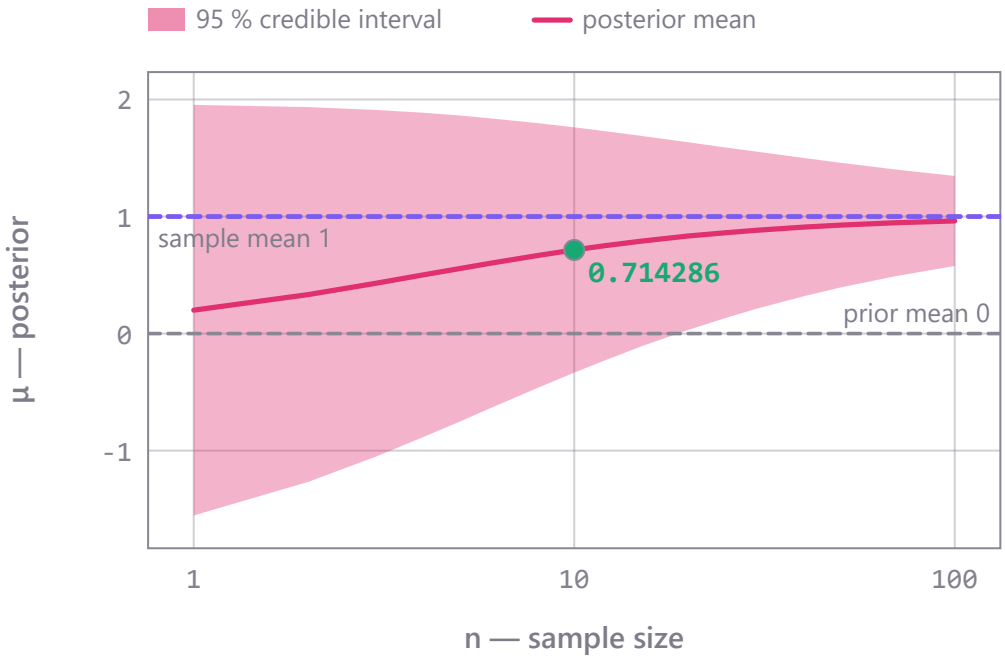
Posterior mean, standard deviation, credible interval and shrinkage factor of a normal mean from the sample mean of n observations with known sigma and a normal prior.

$$\mu_n = B\mu_0 + (1-B)\bar{x},\ B = \frac{\sigma^2/n}{\sigma^2/n + \tau_0^2}, \quad \mu_n \pm z_{(1+c)/2}\,\tau_n$$

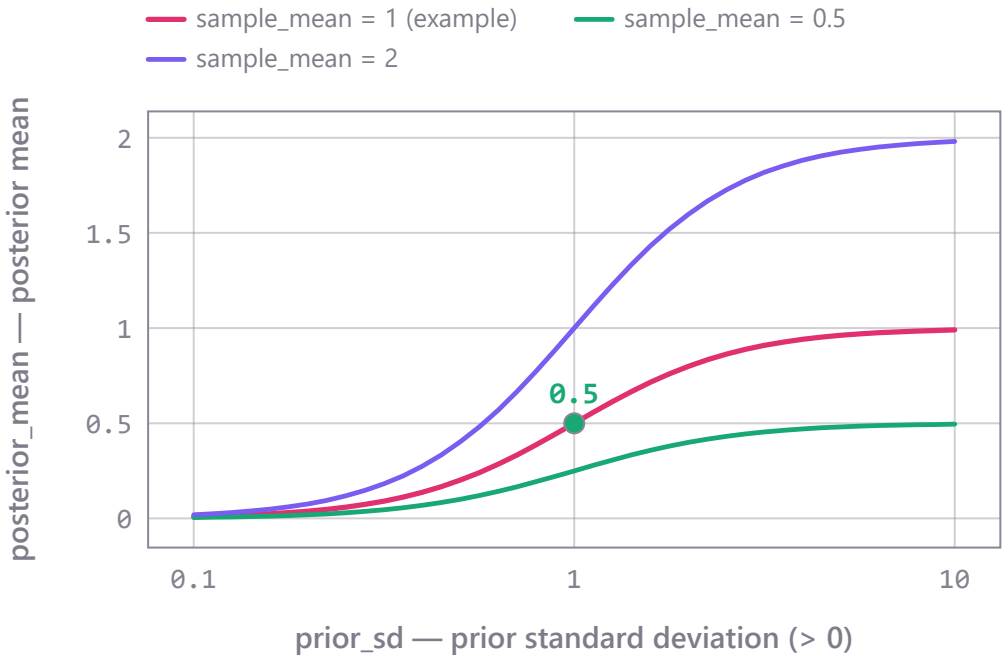
input	meaning	condition
`prior_mean`	prior mean	
`prior_sd`	prior standard deviation (> 0)	
`sigma`	known standard deviation of one observation (> 0)	
`sample_mean`	sample mean	
`n`	sample size (>= 1)	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`posterior_mean`	posterior mean
`posterior_sd`	posterior standard deviation
`lower`	lower credible limit
`upper`	upper credible limit
`shrinkage`	weight on the prior mean $(\sigma^2/n)/(\sigma^2/n + tau0^2)$

- **Method:** normal-normal conjugate update, interval from normal quantiles
- **Assumes:** prior $sd > 0$; $\sigma > 0$; $n \geq 1$; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A normal prior (mean 0, sd 1) and data with σ = 2 and sample mean 1: the package's posterior mean and 95 % credible interval for 17 sample sizes from 1 to 100, moving from the prior towards the data; marked: n = 10, posterior mean = 0.714286, interval [-0.333359, 1.76193].



posterior_mean (posterior mean) against prior_sd (prior standard deviation (> 0)) from 0.1 to 10 on a logarithmic axis. One curve for each value of sample_mean (sample mean): 1, 0.5, 2. Other inputs as in the worked example: prior_mean = 0, sigma = 2, n = 4. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, prior_sd = 1, posterior_mean = 0.5.

```
bayesian_statistics_ai999x.solve('normal_normal_summary', prior_mean=0, prior_sd=1, sigma=2,
sample_mean=1, n=4)
# -> posterior_mean = 0.5, posterior_sd = 0.707107, lower = -0.885904, upper = 1.8859, shrinkage =
0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('normal_normal_summary', prior_mean=100, prior_sd=15, sigma=20,
sample_mean=110, n=25, level=0.95)
# -> posterior_mean = 109.336, posterior_sd = 3.86494, lower = 101.761, upper = 116.911, shrinkage
= 0.06639   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Normal prior with prior mean = 0, prior sd = 1; sample mean = 1 from n = 4 with known sigma = 2. Posterior mean and 95% credible interval?	SOLVED	posterior_mean = 0.5, posterior_sd = 0.707107, lower = -0.885904, upper = 1.8859, shrinkage = 0.5
Prior mean 0, prior standard deviation 2, known sigma = 4, sample mean 1.5, sample size 16: posterior?	SOLVED	posterior_mean = 1.2, posterior_sd = 0.894427, lower = -0.553045, upper = 2.95305, shrinkage = 0.2
Prior mean 0, prior standard deviation 2, known sigma = 4, sample mean 1.5, sample size 0: posterior?	INVALID_INPUT	n: must be at least 1, got 0.

Posterior predictive distribution of a new observation — `normal_posterior_predictive`

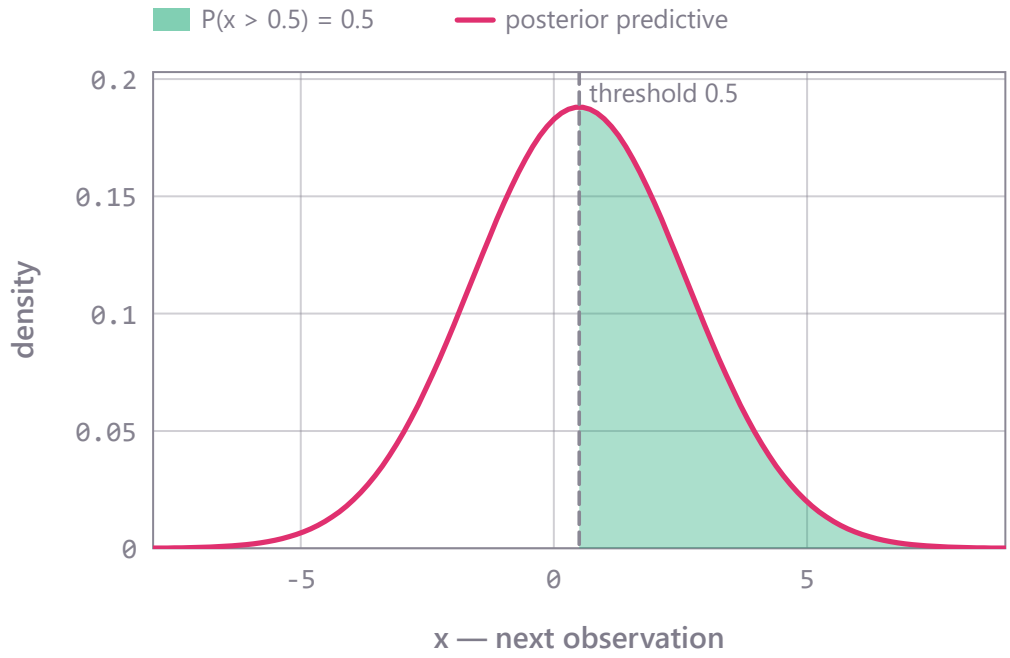
Mean, standard deviation and tail probabilities of the next observation in the normal model with known sigma, after a sample mean of n observations and a normal prior.

$$Y_{new} \mid \bar{x} \sim N\left(\mu_n,\ \tau_n^2 + \sigma^2\right)$$

input	meaning	condition
`prior_mean`	prior mean	
`prior_sd`	prior standard deviation (> 0)	
`sigma`	known standard deviation (> 0)	
`sample_mean`	sample mean	
`n`	sample size (>= 1)	
`threshold`	value c for P(Y > c)	

output	meaning
`predictive_mean`	posterior mean
`predictive_sd`	sqrt(posterior variance + sigma^2)
`prob_above`	P(Y_new > c)
`prob_below`	P(Y_new < c)

- **Method:** predictive variance = posterior variance + sampling variance
- **Assumes:** prior $sd > 0$; $\sigma > 0$; $n \geq 1$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The posterior predictive distribution of the next observation for the worked example (prior mean 0, prior sd 1, σ = 2, sample mean 1 from n = 4): drawn from the package's P(x < t) at 121 thresholds (differenced); returned predictive mean 0.5, sd 2.12132. Shaded: P(x > 0.5) = 0.5 (P(x < 0.5) = 0.5).



predictive_sd (sqrt(posterior variance + sigma^2)) against sigma (known standard deviation (> 0)) from 1 to 3. One curve for each value of prior_sd (prior standard deviation (> 0)): 1, 0.5, 2. Other inputs as in the worked example: prior_mean = 0, sample_mean = 1, n = 4, threshold = 0.5. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, sigma = 2, predictive_sd = 2.12132.

```
bayesian_statistics_ai999x.solve('normal_posterior_predictive', prior_mean=0, prior_sd=1, sigma=2,
sample_mean=1, n=4, threshold=0.5)
# -> predictive_mean = 0.5, predictive_sd = 2.12132, prob_above = 0.5, prob_below = 0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('normal_posterior_predictive', prior_mean=10, prior_sd=3, sigma=2,
sample_mean=12, n=5, threshold=15)
# -> predictive_mean = 11.8367, predictive_sd = 2.17594, prob_above = 0.0730078, prob_below =
0.926992   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Posterior predictive: prior mean = 0, prior sd = 1, known sigma = 2, sample mean = 1, n = 4. Probability that the next observation exceeds 0.5?	SOLVED	predictive_mean = 0.5, predictive_sd = 2.12132, prob_above = 0.5, prob_below = 0.5
Predictive distribution of a new observation with prior mean 10, prior sd 3, sigma = 0, sample mean 12, sample size 5 and threshold = 15.	INVALID_INPUT	Posterior predictive distribution of a new observation: sigma must be positive.

Normal model with unknown mean and variance (normal-inverse-gamma) — `normal_inverse_gamma_posterior`

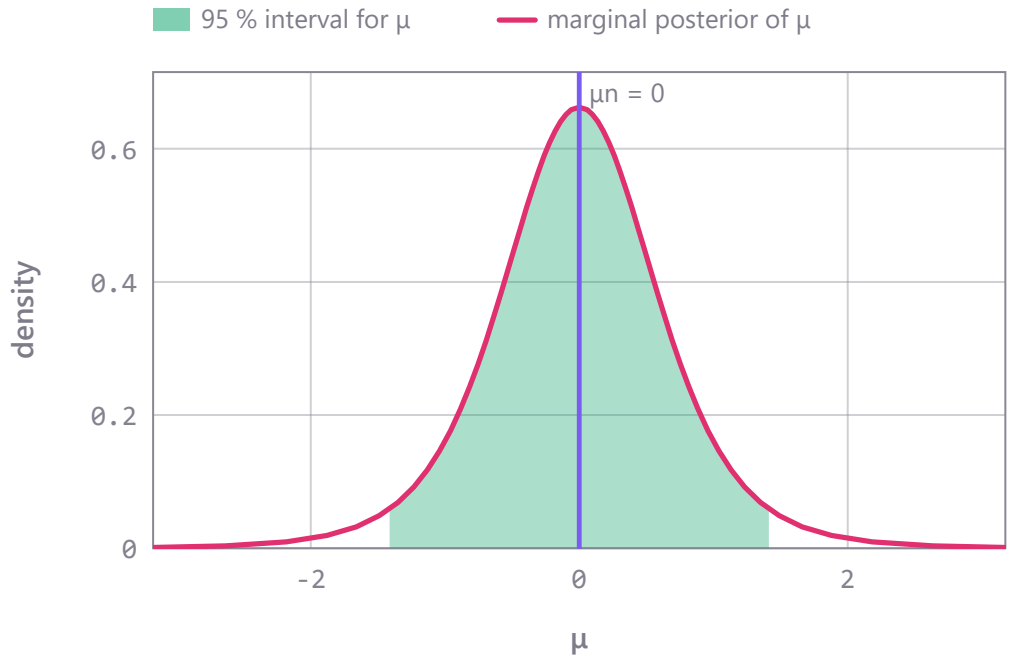
Conjugate normal-inverse-gamma update for normal data with unknown mean and variance: posterior parameters, posterior mean of the variance and the Student-t credible interval for the mean.

$$\beta_n = \beta_0 + \tfrac{1}{2}\sum (x_i - \bar{x})^2 + \frac{\kappa_0 n (\bar{x} - \mu_0)^2}{2(\kappa_0 + n)}, \quad \mu \mid x \sim t_{2\alpha_n}\left(\mu_n,\ \frac{\beta_n}{\alpha_n \kappa_n}\right)$$

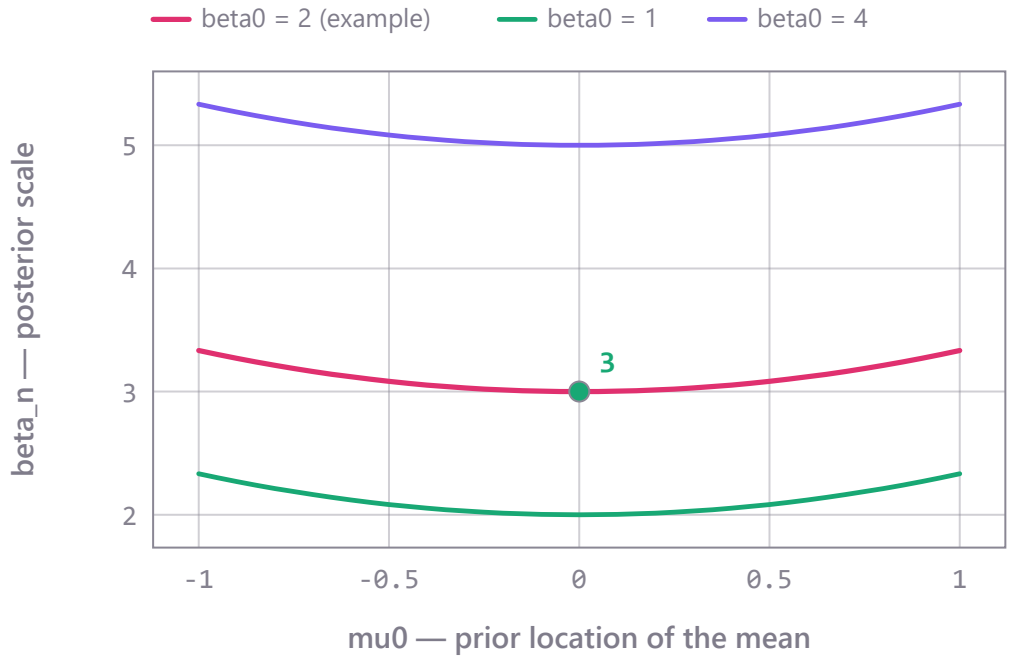
input	meaning	condition
`mu0`	prior location of the mean	
`kappa0`	prior pseudo-observations for the mean (> 0)	
`alpha0`	prior inverse-gamma shape (> 0)	
`beta0`	prior inverse-gamma scale (> 0)	
`data`	the observations	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`mu_n`	posterior location
`kappa_n`	$kappa0 + n$
`alpha_n`	$alpha0 + n/2$
`beta_n`	posterior scale
`sigma2_mean`	posterior mean of σ^2 ($alpha_n > 1$)
`lower`	lower credible limit for the mean
`upper`	upper credible limit for the mean

- **Method:** conjugate update; marginal posterior of the mean is Student t with 2 alpha_n degrees of freedom
- **Assumes:** kappa0, alpha0, $beta0 > 0$; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The marginal posterior of the mean under a normal–inverse-gamma prior (μ0 = 0, κ0 = 1, α0 = 2, β0 = 2) and the data 1, -1: returned μn = 0, κn = 3, αn = 3, βn = 3, posterior mean of σ^2 1.5; each density is drawn from the package (the interval limits at many levels)'s quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Shaded: the 95 % interval [-1.41273, 1.41273].



beta_n (posterior scale) against mu0 (prior location of the mean) from -1 to 1. One curve for each value of beta0 (prior inverse-gamma scale (> 0)): 2, 1, 4. Other inputs as in the worked example: kappa0 = 1, alpha0 = 2, data = 1, -1. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, mu0 = 0, beta_n = 3.

```
bayesian_statistics_ai999x.solve('normal_inverse_gamma_posterior', mu0=0, kappa0=1, alpha0=2, beta0=2, data='1, -1')
# -> mu_n = 0, kappa_n = 3, alpha_n = 3, beta_n = 3, sigma2_mean = 1.5, lower = -1.41273, upper = 1.41273   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('normal_inverse_gamma_posterior', mu0=10, kappa0=0.5, alpha0=3, beta0=4, data=[9, 11, 12])
# -> mu_n = 10.5714, kappa_n = 3.5, alpha_n = 4.5, beta_n = 6.42857, sigma2_mean = 1.83673, lower = 9.12619, upper = 12.0167   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Normal-inverse-gamma posterior: mu0 = 0, kappa0 = 1, alpha0 = 2, beta0 = 2, data = [1, -1]. 95% credible interval for the mean?	SOLVED	mu_n = 0, kappa_n = 3, alpha_n = 3, beta_n = 3, sigma2_mean = 1.5, lower = -1.41273, upper = 1.41273
Unknown mean and variance with mu0 = 10, kappa0 = 0, alpha0 = 3, beta0 = 4 and data = [9, 11, 12].	INVALID_INPUT	Normal model with unknown mean and variance (normal-inverse-gamma): kappa0 must be positive.

Dirichlet-multinomial conjugate posterior — `dirichlet_multinomial_posterior`

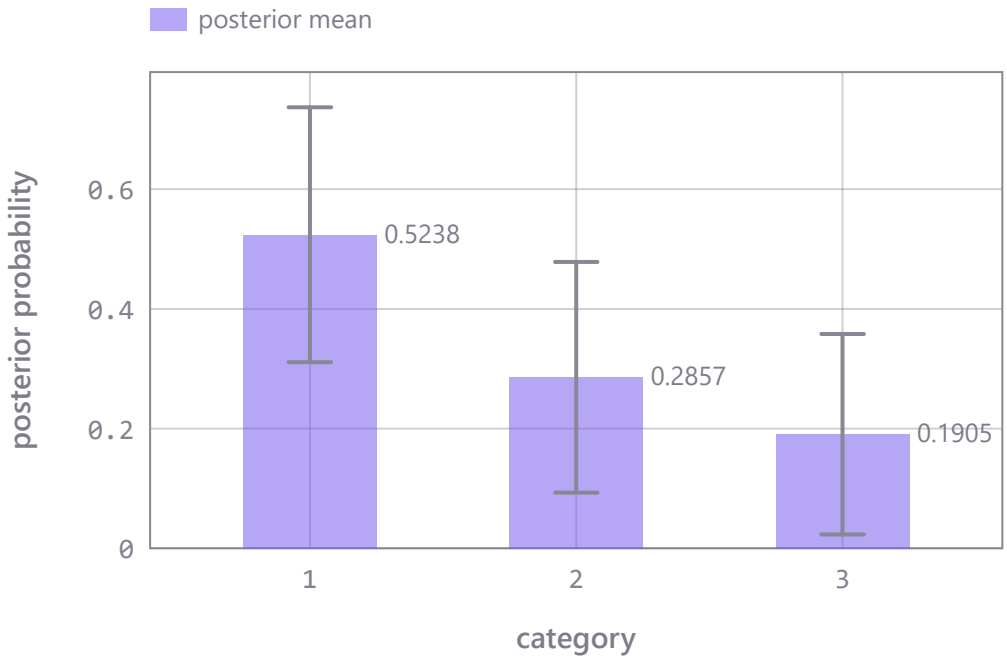
Update a Dirichlet(alpha_1, ..., alpha_k) prior on category probabilities with category counts: posterior concentrations, means and variances.

$$\theta \mid n \sim \mathrm{Dir}(\alpha_1 + n_1, \ldots, \alpha_k + n_k), \quad E[\theta_j] = \frac{\alpha_j + n_j}{\alpha_0 + n}$$

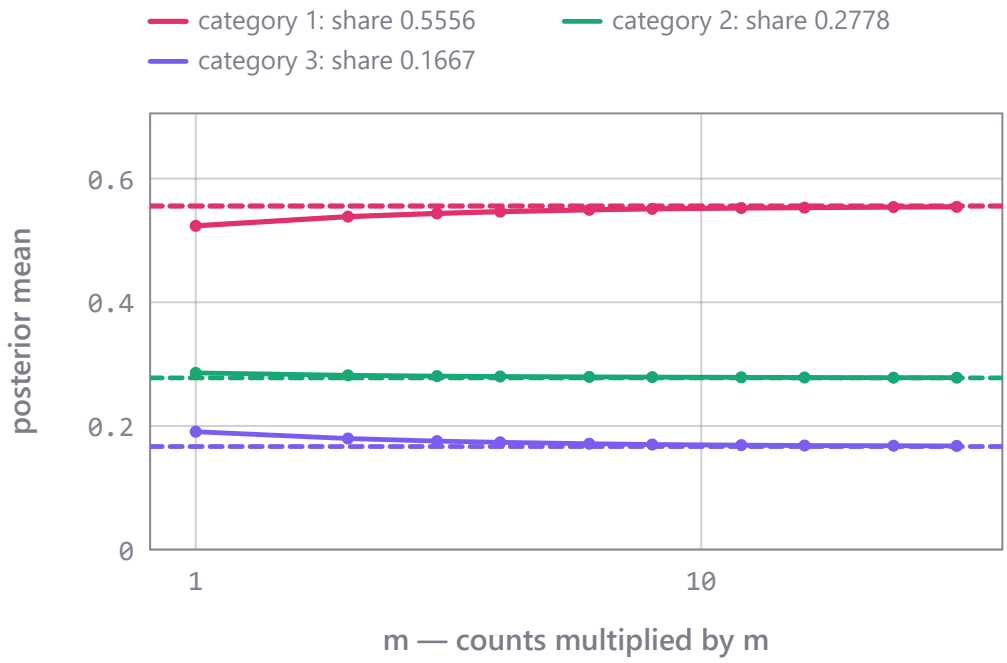
input	meaning	condition
`alpha`	prior concentrations alpha_j (> 0)	
`counts`	observed category counts	

output	meaning
`post_alpha`	alpha_j + n_j
`posterior_means`	posterior means $(alpha_j + n_j)/(alpha_0 + n)$
`posterior_variances`	posterior variances

- **Method:** conjugate update of the Dirichlet prior with the multinomial likelihood
- **Assumes:** $alpha_j > 0$; counts whole numbers >= 0; same length
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A Dirichlet(1, 1, 1) prior over 3 categories updated with the counts 10, 5, 3: the package's posterior Dirichlet(11, 6, 4), posterior means 0.52381, 0.285714, 0.190476; whiskers ±2 posterior standard deviations (0.106479, 0.0963143, 0.083719).



Posterior means of three category probabilities under the Dirichlet(1, 1, 1) prior of the worked example when its counts 10, 5, 3 are multiplied by m = 1 … 32 (inputs chosen), each from the package: with more data the means leave the prior's equal shares and approach the observed shares (dashed lines, the counts divided by their total).

```
bayesian_statistics_ai999x.solve('dirichlet_multinomial_posterior', alpha='1, 1, 1', counts='10, 5, 3')
# -> post_alpha = [11.0, 6.0, 4.0], posterior_means = [0.52380952381, 0.285714285714, 0.190476190476], posterior_variances = [0.0113378684807, 0.00927643784787, 0.00700886415172] (SOLVED)
```

```
bayesian_statistics_ai999x.solve('dirichlet_multinomial_posterior', alpha=[2, 2, 2, 2], counts=[3, 0, 7, 1])
# -> post_alpha = [5.0, 2.0, 9.0, 3.0], posterior_means = [0.263157894737, 0.105263157895, 0.473684210526, 0.157894736842], posterior_variances = [0.00969529085873, 0.00470914127424, 0.0124653739612, 0.00664819944598]   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Dirichlet-multinomial posterior with alpha = [1, 1, 1] and counts = [10, 5, 3]: posterior means?	SOLVED	post_alpha = [11.0, 6.0, 4.0], posterior_means = [0.52380952381, 0.285714285714, 0.190476190476], posterior_variances = [0.0113378684807, 0.00927643784787, 0.00700886415172]
Dirichlet prior concentrations [2, 2, 2, 2] updated with counts [3, 0, 7, 1].	SOLVED	post_alpha = [5.0, 2.0, 9.0, 3.0], posterior_means = [0.263157894737, 0.105263157895, 0.473684210526, 0.157894736842], posterior_variances = [0.00969529085873, 0.00470914127424, 0.0124653739612, 0.00664819944598]
Dirichlet-multinomial posterior with alpha = [1, 1, 1] and counts = [10, 5].	INVALID_INPUT	Dirichlet-multinomial conjugate posterior: the counts must have 3 values (one per entry of the other list).

Dirichlet-multinomial probability of a count vector (evidence) —

`dirichlet_multinomial_evidence`

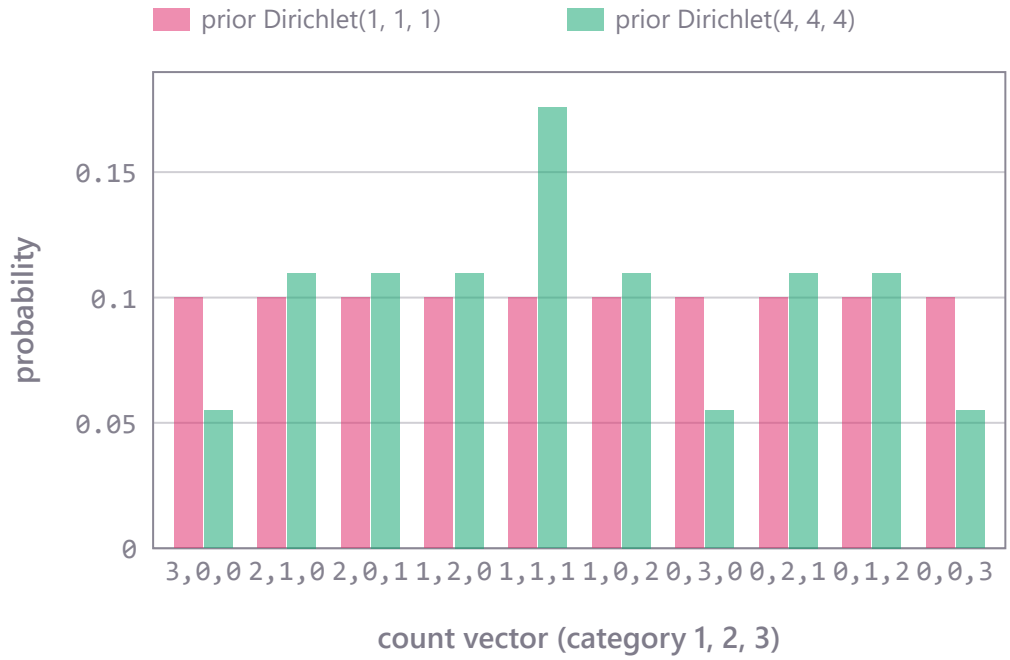
Prior predictive probability (marginal likelihood) of observed category counts when the category probabilities have a Dirichlet prior.

$$P(n \mid \alpha) = \frac{n!\,\Gamma(\alpha_0)}{\Gamma(n+\alpha_0)} \prod_j \frac{\Gamma(n_j+\alpha_j)}{n_j!\,\Gamma(\alpha_j)}$$

input	meaning	condition
`alpha`	prior concentrations (> 0)	
`counts`	category counts	

output	meaning
`probability`	P(counts | alpha)
`log_probability`	its natural log

- **Method:** closed form with log-gamma functions
- **Assumes:** $alpha_j > 0$; counts whole numbers >= 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Dirichlet-multinomial probability of each of the 10 ways to split 3 observations over 3 categories, under the worked example's Dirichlet(1, 1, 1) prior (every split equally likely: 0.1) and a Dirichlet(4, 4, 4) prior (chosen; balanced splits more likely, from 0.0549451 to 0.175824); each probability returned by the package.

```
bayesian_statistics_ai999x.solve('dirichlet_multinomial_evidence', alpha='1, 1, 1', counts='2, 1, 0')
# -> probability = 0.1, log_probability = -2.30259   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('dirichlet_multinomial_evidence', alpha=[1, 1], counts=[3, 3])
# -> probability = 0.142857, log_probability = -1.94591   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Dirichlet-multinomial evidence for counts = [2, 1, 0] with alpha = [1, 1, 1].	SOLVED	probability = 0.1, log_probability = -2.30259
Marginal likelihood of counts [3, 3] under a Dirichlet prior with alpha = [1, 1].	SOLVED	probability = 0.142857, log_probability = -1.94591
Dirichlet-multinomial evidence for counts = [2, 1, 0] with alpha = [1, -1, 1].	INVALID_INPUT	Dirichlet-multinomial probability of a count vector (evidence): Dirichlet concentrations must be positive.

Bayes factor for two simple binomial hypotheses — `bayes_factor_binomial_simple`

Bayes factor BF10 for H1: $p = p_1$ against H0: $p = p_0$ after k successes in n trials, its log10 and the posterior probability of H1 with equal prior odds.

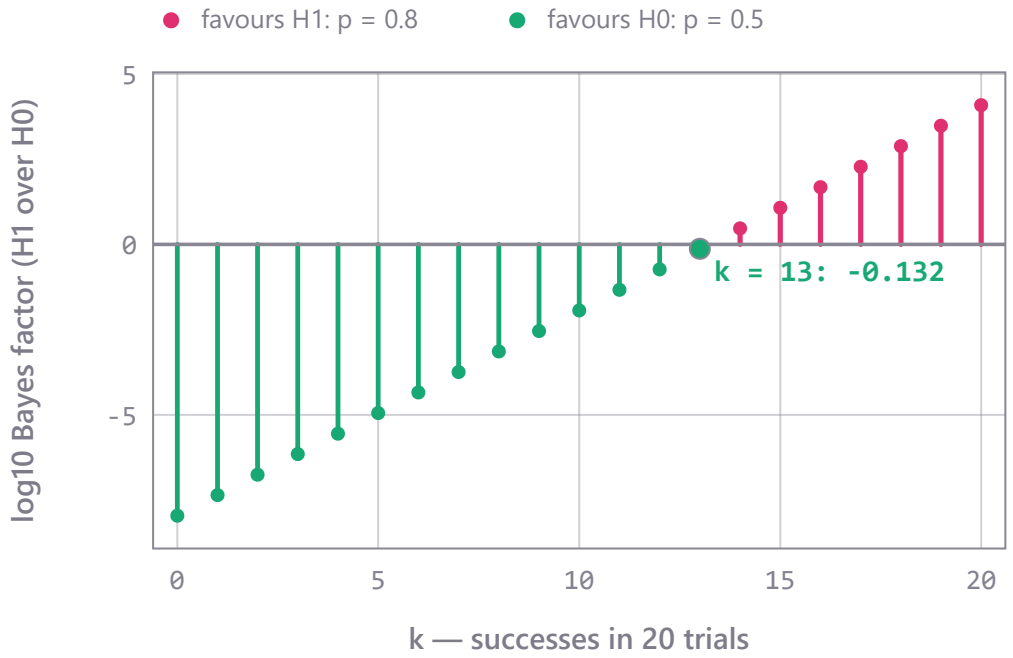
$$BF_{10} = \left(\frac{p_1}{p_0}\right)^{k} \left(\frac{1-p_1}{1-p_0}\right)^{n-k}$$

input	meaning	condition
`successes`	successes k	

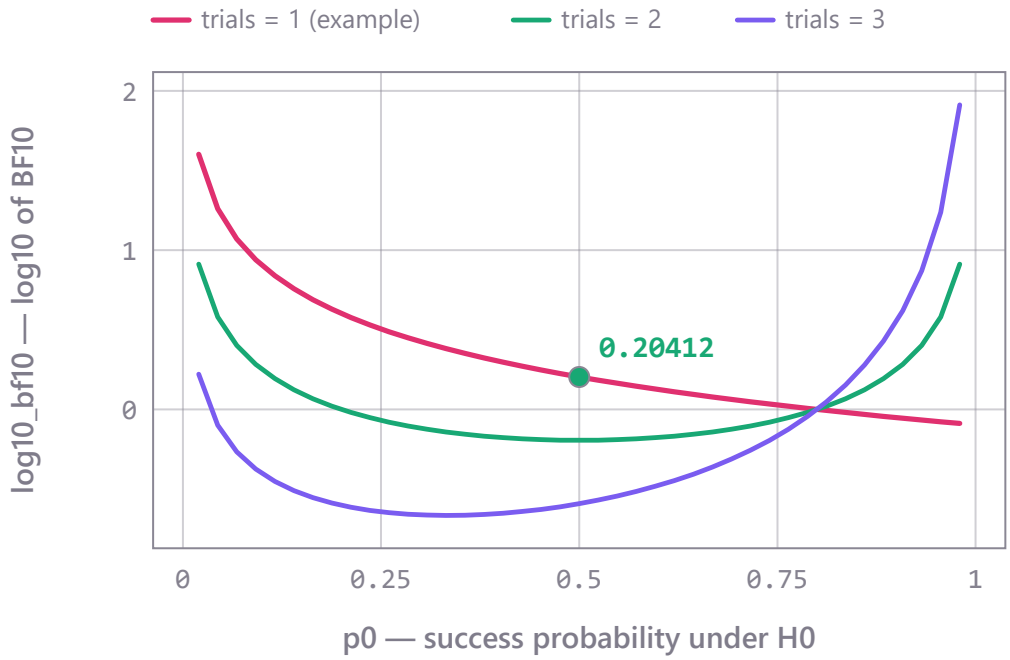
input	meaning	condition
`trials`	trials n (>= 1)	
`p0`	success probability under H0	
`p1`	success probability under H1	

output	meaning
`bf10`	Bayes factor H1 : H0
`log10_bf10`	log10 of BF10
`posterior_h1`	P(H1 | data) with prior odds 1

- **Method:** likelihood ratio of the two simple hypotheses
- **Assumes:** $0 < p_0, p_1 < 1; 0 \leq k \leq n$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Bayes factor of the simple hypotheses p = 0.8 against p = 0.5 for every number of successes in 20 trials (inputs chosen; the worked example has 1 trial), as log10 of the returned factor: each extra success multiplies it by the same amount, so the stems lie on a straight line. Marked: the count nearest to even evidence, k = 13.



log10_bf10 (log10 of BF10) against p0 (success probability under H0) from 0.02 to 0.98. One curve for each value of trials (trials n (>= 1)): 1, 2, 3. Other inputs as in the worked example: successes = 1, p1 = 0.8. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, p0 = 0.5, log10_bf10 = 0.20412.

```
bayesian_statistics_ai999x.solve('bayes_factor_binomial_simple', successes=1, trials=1, p0=0.5,
p1=0.8)
# -> bf10 = 1.6, log10_bf10 = 0.20412, posterior_h1 = 0.615385   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('bayes_factor_binomial_simple', successes=14, trials=20, p0=0.5,
p1=0.7)
# -> bf10 = 5.18442, log10_bf10 = 0.7147, posterior_h1 = 0.838303   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Bayes factor for p1 = 0.7 against p0 = 0.5 after 14 successes in 20 trials.	SOLVED	bf10 = 5.18442, log10_bf10 = 0.7147, posterior_h1 = 0.838303
Compute the Bayes factor with successes = 1, trials = 1, p0 = 0.5, p1 = 0.8.	SOLVED	bf10 = 1.6, log10_bf10 = 0.20412, posterior_h1 = 0.615385
Compute the Bayes factor with successes = 3, trials = 10, p0 = 1.5, p1 = 0.2.	INVALID_INPUT	Bayes factor for two simple binomial hypotheses: p0 must be strictly between 0 and 1.

Bayes factor for a point null on a proportion (beta prior under H1) — `bayes_factor_point_null_binomial`

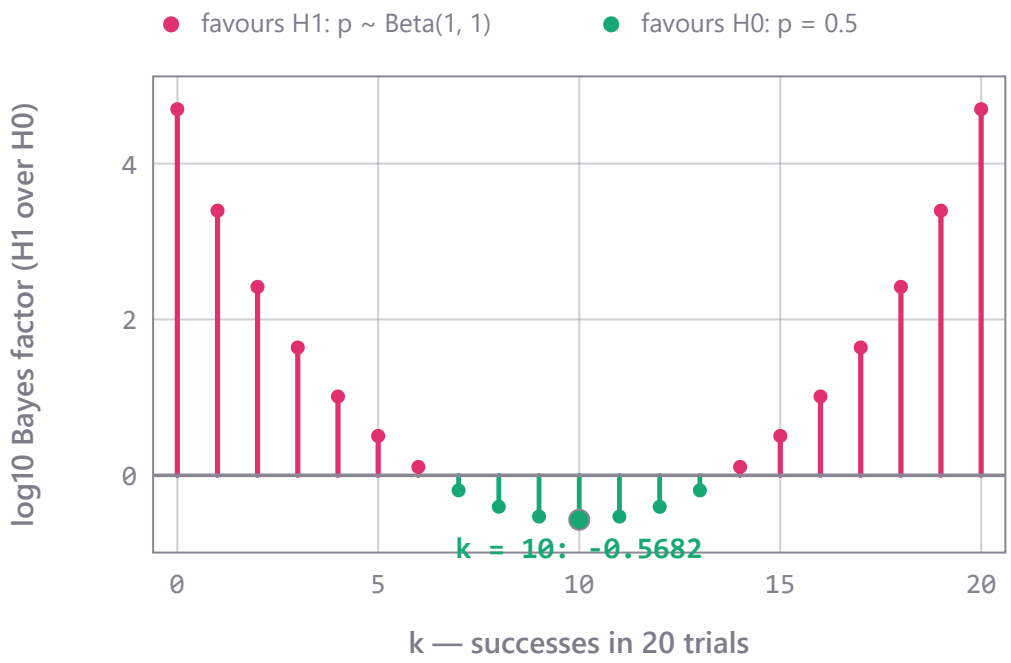
BF10 for H1: p ~ Beta(alpha, beta) against H0: $p = p_0$ after k successes in n trials (default uniform prior), with BF01 and the posterior probability of H1 at equal prior odds.

$$BF_{10} = \frac{B(\alpha + k,\ \beta + n - k) / B(\alpha, \beta)}{p_0^k (1 - p_0)^{n-k}}$$

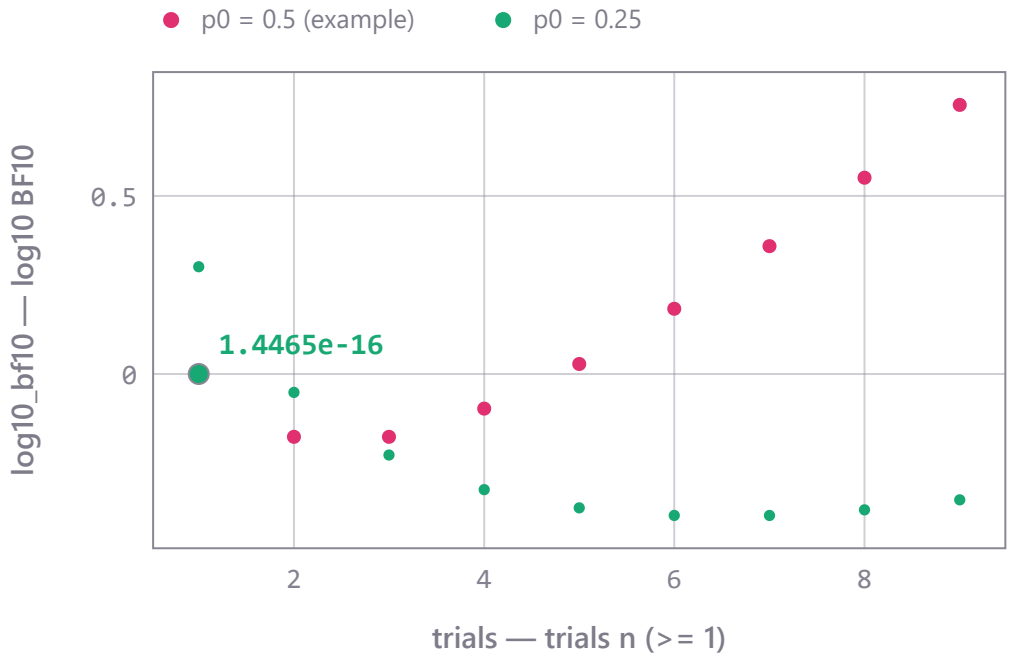
input	meaning	condition
`successes`	successes k	
`trials`	trials n (>= 1)	
`p0`	null value p0	
`alpha`	Beta alpha under H1 (default 1)	
`beta`	Beta beta under H1 (default 1)	

output	meaning
`bf10`	Bayes factor H1 : H0
`bf01`	$1/BF10$
`log10_bf10`	log10 BF10
`posterior_h1`	P(H1 | data) with prior odds 1

- **Method:** ratio of the beta-binomial marginal likelihood to the binomial likelihood at p0
- **Assumes:** $0 < p_0 < 1$; alpha, $\beta > 0$; $0 \leq k \leq n$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Bayes factor of a uniform alternative against the point null p = 0.5 for every number of successes in 20 trials (inputs chosen; the worked example has 1 trial), as log10 of the returned factor: counts near n/2 support the null, but only moderately (marked: k = 10, -0.568218), while extreme counts give strong evidence against it.



log10_bf10 (log10 BF10) against trials (trials n (>= 1)) from 1 to 9. One curve for each value of p0 (null value p0): 0.5, 0.25. Other inputs as in the worked example: successes = 1. 9 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, trials = 1, log10_bf10 = 1.4465e-16.

```
bayesian_statistics_ai999x.solve('bayes_factor_point_null_binomial', successes=1, trials=1, p0=0.5)
# -> bf10 = 1, bf01 = 1, log10_bf10 = 1.44649e-16, posterior_h1 = 0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('bayes_factor_point_null_binomial', successes=2, trials=2, p0=0.5)
# -> bf10 = 1.33333, bf01 = 0.75, log10_bf10 = 0.124939, posterior_h1 = 0.571429   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Bayes factor against the point null p0 = 0.5 with a uniform prior: 2 successes in 2 trials.	SOLVED	bf10 = 1.33333, bf01 = 0.75, log10_bf10 = 0.124939, posterior_h1 = 0.571429
Compute the Bayes factor for p0 = 0.25 with successes = 9, trials = 20, alpha = 2, beta = 2.	SOLVED	bf10 = 2.50393, bf01 = 0.399372, log10_bf10 = 0.398623, posterior_h1 = 0.714606
Compute the Bayes factor for p0 = 0 with successes = 9, trials = 20.	INVALID_INPUT	Bayes factor for a point null on a proportion (beta prior under H1): p0 must be strictly between 0 and 1.

Bayes factor for a normal mean (known variance) — `bayes_factor_normal_mean`

BF10 for H1: mu ~ N(mu0, tau^2) against H0: $mu = mu0$, from the sample mean of n observations with known sigma; BF01 and the posterior probability of H1 at equal prior odds.

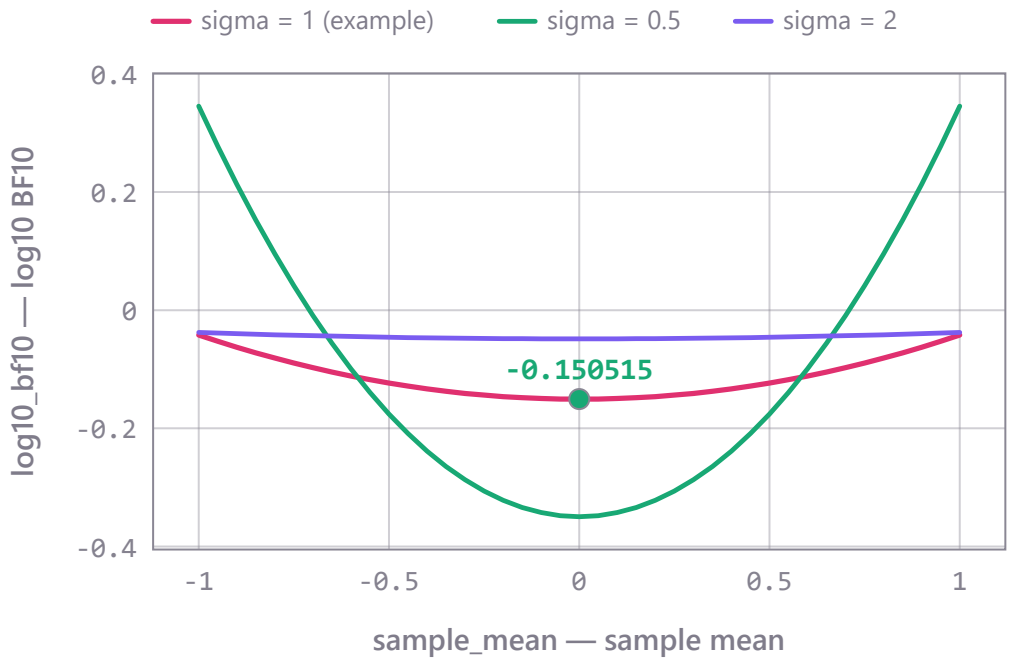
$$BF_{01} = \frac{p(\mu_0 \mid \bar{x}, H_1)}{p(\mu_0 \mid H_1)}$$

input	meaning	condition
`sample_mean`	sample mean	
`n`	sample size (>= 1)	

input	meaning	condition
`sigma`	known standard deviation (> 0)	
`mu0`	null value mu0	
`prior_sd`	prior standard deviation tau of mu under H1 (> 0)	

output	meaning
`bf10`	Bayes factor H1 : H0
`bf01`	$1/BF10$
`log10_bf10`	log10 BF10
`posterior_h1`	P(H1 | data) with prior odds 1

- **Method:** Savage-Dickey density ratio of the posterior and prior at mu0
- **Assumes:** sigma, $\tau > 0$; $n \geq 1$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



log10_bf10 (log10 BF10) against sample_mean (sample mean) from -1 to 1. One curve for each value of sigma (known standard deviation (> 0)): 1, 0.5, 2. Other inputs as in the worked example: n = 1, mu0 = 0, prior_sd = 1. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, sample_mean = 0, log10_bf10 = -0.150515.

```
bayesian_statistics_ai999x.solve('bayes_factor_normal_mean', sample_mean=0, n=1, sigma=1, mu0=0,
prior_sd=1)
# -> bf10 = 0.707107, bf01 = 1.41421, log10_bf10 = -0.150515, posterior_h1 = 0.414214   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('bayes_factor_normal_mean', sample_mean=0.5, n=25, sigma=2, mu0=0,
prior_sd=1)
# -> bf10 = 0.728324, bf01 = 1.37302, log10_bf10 = -0.137675, posterior_h1 = 0.421405   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Bayes factor for a normal mean: sample mean = 0, n = 1, sigma = 1, mu0 = 0, prior sd = 1.	SOLVED	bf10 = 0.707107, bf01 = 1.41421, log10_bf10 = -0.150515, posterior_h1 = 0.414214
Compute the Bayes factor with xbar = 101.2, n = 50, known sigma = 10, mu0 = 100 and tau = 5.	SOLVED	bf10 = 0.379838, bf01 = 2.6327, log10_bf10 = -0.420402, posterior_h1 = 0.275277
Compute the Bayes factor with xbar = 101.2, n = 50, known sigma = -10, mu0 = 100 and tau = 5.	INVALID_INPUT	Bayes factor for a normal mean (known variance): sigma must be positive.

Posterior model probabilities from log evidences — `posterior_model_probabilities`

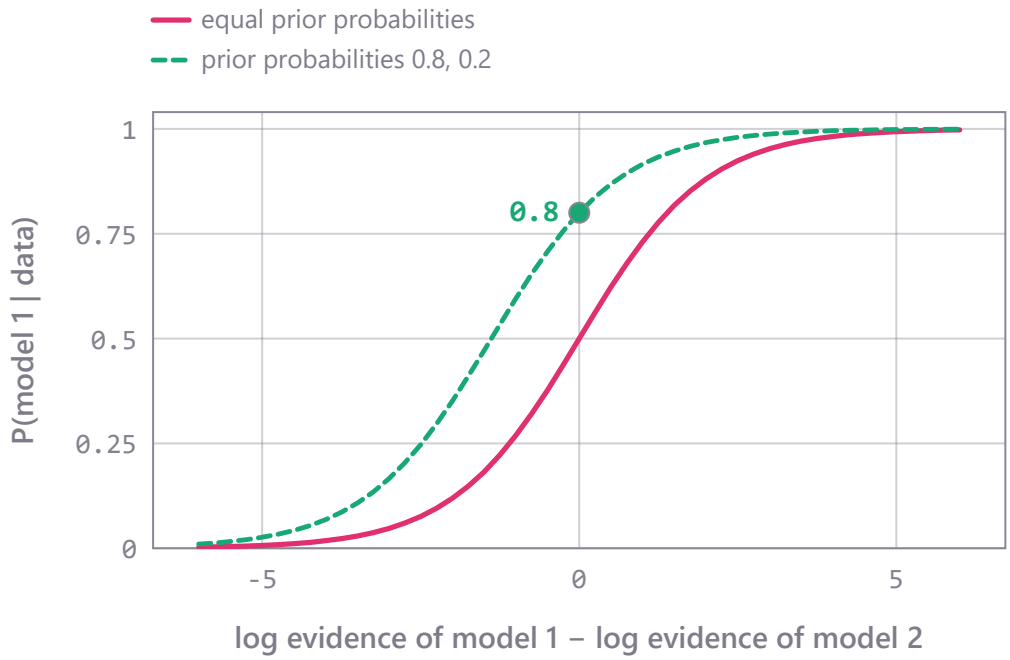
Posterior probabilities of several models from their log marginal likelihoods and prior model probabilities (default equal), computed stably in log space.

$$P(M_i \mid D) = \frac{P(D \mid M_i)P(M_i)}{\sum_j P(D \mid M_j)P(M_j)}$$

input	meaning	condition
`log_evidences`	natural-log marginal likelihoods of the models	
`priors`	prior model probabilities (default equal)	

output	meaning
`posterior`	posterior model probabilities
`best_model`	number of the most probable model
`best_probability`	its posterior probability

- **Method:** log-sum-exp normalization of log prior + log evidence
- **Assumes:** prior probabilities sum to 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The posterior probability of model 1 of two as the difference of their natural-log evidences runs from −6 to 6 (inputs chosen; the worked example has a difference of 0 and equal priors, giving 0.5), for equal priors and for priors 0.8, 0.2, every value from the package. Marked: equal evidence with priors 0.8, 0.2 gives 0.8.

```
bayesian_statistics_ai999x.solve('posterior_model_probabilities', log_evidences='-1, -1')
# -> posterior = [0.5, 0.5], best_model = 1, best_probability = 0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('posterior_model_probabilities', log_evidences=[0,
-1.0986122886681098], priors=[0.5, 0.5])
# -> posterior = [0.75, 0.25], best_model = 1, best_probability = 0.75   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Posterior model probabilities from log marginal likelihoods = [-10, -10, -10].	SOLVED	posterior = [0.333333333333, 0.333333333333, 0.333333333333], best_model = 1, best_probability = 0.333333
Models with log evidences [-100, -101] and prior model probabilities [0.3, 0.7]: posterior probabilities?	SOLVED	posterior = [0.538101526224, 0.461898473776], best_model = 1, best_probability = 0.538102
Models with log evidences [-100, -101] and prior model probabilities [0.3, 0.8]: posterior?	INVALID_INPUT	Posterior model probabilities from log evidences: the prior model probabilities must sum to 1 (they sum to 1.1).

Best action by expected monetary value — `expected_value_decision`

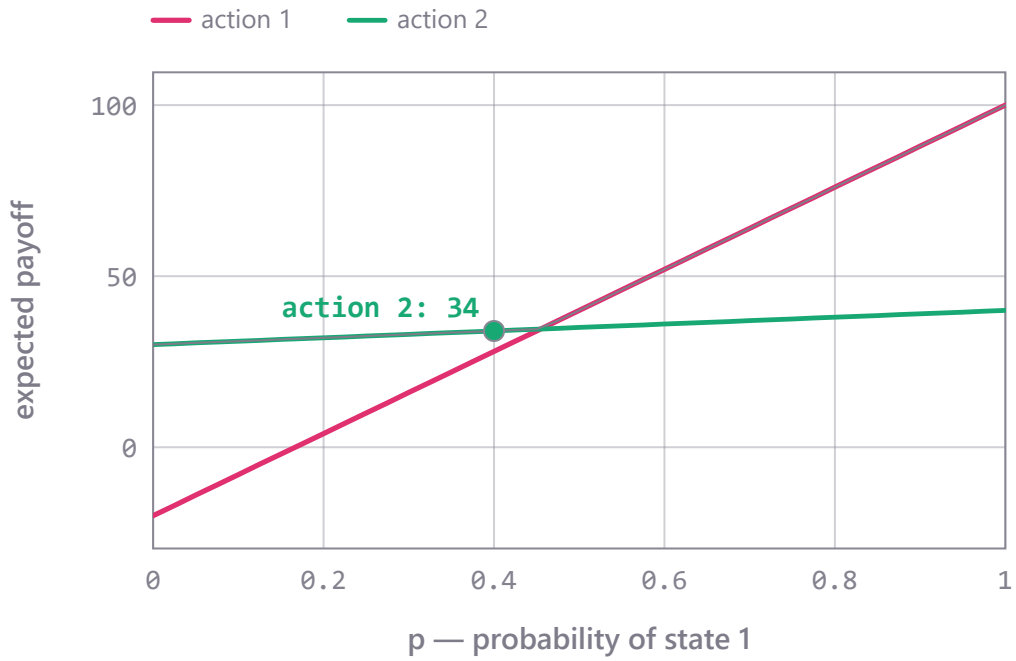
Expected payoff of each action over uncertain states with known probabilities, and the action with the highest expected value.

$$EV(a) = \sum_{s} P(s)\,u(a,s), \quad a^{*} = \arg\max_{a} EV(a)$$

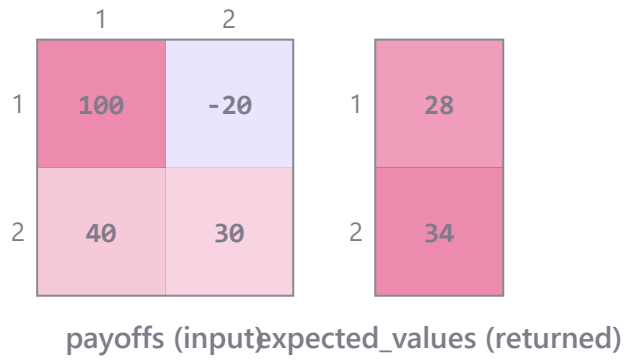
input	meaning	condition
`payoffs`	payoff matrix: one row per action, one column per state	
`probabilities`	state probabilities (sum to 1)	

output	meaning
`expected_values`	expected payoff of each action
`best_action`	number (1-based) of the best action (first if tied)
`best_expected_value`	its expected payoff

- **Method:** expected value of each row
- **Assumes:** probabilities sum to 1; one probability per column
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Expected payoff of each action of the payoff table [[100, -20], [40, 30]] (rows: actions, columns: states) as the probability of state 1 runs from 0 to 1, each from the package; the thin line is the best action's value. Marked: the worked example, probabilities 0.4, 0.6: the returned expected values 28, 34, best action 2 with 34.



The matrices of the worked example as heatmaps (rows and columns numbered; each cell shaded by its magnitude, negative values in the second colour; each panel shaded by its own largest magnitude): payoffs (input); expected_values (returned). Returned with them by bayesian-statistics-ai999x: best_action = 2, best_expected_value = 34. Other inputs: probabilities = 0.4, 0.6.

```
bayesian_statistics_ai999x.solve('expected_value_decision', payoffs='[[100, -20], [40, 30]]',
probabilities='0.4, 0.6')
# -> expected_values = [28.0, 34.0], best_action = 2, best_expected_value = 34   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('expected_value_decision', payoffs=[[5, 1, 0], [3, 3, 3]],
probabilities=[0.2, 0.5, 0.3])
# -> expected_values = [1.5, 3.0], best_action = 2, best_expected_value = 3   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Expected value decision: payoffs = [[100, -20], [40, 30]], probabilities = [0.4, 0.6]. Best action?	SOLVED	expected_values = [28.0, 34.0], best_action = 2, best_expected_value = 34
Which action maximizes expected monetary value with payoff matrix [[5, 1, 0], [3, 3, 3]] and state probabilities [0.2, 0.5, 0.3]?	SOLVED	expected_values = [1.5, 3.0], best_action = 2, best_expected_value = 3
Expected value decision: payoffs = [[100, -20], [40, 30]], probabilities = [0.4, 0.4].	INVALID_INPUT	Best action by expected monetary value: the state probabilities must sum to 1 (they sum to 0.8).
Expected value decision: payoffs = [[100, -20], [40, 30]]. Best action?	NEEDS_INPUT	expected_value_decision needs probabilities (state probabilities (sum to 1)). Example: solve('expected_value_decision',

Decision under uncertainty: maximin, maximax, Laplace, minimax regret, Hurwicz — `decision_criteria_under_uncertainty`

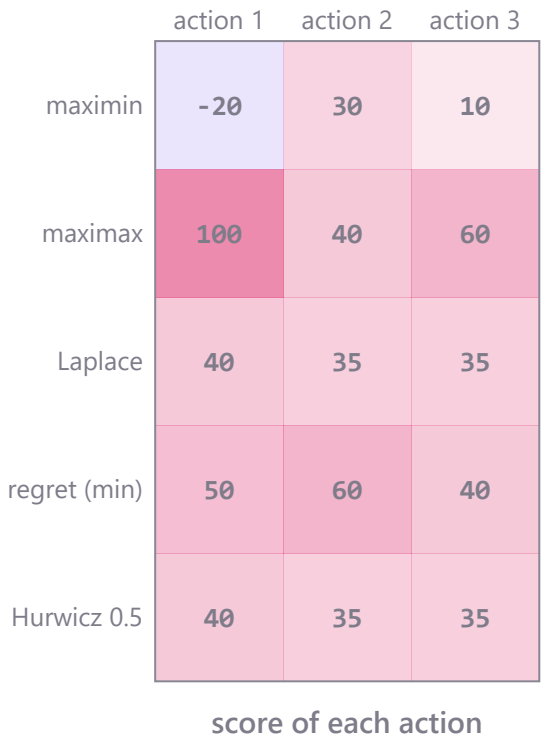
The best action and its criterion value when state probabilities are unknown, by a stated criterion.

$$\textbf{maximin: } \max_{a} \min_{s} u, \quad \text{regret } r(a,s) = \max_{a'} u(a',s) - u(a,s)$$

input	meaning	condition
`payoffs`	payoff matrix: one row per action, one column per state	
`criterion`	maximin, maximax, laplace, minimax regret or hurwicz	
`optimism`	Hurwicz optimism coefficient alpha (Hurwicz only)	

output	meaning
`scores`	criterion score of each action (maximum regret for minimax regret)
`best_action`	number of the chosen action (first if tied)
`criterion_value`	criterion value of the chosen action

- **Method:** row minima / maxima / means / Hurwicz mixtures, or maximum regret per row
- **Assumes:** $0 \leq \alpha \leq 1$ for Hurwicz
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The score of each action of the payoff table [[100, -20], [40, 30], [60, 10]] under five criteria (Hurwicz optimism 0.5, chosen), every score from the package; the regret row is minimised, the others maximised. Best action: maximin: action 2; maximax: action 1; laplace: action 1; minimax regret: action 3; hurwicz: action 1.

```
bayesian_statistics_ai999x.solve('decision_criteria_under_uncertainty', payoffs='[[100, -20], [40, 30], [60, 10]]', criterion='maximin')
# -> scores = [-20.0, 30.0, 10.0], best_action = 2, criterion_value = 30   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('decision_criteria_under_uncertainty', payoffs=[[5, 1, 0], [3, 3, 3]], criterion='minimax regret')
# -> scores = [3.0, 2.0], best_action = 2, criterion_value = 2   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Use the maximin criterion with payoffs = [[100, -20], [40, 30], [60, 10]].	SOLVED	scores = [-20.0, 30.0, 10.0], best_action = 2, criterion_value = 30
Minimax regret decision for payoff matrix [[5, 1, 0], [3, 3, 3]].	SOLVED	scores = [3.0, 2.0], best_action = 2, criterion_value = 2
Hurwicz criterion with optimism = 0.6 and payoffs = [[100, -20], [40, 30]].	SOLVED	scores = [52.0, 36.0], best_action = 1, criterion_value = 52
Hurwicz criterion with payoffs = [[100, -20], [40, 30]].	INVALID_INPUT	Decision under uncertainty: maximin, maximax, Laplace, minimax regret, Hurwicz: the Hurwicz criterion needs the optimism

Expected value of perfect information (EVPI) — `expected_value_of_perfect_information`

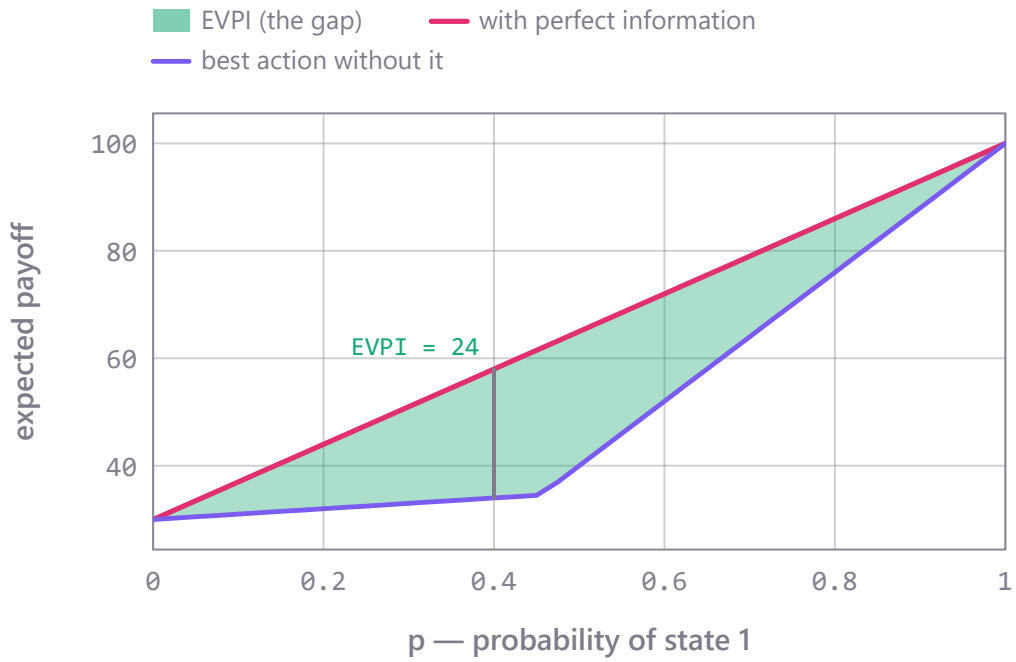
The most a decision maker should pay to learn the true state before acting: expected value with perfect information minus the best expected value without it.

$$EVPI = \sum_s P(s) \max_a u(a,s) - \max_a \sum_s P(s)\, u(a,s)$$

input	meaning	condition
`payoffs`	payoff matrix: one row per action, one column per state	
`probabilities`	state probabilities (sum to 1)	

output	meaning
`evpi`	expected value of perfect information
`ev_with_perfect_information`	expected value of the best action in each state
`best_expected_value`	best expected value without information

- **Method:** E[max over actions] - max over actions of E
- **Assumes:** probabilities sum to 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



For the payoff table [[100, -20], [40, 30]]: the expected payoff with perfect information about the state and the best expected payoff without it, over the probability of state 1 (41 runs each); their gap is the expected value of perfect information. Marked: probabilities 0.4, 0.6, returned values 58 and 34, EVPI = 24.

```
bayesian_statistics_ai999x.solve('expected_value_of_perfect_information', payoffs='[[100, -20], [40, 30]]', probabilities='0.4, 0.6')
# -> evpi = 24, ev_with_perfect_information = 58, best_expected_value = 34   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('expected_value_of_perfect_information', payoffs=[[5, 1, 0], [3, 3, 3]], probabilities=[0.2, 0.5, 0.3])
# -> evpi = 0.4, ev_with_perfect_information = 3.4, best_expected_value = 3   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Expected value of perfect information: payoffs = [[100, -20], [40, 30]], probabilities = [0.4, 0.6].	SOLVED	evpi = 24, ev_with_perfect_information = 58, best_expected_value = 34
Find the EVPI for payoff matrix [[5, 1, 0], [3, 3, 3]] and state probabilities [0.2, 0.5, 0.3].	SOLVED	evpi = 0.4, ev_with_perfect_information = 3.4, best_expected_value = 3
Find the EVPI for payoff matrix [[5, 1, 0], [3, 3, 3]] and state probabilities [0.2, 0.8].	INVALID_INPUT	Expected value of perfect information (EVPI): the state probabilities must have 3 values (one per entry of the other lis

Expected value of sample information (EVSI) — `expected_value_of_sample_information`

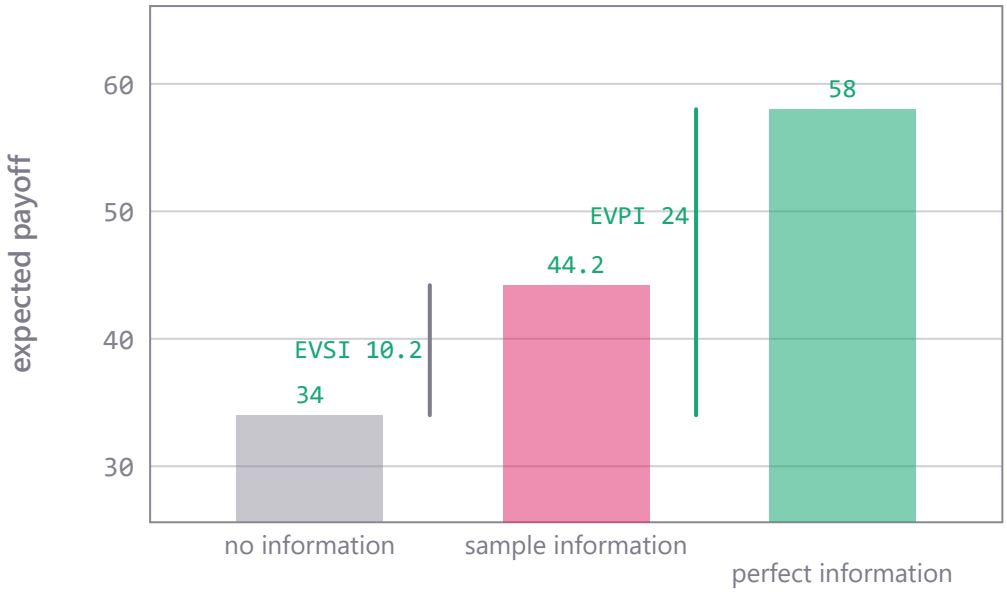
Value of an imperfect test before acting: for each possible signal, update the state probabilities by Bayes' theorem, choose the best action, and compare the expected value with that of acting now.

$$EVSI = \sum_{z} P(z) \max_{a} \sum_{s} P(s \mid z)\, u(a, s) - \max_{a} \sum_{s} P(s)\, u(a, s)$$

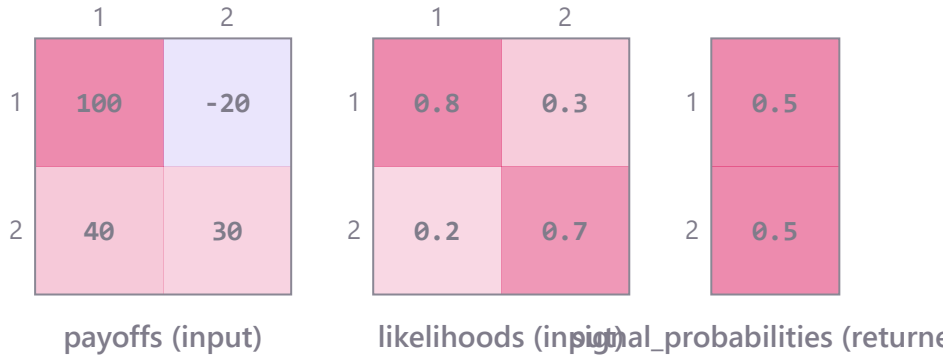
input	meaning	condition
`payoffs`	payoff matrix: one row per action, one column per state	
`priors`	prior state probabilities (sum to 1)	
`likelihoods`	likelihood matrix P(signal | state): one row per signal, one column per state	

output	meaning
`evsi`	expected value of sample information
`ev_with_information`	expected value with the test
`ev_without_information`	best expected value without the test
`signal_probabilities`	probability of each signal
`best_action_per_signal`	best action after each signal

- **Method:** preposterior analysis: Bayes' theorem per signal, best action per posterior, weighted by P(signal)
- **Assumes:** priors sum to 1; each likelihood column sums to 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



For the payoff table [[100, -20], [40, 30]], priors 0.4, 0.6 and a signal with likelihoods [[0.8, 0.3], [0.2, 0.7]] (the worked example): the returned expected payoff without information 34, with the sample information 44.2 (EVSI = 10.2; signal probabilities 0.5, 0.5, best action per signal [1, 2]) and, from the perfect-information capability, 58 (EVPI = 24): sample information is worth at most the perfect kind.



The matrices of the worked example as heatmaps (rows and columns numbered; each cell shaded by its magnitude, negative values in the second colour; each panel shaded by its own largest magnitude): payoffs (input); likelihoods (input); signal_probabilities (returned). Returned with them by bayesian-statistics-ai999x: evsi = 10.2, ev_with_information = 44.2, ev_without_information = 34. Other inputs: priors = 0.4, 0.6.

```
bayesian_statistics_ai999x.solve('expected_value_of_sample_information', payoffs='[[100, -20], [40, 30]]', priors='0.4, 0.6', likelihoods='[[0.8, 0.3], [0.2, 0.7]]')
# -> evsi = 10.2, ev_with_information = 44.2, ev_without_information = 34, signal_probabilities = [0.5, 0.5], best_action_per_signal = [1, 2]   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('expected_value_of_sample_information', payoffs=[[50, -10], [0, 0]], priors=[0.5, 0.5], likelihoods=[[0.9, 0.2], [0.1, 0.8]])
# -> evsi = 1.5, ev_with_information = 21.5, ev_without_information = 20, signal_probabilities = [0.55, 0.45], best_action_per_signal = [1, 2]   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Find the EVSI with payoff matrix [[50, -10], [0, 0]], prior probabilities [0.5, 0.5] and likelihood matrix [[0.9, 0.2], [0.1, 0.8]].	SOLVED	evsi = 1.5, ev_with_information = 21.5, ev_without_information = 20, signal_probabilities = [0.55, 0.45], best_action_per_signal = [1, 2]
Expected value of sample information: payoffs = [[100, -20], [40, 30]], priors = [0.4, 0.6], likelihoods = [[0.8, 0.3], [0.3, 0.7]].	INVALID_INPUT	Expected value of sample information (EVSI): each column of the likelihood matrix (P(signal | state)) must sum to 1.

Bayes action: minimum posterior expected loss — `bayes_action_minimum_expected_loss`

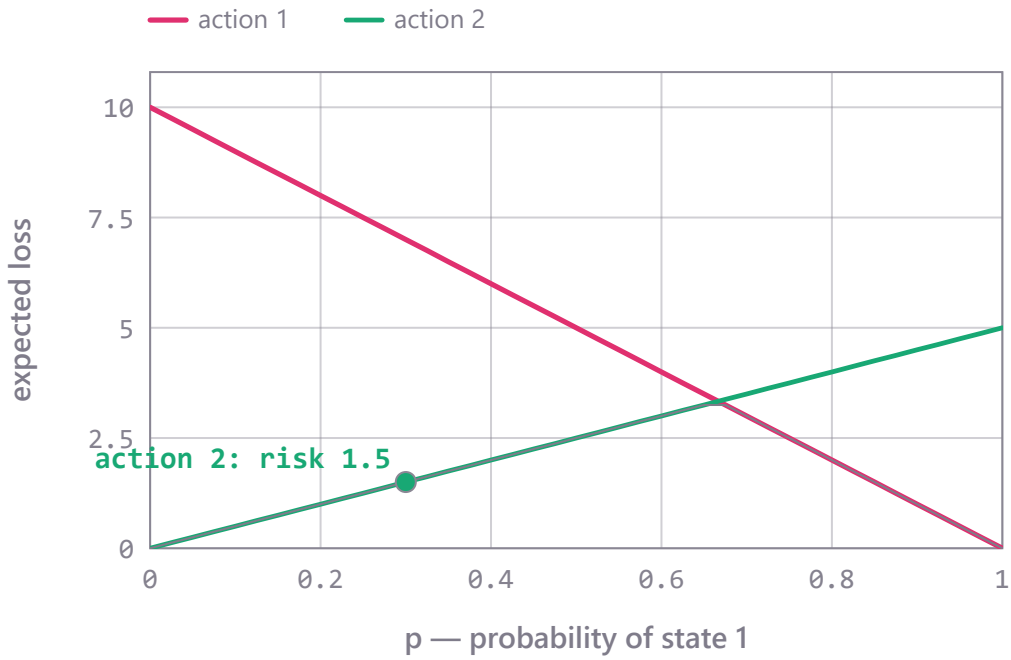
Posterior expected loss of each action from a loss matrix and (posterior) state probabilities; the Bayes action minimizes it.

$$\rho(a) = \sum_{s} P(s \mid x)\, L(a, s), \quad a^* = \arg\min_{a} \rho(a)$$

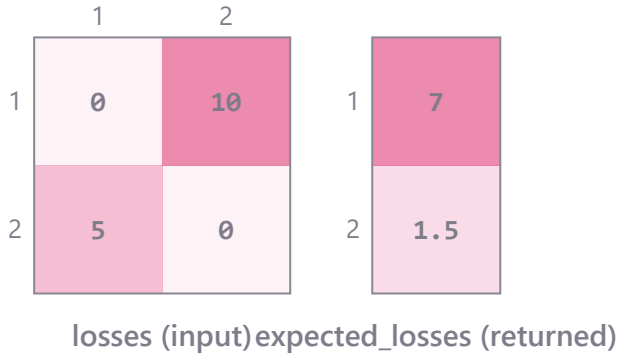
input	meaning	condition
`losses`	loss matrix: one row per action, one column per state	
`probabilities`	posterior state probabilities (sum to 1)	

output	meaning
`expected_losses`	expected loss of each action
`bayes_action`	action with least expected loss
`bayes_risk`	its expected loss

- **Method:** expected loss per row, minimum
- **Assumes:** probabilities sum to 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Expected loss of each action of the loss table [[0, 10], [5, 0]] (rows: actions, columns: states) over the probability of state 1, 41 runs; the thin line is the smallest (the Bayes risk). Marked: the worked example, probabilities 0.3, 0.7: returned expected losses 7, 1.5, Bayes action 2, Bayes risk 1.5.



The matrices of the worked example as heatmaps (rows and columns numbered; each cell shaded by its magnitude, negative values in the second colour; each panel shaded by its own largest magnitude): losses (input); expected_losses (returned). Returned with them by bayesian-statistics-ai999x: bayes_action = 2, bayes_risk = 1.5. Other inputs: probabilities = 0.3, 0.7.

```
bayesian_statistics_ai999x.solve('bayes_action_minimum_expected_loss', losses='[[0, 10], [5, 0]]',
probabilities='0.3, 0.7')
# -> expected_losses = [7.0, 1.5], bayes_action = 2, bayes_risk = 1.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('bayes_action_minimum_expected_loss', losses=[[0, 4, 8], [2, 2,
2]], probabilities=[0.5, 0.3, 0.2])
# -> expected_losses = [2.8, 2.0], bayes_action = 2, bayes_risk = 2   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Bayes action: losses = [[0, 10], [5, 0]], probabilities = [0.3, 0.7]. Minimum expected loss?	SOLVED	expected_losses = [7.0, 1.5], bayes_action = 2, bayes_risk = 1.5

question	outcome	result
Which action minimizes the expected loss with loss matrix [[0, 4, 8], [2, 2, 2]] and probabilities [0.5, 0.3, 0.2]?	SOLVED	expected_losses = [2.8, 2.0], bayes_action = 2, bayes_risk = 2
Bayes action: losses = [[0, 10], [5, 0]], probabilities = [0.3, 0.9].	INVALID_INPUT	Bayes action: minimum posterior expected loss: the state probabilities must sum to 1 (they sum to 1.2).

Expected utility and certainty equivalent (exponential utility) —

`certainty_equivalent_exponential_utility`

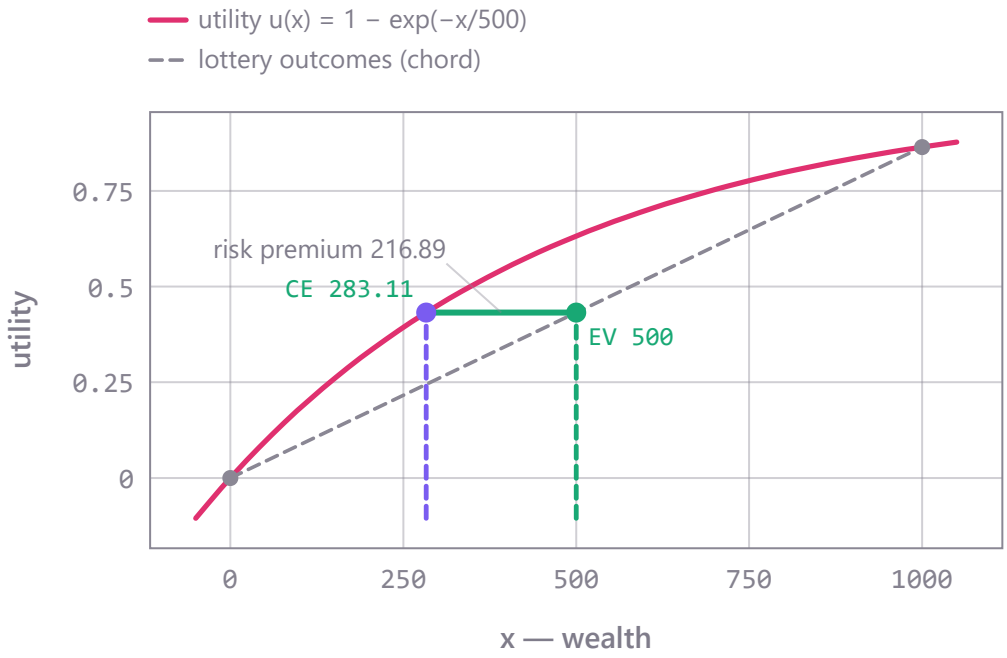
Expected utility, certainty equivalent and risk premium of a lottery for a risk-averse decision maker with exponential utility u(x) = 1 - exp(-x/R).

$$CE = -R \ln \sum_i p_i e^{-x_i/R}$$

input	meaning	condition
`outcomes`	monetary outcomes	
`probabilities`	their probabilities (sum to 1)	
`risk_tolerance`	risk tolerance R (> 0)	

output	meaning
`expected_utility`	E[$1 - exp$(-x/R)]
`certainty_equivalent`	-R ln E[exp(-x/R)]
`risk_premium`	expected value - certainty equivalent
`expected_value`	expected value

- **Method:** log-sum-exp evaluation of E[exp(-x/R)]
- **Assumes:** probabilities sum to 1; $R > 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A lottery with outcomes 1000, 0 and probabilities 0.5, 0.5 (the worked example): the utility is drawn from the package's expected utility of a sure outcome at 61 values. The returned expected utility 0.432332 sits on the chord at the expected value 500; the sure amount with the same utility is the returned certainty equivalent 283.11, and the gap is the returned risk premium 216.89.



certainty_equivalent (-R ln E[exp(-x/R)]) against risk_tolerance (risk tolerance R (> 0)) from 50 to 5000 on a logarithmic axis. Other inputs as in the worked example: outcomes = 1000, 0, probabilities = 0.5, 0.5. 41 points computed by bayesian-statistics-ai999x; marked: the worked example, risk_tolerance = 500, certainty_equivalent = 283.11.

```
bayesian_statistics_ai999x.solve('certainty_equivalent_exponential_utility', outcomes='1000, 0',
probabilities='0.5, 0.5', risk_tolerance=500)
# -> expected_utility = 0.432332, certainty_equivalent = 283.11, risk_premium = 216.89,
expected_value = 500   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('certainty_equivalent_exponential_utility', outcomes=[200, 50,
-100], probabilities=[0.3, 0.5, 0.2], risk_tolerance=250)
# -> expected_utility = 0.157471, certainty_equivalent = 42.8368, risk_premium = 22.1632,
expected_value = 65   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Exponential utility with risk tolerance = 500: outcomes = [1000, 0], probabilities = [0.5, 0.5]. Certainty equivalent?	SOLVED	expected_utility = 0.432332, certainty_equivalent = 283.11, risk_premium = 216.89, expected_value = 500
Exponential utility with risk tolerance = -5: outcomes = [1000, 0], probabilities = [0.5, 0.5].	INVALID_INPUT	Expected utility and certainty equivalent (exponential utility): the risk tolerance must be positive.

Certainty equivalent with logarithmic utility — `certainty_equivalent_log_utility`

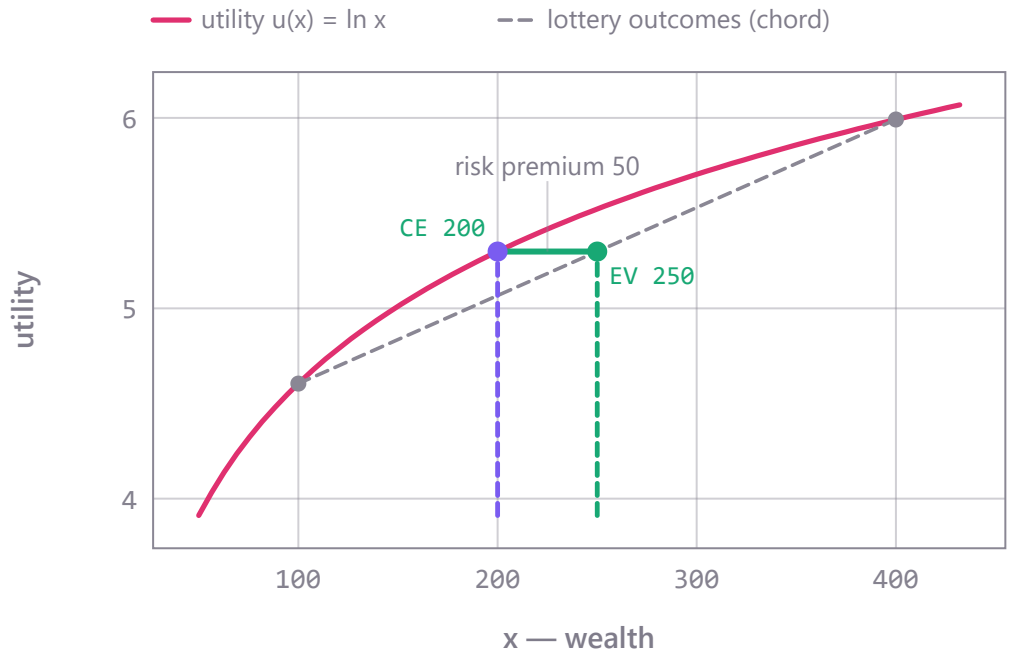
Expected log utility, certainty equivalent (the weighted geometric mean) and risk premium of a lottery over positive wealth levels.

$$CE = \exp\left(\sum_i p_i \ln x_i\right) = \prod_i x_i^{p_i}$$

input	meaning	condition
`outcomes`	positive wealth levels	
`probabilities`	their probabilities (sum to 1)	

output	meaning
`expected_log_utility`	E[ln x]
`certainty_equivalent`	exp(E[ln x])
`risk_premium`	E[x] - CE
`expected_value`	E[x]

- **Method:** weighted geometric mean
- **Assumes:** outcomes > 0; probabilities sum to 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A lottery with outcomes 100, 400 and probabilities 0.5, 0.5 (the worked example): the utility ln x is drawn from the package's expected utility of a sure outcome at 61 values. The returned expected utility 5.29832 sits on the chord at the expected value 250; the sure amount with the same utility is the returned certainty equivalent 200, and the gap is the returned risk premium 50.

```
bayesian_statistics_ai999x.solve('certainty_equivalent_log_utility', outcomes='100, 400',
probabilities='0.5, 0.5')
# -> expected_log_utility = 5.29832, certainty_equivalent = 200, risk_premium = 50, expected_value
= 250   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('certainty_equivalent_log_utility', outcomes=[50, 100, 200],
probabilities=[0.25, 0.5, 0.25])
# -> expected_log_utility = 4.60517, certainty_equivalent = 100, risk_premium = 12.5,
expected_value = 112.5   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Logarithmic utility: outcomes = [100, 400], probabilities = [0.5, 0.5]. Certainty equivalent?	SOLVED	expected_log_utility = 5.29832, certainty_equivalent = 200, risk_premium = 50, expected_value = 250
Find the certainty equivalent under log utility of wealth levels [50, 100, 200] with probabilities [0.25, 0.5, 0.25].	SOLVED	expected_log_utility = 4.60517, certainty_equivalent = 100, risk_premium = 12.5, expected_value = 112.5
Logarithmic utility: outcomes = [-100, 400], probabilities = [0.5, 0.5].	INVALID_INPUT	Certainty equivalent with logarithmic utility: logarithmic utility needs positive outcomes (wealth levels).

Break-even (indifference) probability of two actions — `break_even_probability`

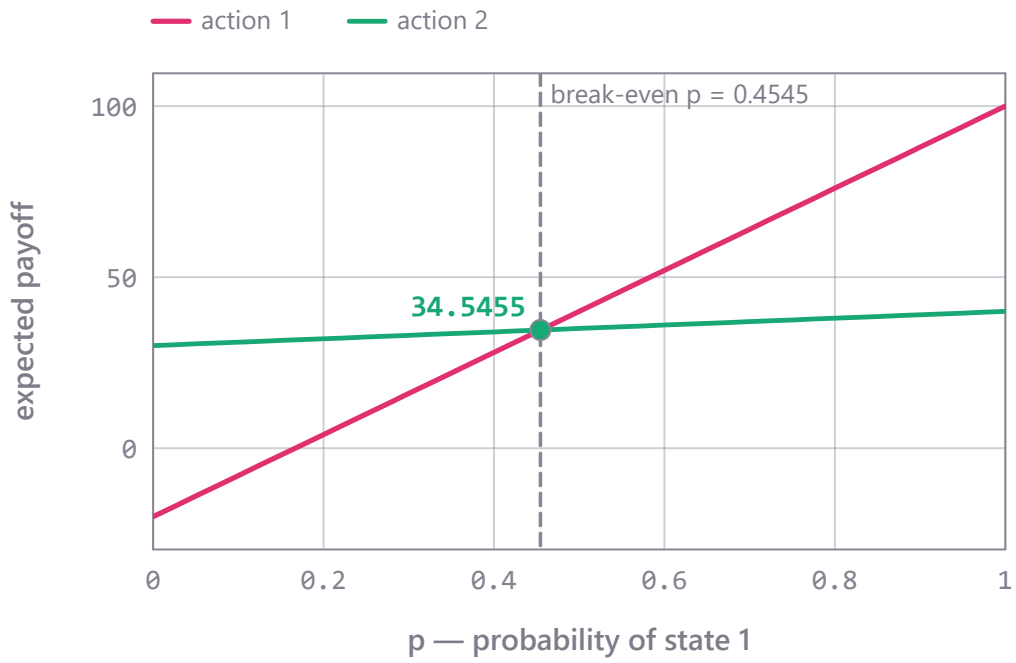
The probability of state 1 at which two actions have equal expected value, and which action is preferred above it.

$$p^* = \frac{u_{22} - u_{12}}{(u_{11} - u_{12}) - (u_{21} - u_{22})}$$

input	meaning	condition
`payoffs`	2 x 2 payoff matrix [[action 1 in state 1, in state 2], [action 2 ...]]	

output	meaning
`break_even`	probability of state 1 at which the expected values are equal
`preferred_above`	action preferred when P(state 1) exceeds it
`expected_value_at_break_even`	the common expected value there

- **Method:** intersection of the two expected-value lines
- **Assumes:** a crossing strictly between 0 and 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The two actions of the payoff table [[100, -20], [40, 30]]: their expected payoffs over the probability of state 1 (from the package's expected-value capability, 41 values of p) cross at the returned break-even probability 0.454545, where both are worth 34.5455; action 1 is preferred above it.

```
bayesian_statistics_ai999x.solve('break_even_probability', payoffs='[[100, -20], [40, 30]]')
# -> break_even = 0.454545, preferred_above = 1, expected_value_at_break_even = 34.5455   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('break_even_probability', payoffs=[[50, 0], [20, 10]])
# -> break_even = 0.25, preferred_above = 1, expected_value_at_break_even = 12.5   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Break-even probability for payoffs = [[100, -20], [40, 30]].	SOLVED	break_even = 0.454545, preferred_above = 1, expected_value_at_break_even = 34.5455
Break-even probability for payoffs = [[100, 50], [40, 30]].	INVALID_INPUT	Break-even (indifference) probability of two actions: one action is at least as good for every probability: there is no

Decision threshold probability (treat / do not treat) — `treatment_threshold`

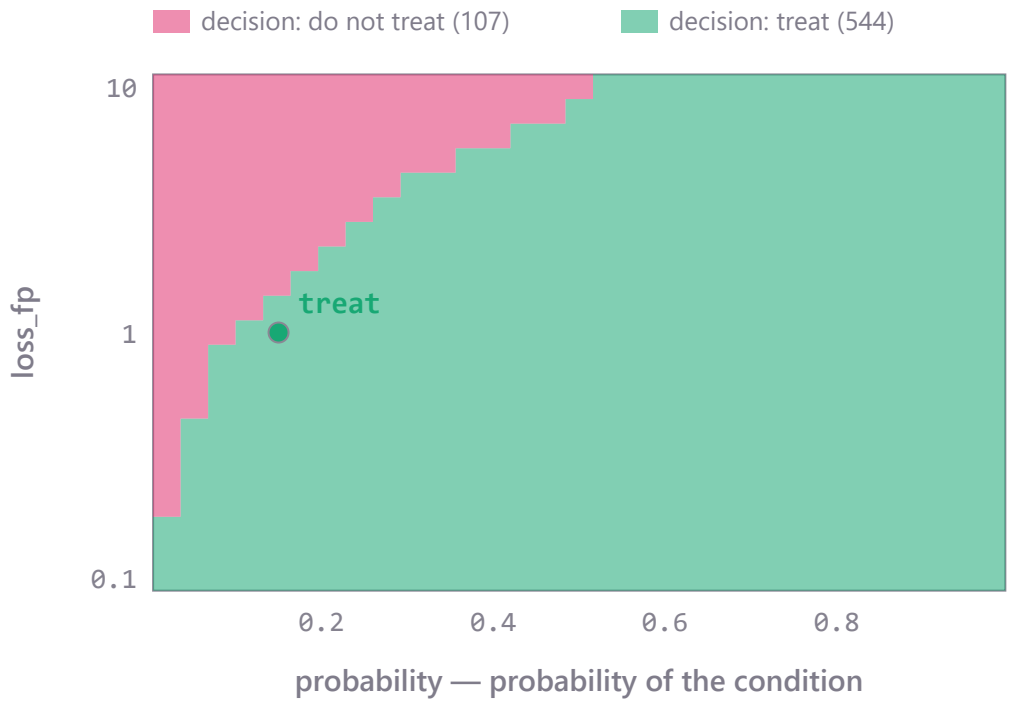
The probability above which acting (treating) has lower expected loss than not acting, from the loss of a false positive and of a false negative; the decision for a given probability.

$$p^* = \frac{L_{FP}}{L_{FP} + L_{FN}}$$

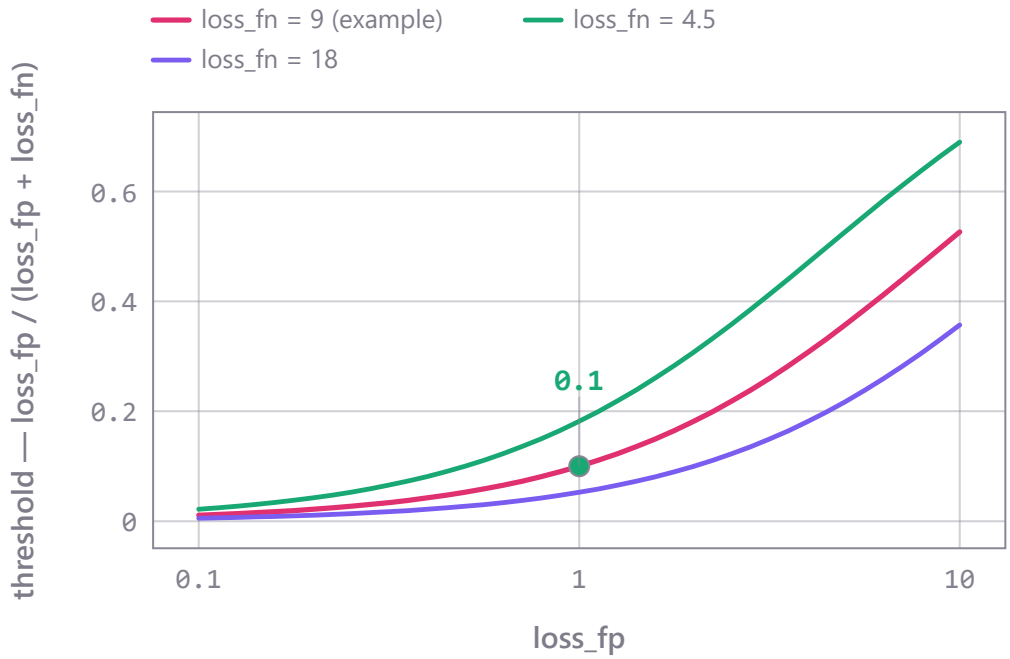
input	meaning	condition
`loss_fp`	loss of acting when the condition is absent (false positive)	
`loss_fn`	loss of not acting when the condition is present (false negative)	
`probability`	probability of the condition (optional)	

output	meaning
`threshold`	loss_fp / (loss_fp + loss_fn)
`decision`	treat / do not treat (if p given)

- **Method:** equal expected losses
- **Assumes:** losses > 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



treat / do not treat (if p given) (decision) over probability (probability of the condition (optional)) from 0.02 to 0.98 and loss_fp (loss of acting when the condition is absent (false positive)) from 0.1 to 10 (logarithmic vertical axis): each of the 31 × 21 cells is one run of bayesian-statistics-ai999x, coloured by the verdict it returned — do not treat: 107 cells, treat: 544 cells. Other inputs as in the worked example: loss_fn = 9. Marked: the worked example, probability = 0.15, loss_fp = 1 → treat.



threshold (loss_fp / (loss_fp + loss_fn)) against loss_fp (loss of acting when the condition is absent (false positive)) from 0.1 to 10 on a logarithmic axis. One curve for each value of loss_fn (loss of not acting when the condition is present (false negative)): 9, 4.5, 18. Other inputs as in the worked example: probability = 0.15. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, loss_fp = 1, threshold = 0.1.

```
bayesian_statistics_ai999x.solve('treatment_threshold', loss_fp=1, loss_fn=9, probability=0.15)
# -> threshold = 0.1, decision = treat   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('treatment_threshold', loss_fp=20, loss_fn=80)
# -> threshold = 0.2, decision = none   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Treatment threshold with loss_fp = 1 and loss_fn = 9; the probability of disease is 0.05. Treat?	SOLVED	threshold = 0.1, decision = do not treat
Decision threshold when a false-positive loss is 20 and a false-negative loss is 80.	SOLVED	threshold = 0.2, decision = none
Treatment threshold with loss_fp = -1 and loss_fn = 9.	INVALID_INPUT	Decision threshold probability (treat / do not treat): the loss of a false positive must be positive.

Uncertainty of a sum or linear combination (independent inputs) — `uncertainty_sum`

Value and combined standard uncertainty of y = sum c_i x_i for independent inputs (default $c_i = 1$; use -1 for a difference).

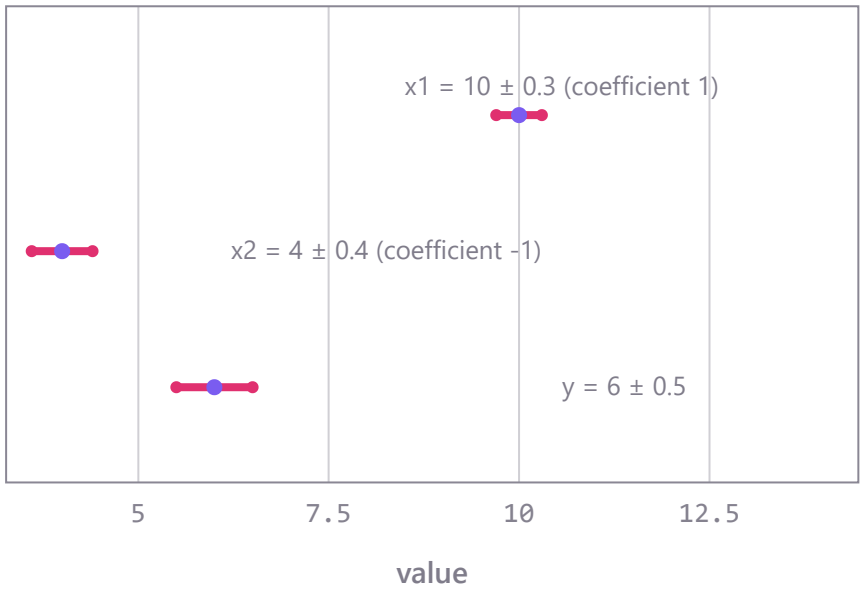
$$u_y = \sqrt{\sum_i c_i^2 u_i^2}$$

input	meaning	condition
`values`	input values x_i	

input	meaning	condition
`uncertainties`	standard uncertainties u_i (>= 0)	
`coefficients`	coefficients c_i (default all 1)	

output	meaning
`value`	y = sum c_i x_i
`uncertainty`	$u_y = sqrt$(sum (c_i u_i)^2)

- **Method:** root-sum-of-squares (GUM, first order)
- **Assumes:** independent inputs; $u_i \geq 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A weighted sum y = Σ c_i x_i of the inputs 10, 4 with standard uncertainties 0.3, 0.4 and coefficients 1, -1 (the worked example), each drawn as value ± uncertainty: the returned y = 6 ± 0.5 — uncertainties add in quadrature, so the result is less uncertain than their plain sum.

```
bayesian_statistics_ai999x.solve('uncertainty_sum', values='10, 4', uncertainties='0.3, 0.4',
coefficients='1, -1')
# -> value = 6, uncertainty = 0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('uncertainty_sum', values=[12.3, 4.5, 7.1], uncertainties=[0.2,
0.1, 0.3])
# -> value = 23.9, uncertainty = 0.374166   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Uncertainty of a sum: values = [12.3, 4.5, 7.1], uncertainties = [0.2, 0.1, 0.3].	SOLVED	value = 23.9, uncertainty = 0.374166
Linear combination with values [10, 4], uncertainties [0.3, 0.4] and coefficients [1, -1]: uncertainty of the difference?	SOLVED	value = 6, uncertainty = 0.5

question	outcome	result
Uncertainty of a sum: values = [12.3, 4.5], uncertainties = [0.2, -0.1].	INVALID_INPUT	Uncertainty of a sum or linear combination (independent inputs): standard uncertainties must not be negative.

Uncertainty of a product and a quotient — `uncertainty_product_quotient`

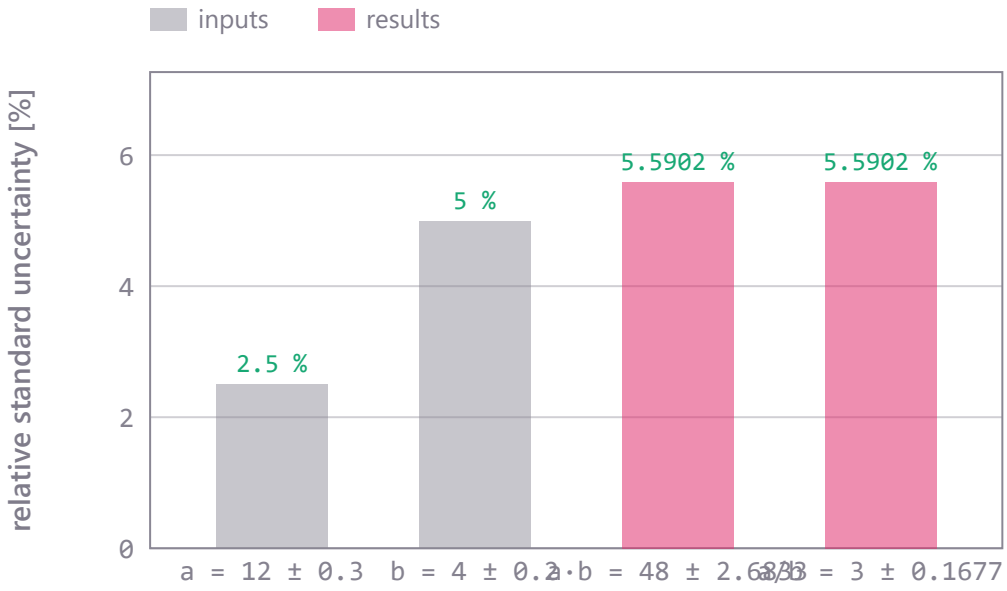
Value and first-order standard uncertainty of ab and a/b for independent a and b.

$$\frac{u_{ab}}{|ab|} = \frac{u_{a/b}}{|a/b|} = \sqrt{\left(\frac{u_a}{a}\right)^2 + \left(\frac{u_b}{b}\right)^2}$$

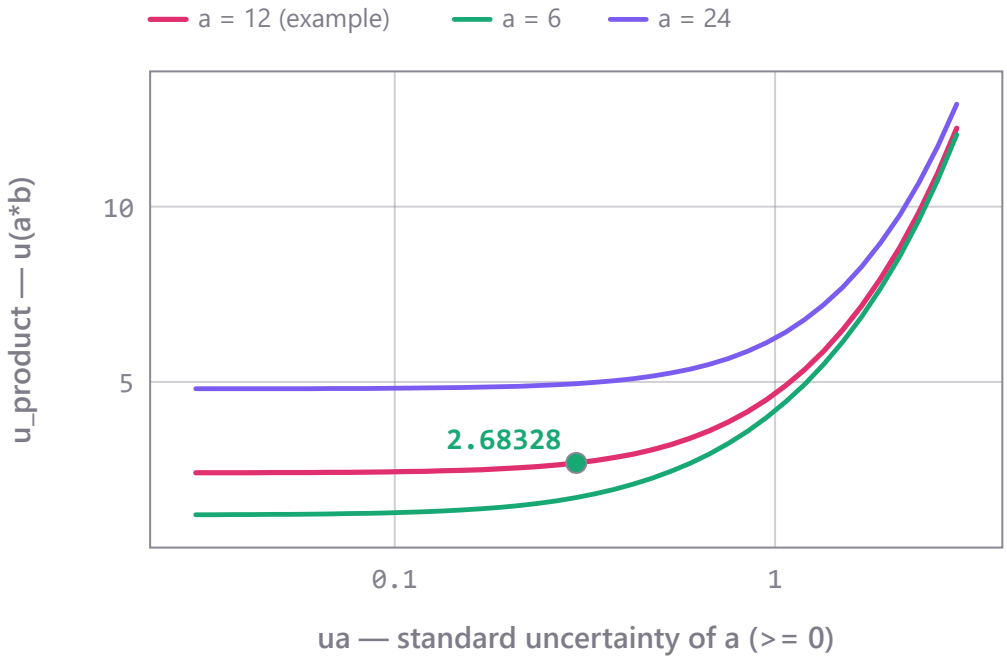
input	meaning	condition
`a`	value of a	
`ua`	standard uncertainty of a (>= 0)	
`b`	value of b (non-zero)	
`ub`	standard uncertainty of b (>= 0)	

output	meaning
`product`	ab
`u_product`	u(a*b)
`quotient`	a/b
`u_quotient`	u(a/b)

- **Method:** first-order propagation (relative uncertainties add in quadrature)
- **Assumes:** independent a and b; $b \neq 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Relative uncertainties for a = 12 ± 0.3 and b = 4 ± 0.2 (the worked example): the returned product 48 ± 2.68328 and quotient 3 ± 0.167705. Each bar is an uncertainty divided by its value (derived): product and quotient share the same relative uncertainty, the quadrature sum of the inputs'.



u_product (u(a*b)) against ua (standard uncertainty of a (>= 0)) from 0.03 to 3 on a logarithmic axis. One curve for each value of a (value of a): 12, 6, 24. Other inputs as in the worked example: b = 4.0, ub = 0.2. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, ua = 0.3, u_product = 2.68328.

```
bayesian_statistics_ai999x.solve('uncertainty_product_quotient', a=12.0, ua=0.3, b=4.0, ub=0.2)
# -> product = 48, u_product = 2.68328, quotient = 3, u_quotient = 0.167705   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('uncertainty_product_quotient', a=5, ua=0.1, b=2, ub=0.05)
# -> product = 10, u_product = 0.320156, quotient = 2.5, u_quotient = 0.0800391   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Uncertainty of the product and quotient: a = 12.0, ua = 0.3, b = 4.0, ub = 0.2.	SOLVED	product = 48, u_product = 2.68328, quotient = 3, u_quotient = 0.167705
Find the uncertainty of the product of a = 5 with u(a) = 0.1 and b = 2 with u(b) = 0.05.	SOLVED	product = 10, u_product = 0.320156, quotient = 2.5, u_quotient = 0.0800391
Uncertainty of the product and quotient: a = 12.0, ua = 0.3, b = 0, ub = 0.2.	INVALID_INPUT	Uncertainty of a product and a quotient: b must not be 0 (the quotient a/b).

Uncertainty of a product and quotient with correlation — `uncertainty_product_correlated`

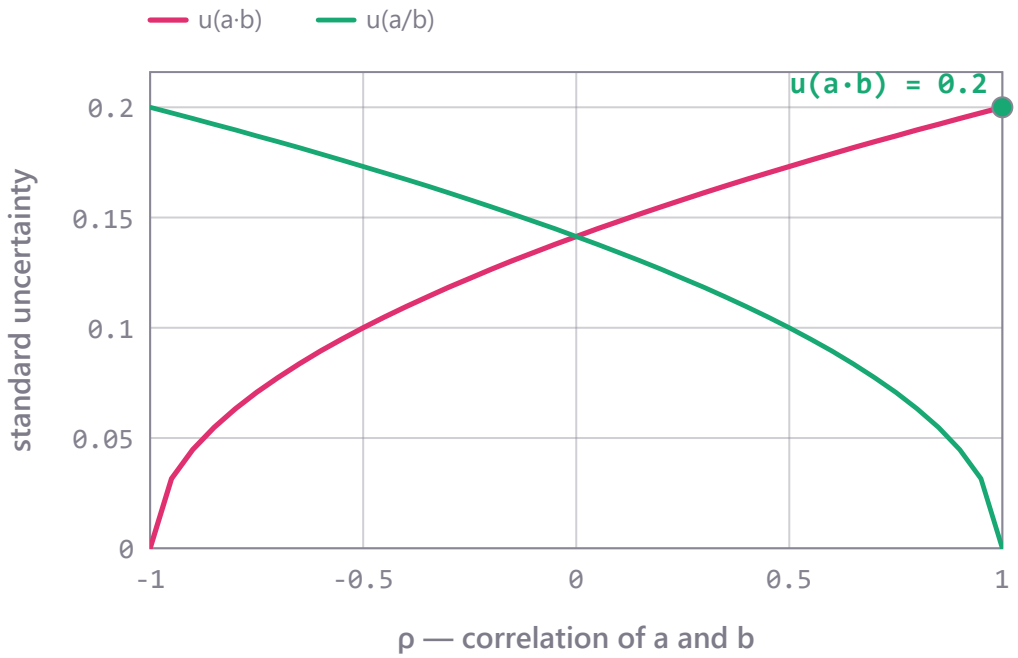
First-order standard uncertainty of ab and a/b when a and b are correlated with coefficient rho.

$$u_{ab}^2 = b^2 u_a^2 + a^2 u_b^2 + 2ab\,\rho\,u_a u_b$$

input	meaning	condition
`a`	value of a	
`ua`	standard uncertainty of a	
`b`	value of b (non-zero)	
`ub`	standard uncertainty of b	
`rho`	correlation coefficient of a and b (-1 to 1)	

output	meaning
`product`	ab
`u_product`	u(a*b)
`quotient`	a/b
`u_quotient`	u(a/b)

- **Method:** first-order propagation with the covariance term
- **Assumes:** $-1 \le \rho \le 1$; $b \neq 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Uncertainties of the product and the quotient of a = 1 ± 0.1 and b = 1 ± 0.1 as their correlation runs from −1 to 1, each from the package: positive correlation inflates the product's uncertainty and cancels the quotient's. Marked: the worked example, ρ = 1: u(a·b) = 0.2, u(a/b) = 0.

```
bayesian_statistics_ai999x.solve('uncertainty_product_correlated', a=1, ua=0.1, b=1, ub=0.1, rho=1)
# -> product = 1, u_product = 0.2, quotient = 1, u_quotient = 0   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('uncertainty_product_correlated', a=12.0, ua=0.3, b=4.0, ub=0.2,
rho=0.5)
# -> product = 48, u_product = 3.1749, quotient = 3, u_quotient = 0.129904   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Uncertainty of the product with correlation: a = 1, ua = 0.1, b = 1, ub = 0.1, rho = -1.	SOLVED	product = 1, u_product = 0, quotient = 1, u_quotient = 0.2
Uncertainty of the product with correlation: a = 1, ua = 0.1, b = 1, ub = 0.1, rho = 1.5.	INVALID_INPUT	Uncertainty of a product and quotient with correlation: a correlation coefficient lies between -1 and 1.

Uncertainty of a power x^n — `uncertainty_power`

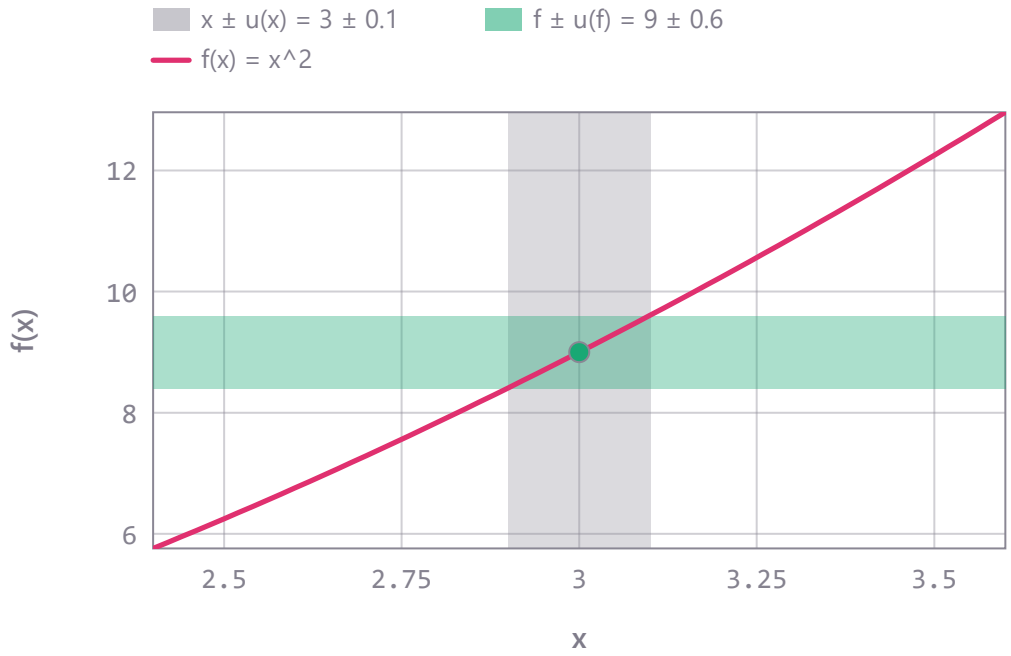
Value, standard uncertainty and relative uncertainty of $y = x^n$ (first order).

$$\frac{u_y}{|y|} = |n| \, \frac{u_x}{|x|}$$

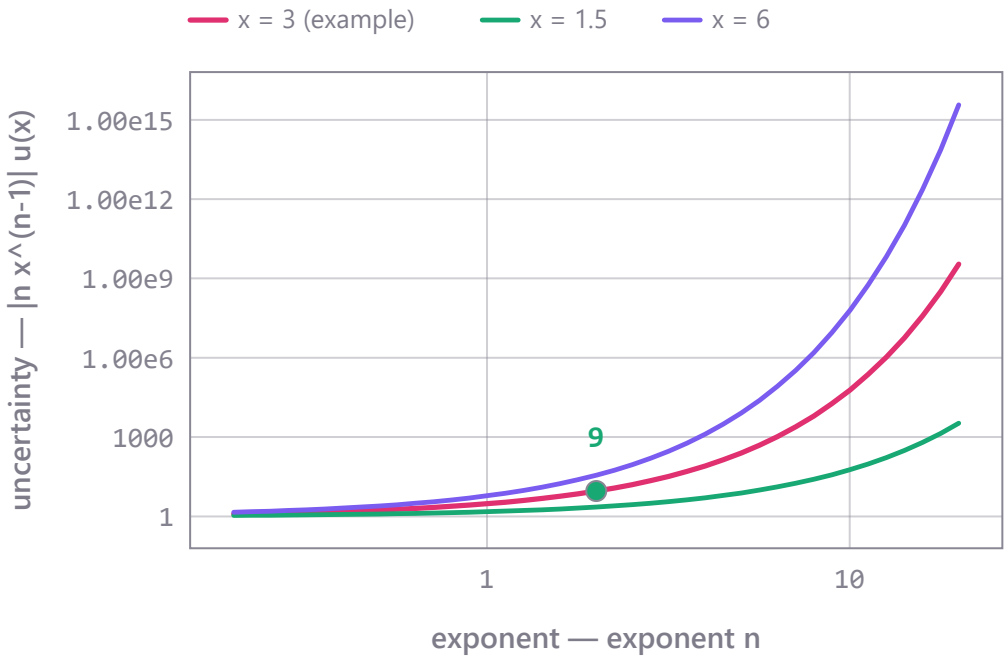
input	meaning	condition
`x`	value of x	
`ux`	standard uncertainty of x (>= 0)	
`exponent`	exponent n	

output	meaning
`value`	x^n
`uncertainty`	|n x^{n-1}| u(x)
`relative_uncertainty`	|n| u(x)/|x|

- **Method:** first-order propagation
- **Assumes:** $x > 0$ for a non-integer exponent; $x \neq 0$ for $n < 1$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Uncertainty of x^2 (the worked example): f is drawn from the package's value at 61 points around x = 3; the vertical band is the stated x ± u(x) = 3 ± 0.1, the horizontal band the returned f = 9 ± 0.6. The slope of f at x carries one band into the other. Returned relative uncertainty 0.0666667.



uncertainty (|n x^(n-1)| u(x)) against exponent (exponent n) from 0.2 to 20 on a logarithmic axis, logarithmic output axis. One curve for each value of x (value of x): 3, 1.5, 6. Other inputs as in the worked example: ux = 0.1. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, exponent = 2, uncertainty = 9.

```
bayesian_statistics_ai999x.solve('uncertainty_power', x=3.0, ux=0.1, exponent=2)
# -> value = 9, uncertainty = 0.6, relative_uncertainty = 0.0666667   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('uncertainty_power', x=16, ux=0.4, exponent=0.5)
# -> value = 4, uncertainty = 0.05, relative_uncertainty = 0.0125   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Uncertainty of a power: x = 3.0, ux = 0.1, exponent = 2.	SOLVED	value = 9, uncertainty = 0.6, relative_uncertainty = 0.0666667
Find the uncertainty of x^n with x = 16, u(x) = 0.4 and n = 0.5.	SOLVED	value = 4, uncertainty = 0.05, relative_uncertainty = 0.0125
Uncertainty of a power: x = -3.0, ux = 0.1, exponent = 0.5.	INVALID_INPUT	Uncertainty of a power x^n: a negative x needs a whole-number exponent.

Uncertainty of a function of one variable — `uncertainty_function`

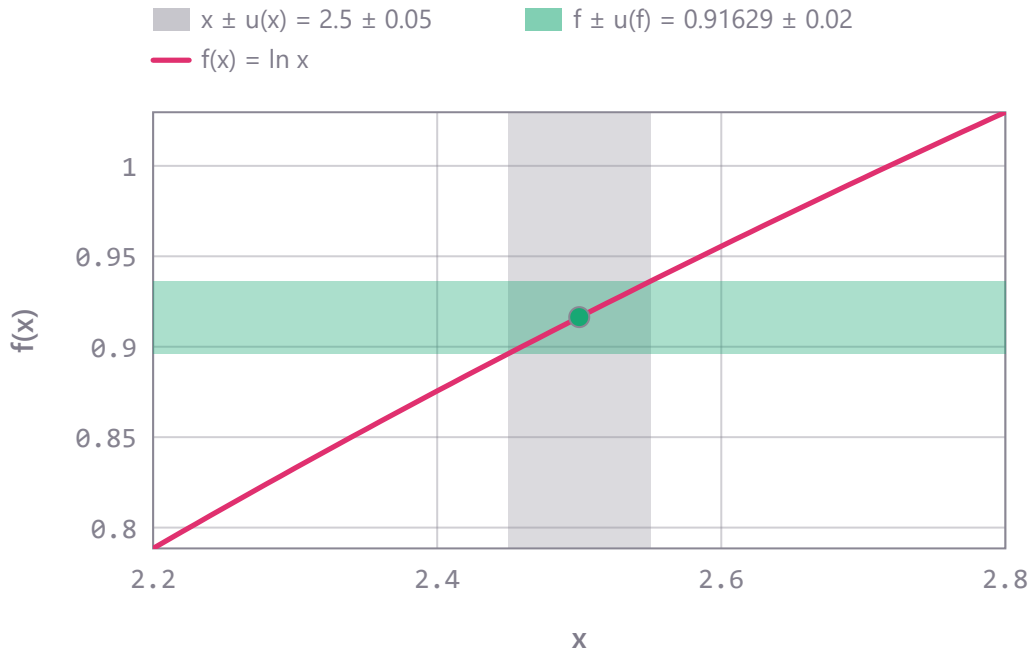
Value and first-order standard uncertainty of f(x) for exp, ln, log10, sqrt, sin, cos, tan (x in radians) or $1/x$.

$$u_f = |f'(x)|u_x$$

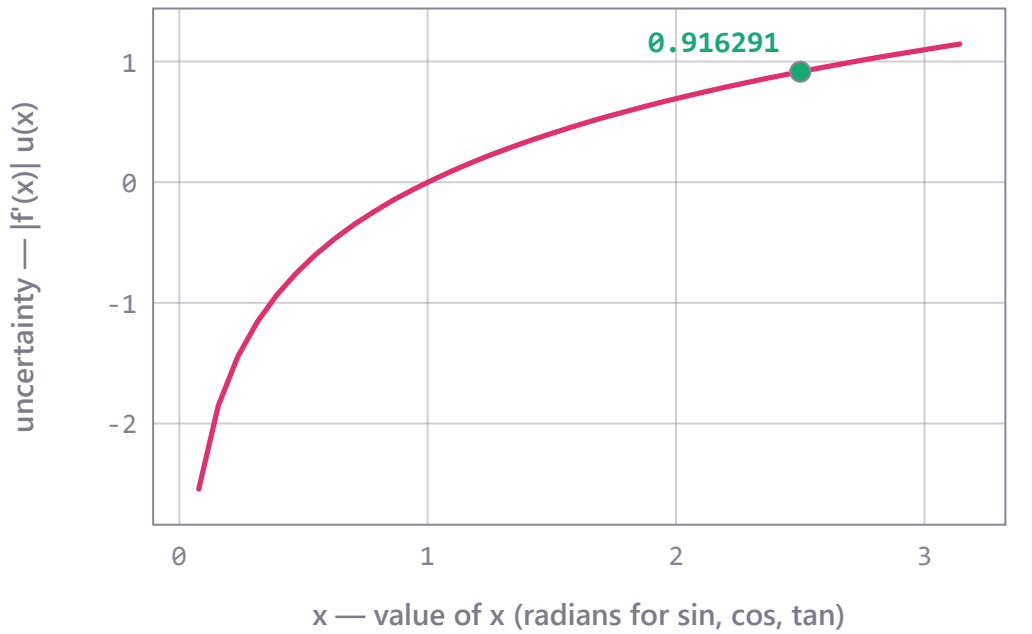
input	meaning	condition
`x`	value of x (radians for sin, cos, tan)	
`ux`	standard uncertainty of x	
`function`	exp, ln, log10, sqrt, sin, cos, tan or reciprocal	

output	meaning
`value`	f(x)
`uncertainty`	|f'(x)| u(x)

- **Method:** first-order propagation u_f = |f'(x)| u_x
- **Assumes:** x inside the function's domain
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Uncertainty through the function ln (the worked example): f is drawn from the package's value at 61 points around x = 2.5; the vertical band is the stated x ± u(x) = 2.5 ± 0.05, the horizontal band the returned f = 0.916291 ± 0.02. The slope of f at x carries one band into the other.



uncertainty (|f'(x)| u(x)) against x (value of x (radians for sin, cos, tan)) from 0 to 3.14159. Other inputs as in the worked example: ux = 0.05, function = ln. 40 points computed by bayesian-statistics-ai999x; marked: the worked example, x = 2.5, uncertainty = 0.916291.

```
bayesian_statistics_ai999x.solve('uncertainty_function', x=2.5, ux=0.05, function='ln')
# -> value = 0.916291, uncertainty = 0.02   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('uncertainty_function', x=0.0, ux=0.01, function='exp')
# -> value = 1, uncertainty = 0.01   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Uncertainty of a function: function = ln, x = 2.5, ux = 0.05.	SOLVED	value = 0.916291, uncertainty = 0.02
Uncertainty of a function: function = sqrt, x = 4, ux = 0.2.	SOLVED	value = 2, uncertainty = 0.05
Uncertainty of a function: function = ln, x = -2.5, ux = 0.05.	INVALID_INPUT	Uncertainty of a function of one variable: ln needs x > 0.

Uncertainty of a power-law product y = k x1^p1 x2^p2 … — `uncertainty_power_law`

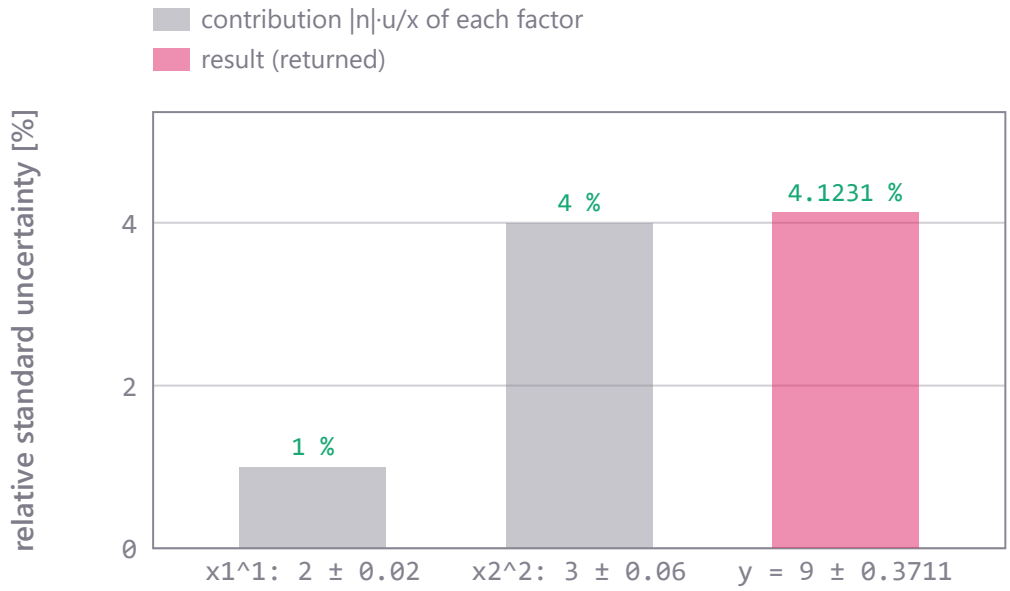
Value and first-order uncertainty of a product of powers of independent positive inputs (relative uncertainties weighted by the exponents add in quadrature).

$$\frac{u_y}{|y|} = \sqrt{\sum_i \left(p_i \frac{u_i}{x_i} \right)^2}$$

input	meaning	condition
`coefficient`	exact constant k	
`values`	positive input values x_i	
`exponents`	exponents p_i	
`uncertainties`	standard uncertainties u_i	

output	meaning
`value`	y
`uncertainty`	u_y
`relative_uncertainty`	u_y/|y|

- **Method:** first-order propagation in relative form
- **Assumes:** $x_i > 0$; independent inputs
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The power law y = 0.5 · Π x_i^n_i with values 2, 3, exponents 1, 2 and uncertainties 0.02, 0.06 (the worked example): each factor contributes |n|·u/x to the relative uncertainty (derived), and the returned relative uncertainty of y, 0.0412311, is their quadrature sum (y = 9 ± 0.37108).

```
bayesian_statistics_ai999x.solve('uncertainty_power_law', coefficient=0.5, values='2, 3',
exponents='1, 2', uncertainties='0.02, 0.06')
# -> value = 9, uncertainty = 0.37108, relative_uncertainty = 0.0412311   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('uncertainty_power_law', coefficient=1, values=[10, 4], exponents=
[2, -1], uncertainties=[0.1, 0.08])
# -> value = 25, uncertainty = 0.707107, relative_uncertainty = 0.0282843   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Uncertainty of a power law: coefficient = 0.5, values = [2.0, 3.0], exponents = [1, 2], uncertainties = [0.02, 0.06].	SOLVED	value = 9, uncertainty = 0.37108, relative_uncertainty = 0.0412311
Uncertainty of a power law: coefficient = 1, values = [-2.0, 3.0], exponents = [1, 2], uncertainties = [0.02, 0.06].	INVALID_INPUT	Uncertainty of a power-law product y = k x1^p1 x2^p2 ...: a power law needs positive values.

Inverse-variance weighted mean of measurements — `inverse_variance_weighted_mean`

Combine measurements of one quantity with different standard uncertainties: weighted mean, its uncertainty, chi-square of consistency and Birge ratio (equal to the posterior under a flat prior).

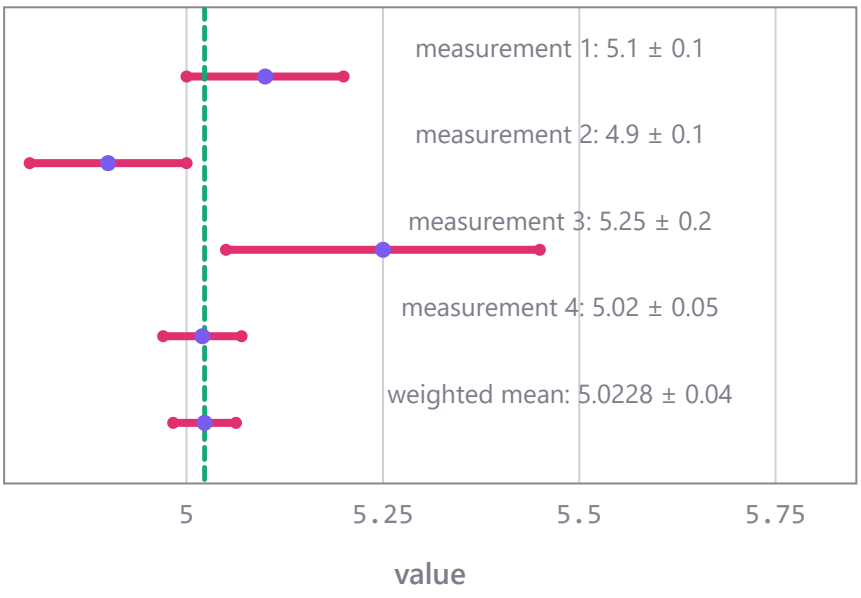
$$\bar{x}_w = \frac{\sum x_i/u_i^2}{\sum 1/u_i^2}, \quad u(\bar{x}_w) = \left(\sum 1/u_i^2\right)^{-1/2}$$

input	meaning	condition
`values`	measured values	

input	meaning	condition
`uncertainties`	their standard uncertainties (> 0)	

output	meaning
`weighted_mean`	sum(x_i/u_i^2)/sum(1/u_i^2)
`uncertainty`	$1/sqrt$(sum 1/u_i^2)
`chi_square`	sum((x_i - mean)^2/u_i^2)
`birge_ratio`	sqrt($\chi^2/(n-1)$)

- **Method:** inverse-variance weighting
- **Assumes:** $u_i > 0$; $n \geq 2$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Four measurements 5.1, 4.9, 5.25, 5.02 with standard uncertainties 0.1, 0.1, 0.2, 0.05 (inputs chosen; the worked example has two), each drawn as value ± uncertainty, and the returned inverse-variance weighted mean 5.0228 ± 0.04 (bottom row and vertical line): the most precise measurement pulls hardest. Returned χ^2 = 3.3976, Birge ratio 1.06421.

```
bayesian_statistics_ai999x.solve('inverse_variance_weighted_mean', values='5.1, 4.9',
uncertainties='0.1, 0.1')
# -> weighted_mean = 5, uncertainty = 0.0707107, chi_square = 2, birge_ratio = 1.41421   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('inverse_variance_weighted_mean', values=[10.2, 9.8, 10.5],
uncertainties=[0.2, 0.1, 0.4])
# -> weighted_mean = 9.90952, uncertainty = 0.0872872, chi_square = 5.4881, birge_ratio = 1.65652
(SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Combine the measurements [5.1, 4.9] with uncertainties [0.1, 0.1]: weighted mean?	SOLVED	weighted_mean = 5, uncertainty = 0.0707107, chi_square = 2, birge_ratio = 1.41421

question	outcome	result
Inverse-variance weighted mean of values = [10.2, 9.8] with uncertainties = [0.2, 0].	INVALID_INPUT	Inverse-variance weighted mean of measurements: uncertainties must be positive.

Uncertainty of a linear combination of correlated inputs — `uncertainty_linear_correlated`

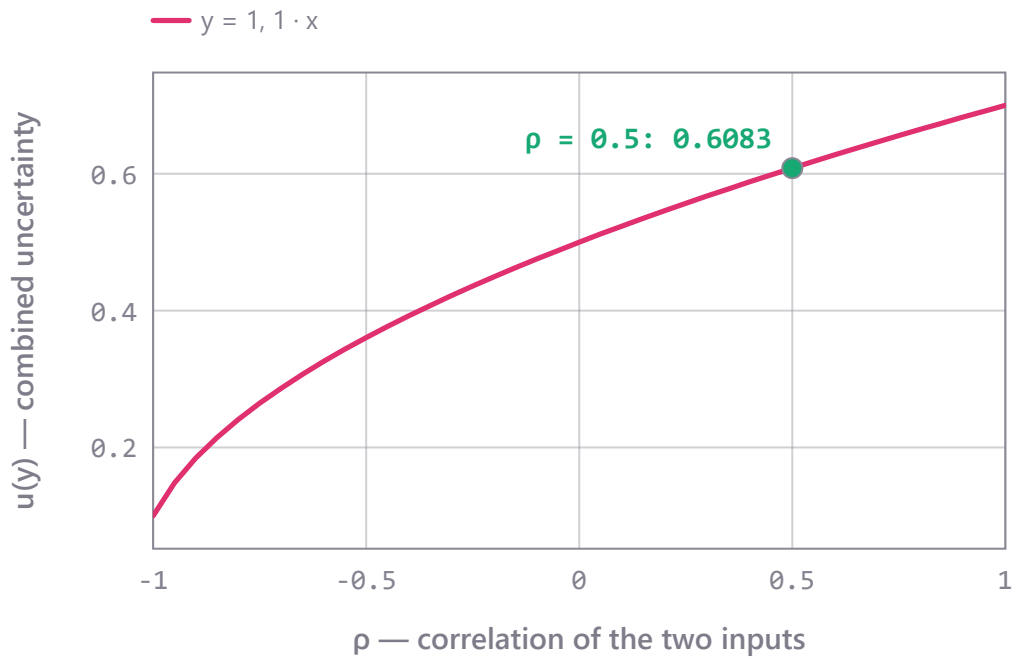
Combined standard uncertainty of y = sum c_i x_i when the inputs are correlated (correlation matrix), u_y^2 = sum_ij c_i c_j u_i u_j r_ij.

$$u_y^2 = \sum_i \sum_j c_i c_j u_i u_j r_{ij}$$

input	meaning	condition
`values`	input values	
`uncertainties`	standard uncertainties	
`correlation`	correlation matrix (ones on the diagonal, positive semi-definite)	
`coefficients`	coefficients c_i (default 1)	

output	meaning
`value`	y
`uncertainty`	u_y

- **Method:** quadratic form c^T Sigma c (GUM eq. 13)
- **Assumes:** valid correlation matrix
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The combined standard uncertainty of y = c1 x1 + c2 x2 for values 10, 20, uncertainties 0.3, 0.4 and coefficients 1, 1 (the worked example) as the correlation of the inputs runs from −1 to 1, each from the package: from 0.1 to 0.7. Marked: the worked example's ρ = 0.5, u(y) = 0.608276 (y = 30).

```
bayesian_statistics_ai999x.solve('uncertainty_linear_correlated', values='10, 20',
uncertainties='0.3, 0.4', correlation='[[1, 0.5], [0.5, 1]]', coefficients='1, 1')
# -> value = 30, uncertainty = 0.608276   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('uncertainty_linear_correlated', values=[1, 2], uncertainties=
[0.1, 0.1], correlation=[[1, 1], [1, 1]], coefficients=[1, -1])
# -> value = -1, uncertainty = 0   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Uncertainty of a linear combination with correlation matrix [[1, 0.5], [0.5, 1]], values = [10, 20], uncertainties = [0.3, 0.4], coefficients = [1, 1].	SOLVED	value = 30, uncertainty = 0.608276
Uncertainty of a linear combination with correlation matrix [[1, 1.5], [1.5, 1]], values = [10, 20], uncertainties = [0.3, 0.4].	INVALID_INPUT	Uncertainty of a linear combination of correlated inputs: a correlation matrix has ones on the diagonal and entries betw

Covariance propagation through a linear map y = A x — `covariance_propagation_linear`

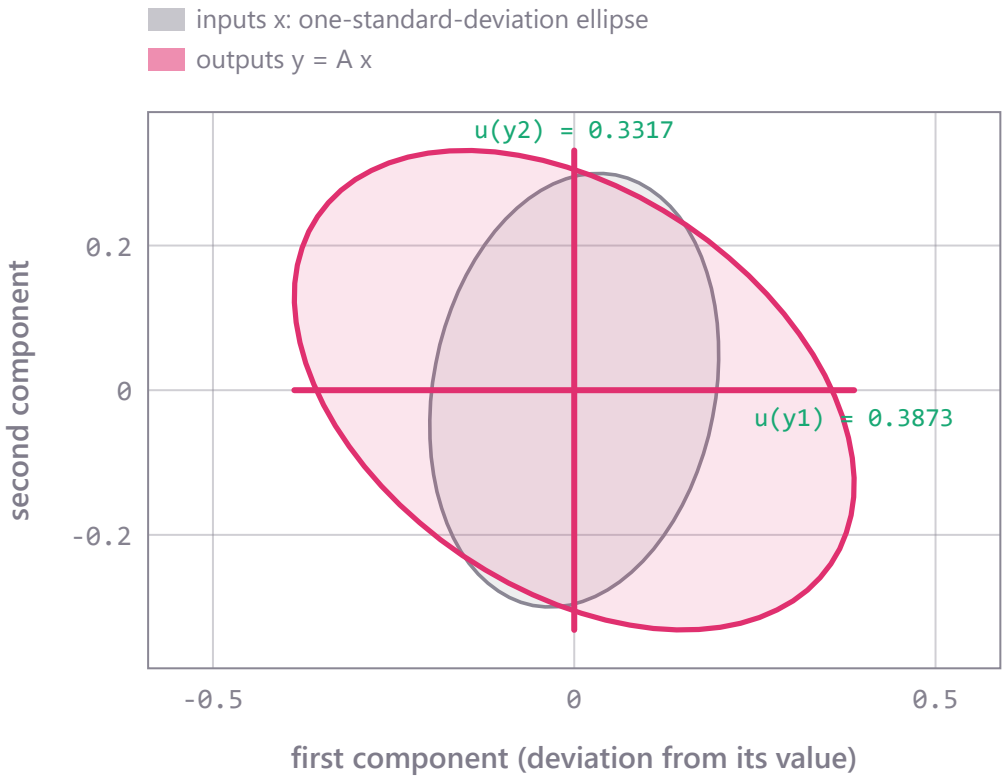
Covariance matrix of the outputs $y = A$ x and their standard uncertainties, given the covariance matrix of x (first-order propagation with a sensitivity matrix).

$$\Sigma_y = A\,\Sigma_x\,A^{\mathsf{T}}$$

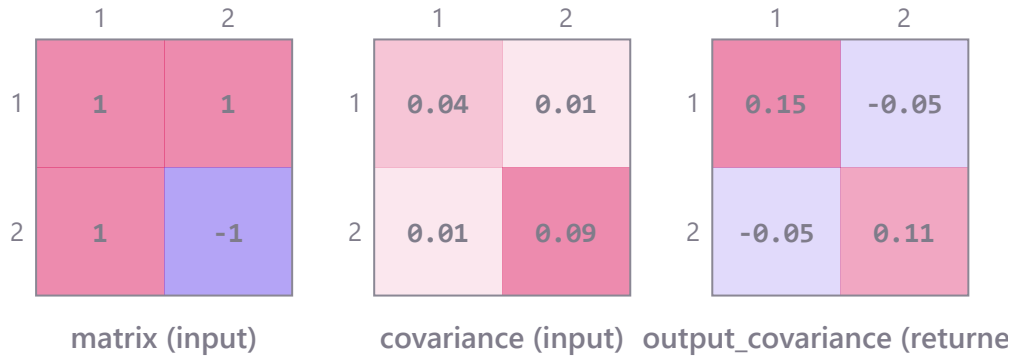
input	meaning	condition
`matrix`	sensitivity (Jacobian) matrix A: one row per output	
`covariance`	covariance matrix of the inputs (symmetric, positive semi-definite)	

output	meaning
`output_covariance`	A Sigma A^T
`uncertainties`	square roots of its diagonal

- **Method:** matrix product A Sigma A^T
- **Assumes:** covariance symmetric positive semi-definite; columns of A = inputs
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Linear propagation of a covariance: the input covariance [[0.04, 0.01], [0.01, 0.09]] through the matrix A = [[1, 1], [1, -1]] (the worked example) gives the returned output covariance [['0.15', '-0.05'], ['-0.05', '0.11']] and uncertainties 0.3873, 0.3317 (thick bars). Each ellipse is the one-standard-deviation contour of its covariance matrix (axes from its eigenvalues, derived); the tilt shows the correlation.



The matrices of the worked example as heatmaps (rows and columns numbered; each cell shaded by its magnitude, negative values in the second colour; each panel shaded by its own largest magnitude): matrix (input); covariance (input); output_covariance (returned).

```
bayesian_statistics_ai999x.solve('covariance_propagation_linear', matrix='[[1, 1], [1, -1]]',
covariance='[[0.04, 0.01], [0.01, 0.09]]')
# -> output_covariance = [[0.15, -0.05], [-0.05, 0.11]], uncertainties = [0.387298334621,
0.331662479036]   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('covariance_propagation_linear', matrix=[[2, 1]], covariance=[[1,
0], [0, 4]])
# -> output_covariance = [[8.0]], uncertainties = [2.82842712475]   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Covariance propagation: matrix = [[1, 1], [1, -1]], covariance = [[0.04, 0.01], [0.01, 0.09]].	SOLVED	output_covariance = [[0.15, -0.05], [-0.05, 0.11]], uncertainties = [0.387298334621, 0.331662479036]
Propagate the covariance matrix [[1, 2], [2, 1]] through the sensitivity matrix [[2, 1]].	INVALID_INPUT	Covariance propagation through a linear map y = A x: the covariance matrix must be positive semi-definite.

Effective degrees of freedom and expanded uncertainty (Welch-Satterthwaite) — `welch_satterthwaite_expanded_uncertainty`

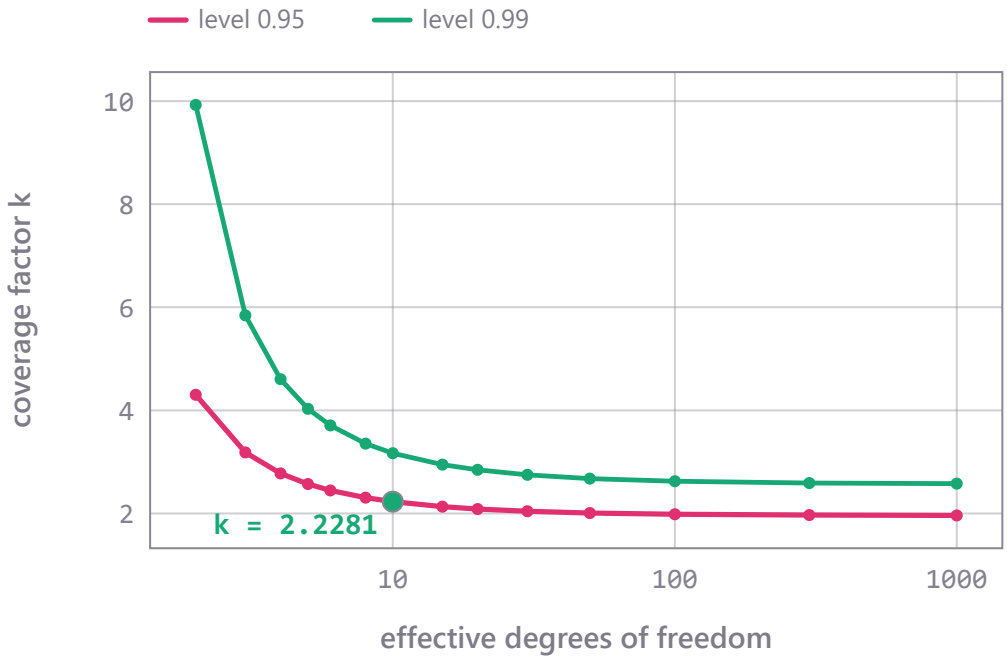
Combined standard uncertainty, Welch-Satterthwaite effective degrees of freedom, the Student-t coverage factor and the expanded uncertainty at a stated coverage probability.

$$\nu_{eff} = \frac{u_c^4}{\sum_i u_i^4/\nu_i}, \quad U = t_{(1+p)/2,\,\nu_{eff}}\, u_c$$

input	meaning	condition
`uncertainties`	uncertainty components c_i u_i (> 0)	
`dof`	degrees of freedom of each component (> 0)	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`combined_uncertainty`	u_c
`effective_dof`	nu_eff
`coverage_factor`	Missing superscript or subscript argument{$(1+p)/2$, nu_eff}
`expanded_uncertainty`	$U = k$ u_c

- **Method:** Welch-Satterthwaite formula and the Student-t quantile at non-integer degrees of freedom
- **Assumes:** components and degrees of freedom > 0; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The coverage factor k that turns a combined standard uncertainty into an expanded one, against the degrees of freedom (one component, 14 values from 2 to 1000, levels 0.95 and 0.99; each from the package): few degrees of freedom need a large k. Marked: the worked example, 10 degrees of freedom at level 0.95: k = 2.22814, expanded uncertainty 2.22814.



expanded_uncertainty (U = k u_c) against level (credibility level as a fraction (0.95 for 95%; default 0.95)) from 0.02 to 0.98. Other inputs as in the worked example: uncertainties = 1, dof = 10. 41 points computed by bayesian-statistics-ai999x; marked: the worked example, level = 0.95, expanded_uncertainty = 2.22814.

```
bayesian_statistics_ai999x.solve('welch_satterthwaite_expanded_uncertainty', uncertainties='1',
dof='10', level=0.95)
# -> combined_uncertainty = 1, effective_dof = 10, coverage_factor = 2.22814, expanded_uncertainty
= 2.22814   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('welch_satterthwaite_expanded_uncertainty', uncertainties=[0.3, 0.2], dof=[4, 9])
# -> combined_uncertainty = 0.360555, effective_dof = 7.67213, coverage_factor = 2.32327, expanded_uncertainty = 0.837667   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Welch-Satterthwaite: uncertainties = [1], degrees of freedom = [10]. Expanded uncertainty at 95% coverage?	SOLVED	combined_uncertainty = 1, effective_dof = 10, coverage_factor = 2.22814, expanded_uncertainty = 2.22814
Expanded uncertainty with components [0.1, 0.1], dof = [5, 0] and level = 0.95.	INVALID_INPUT	Effective degrees of freedom and expanded uncertainty (Welch-Satterthwaite): uncertainty components and degrees of freed

Type B standard uncertainty from a stated half-width — `type_b_standard_uncertainty`

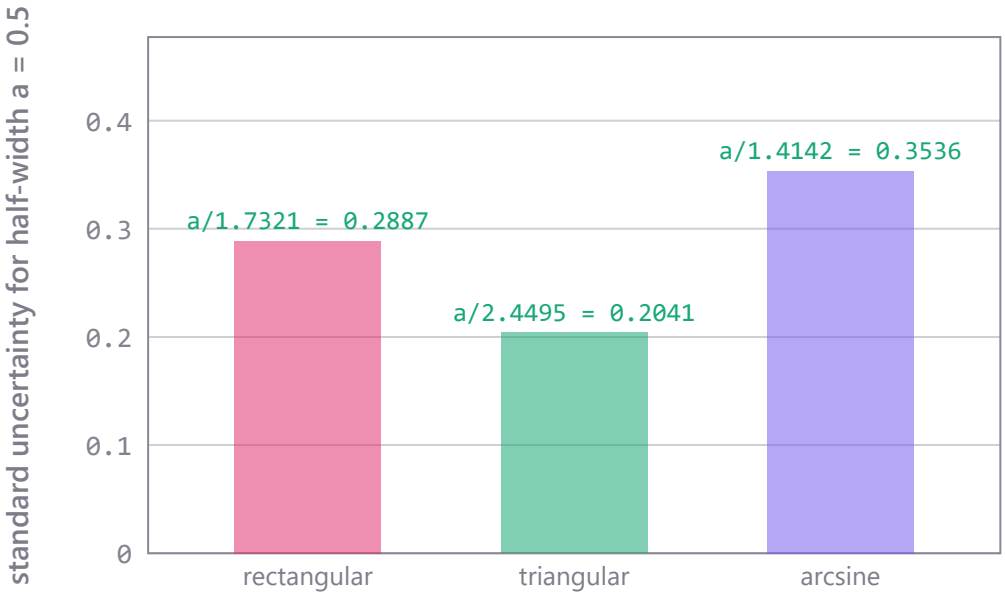
Standard uncertainty of a quantity known only to lie within +/- a, for a rectangular, triangular or U-shaped (arcsine) distribution.

$$u = \frac{a}{\sqrt{3}}\ (\text{rect.}),\ \frac{a}{\sqrt{6}}\ (\text{tri.}),\ \frac{a}{\sqrt{2}}\ (\text{arcsine})$$

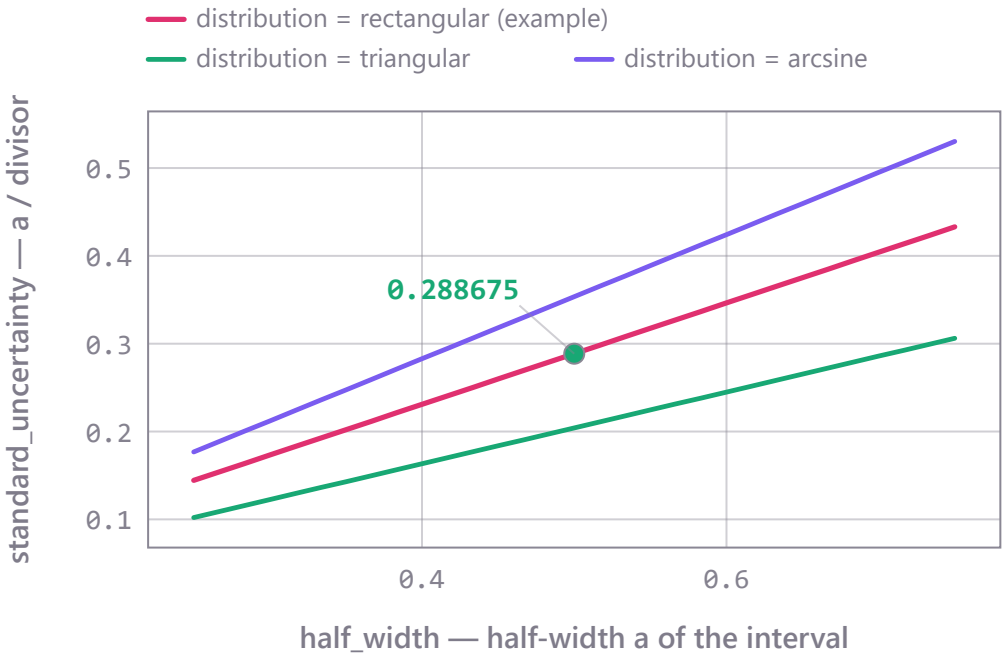
input	meaning	condition
`half_width`	half-width a of the interval (> 0)	
`distribution`	rectangular, triangular or arcsine (U-shaped)	

output	meaning
`standard_uncertainty`	a / divisor
`divisor`	sqrt(3), sqrt(6) or sqrt(2)

- **Method:** standard deviation of the assumed distribution
- **Assumes:** $a > 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The standard uncertainty of a quantity known only to lie within ±0.5 (the worked example's half-width) for each assumed distribution, as returned: rectangular: divisor 1.73205, u = 0.288675; triangular: divisor 2.44949, u = 0.204124; arcsine: divisor 1.41421, u = 0.353553. A distribution that piles up at the edges gives the largest uncertainty.



standard_uncertainty (a / divisor) against half_width (half-width a of the interval (> 0)) from 0.25 to 0.75. One curve for each value of distribution (rectangular, triangular or arcsine (U-shaped)): rectangular, triangular, arcsine. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, half_width = 0.5, standard_uncertainty = 0.288675.

```
bayesian_statistics_ai999x.solve('type_b_standard_uncertainty', half_width=0.5,
distribution='rectangular')
# -> standard_uncertainty = 0.288675, divisor = 1.73205   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('type_b_standard_uncertainty', half_width=0.3,
distribution='triangular')
# -> standard_uncertainty = 0.122474, divisor = 2.44949   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Type B uncertainty: half-width = 0.5 with a rectangular distribution.	SOLVED	standard_uncertainty = 0.288675, divisor = 1.73205
Type B uncertainty: half-width = 2 with a U-shaped distribution.	SOLVED	standard_uncertainty = 1.41421, divisor = 1.41421
Type B uncertainty: half-width = -0.5 with a rectangular distribution.	INVALID_INPUT	Type B standard uncertainty from a stated half-width: the half-width must be positive.

Exact versus first-order propagation through exp(X) — `lognormal_propagation`

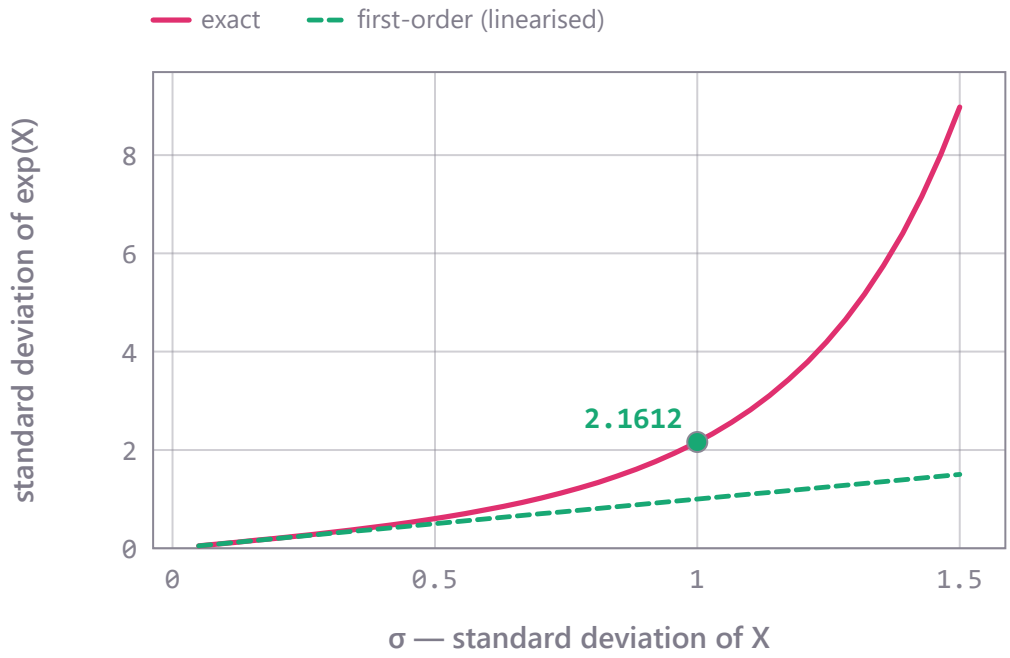
Mean and standard deviation of $Y = exp$(X) for normal X ~ N(mu, sigma^2), exactly (lognormal) and by first-order propagation, showing the bias of the linear approximation.

$$E[e^X] = e^{\mu+\sigma^2/2}, \quad \mathrm{SD}[e^X] = e^{\mu+\sigma^2/2}\sqrt{e^{\sigma^2}-1}$$

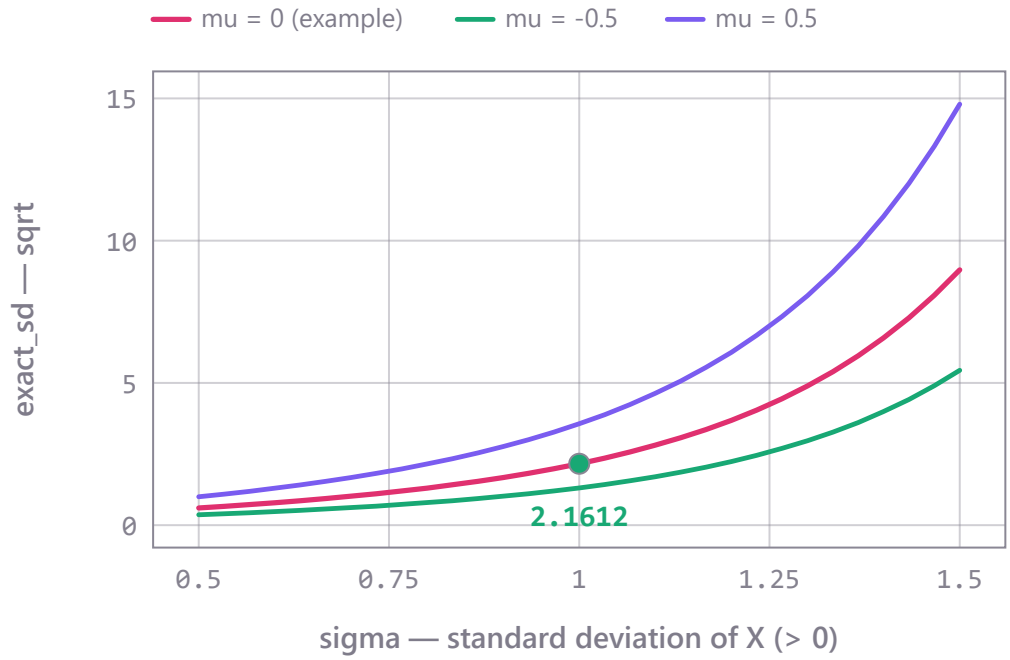
input	meaning	condition
`mu`	mean of X	
`sigma`	standard deviation of X (> 0)	

output	meaning
`exact_mean`	exp(mu + sigma^2/2)
`exact_sd`	sqrt(exp(sigma^2) - 1) exp(mu + sigma^2/2)
`first_order_mean`	exp(mu)
`first_order_sd`	exp(mu) sigma

- **Method:** lognormal moments; first-order Taylor propagation
- **Assumes:** $\sigma > 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The standard deviation of Y = exp(X) for normal X with mean μ = 0 as σ runs from 0.05 to 1.5: exact and first-order (linearised) values, each from the package; they agree for small σ and part as it grows. Marked: the worked example σ = 1: exact 2.1612, first order 1 (means 1.64872 and 1).



exact_sd (sqrt(exp(sigma^2) - 1) exp(mu + sigma^2/2)) against sigma (standard deviation of X (> 0)) from 0.5 to 1.5. One curve for each value of mu (mean of X): 0, -0.5, 0.5. 31 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, sigma = 1, exact_sd = 2.1612.

```
bayesian_statistics_ai999x.solve('lognormal_propagation', mu=0, sigma=1)
# -> exact_mean = 1.64872, exact_sd = 2.1612, first_order_mean = 1, first_order_sd = 1   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('lognormal_propagation', mu=1, sigma=0.5)
# -> exact_mean = 3.08022, exact_sd = 1.64157, first_order_mean = 2.71828, first_order_sd = 1.35914
(SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Lognormal propagation with mu = 0 and sigma = 1.	SOLVED	exact_mean = 1.64872, exact_sd = 2.1612, first_order_mean = 1, first_order_sd = 1
Lognormal propagation with mu = 0 and sigma = -1.	INVALID_INPUT	Exact versus first-order propagation through exp(X): sigma must be positive.

Exact mean and standard deviation of a product of independent normals — `product_of_normals_exact_moments`

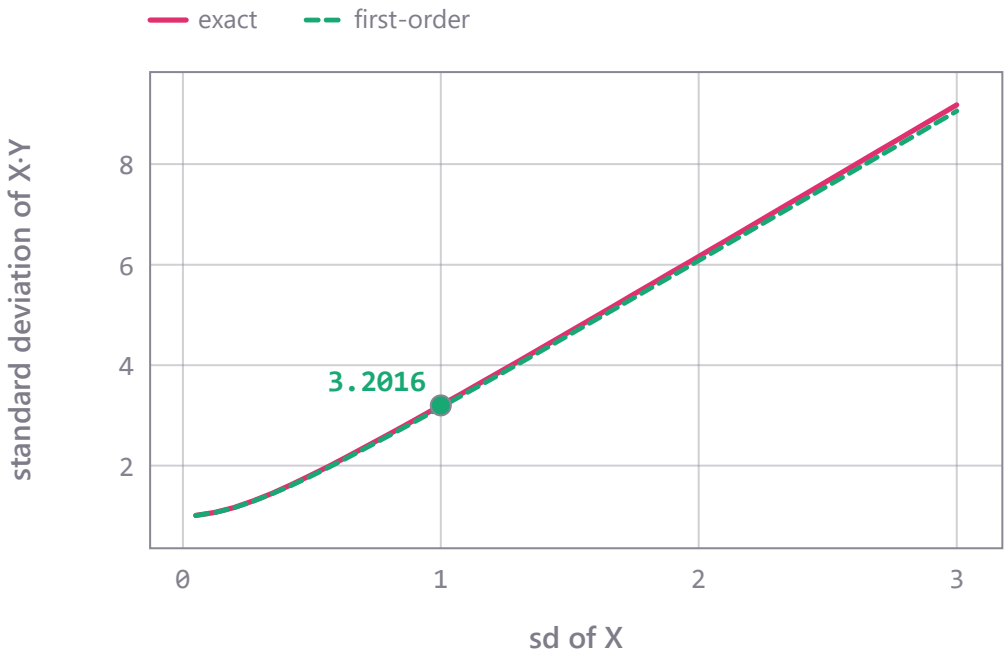
Exact moments of XY for independent normal X and Y, compared with first-order propagation (which misses the sd_x sd_y term).

$$\mathrm{Var}(XY) = \mu_x^2\sigma_y^2 + \mu_y^2\sigma_x^2 + \sigma_x^2\sigma_y^2$$

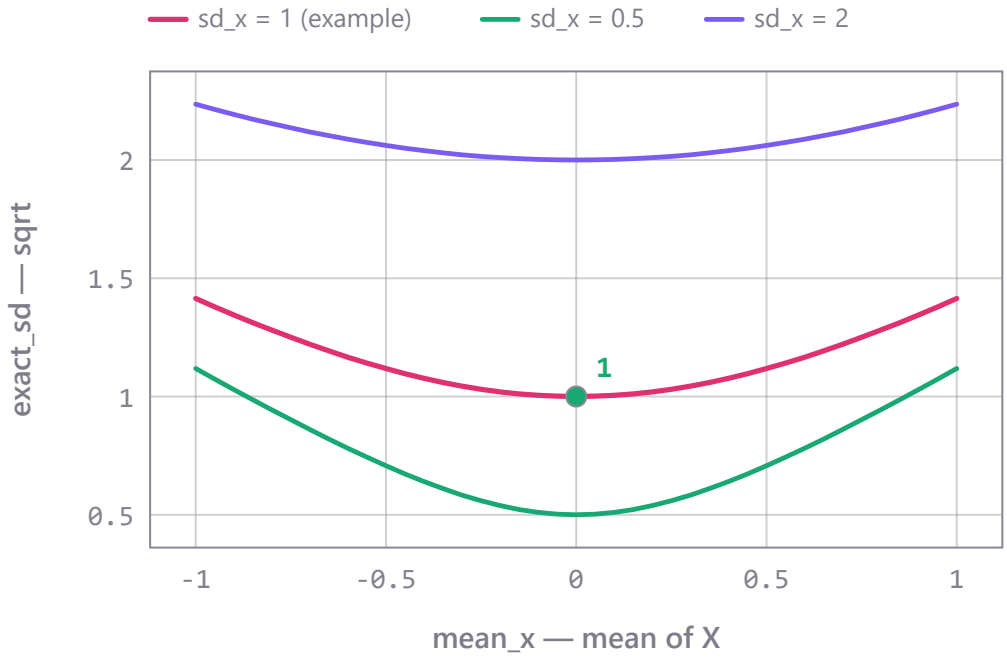
input	meaning	condition
`mean_x`	mean of X	
`sd_x`	sd of X (> 0)	
`mean_y`	mean of Y	
`sd_y`	sd of Y (> 0)	

output	meaning
`exact_mean`	mean_x mean_y
`exact_sd`	sqrt($mx^2\ sy^2 + my^2\ sx^2 + sx^2$ sy^2)
`first_order_sd`	sqrt($mx^2\ sy^2 + my^2$ sx^2)

- **Method:** exact second moments of a product
- **Assumes:** independent normal inputs
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The standard deviation of the product of independent normals X (mean 2) and Y (mean 3, sd 0.5) as the sd of X runs from 0.05 to 3 (inputs chosen; the worked example has zero means, where the first-order value vanishes): exact and first-order values, each from the package. Marked: sd of X = 1, exact 3.20156, first order 3.16228.



exact_sd (sqrt(mx^2 sy^2 + my^2 sx^2 + sx^2 sy^2)) against mean_x (mean of X) from -1 to 1. One curve for each value of sd_x (sd of X (> 0)): 1, 0.5, 2. Other inputs as in the worked example: mean_y = 0, sd_y = 1. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, mean_x = 0, exact_sd = 1.

```
bayesian_statistics_ai999x.solve('product_of_normals_exact_moments', mean_x=0, sd_x=1, mean_y=0,
sd_y=1)
# -> exact_mean = 0, exact_sd = 1, first_order_sd = 0   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('product_of_normals_exact_moments', mean_x=10, sd_x=1, mean_y=5,
sd_y=2)
# -> exact_mean = 50, exact_sd = 20.7123, first_order_sd = 20.6155   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Exact moments of the product of independent normals: mean_x = 10, sd_x = 1, mean_y = 5, sd_y = 2.	SOLVED	exact_mean = 50, exact_sd = 20.7123, first_order_sd = 20.6155
Exact moments of the product of independent normals: mean_x = 10, sd_x = 0, mean_y = 5, sd_y = 2.	INVALID_INPUT	Exact mean and standard deviation of a product of independent normals: sd_x must be positive.

Number of Monte Carlo draws for a target accuracy — `monte_carlo_sample_size`

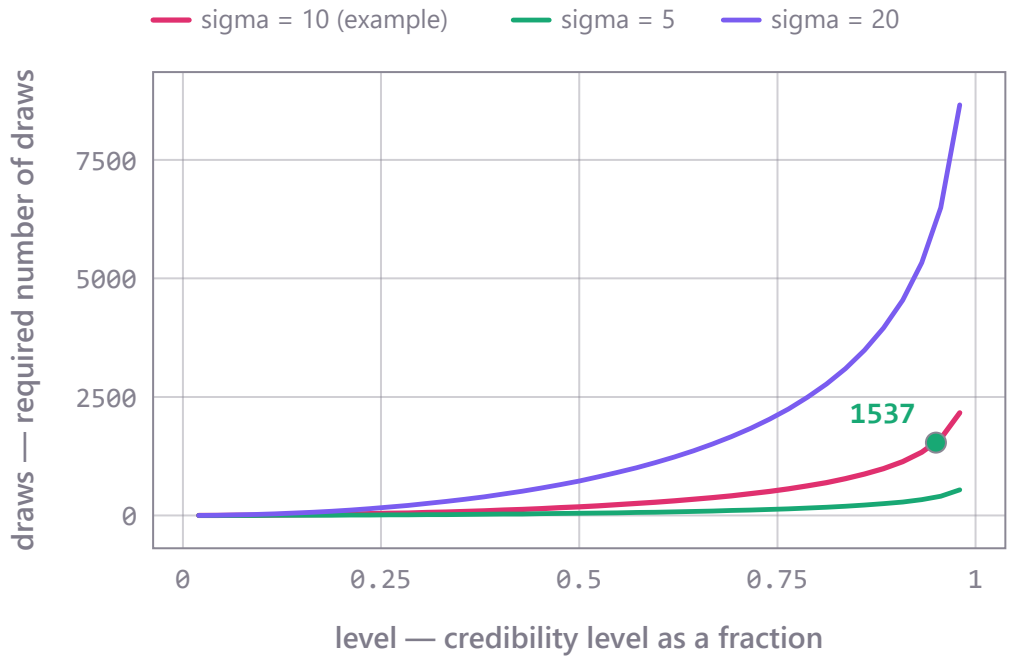
How many independent draws make the half-width of the confidence interval of a Monte Carlo mean at most the stated value, given the standard deviation of one draw.

$$n = \left\lceil \left(\frac{z_{(1+c)/2}\,\sigma}{E} \right)^2 \right\rceil$$

input	meaning	condition
`sigma`	standard deviation of one draw (> 0)	
`half_width`	target half-width (> 0)	
`level`	credibility level as a fraction (0.95 for 95%; default 0.95)	

output	meaning
`draws`	required number of draws (rounded up)
`z`	normal quantile z_{$(1+c)/2$}
`standard_error_target`	target standard error half_width/z

- **Method:** n = (z sigma / E)^2, rounded up
- **Assumes:** sigma, half-width > 0; 0 < level < 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



draws (required number of draws (rounded up)) against level (credibility level as a fraction (0.95 for 95%; default 0.95)) from 0.02 to 0.98. One curve for each value of sigma (standard deviation of one draw (> 0)): 10, 5, 20. Other inputs as in the worked example: half_width = 0.5. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, level = 0.95, draws = 1537.

```
bayesian_statistics_ai999x.solve('monte_carlo_sample_size', sigma=10, half_width=0.5, level=0.95)
# -> draws = 1537, z = 1.95996, standard_error_target = 0.255107   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('monte_carlo_sample_size', sigma=1, half_width=0.1, level=0.9)
# -> draws = 271, z = 1.64485, standard_error_target = 0.0607957   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Monte Carlo sample size: standard deviation 10, half-width 0.5, 95% confidence. How many draws?	SOLVED	draws = 1537, z = 1.95996, standard_error_target = 0.255107
Monte Carlo sample size: standard deviation 10, half-width 0, 95% confidence.	INVALID_INPUT	Number of Monte Carlo draws for a target accuracy: the half-width must be positive.

Monte Carlo estimate of an integral from uniform samples — `monte_carlo_integration`

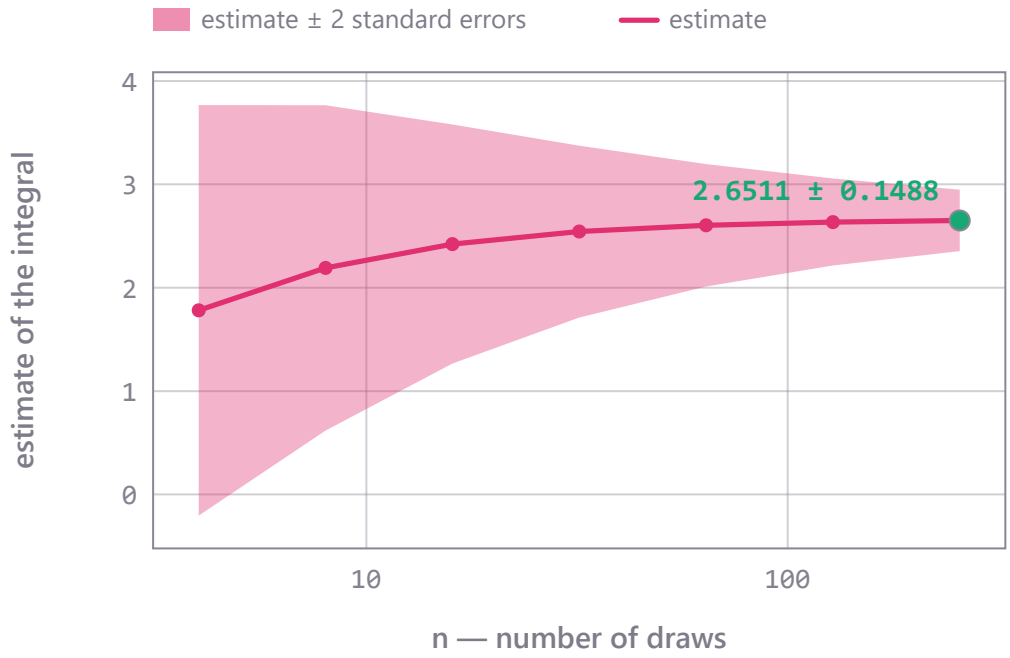
Estimate of the integral of f over [a, b] from f evaluated at n points drawn uniformly on [a, b], with its Monte Carlo standard error.

$$\hat{I} = \frac{b-a}{n}\sum_i f(x_i), \quad SE = (b-a)\frac{s_f}{\sqrt{n}}$$

input	meaning	condition
`values`	f(x_i) at uniform sample points	
`lower`	lower limit a	
`upper`	upper limit b (> a)	

output	meaning
`estimate`	(b - a) times the mean of f
`standard_error`	(b - a) s / sqrt(n)
`n`	number of samples

- **Method:** sample mean of f scaled by the interval length
- **Assumes:** $b > a$; $n \geq 2$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Monte Carlo estimates of the integral of f(x) = x^2 over [0, 2] (inputs chosen: the values f at x = 2·h_i, where h_i are the package's first 256 one-dimensional Halton points) from the first n = 4, 8, 16, 32, 64, 128, 256 values, each estimate and standard error returned by the package; the band narrows like 1/√n. Marked: n = 256, 2.65106 ± 0.148815.

```
bayesian_statistics_ai999x.solve('monte_carlo_integration', values='1, 3', lower=0, upper=2)
# -> estimate = 4, standard_error = 2, n = 2   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('monte_carlo_integration', values=[3, 5, 4, 6], lower=1, upper=4)
# -> estimate = 13.5, standard_error = 1.93649, n = 4   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Monte Carlo integration with values = [1, 3]: lower = 0, upper = 2.	SOLVED	estimate = 4, standard_error = 2, n = 2
Monte Carlo integration with values = [1, 3]: lower = 2, upper = 0.	INVALID_INPUT	Monte Carlo estimate of an integral from uniform samples: the upper limit must exceed the lower limit.

Importance sampling estimate and effective sample size — `importance_sampling_estimate`

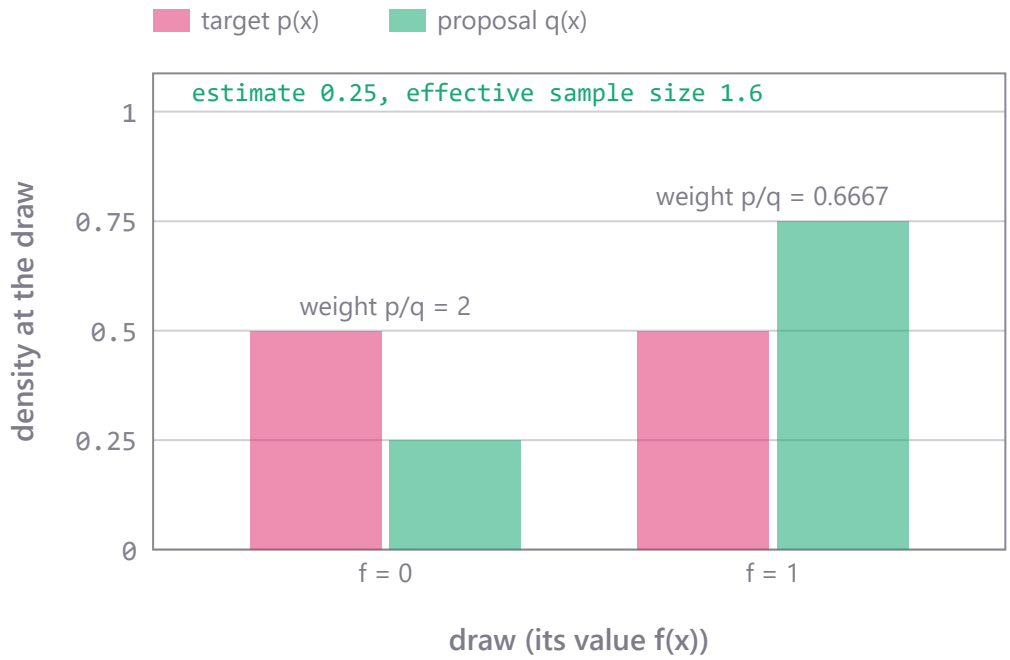
Self-normalized and plain importance-sampling estimates of E_p[f] from draws of a proposal q, with the effective sample size of the weights p/q.

$$\hat{\mu} = \frac{\sum_i w_i f(x_i)}{\sum_i w_i},\ w_i = \frac{p(x_i)}{q(x_i)}, \quad ESS = \frac{(\sum w_i)^2}{\sum w_i^2}$$

input	meaning	condition
`values`	f(x_i) at the draws	
`target_density`	target density p(x_i) (or unnormalized)	
`proposal_density`	proposal density q(x_i) > 0	

output	meaning
`estimate`	self-normalized estimate sum w f / sum w
`unnormalized_estimate`	mean of w f (needs normalized densities)
`effective_sample_size`	(sum w)^2 / sum w^2

- **Method:** weights $w = p/q$
- **Assumes:** $q > 0$; $p \geq 0$, not all zero
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Importance sampling with draws whose values are 0, 1, target densities 0.5, 0.5 and proposal densities 0.25, 0.75 (the worked example): the importance weights p/q (derived, written over the bars) reweight the draws; returned self-normalised estimate 0.25, unnormalised 0.333333, effective sample size 1.6.

```
bayesian_statistics_ai999x.solve('importance_sampling_estimate', values='0, 1',
target_density='0.5, 0.5', proposal_density='0.25, 0.75')
# -> estimate = 0.25, unnormalized_estimate = 0.333333, effective_sample_size = 1.6   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('importance_sampling_estimate', values=[1, 2, 3], target_density=
[0.2, 0.5, 0.3], proposal_density=[0.4, 0.4, 0.2])
# -> estimate = 2.30769, unnormalized_estimate = 2.5, effective_sample_size = 2.6   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Importance sampling estimate with f values [0, 1], target density [0.5, 0.5] and proposal density [0.25, 0.75].	SOLVED	estimate = 0.25, unnormalized_estimate = 0.333333, effective_sample_size = 1.6
Importance sampling: values = [1, 2], target densities = [0.2, 0.5], proposal densities = [0.4, 0].	INVALID_INPUT	Importance sampling estimate and effective sample size: proposal densities must be positive and target densities non-neg

Control-variate Monte Carlo estimate — `control_variate_estimate`

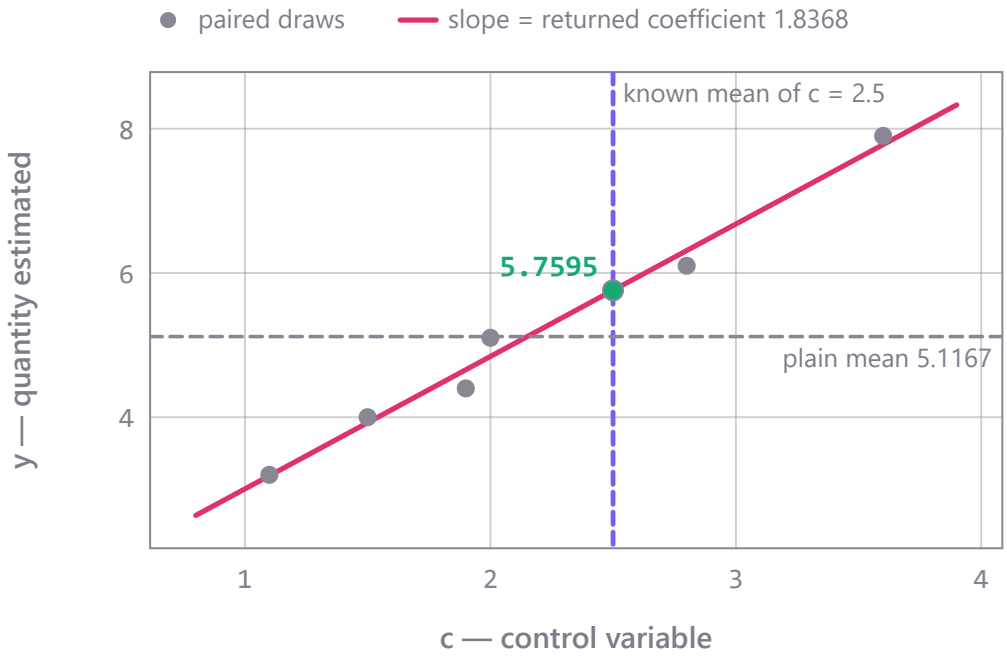
Variance-reduced estimate of E[Y] from paired draws (y_i, c_i) of the output and a control variate with known mean: optimal coefficient, estimate and the fraction of variance left $(1 - r^2)$.

$$\hat{\mu}_{cv} = \bar{y} - \hat{b}\,(\bar{c} - \mu_c), \quad \hat{b} = \frac{\sum (y_i - \bar{y})(c_i - \bar{c})}{\sum (c_i - \bar{c})^2}$$

input	meaning	condition
`y`	simulated outputs y_i	
`c`	control values c_i (same draws)	
`control_mean`	known mean of the control	

output	meaning
`estimate`	mean(y) - b (mean(c) - mu_c)
`coefficient`	b = cov(y, c)/var(c)
`plain_estimate`	mean(y)
`variance_reduction`	variance factor $1 - r^2$

- **Method:** regression (optimal) control-variate coefficient
- **Assumes:** outputs and control vary; $n \geq 3$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



A control-variate estimate for paired draws y = 3.2, 5.1, 4.4, 7.9, 6.1, 4.0 and c = 1.1, 2.0, 1.9, 3.6, 2.8, 1.5 with the known mean of c = 2.5 (inputs chosen; the worked example's y and c are identical): the line through the sample means ($\bar{c}$ = 2.15, $\bar{y}$ = 5.1167, derived) with the returned slope 1.83676 is read at the known mean, giving the returned estimate 5.75953 instead of the plain mean 5.11667; returned variance reduction 0.0140049.

```
bayesian_statistics_ai999x.solve('control_variate_estimate', y='1, 2, 3', c='1, 2, 3',
control_mean=2.5)
# -> estimate = 2.5, coefficient = 1, plain_estimate = 2, variance_reduction = 0   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('control_variate_estimate', y=[2.1, 3.9, 6.2, 8.1], c=[1, 2, 3,
4], control_mean=2.5)
# -> estimate = 5.075, coefficient = 2.03, plain_estimate = 5.075, variance_reduction = 0.00208258
(SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Control variate: outputs = [1, 2, 3], controls = [1, 2, 3], known mean = 2.5.	SOLVED	estimate = 2.5, coefficient = 1, plain_estimate = 2, variance_reduction = 0
Control variate: outputs = [1, 2, 3], controls = [1, 1, 1], known mean = 2.5.	INVALID_INPUT	Control-variate Monte Carlo estimate: the outputs and the control must both vary.

Stratified sampling estimate and standard error — `stratified_sampling_estimate`

Estimate of a mean from strata with known weights, the sample mean, standard deviation and size of each stratum, and its standard error.

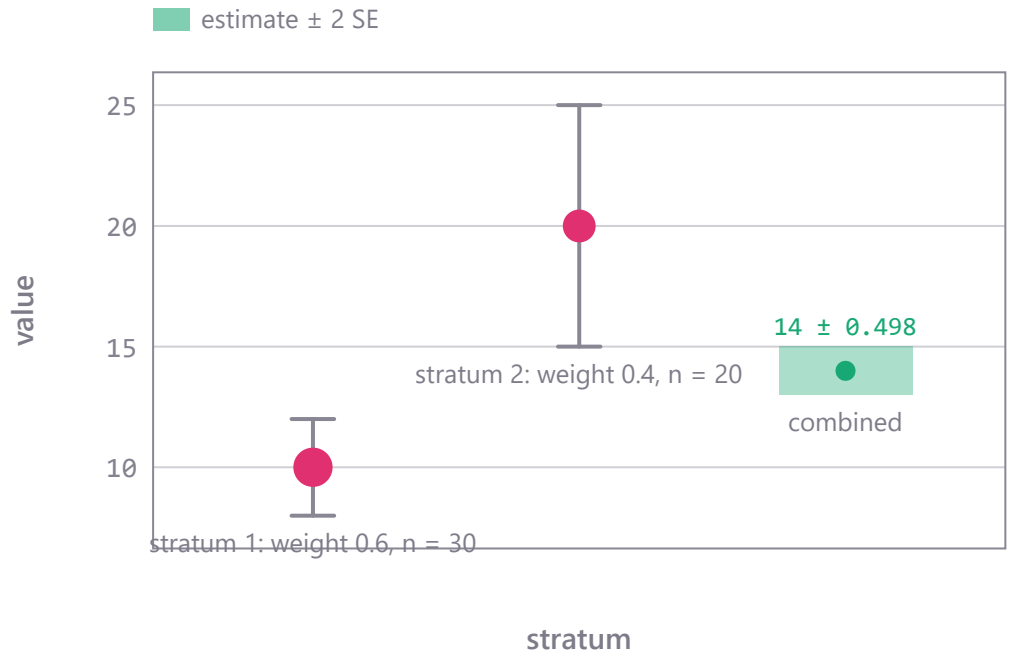
$$\hat{\mu} = \sum_h W_h \bar{y}_h, \quad SE = \sqrt{\sum_h W_h^2 s_h^2 / n_h}$$

input	meaning	condition
`weights`	stratum weights (sum to 1)	

input	meaning	condition
`means`	stratum sample means	
`sds`	stratum sample standard deviations	
`sizes`	stratum sample sizes	

output	meaning
`estimate`	sum W_h mean_h
`standard_error`	sqrt(sum W_h^2 s_h^2 / n_h)

- **Method:** weighted stratum means
- **Assumes:** weights sum to 1; sizes >= 1
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Stratified sampling with weights 0.6, 0.4, stratum means 10, 20, standard deviations 2, 5 and sample sizes 30, 20 (the worked example): each stratum's mean ± sd (dot size by weight), and the returned overall estimate 14 with standard error 0.497996 (band ± 2 SE).

```
bayesian_statistics_ai999x.solve('stratified_sampling_estimate', weights='0.6, 0.4', means='10, 20', sds='2, 5', sizes='30, 20')
# -> estimate = 14, standard_error = 0.497996   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('stratified_sampling_estimate', weights=[0.5, 0.5], means=[1, 3], sds=[1, 1], sizes=[4, 4])
# -> estimate = 2, standard_error = 0.353553   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Stratified sampling: weights = [0.6, 0.4], means = [10, 20], sds = [2, 5], sizes = [30, 20].	SOLVED	estimate = 14, standard_error = 0.497996

question	outcome	result
Stratified sampling: weights = [0.6, 0.6], means = [10, 20], sds = [2, 5], sizes = [30, 20].	INVALID_INPUT	Stratified sampling estimate and standard error: the stratum weights must sum to 1 (they sum to 1.2).

Point of the Halton low-discrepancy sequence — `halton_point`

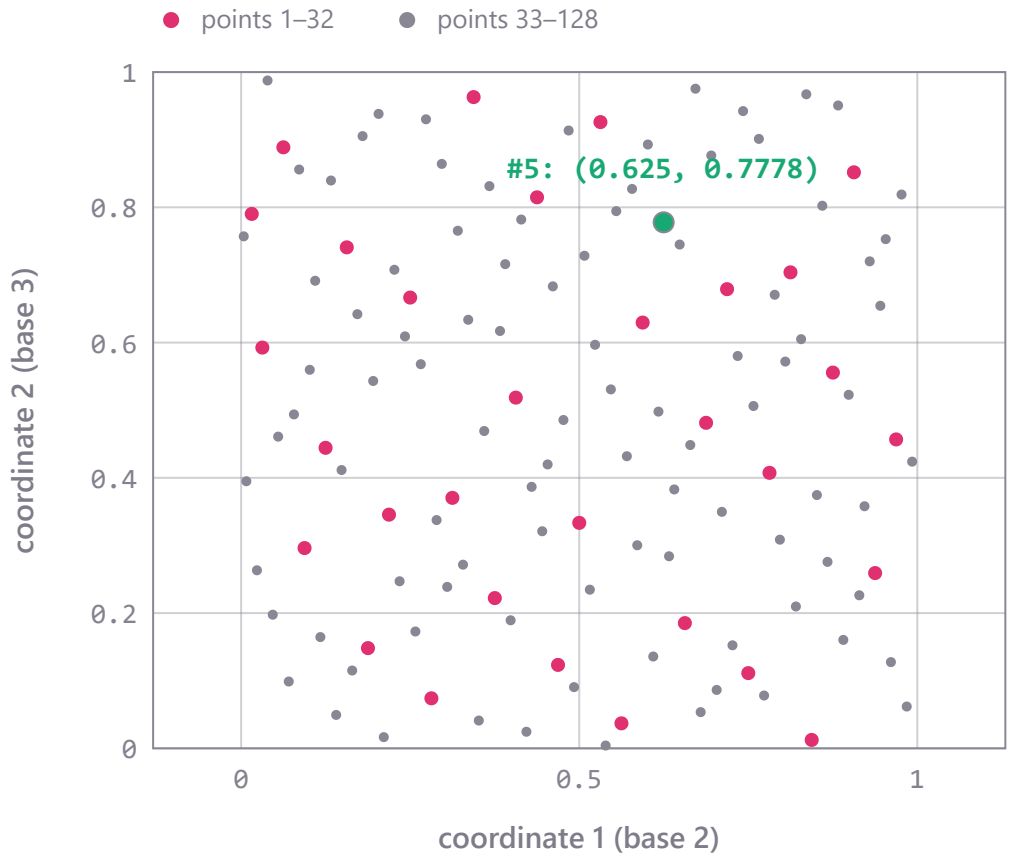
Coordinates of the point with a given index in the (unscrambled) Halton sequence of a given dimension: the radical inverses of the index in the first prime bases (index 0 is the origin).

$$\phi_b(i) = \sum_k d_k \, b^{-k-1} \;\text{ for } i = \sum_k d_k b^k$$

input	meaning	condition
`index`	index of the point (0, 1, 2, ...)	
`dimension`	dimension (1 to 20)	

output	meaning
`point`	coordinates in [0, 1)
`bases`	prime bases used

- **Method:** radical inverse: reverse the base-b digits of the index behind the radix point
- **Assumes:** index >= 0; 1 <= dimension <= 20
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The first 128 points of the two-dimensional Halton sequence (bases 2 and 3), each returned by the package: they fill the unit square evenly, without the clumps and gaps of random points. Marked: index 5 (the worked example's index), point (0.625, 0.777778).

```
bayesian_statistics_ai999x.solve('halton_point', index=5, dimension=3)
# -> point = [0.625, 0.777777777778, 0.04], bases = [2.0, 3.0, 5.0]   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('halton_point', index=1, dimension=2)
# -> point = [0.5, 0.333333333333], bases = [2.0, 3.0]   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Halton sequence point with index = 5 and dimension = 3.	SOLVED	point = [0.625, 0.777777777778, 0.04], bases = [2.0, 3.0, 5.0]
Halton sequence point with index = 5 and dimension = 30.	INVALID_INPUT	Point of the Halton low-discrepancy sequence: dimension must be at most 20.

Sobol sensitivity indices of the Ishigami function — `sobol_indices_ishigami`

Exact first-order and total Sobol indices and the variance of the Ishigami test function $f = sin\, x_1 + a\, sin^2\, x_2 + b\, x_3^4$ sin x1 with x_i uniform on [-pi, pi].

$$V = \frac{a^2}{8} + \frac{b\pi^4}{5} + \frac{b^2\pi^8}{18} + \frac{1}{2},\ V_1 = \frac{1}{2}\left(1 + \frac{b\pi^4}{5}\right)^2,\ V_2 = \frac{a^2}{8},\ V_{13} = \frac{8b^2\pi^8}{225}$$

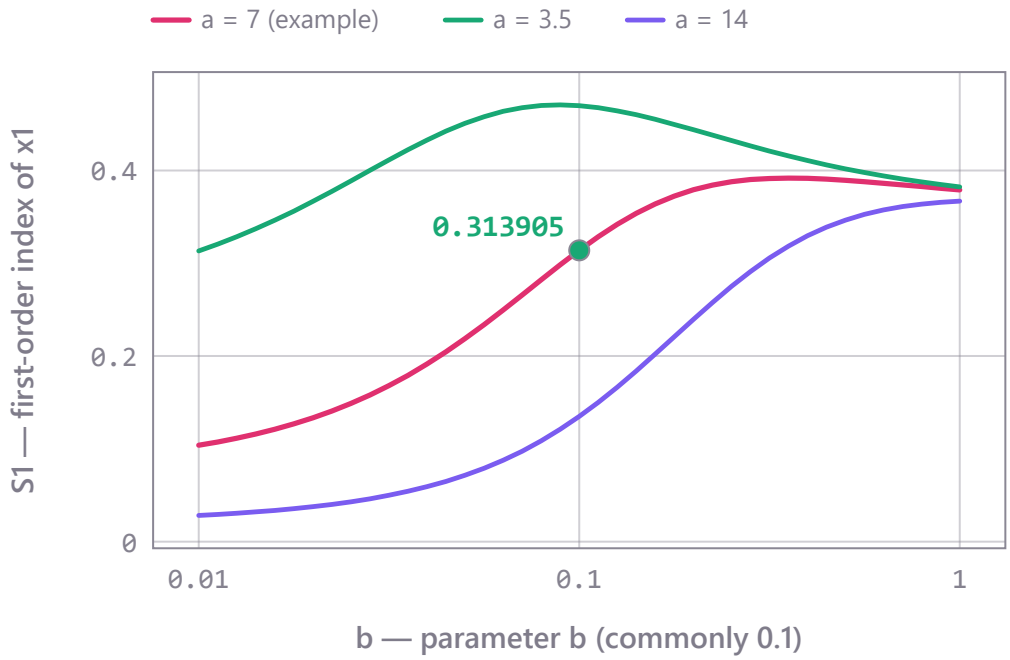
input	meaning	condition
`a`	parameter a (commonly 7)	
`b`	parameter b (commonly 0.1)	

output	meaning
`S1`	first-order index of x1
`S2`	first-order index of x2
`S3`	first-order index of x3
`ST1`	total index of x1
`ST2`	total index of x2
`ST3`	total index of x3
`variance`	variance of f

- **Method:** closed-form variance decomposition
- **Assumes:** inputs independent uniform on [-pi, pi]
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



First- and total-order Sobol' indices of the Ishigami function with a = 7, b = 0.1 (the worked example), as returned: S = 0.3139, 0.4424, 0, ST = 0.5576, 0.4424, 0.2437, variance 13.8446. x3 has no effect alone (S3 = 0) but acts through its interaction with x1.



S1 (first-order index of x1) against b (parameter b (commonly 0.1)) from 0.01 to 1 on a logarithmic axis. One curve for each value of a (parameter a (commonly 7)): 7, 3.5, 14. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, b = 0.1, S1 = 0.313905.

```
bayesian_statistics_ai999x.solve('sobol_indices_ishigami', a=7, b=0.1)
# -> S1 = 0.313905, S2 = 0.442411, S3 = 0, ST1 = 0.557589, ST2 = 0.442411, ST3 = 0.243684, variance
= 13.8446   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('sobol_indices_ishigami', a=0, b=0)
# -> S1 = 1, S2 = 0, S3 = 0, ST1 = 1, ST2 = 0, ST3 = 0, variance = 0.5   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Sobol indices of the Ishigami function with a = 7 and b = 0.1.	SOLVED	S1 = 0.313905, S2 = 0.442411, S3 = 0, ST1 = 0.557589, ST2 = 0.442411, ST3 = 0.243684, variance = 13.8446
Sobol indices of the Ishigami function with a = 7.	NEEDS_INPUT	sobol_indices_ishigami needs b (parameter b (commonly 0.1)). Example: solve('sobol_indices_ishigami', a=7, b=0.1).

Sobol indices of Sobol's G-function — `sobol_indices_g_function`

Exact first-order and total Sobol indices and the variance of the G-function prod_i (|4 $x_i - 2$| + a_i)/(1 + a_i), x_i uniform on [0, 1].

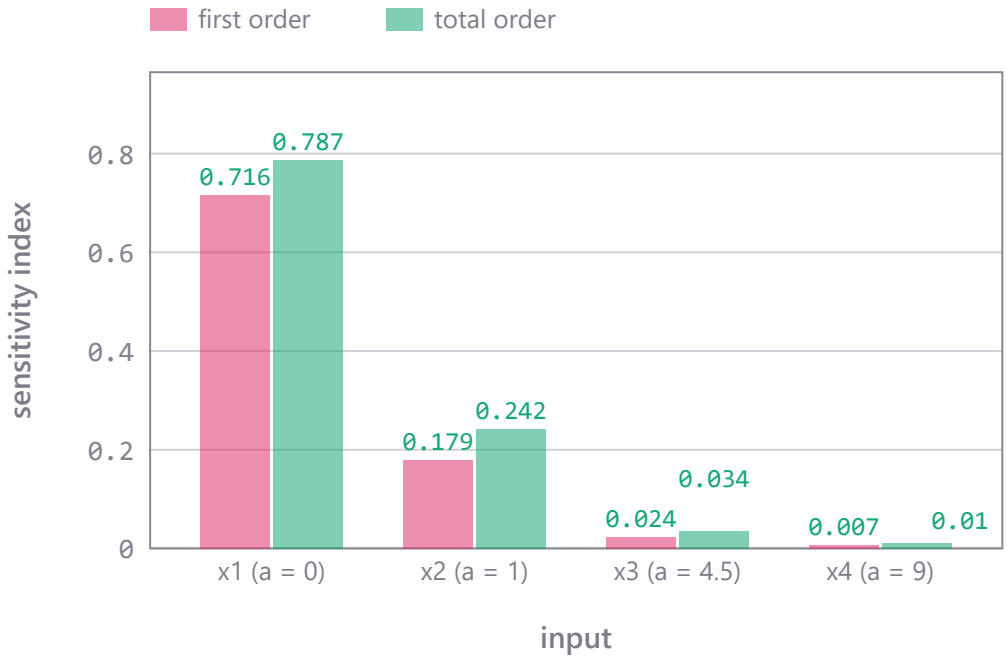
$$V_i = \frac{1}{3(1+a_i)^2},\ V = \prod_i (1+V_i) - 1,\ S_i = \frac{V_i}{V},\ S_{Ti} = \frac{V_i \prod_{j \neq i}(1+V_j)}{V}$$

input	meaning	condition
`a`	parameters $a_i \geq 0$ (small a_i = important input)	

output	meaning
`first_order`	first-order indices S_i

output	meaning
`total_order`	total indices ST_i
`variance`	variance of the function

- **Method:** partial variances V_i = 1/(3(1 + a_i)^2) of the product form
- **Assumes:** $a_i \geq 0$; independent uniform inputs
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Sobol' indices of the g-function with importance parameters a = 0, 1, 4.5, 9 (chosen; the worked example has a single input, which then carries everything): returned first-order indices 0.7165, 0.1791, 0.0237, 0.0072 and total-order indices 0.7874, 0.2423, 0.0343, 0.0105, output variance 0.465229. A larger a makes an input less important; total exceeds first order by the interactions.

```
bayesian_statistics_ai999x.solve('sobol_indices_g_function', a='0')
# -> first_order = [1.0], total_order = [1.0], variance = 0.333333   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('sobol_indices_g_function', a=[0, 0])
# -> first_order = [0.428571428571, 0.428571428571], total_order = [0.571428571429,
0.571428571429], variance = 0.777778   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Sobol indices of the G-function with a = [0, 0].	SOLVED	first_order = [0.428571428571, 0.428571428571], total_order = [0.571428571429, 0.571428571429], variance = 0.777778
Sobol indices of the G-function with a = [0, -1].	INVALID_INPUT	Sobol indices of Sobol's G-function: the G-function parameters must be >= 0.

Sobol indices of a linear (additive) model — `sobol_indices_linear_model`

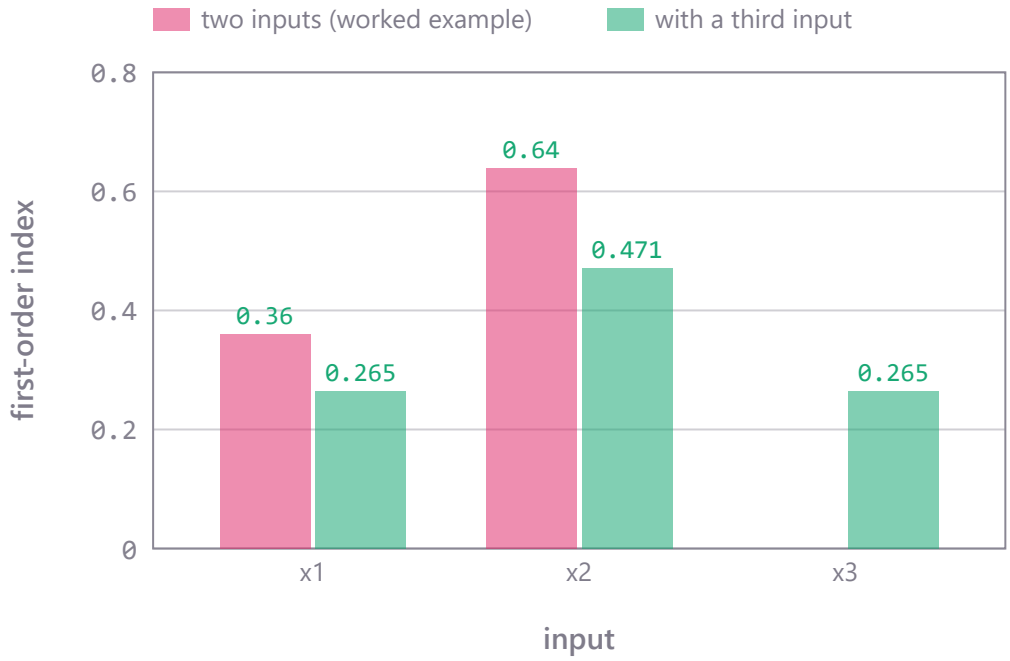
First-order Sobol indices of y = sum a_i x_i with independent inputs of given standard deviations (S_i = a_i^2 sigma_i^2 / Var y; they sum to 1).

$$S_i = \frac{a_i^2 \sigma_i^2}{\sum_j a_j^2 \sigma_j^2}$$

input	meaning	condition
`coefficients`	model coefficients a_i	
`sds`	input standard deviations sigma_i (> 0)	

output	meaning
`first_order`	first-order (= total) indices
`variance`	Var y

- **Method:** variance shares of the terms
- **Assumes:** independent inputs; $sigma_i > 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



First-order Sobol' indices of a linear model y = Σ a_i x_i with independent inputs: the worked example (coefficients 1, 1, sds 3, 4) returns 0.36, 0.64 (variance 25); adding a third input with coefficient 2 and sd 1.5 (chosen) returns 0.2647, 0.4706, 0.2647 (variance 34). Each index is that input's share of the output variance.

```
bayesian_statistics_ai999x.solve('sobol_indices_linear_model', coefficients='1, 1', sds='3, 4')
# -> first_order = [0.36, 0.64], variance = 25   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('sobol_indices_linear_model', coefficients=[2, 1, 0.5], sds=[1, 2, 4])
# -> first_order = [0.333333333333, 0.333333333333, 0.333333333333], variance = 12   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Sobol indices of a linear model with coefficients = [2, 1, 0.5] and sds = [1, 2, 4].	SOLVED	first_order = [0.333333333333, 0.333333333333, 0.333333333333], variance = 12
Sobol indices of a linear model with coefficients = [2, 1] and sds = [1, 0].	INVALID_INPUT	Sobol indices of a linear (additive) model: input standard deviations must be positive.

Sobol indices of a model with an interaction term — `sobol_indices_interaction_model`

First-order, interaction and total Sobol indices of $y = a_1\,x_1 + a_2\,x_2 + a12$ x1 x2 with independent normal inputs x1 ~ N(mean1, sd1^2), x2 ~ N(mean2, sd2^2).

$$V_1 = (a_1 + a_{12}\mu_2)^2\sigma_1^2,\; V_2 = (a_2 + a_{12}\mu_1)^2\sigma_2^2,\; V_{12} = a_{12}^2\sigma_1^2\sigma_2^2$$

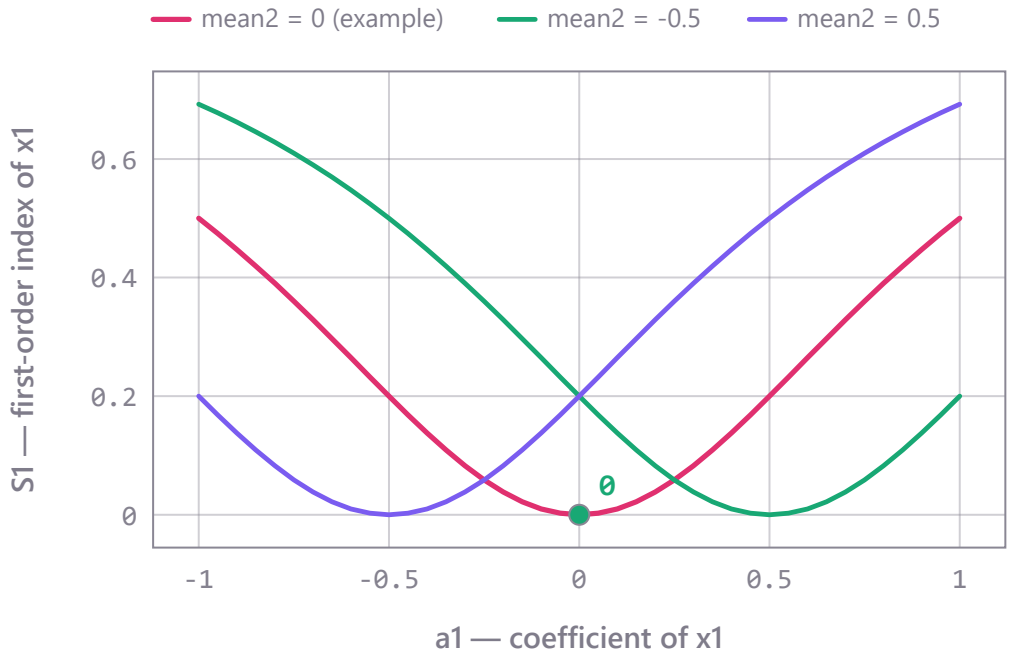
input	meaning	condition
`a1`	coefficient of x1	
`a2`	coefficient of x2	
`a12`	interaction coefficient	
`mean1`	mean of x1	
`sd1`	sd of x1 (> 0)	
`mean2`	mean of x2	
`sd2`	sd of x2 (> 0)	

output	meaning
`S1`	first-order index of x1
`S2`	first-order index of x2
`S12`	interaction index
`ST1`	total index of x1
`ST2`	total index of x2
`variance`	Var y

- **Method:** ANOVA decomposition of the bilinear polynomial
- **Assumes:** independent normal inputs
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Sobol' indices of y = a1 x1 + a2 x2 + a12 x1 x2 with a1 = 1, a2 = 0.5, a12 = 0.8 and standard normal inputs (chosen; the worked example has only the interaction term, which then carries everything): returned S1 = 0.529101, S2 = 0.132275, S12 = 0.338624, ST1 = 0.867725, ST2 = 0.470899, variance 1.89. Each total index adds the interaction share to the first-order one (the interaction bar repeats S12).



S1 (first-order index of x1) against a1 (coefficient of x1) from -1 to 1. One curve for each value of mean2 (mean of x2): 0, -0.5, 0.5. Other inputs as in the worked example: a2 = 0, a12 = 1, mean1 = 0, sd1 = 1, sd2 = 1. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, a1 = 0, S1 = 0.

```
bayesian_statistics_ai999x.solve('sobol_indices_interaction_model', a1=0, a2=0, a12=1, mean1=0,
sd1=1, mean2=0, sd2=1)
# -> S1 = 0, S2 = 0, S12 = 1, ST1 = 1, ST2 = 1, variance = 1   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('sobol_indices_interaction_model', a1=1, a2=2, a12=0.5, mean1=0, sd1=1, mean2=1, sd2=0.5)
# -> S1 = 0.679245, S2 = 0.301887, S12 = 0.0188679, ST1 = 0.698113, ST2 = 0.320755, variance = 3.3125   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Sobol indices of the interaction model with a1 = 0, a2 = 0, a12 = 1, mean1 = 0, sd1 = 1, mean2 = 0, sd2 = 1.	SOLVED	S1 = 0, S2 = 0, S12 = 1, ST1 = 1, ST2 = 1, variance = 1
Sobol indices of the interaction model with a1 = 1, a2 = 2, a12 = 0.5, mean1 = 0, sd1 = 0, mean2 = 1, sd2 = 0.5.	INVALID_INPUT	Sobol indices of a model with an interaction term: sd1 must be positive.

Morris elementary-effects screening — `morris_elementary_effects`

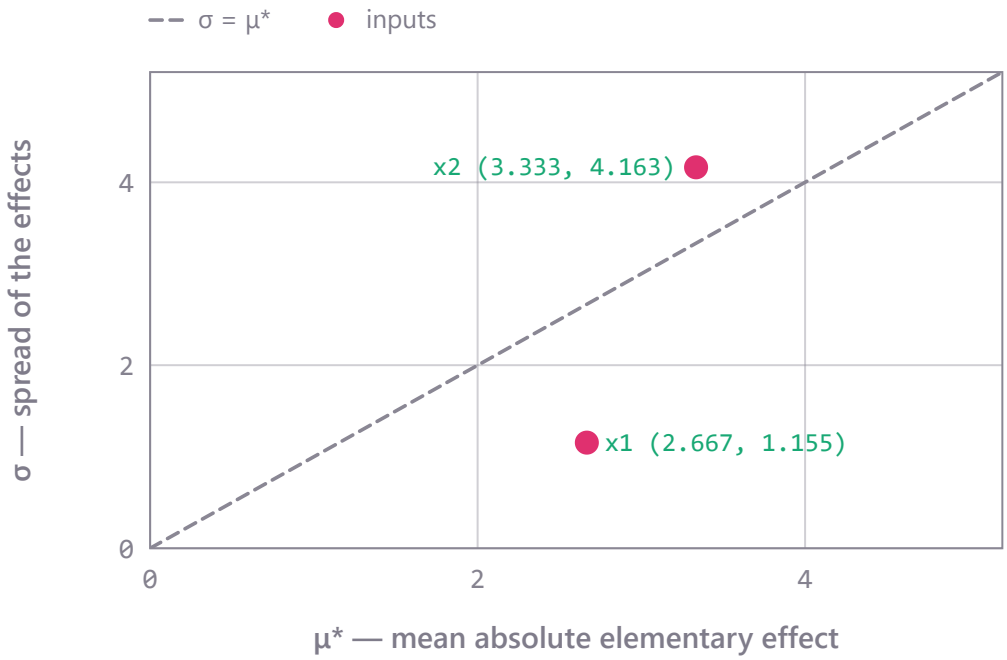
Mean, mean absolute value (mu*) and standard deviation of the elementary effects from one-at-a-time trajectories: in each row, output $j+1$ is the model at the previous point with input j raised by delta.

$$EE_j = \frac{f(x + \Delta e_j) - f(x)}{\Delta}, \quad \mu_j^* = \frac{1}{r}\sum_t |EE_j^{(t)}|$$

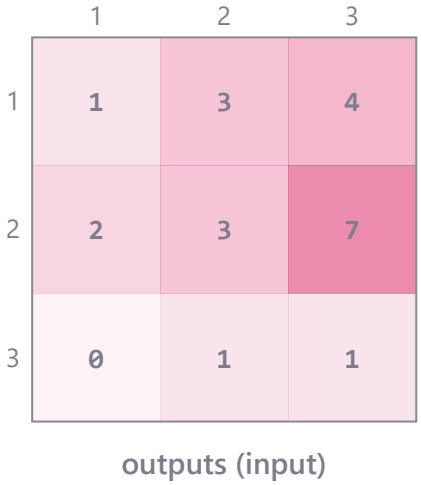
input	meaning	condition
`outputs`	model outputs: one row per trajectory, $k+1$ values (base point, then input 1, 2, ... raised in turn)	
`delta`	step delta (> 0)	

output	meaning
`mu`	mean elementary effect of each input
`mu_star`	mean absolute elementary effect
`sigma`	standard deviation of the elementary effects

- **Method:** elementary effects EE_ij = (y_{i,$j+1$} - y_{ij})/delta
- **Assumes:** at least 2 trajectories; $\delta > 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The Morris screening plane for the elementary-effect outputs [[1, 3, 4], [2, 3, 7], [0, 1, 1]] with step Δ = 0.5 (the worked example): each input at its returned μ* (importance) and σ (spread, a sign of interactions or non-linearity): μ* = 2.6667, 3.3333, σ = 1.1547, 4.1633, μ = 2.6667, 3.3333. Above the dashed line the spread exceeds the mean effect.



The matrices of the worked example as heatmaps (rows and columns numbered; each cell shaded by its magnitude, negative values in the second colour; each panel shaded by its own largest magnitude): outputs (input). Other inputs: delta = 0.5.

```
bayesian_statistics_ai999x.solve('morris_elementary_effects', outputs='[[1, 3, 4], [2, 3, 7], [0, 1, 1]]', delta=0.5)
# -> mu = [2.66666666667, 3.33333333333], mu_star = [2.66666666667, 3.33333333333], sigma = [1.15470053838, 4.16333199893]   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('morris_elementary_effects', outputs=[[10, 12], [11, 10]], delta=1)
# -> mu = [0.5], mu_star = [1.5], sigma = [2.12132034356]   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Elementary effects screening with model outputs [[10, 12], [11, 10]] and step 1.	SOLVED	mu = [0.5], mu_star = [1.5], sigma = [2.12132034356]
Morris elementary effects: outputs = [[1, 3, 4]], delta = 0.5.	INVALID_INPUT	Morris elementary-effects screening: give at least 2 trajectories, each with k + 1 >= 2 outputs.

Kullback-Leibler divergence between two normal distributions — `kl_divergence_normal`

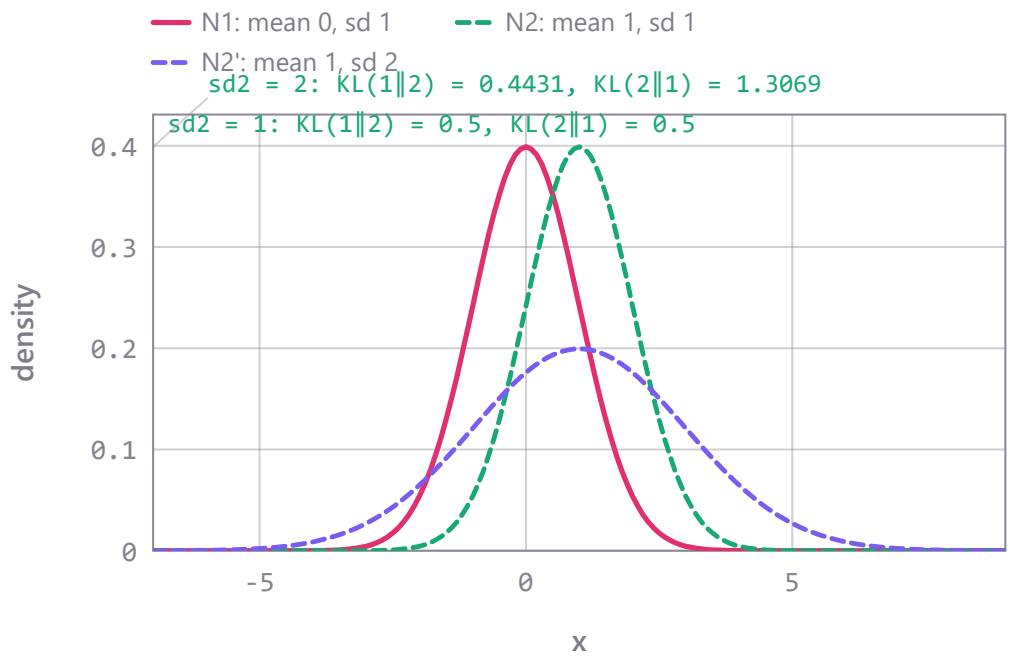
KL(N1 || N2) and KL(N2 || N1), e.g. the information gained from a normal prior to a normal posterior.

$$KL(N_1 \| N_2) = \ln \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2}$$

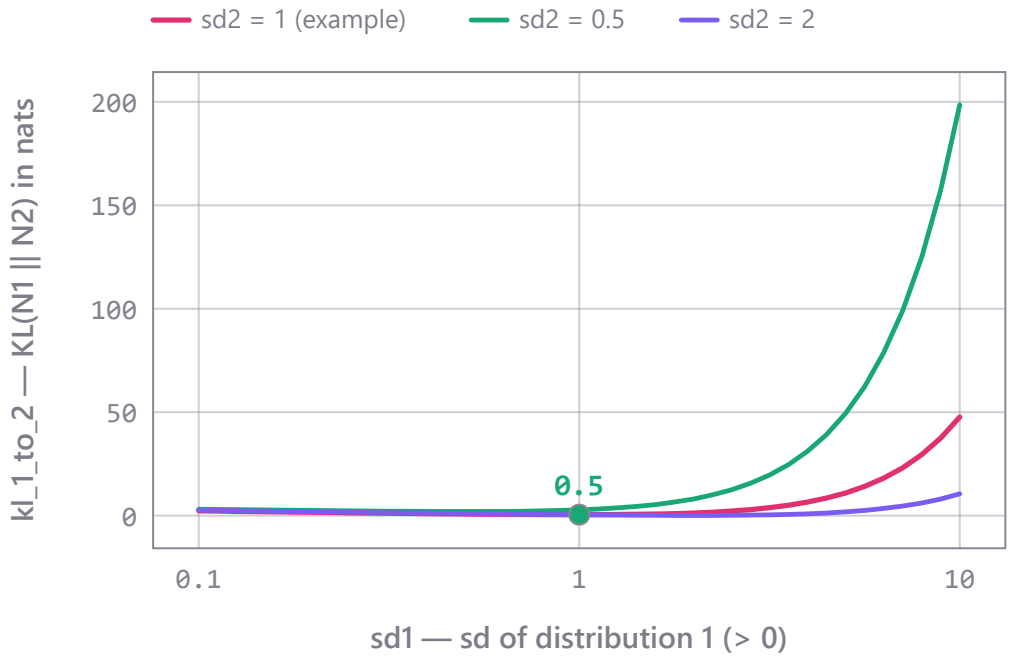
input	meaning	condition
`mean1`	mean of distribution 1	
`sd1`	sd of distribution 1 (> 0)	
`mean2`	mean of distribution 2	
`sd2`	sd of distribution 2 (> 0)	

output	meaning
`kl_1_to_2`	KL(N1 || N2) in nats
`kl_2_to_1`	KL(N2 || N1) in nats

- **Method:** closed form for two normals
- **Assumes:** sd1, $sd2 > 0$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Normal densities (drawn from their stated means and standard deviations): N1 = N(0, 1^2) and N2 = N(1, 1^2) of the worked example, and a wider N2' with sd 2 (chosen). Returned divergences in nats: sd2 = 1: KL(1||2) = 0.5, KL(2||1) = 0.5; sd2 = 2: KL(1||2) = 0.443147, KL(2||1) = 1.30685.



kl_1_to_2 (KL(N1 || N2) in nats) against sd1 (sd of distribution 1 (> 0)) from 0.1 to 10 on a logarithmic axis. One curve for each value of sd2 (sd of distribution 2 (> 0)): 1, 0.5, 2. Other inputs as in the worked example: mean1 = 0, mean2 = 1. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, sd1 = 1, kl_1_to_2 = 0.5.

```
bayesian_statistics_ai999x.solve('kl_divergence_normal', mean1=0, sd1=1, mean2=1, sd2=1)
# -> kl_1_to_2 = 0.5, kl_2_to_1 = 0.5   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('kl_divergence_normal', mean1=0, sd1=1, mean2=1, sd2=2)
# -> kl_1_to_2 = 0.443147, kl_2_to_1 = 1.30685   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
KL divergence between normal distributions: mean1 = 0, sd1 = 1, mean2 = 1, sd2 = 2.	SOLVED	kl_1_to_2 = 0.443147, kl_2_to_1 = 1.30685
KL divergence between normal distributions: mean1 = 0, sd1 = 1, mean2 = 1, sd2 = 0.	INVALID_INPUT	Kullback-Leibler divergence between two normal distributions: sd2 must be positive.

Kullback-Leibler divergence between two beta distributions — `kl_divergence_beta`

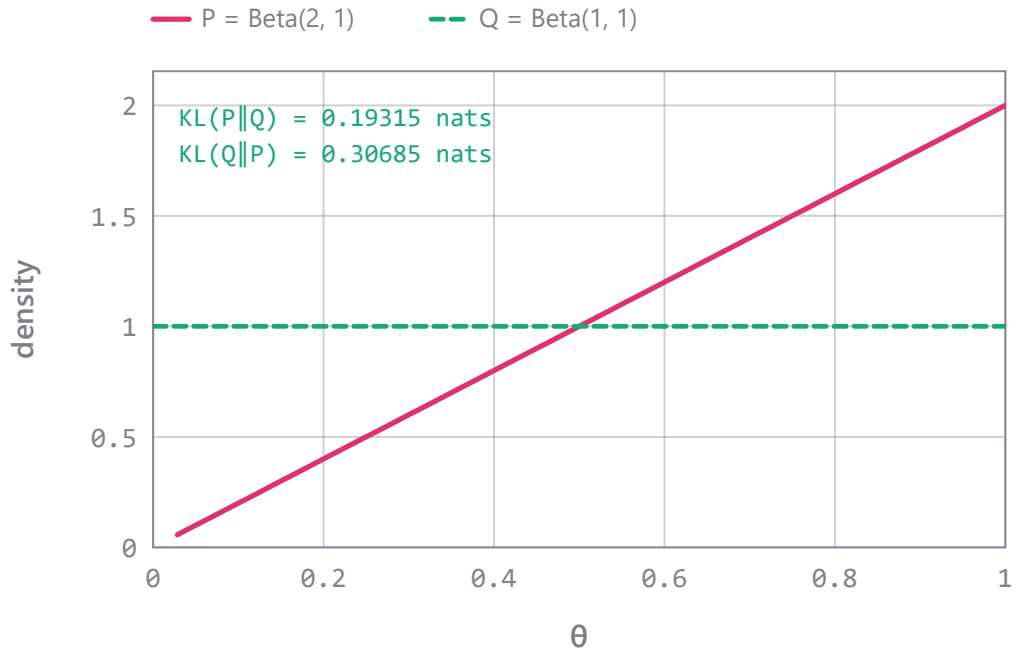
KL(Beta(alpha1, beta1) || Beta(alpha2, beta2)): for example the information gained from a beta prior (distribution 2) to a beta posterior (distribution 1).

$$KL = \ln \frac{B(\alpha_2, \beta_2)}{B(\alpha_1, \beta_1)} + (\alpha_1 - \alpha_2)\psi(\alpha_1) + (\beta_1 - \beta_2)\psi(\beta_1) + (\alpha_2 - \alpha_1 + \beta_2 - \beta_1)\psi(\alpha_1 + \beta_1)$$

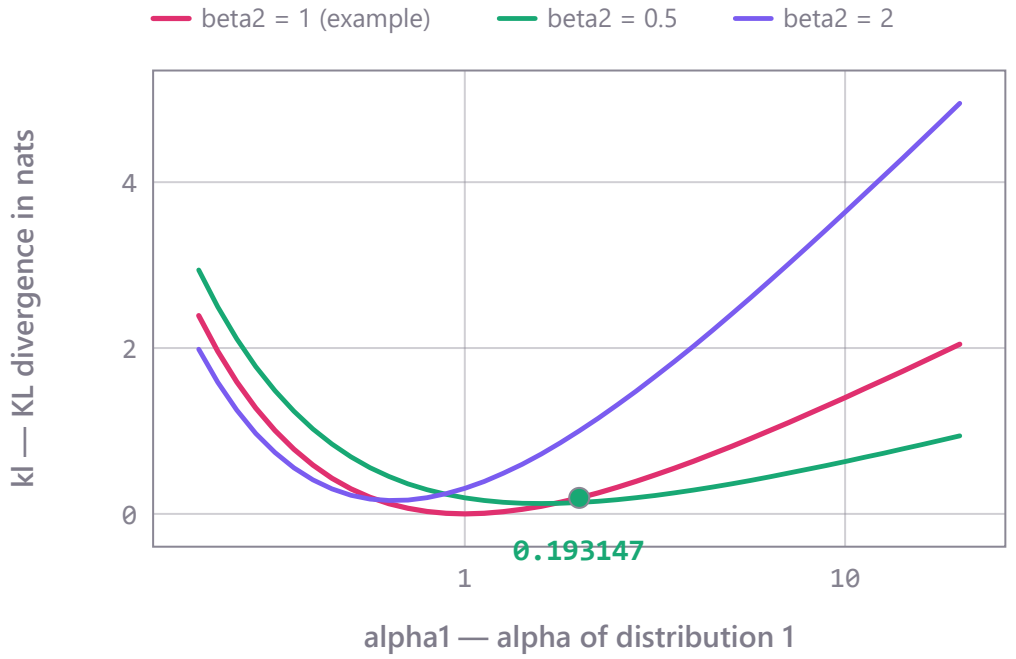
input	meaning	condition
`alpha1`	alpha of distribution 1	
`beta1`	beta of distribution 1	
`alpha2`	alpha of distribution 2	
`beta2`	beta of distribution 2	

output	meaning
`kl`	KL divergence in nats

- **Method:** closed form with log-beta and digamma functions
- **Assumes:** all parameters > 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



The two Beta densities of the worked example; each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Marked: the returned divergence KL(P‖Q) = 0.193147 nats and, with the two swapped, KL(Q‖P) = 0.306853 nats — the divergence is not symmetric.



kl (KL divergence in nats) against alpha1 (alpha of distribution 1) from 0.2 to 20 on a logarithmic axis. One curve for each value of beta2 (beta of distribution 2): 1, 0.5, 2. Other inputs as in the worked example: beta1 = 1, alpha2 = 1. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, alpha1 = 2, kl = 0.193147.

```
bayesian_statistics_ai999x.solve('kl_divergence_beta', alpha1=2, beta1=1, alpha2=1, beta2=1)
# -> kl = 0.193147   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('kl_divergence_beta', alpha1=8, beta1=4, alpha2=1, beta2=1)
# -> kl = 0.636614   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
KL divergence between beta distributions: alpha1 = 2, beta1 = 1, alpha2 = 1, beta2 = 1.	SOLVED	kl = 0.193147
KL divergence between beta distributions: alpha1 = 2, beta1 = 0, alpha2 = 1, beta2 = 1.	INVALID_INPUT	Kullback-Leibler divergence between two beta distributions: beta1 must be positive.

Gauss-Hermite rule for expectations under a normal distribution — `gauss_hermite_rule`

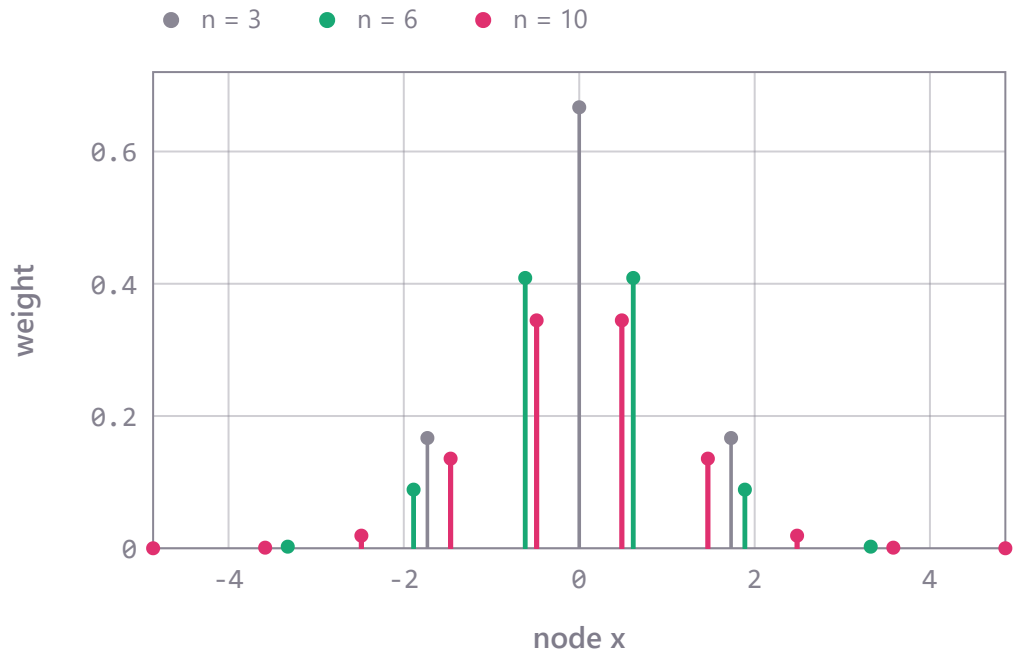
Nodes and weights of the n-point Gauss-Hermite rule for E[f(Z)], Z standard normal (probabilists' Hermite polynomials; weights sum to 1; exact for polynomials of degree up to 2n - 1). For N(mu, sigma^2) use mu + sigma x_i.

$$E[f(Z)] \approx \sum_{i=1}^{n} w_i f(x_i), \quad He_n(x_i) = 0, \quad w_i = \frac{n!}{n^2\, He_{n-1}(x_i)^2}$$

input	meaning	condition
`n`	number of nodes (1 to 40)	

output	meaning
`nodes`	nodes x_i (ascending)
`weights`	weights w_i (sum to 1)

- **Method:** Golub-Welsch: eigenvalues of the Jacobi matrix (implicit QL), Newton-polished on He_n; weights n!/(n^2 He_{$n-1$}(x_i)^2)
- **Assumes:** $1 \le n \le 40$
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Gauss–Hermite rules with n = 3, 6 and 10 nodes (inputs chosen; the worked example has n = 2): each node with its weight as a stem, all returned by the package (n = 3: nodes ±1.7321 at most, largest weight 0.6667, weights summing to 1; n = 6: nodes ±3.3243 at most, largest weight 0.4088, weights summing to 1; n = 10: nodes ±4.8595 at most, largest weight 0.3446, weights summing to 1, the sum being derived). Wider rules reach further into the tails with ever smaller weights.

```
bayesian_statistics_ai999x.solve('gauss_hermite_rule', n=2)
# -> nodes = [-1.0, 1.0], weights = [0.5, 0.5]   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('gauss_hermite_rule', n=3)
# -> nodes = [-1.73205080757, 0.0, 1.73205080757], weights = [0.166666666667, 0.666666666667,
0.166666666667]   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Gauss-Hermite rule with n = 3 nodes for a normal expectation.	SOLVED	nodes = [-1.73205080757, 0.0, 1.73205080757], weights = [0.166666666667, 0.666666666667, 0.166666666667]
Gauss-Hermite rule with n = 0 nodes.	INVALID_INPUT	n: must be at least 1, got 0.

Exact mean and variance of a polynomial of a normal variable — `polynomial_of_normal_moments`

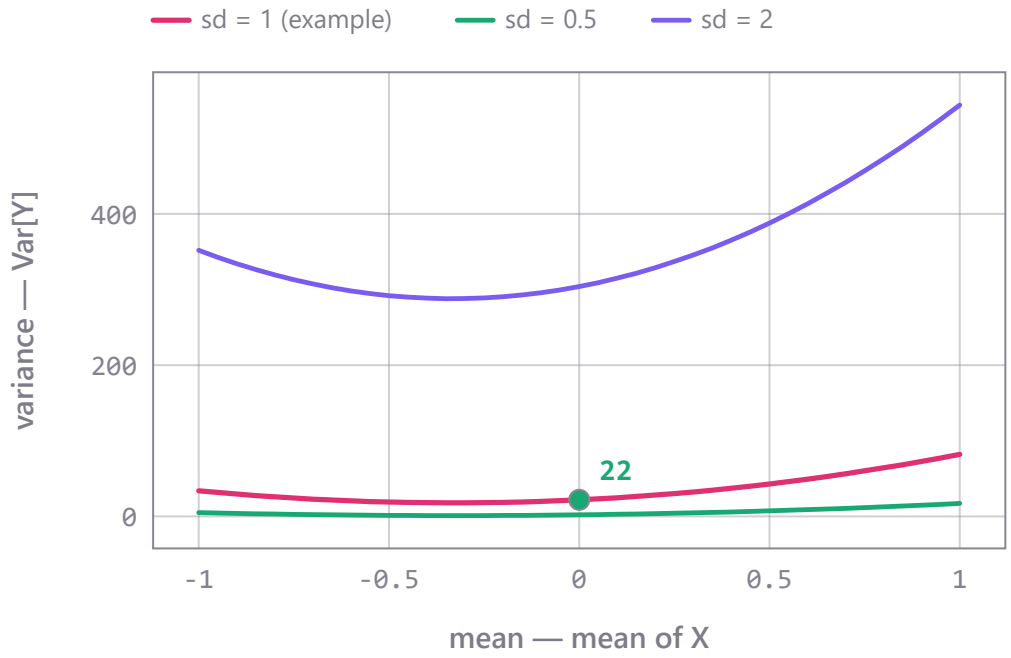
Exact mean, variance and standard deviation of $Y = c_0 + c_1\,X + \ldots + c_6\,X^6$ for X ~ N(mean, sd^2) (nonlinear uncertainty propagation without linearization).

$$E[Z^{2m}] = (2m-1)!!, \quad X = \mu + \sigma Z$$

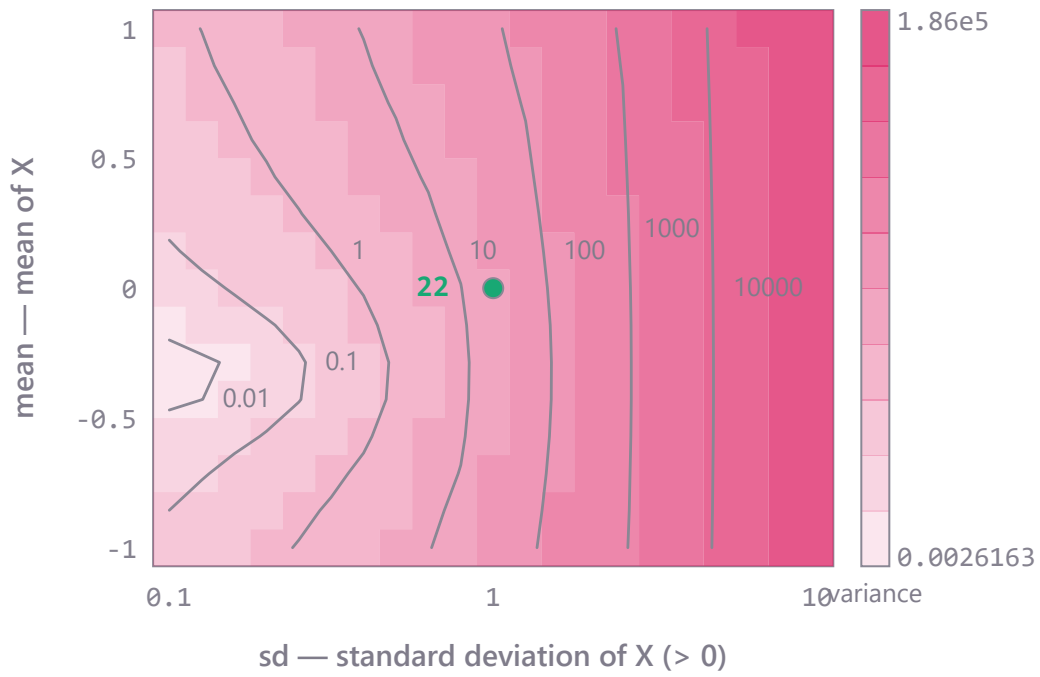
input	meaning	condition
`coefficients`	coefficients c0, c1, ... (constant first, up to degree 6)	
`mean`	mean of X	
`sd`	standard deviation of X (> 0)	

output	meaning
`mean`	E[Y]
`variance`	Var[Y]
`sd`	sd of Y

- **Method:** substitute X = mean + sd Z and use the normal moments E[Z^k] = $(k-1)$!! for even k
- **Assumes:** $sd > 0$; degree <= 6
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



variance (Var[Y]) against mean (mean of X) from -1 to 1. One curve for each value of sd (standard deviation of X (> 0)): 1, 0.5, 2. Other inputs as in the worked example: coefficients = 1, 2, 3. 41 points per curve computed by bayesian-statistics-ai999x; marked: the worked example, mean = 0, variance = 22.



variance (Var[Y]) over sd (standard deviation of X (> 0)) from 0.1 to 10 and mean (mean of X) from -1 to 1 (logarithmic horizontal axis): each of the 21 × 15 cells is one run of bayesian-statistics-ai999x, shaded by its value (light 0.00261633 to dark 186400, logarithmic shading); iso-lines at variance = 0.01, 0.1, 1, 10, 100, 1000, 10000. Other inputs as in the worked example: coefficients = 1, 2, 3. Marked: the worked example, sd = 1, mean = 0, variance = 22.

```
bayesian_statistics_ai999x.solve('polynomial_of_normal_moments', coefficients='1, 2, 3', mean=0, sd=1)
# -> mean = 4, variance = 22, sd = 4.69042   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('polynomial_of_normal_moments', coefficients=[0, 0, 1], mean=2, sd=0.5)
# -> mean = 4.25, variance = 4.125, sd = 2.03101   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Mean and variance of the polynomial with coefficients = [1, 2, 3] of a normal X with mean = 0, sd = 1.	SOLVED	mean = 4, variance = 22, sd = 4.69042
Mean and variance of the polynomial with coefficients = [1, 2, 3] of a normal X with mean = 0, sd = 0.	INVALID_INPUT	Exact mean and variance of a polynomial of a normal variable: the standard deviation must be positive.

Probability that one Poisson rate exceeds another (gamma posteriors) — `probability_rate_greater_gamma`

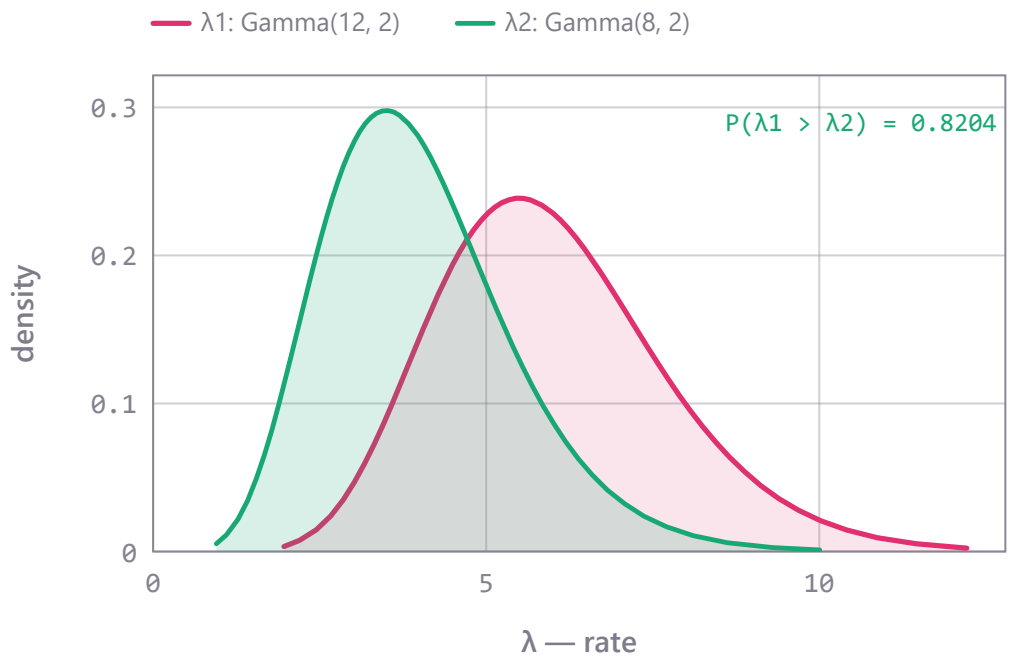
P(lambda_1 > lambda_2) for independent Gamma(shape, rate) posteriors of two Poisson rates, and the ratio of their posterior means.

$$P(\lambda_1 > \lambda_2) = 1 - I_{b_1/(b_1+b_2)}(a_1, a_2)$$

input	meaning	condition
`shape1`	posterior shape of rate 1	
`rate1`	posterior rate parameter of rate 1	
`shape2`	posterior shape of rate 2	
`rate2`	posterior rate parameter of rate 2	

output	meaning
`prob_1_greater`	P(lambda_1 > lambda_2)
`prob_2_greater`	P(lambda_2 > lambda_1)
`rate_ratio_mean`	ratio of posterior means

- **Method:** lambda_1 b_1 / (lambda_1 $b_1 + lambda_2$ b_2) ~ Beta(shape1, shape2): one regularized incomplete beta
- **Assumes:** all parameters > 0
- **Evidence:** certified procedure (gates: conditions, control, examples, fast_enough, questions, reference, sensitivity, templates)



Two gamma posteriors for rates (inputs chosen: shape 12, rate 2 and shape 8, rate 2); each density is drawn from the package's quantiles at 60 cumulative probabilities (the probability of each small interval divided by its width). Marked: the returned P(λ1 > λ2) = 0.820358, and the posterior mean of λ1/λ2 = 1.5.

```
bayesian_statistics_ai999x.solve('probability_rate_greater_gamma', shape1=5, rate1=1, shape2=5,
rate2=1)
# -> prob_1_greater = 0.5, prob_2_greater = 0.5, rate_ratio_mean = 1   (SOLVED)
```

```
bayesian_statistics_ai999x.solve('probability_rate_greater_gamma', shape1=1, rate1=1, shape2=1,
rate2=2)
# -> prob_1_greater = 0.666667, prob_2_greater = 0.333333, rate_ratio_mean = 2   (SOLVED)
```

Questions in words (documented forms), as executed:

question	outcome	result
Compare gamma posteriors shape1 = 5, rate1 = 1, shape2 = 5, rate2 = 1.	SOLVED	prob_1_greater = 0.5, prob_2_greater = 0.5, rate_ratio_mean = 1
Compare gamma posteriors shape1 = 1, rate1 = 1, shape2 = 1, rate2 = 2.	SOLVED	prob_1_greater = 0.666667, prob_2_greater = 0.333333, rate_ratio_mean = 2
Compare gamma posteriors shape1 = 5, rate1 = 0, shape2 = 5, rate2 = 1.	INVALID_INPUT	Probability that one Poisson rate exceeds another (gamma posteriors): rate1 must be positive.